

Computer Vision

Bogdan Alexe

bogdan.alexe@fmi.unibuc.ro

University of Bucharest, 2nd semester, 2023-2024

Course structure

1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

4. Object Recognition: high – level vision

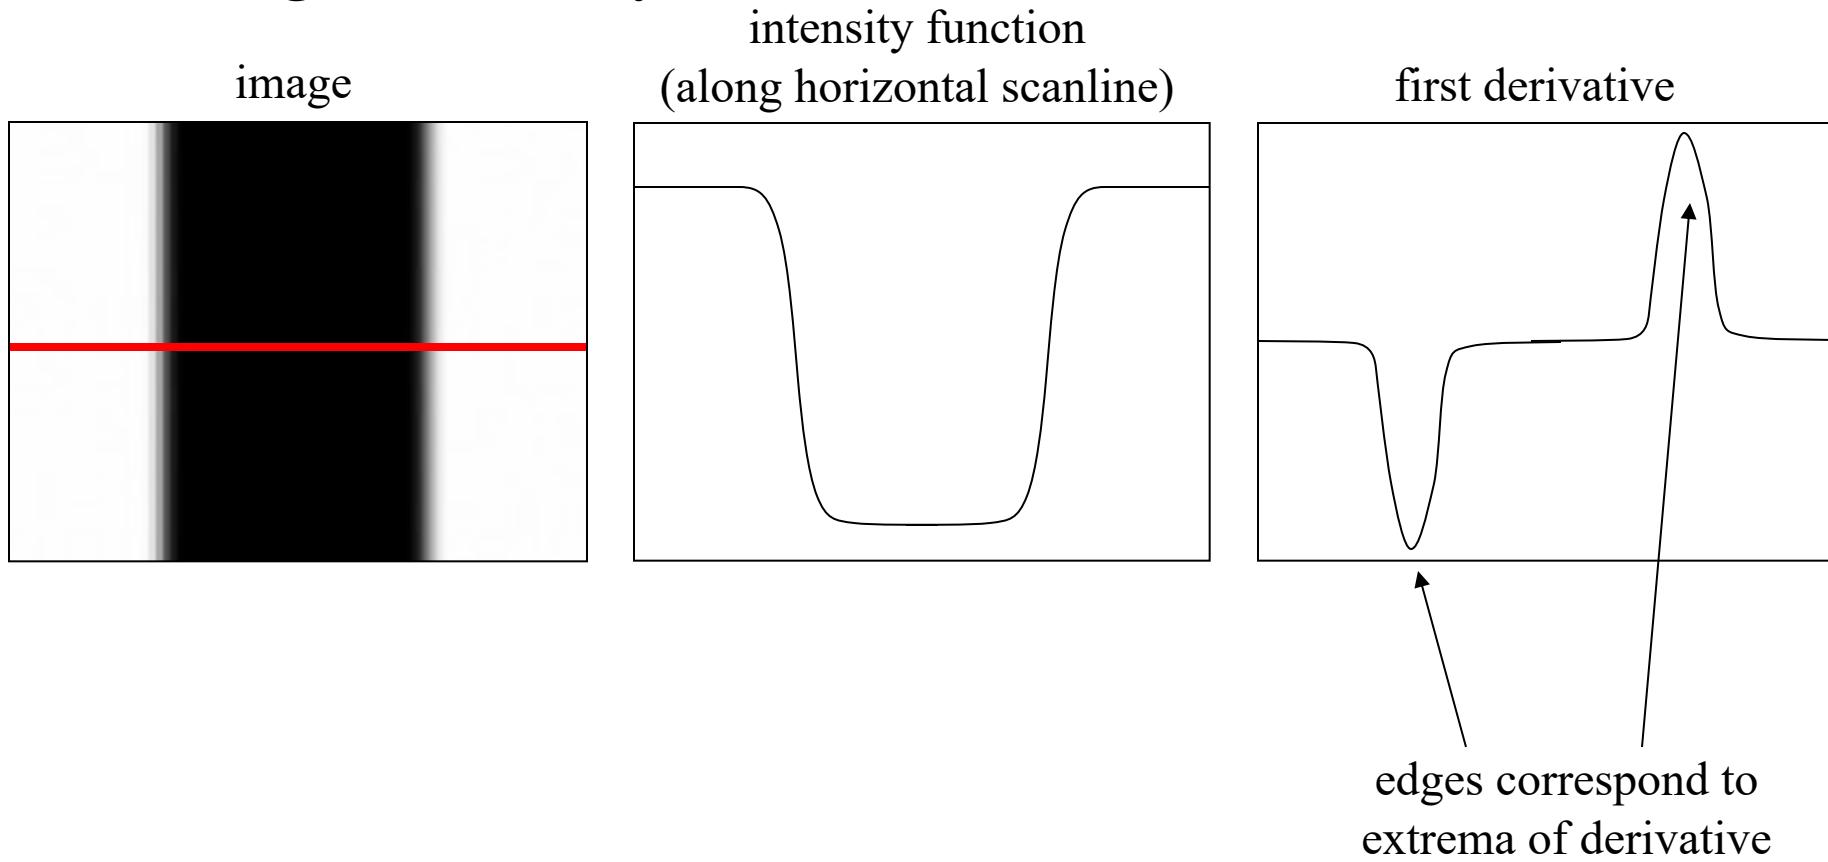
Object classification, object detection, part based models, bovw models

5. Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

Characterizing edges

- An edge is a place of rapid change in the image intensity function



Derivatives with convolution

For 2D function, $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

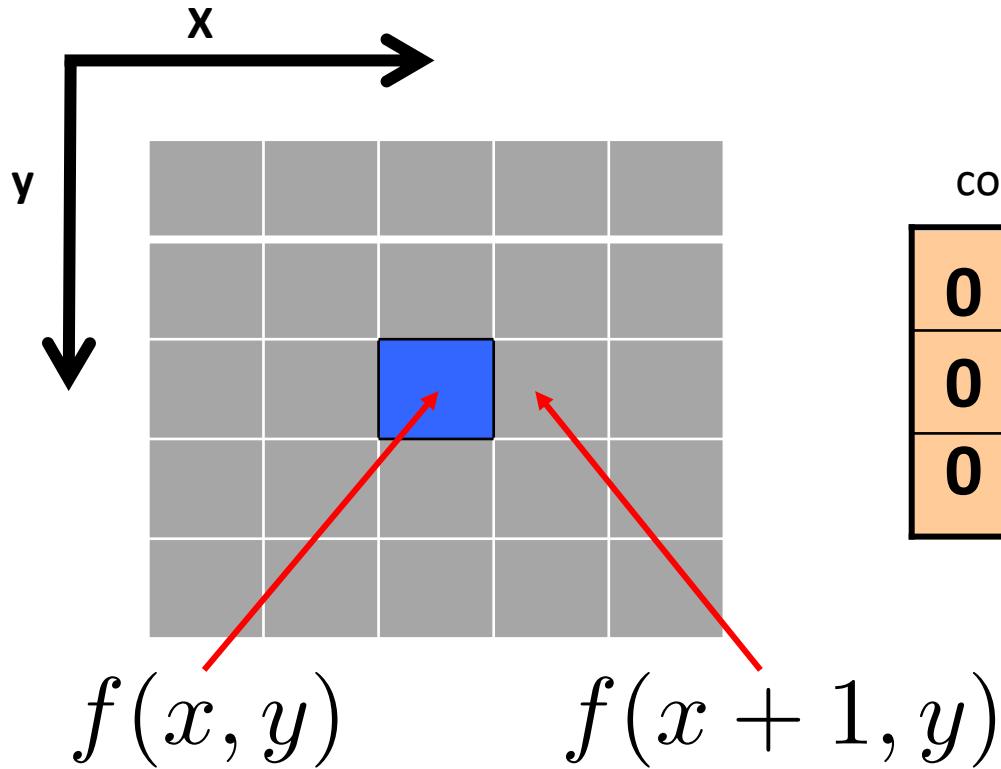
For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

To implement above as convolution, what would be the associated filter?

Derivatives with convolution

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$



correlation			convolution		
0	?	0	1	-1	
0	?	1	0	0	
0	0	0			

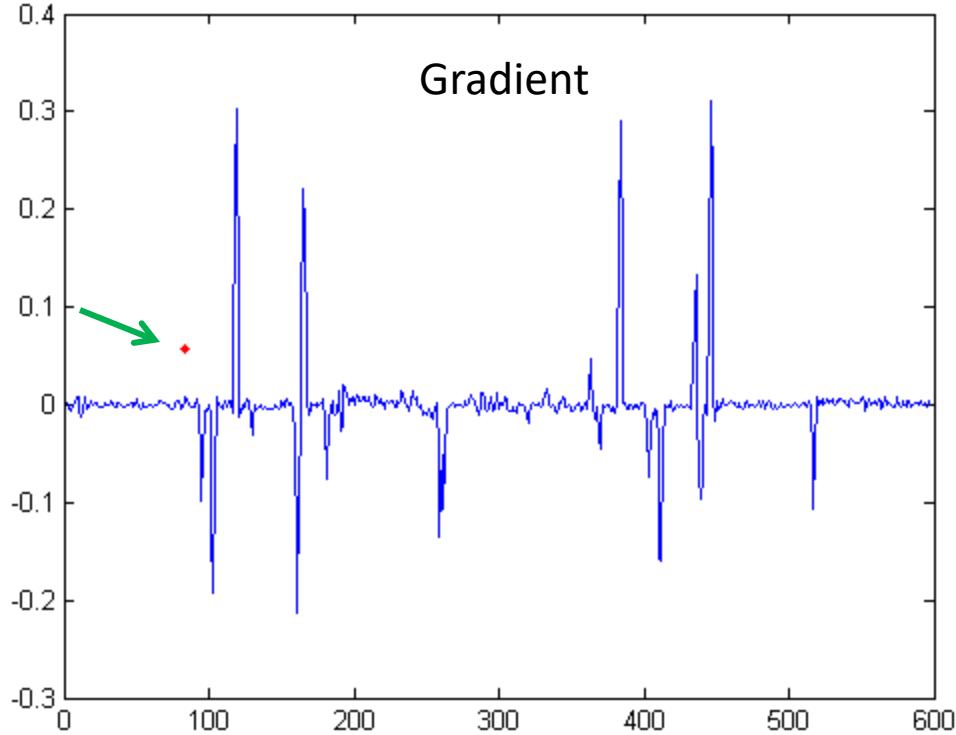
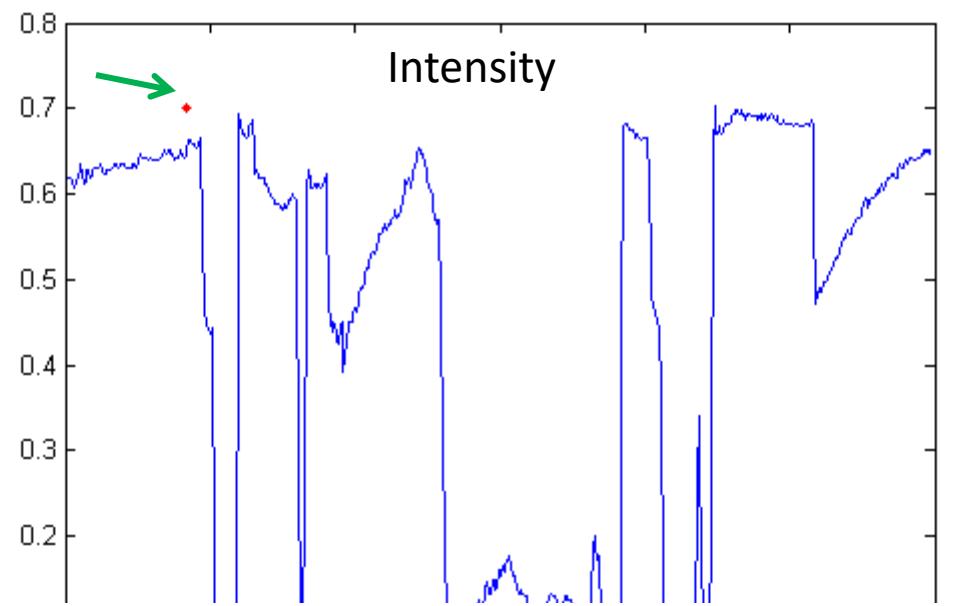
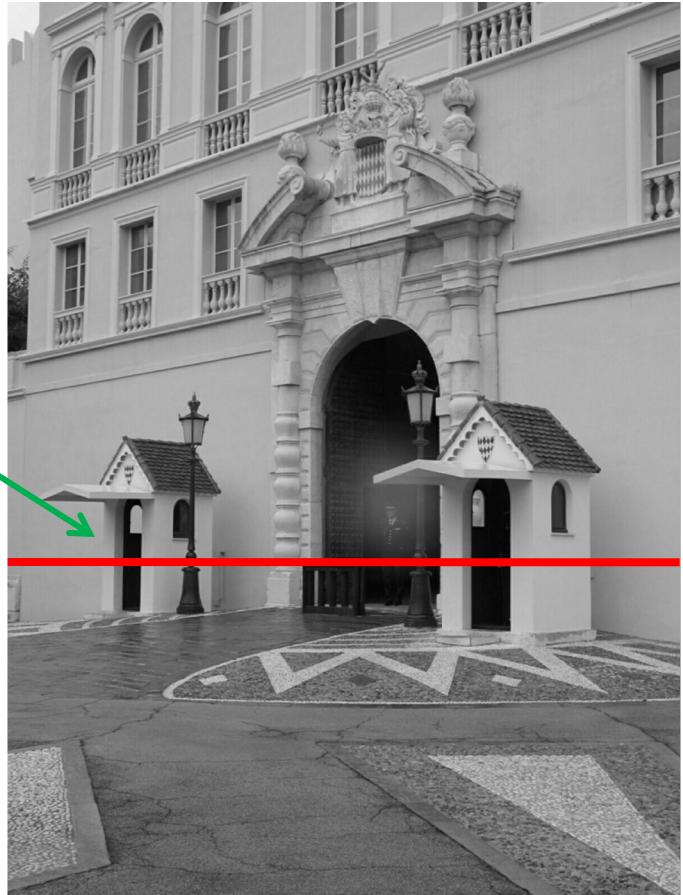
Other filters for finite differences

Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

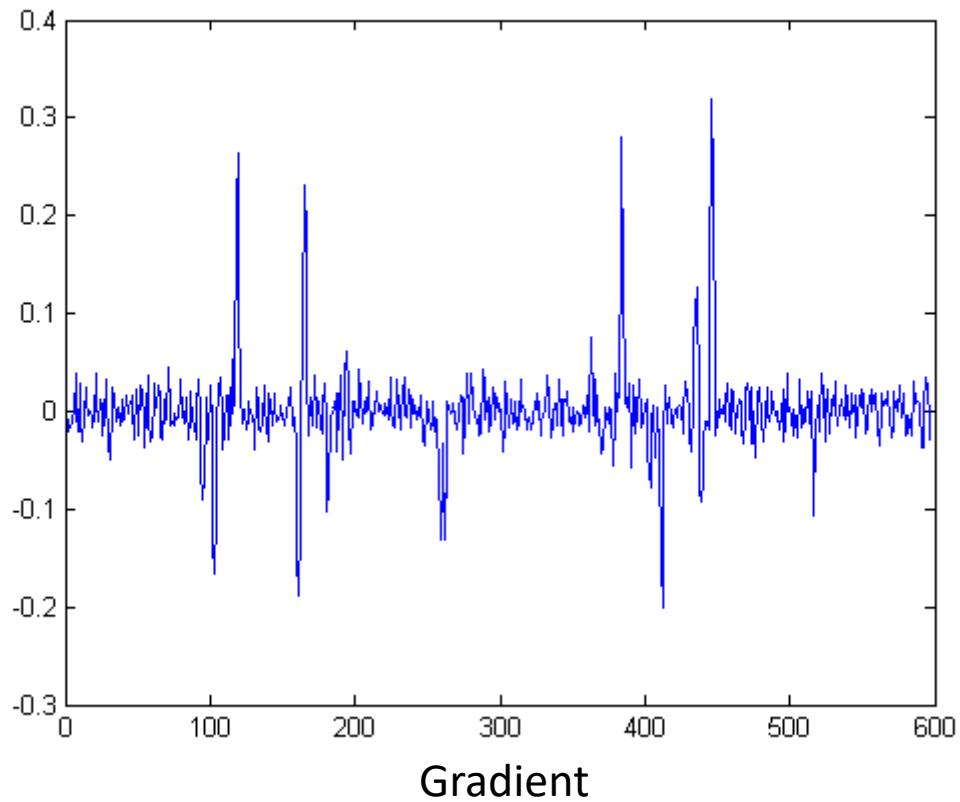
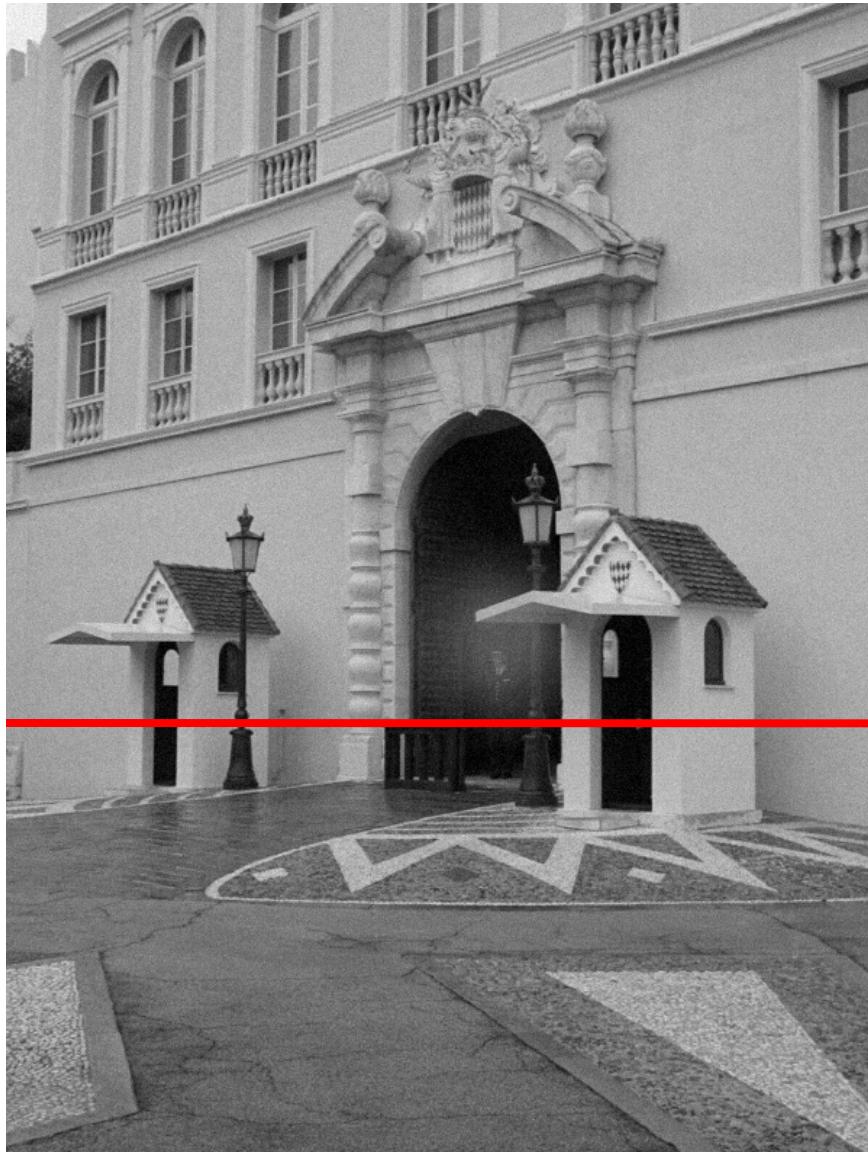
Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

Intensity profile

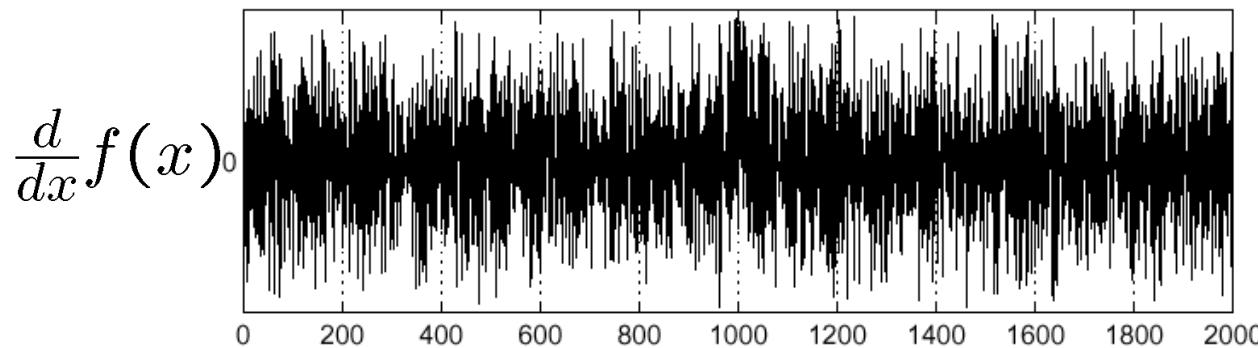
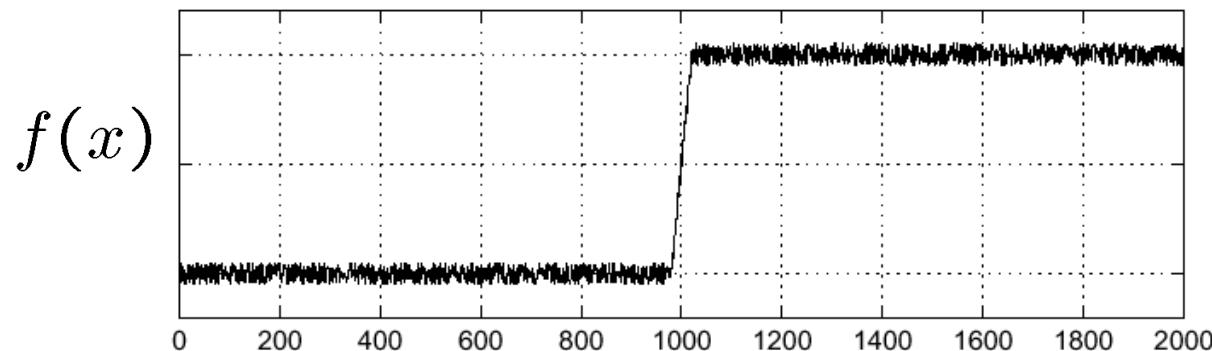


With a little Gaussian noise



Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal

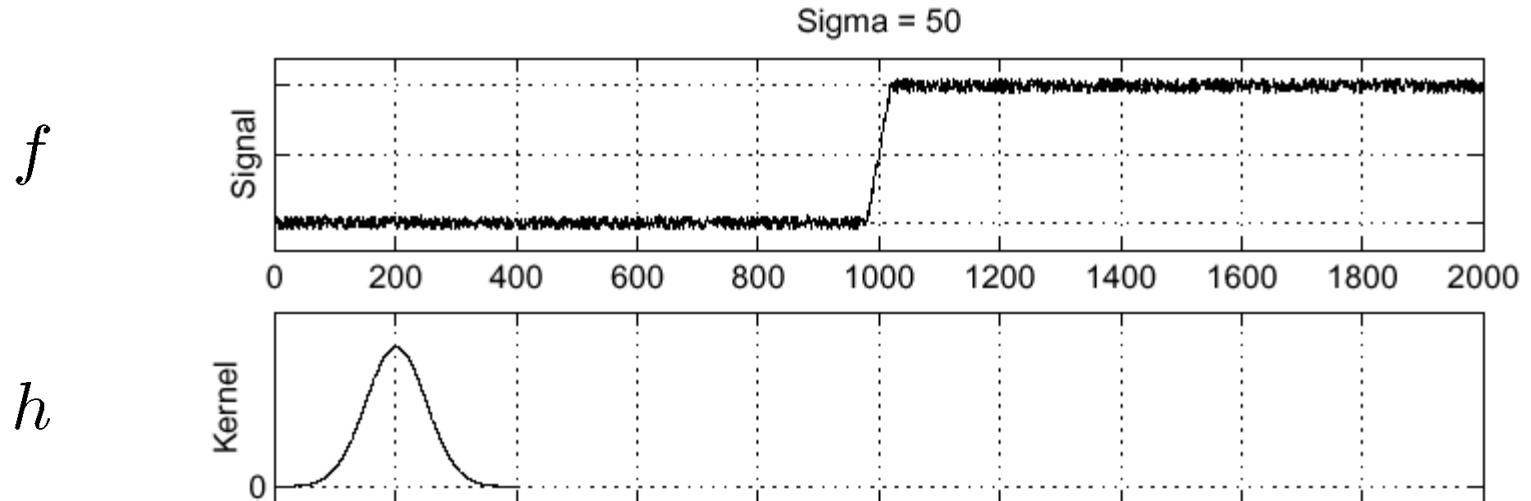


Where is the edge?

Effects of noise

- Difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbors
 - Generally, the larger the noise the stronger the response
- What can we do about it?

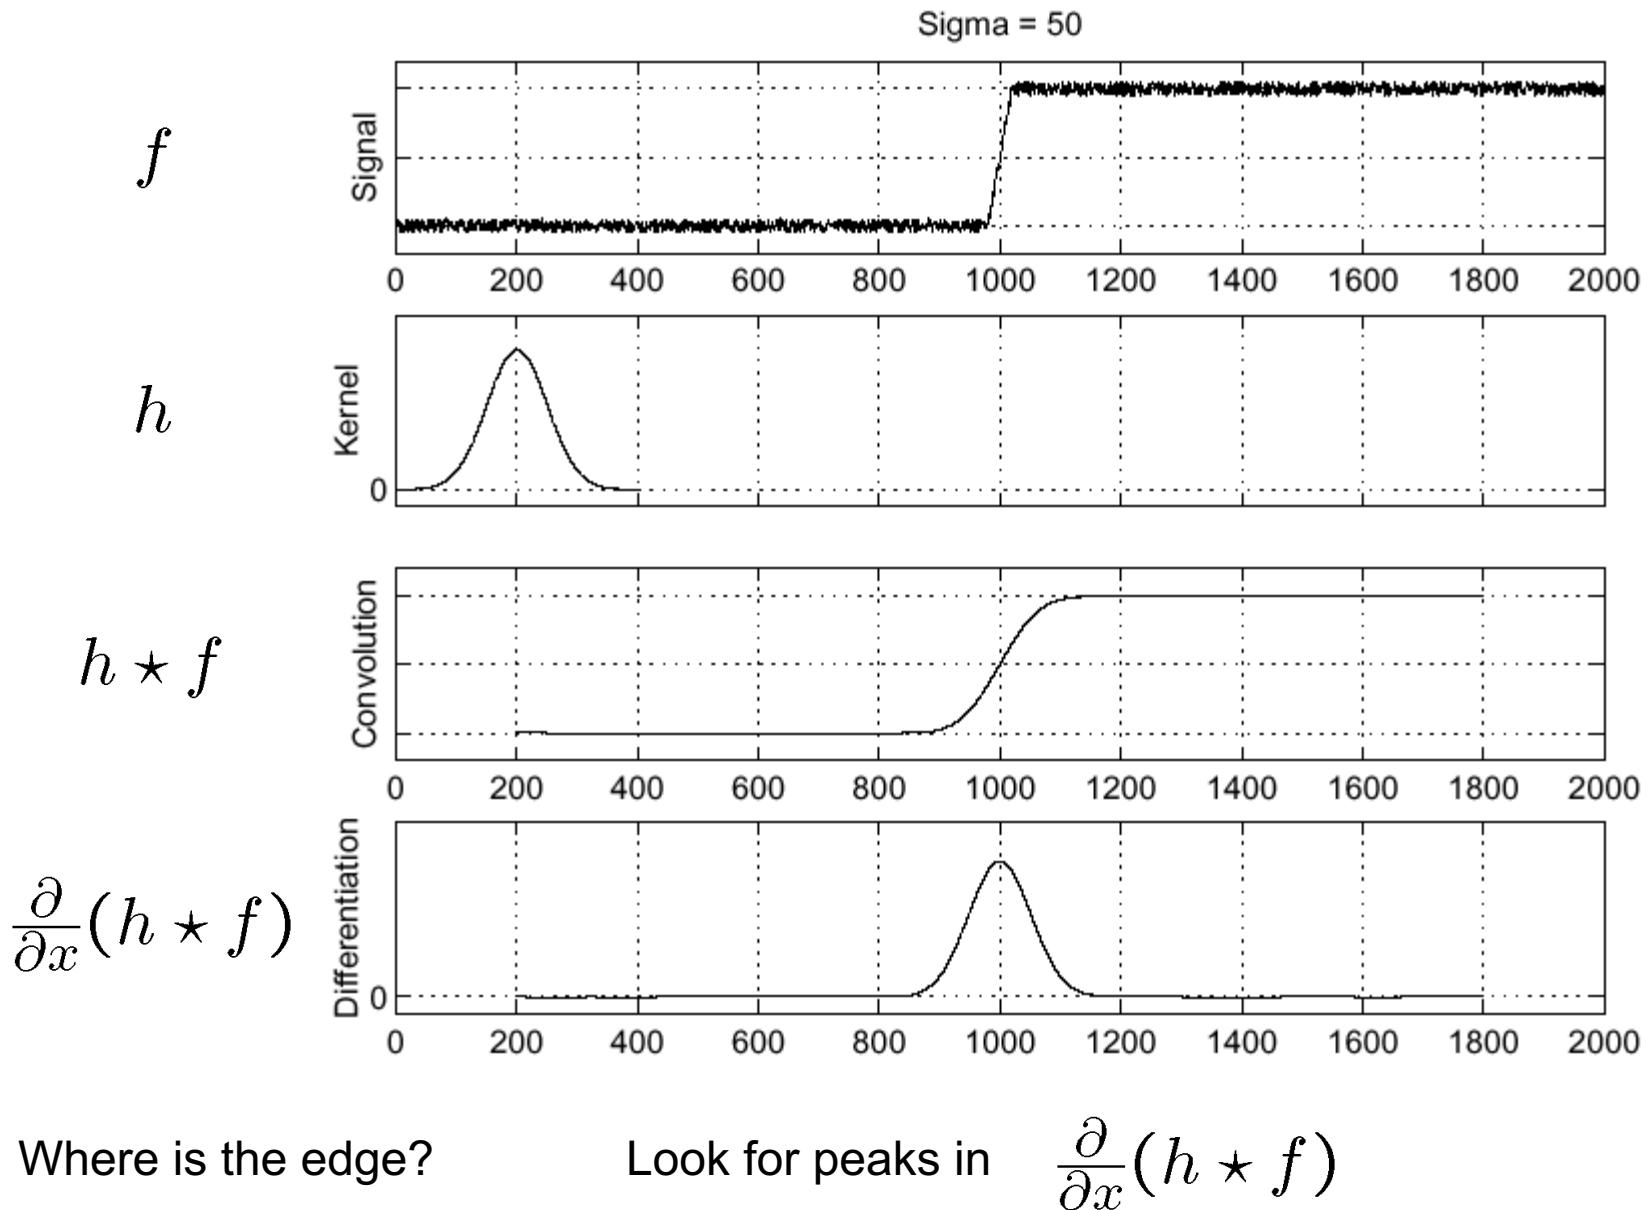
Solution: smooth first



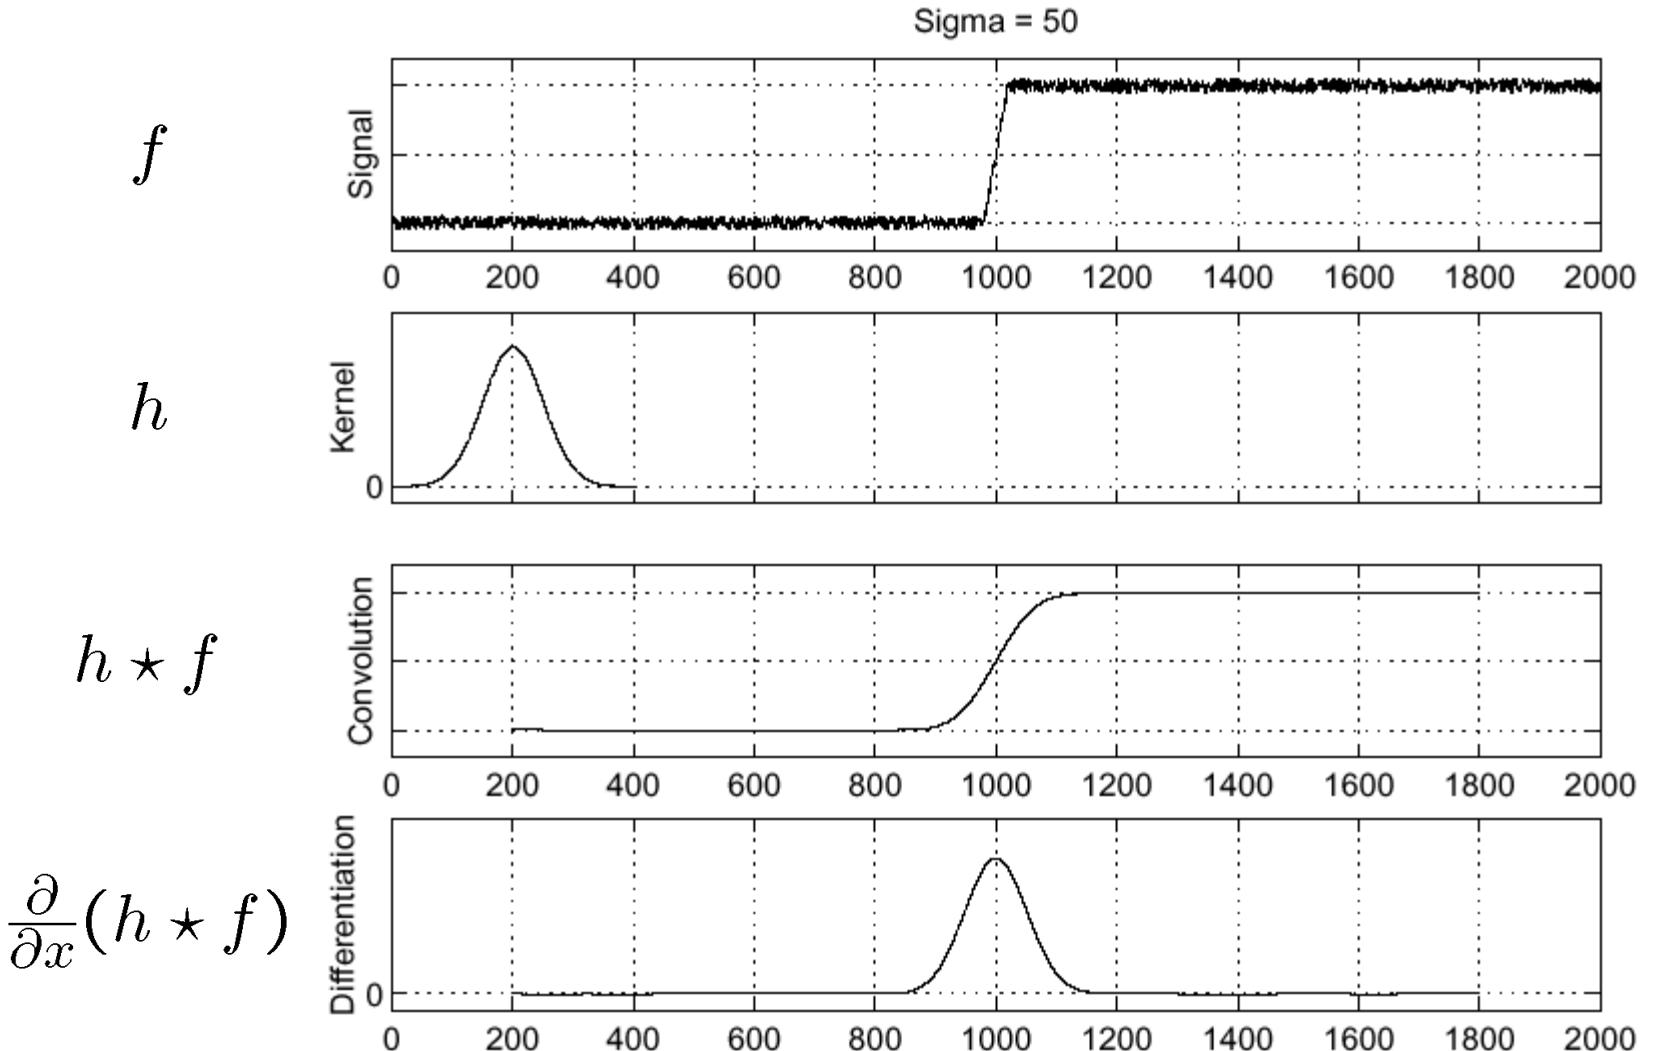
Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Solution: smooth first



Solution: smooth first



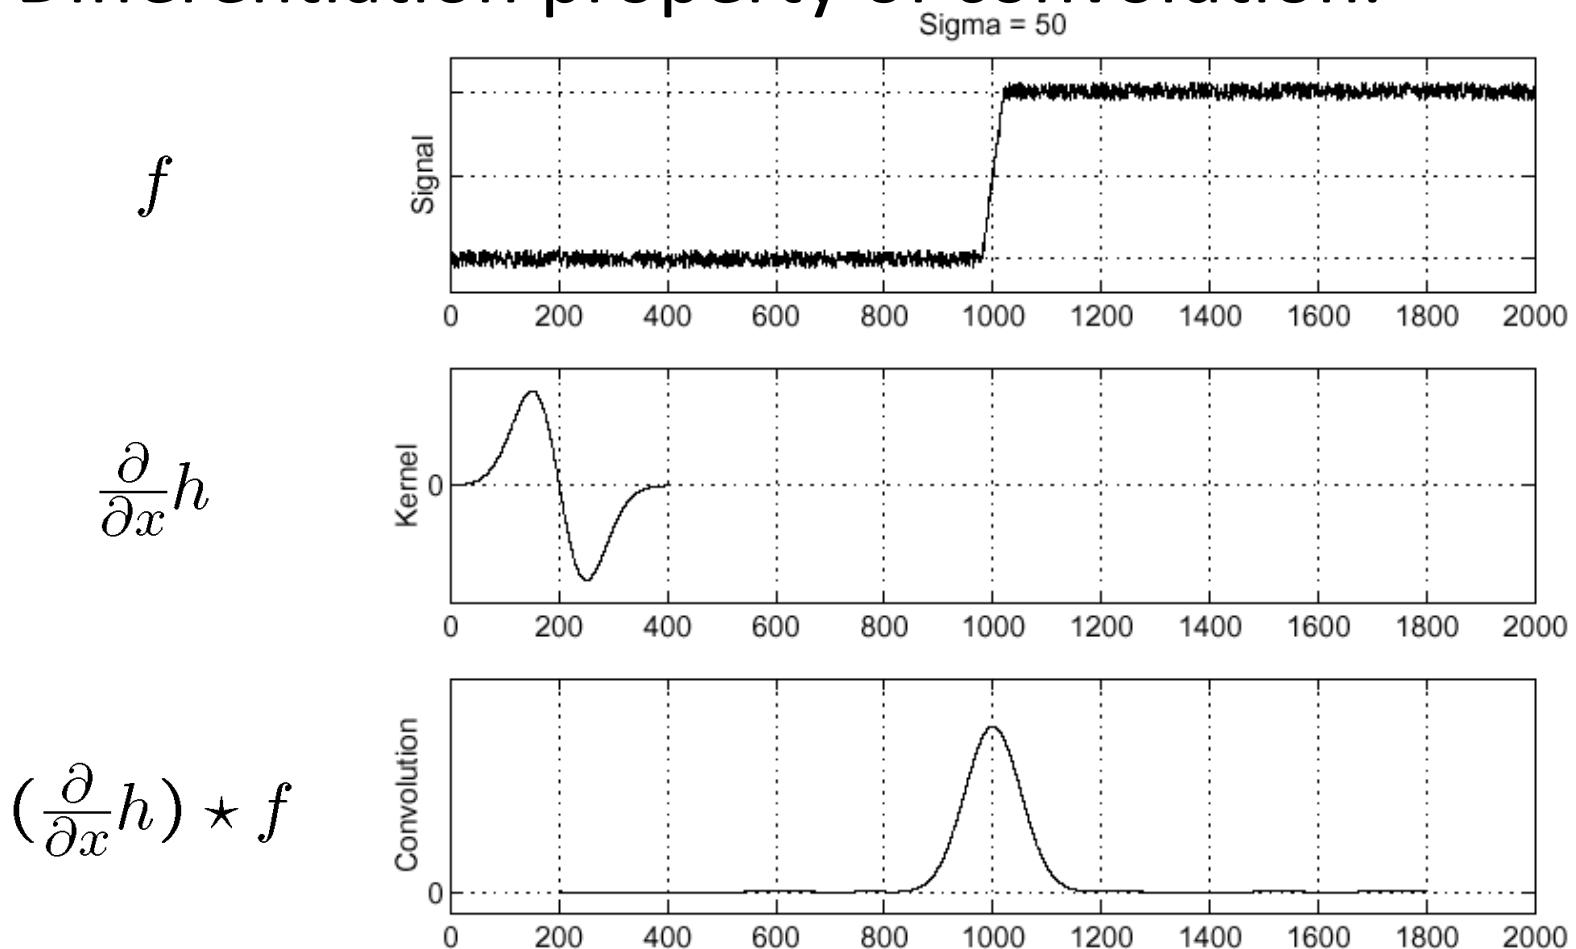
Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h \star f)$ = zero's of the $\frac{\partial^2}{\partial x^2}(h \star f)$ function

Derivative theorem of convolution

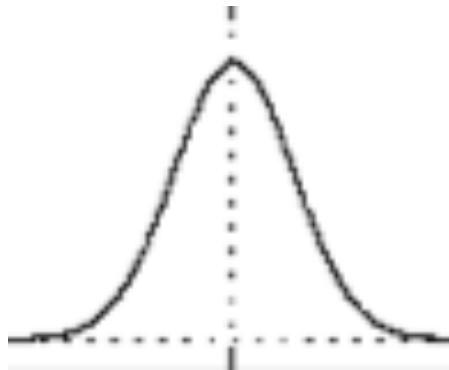
$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

Differentiation property of convolution.

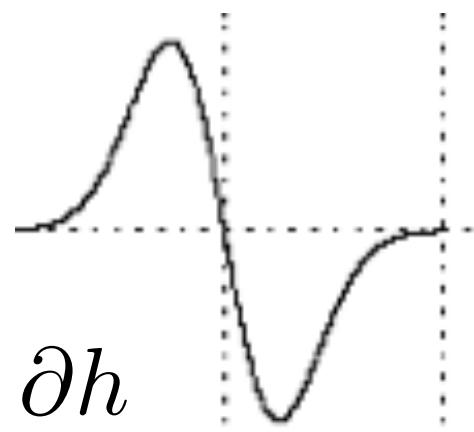


1D

-1	1
----	---



=



d_x

h

$\frac{\partial h}{\partial x}$

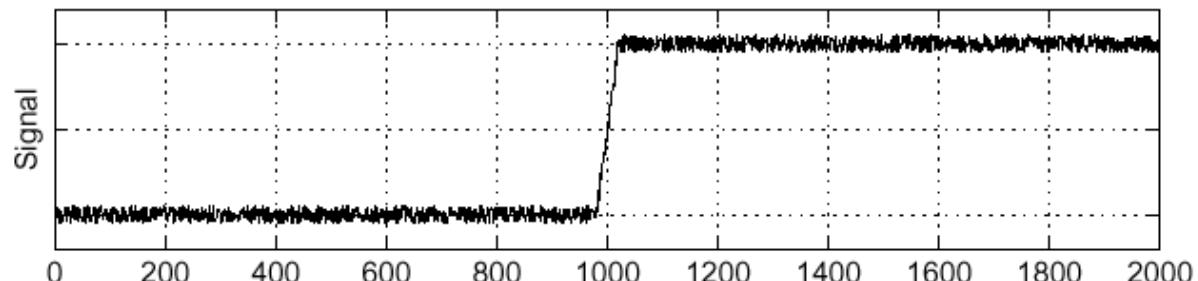
Laplacian of a Gaussian – 1D

Consider

$$\frac{\partial^2}{\partial x^2}(h \star f)$$

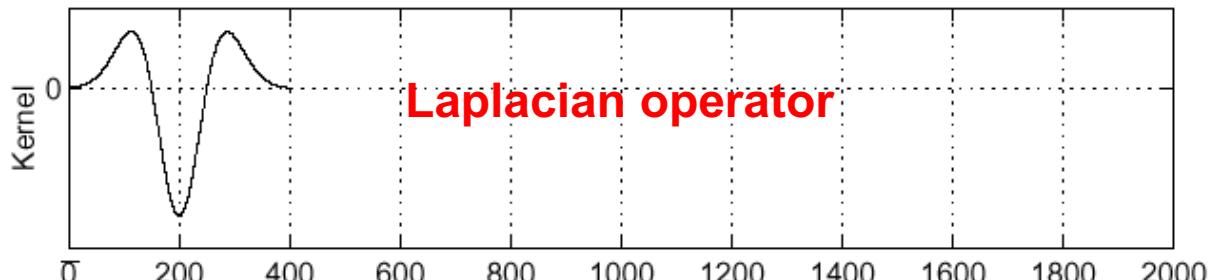
f

Sigma = 50

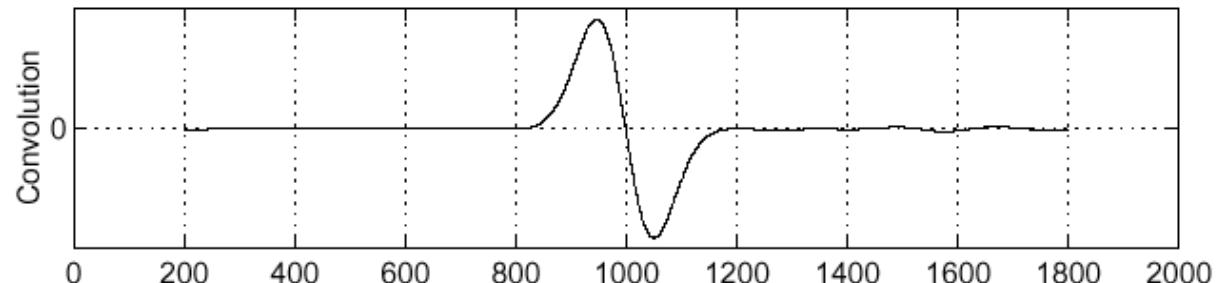


$$\frac{\partial^2}{\partial x^2} h$$

Laplacian operator



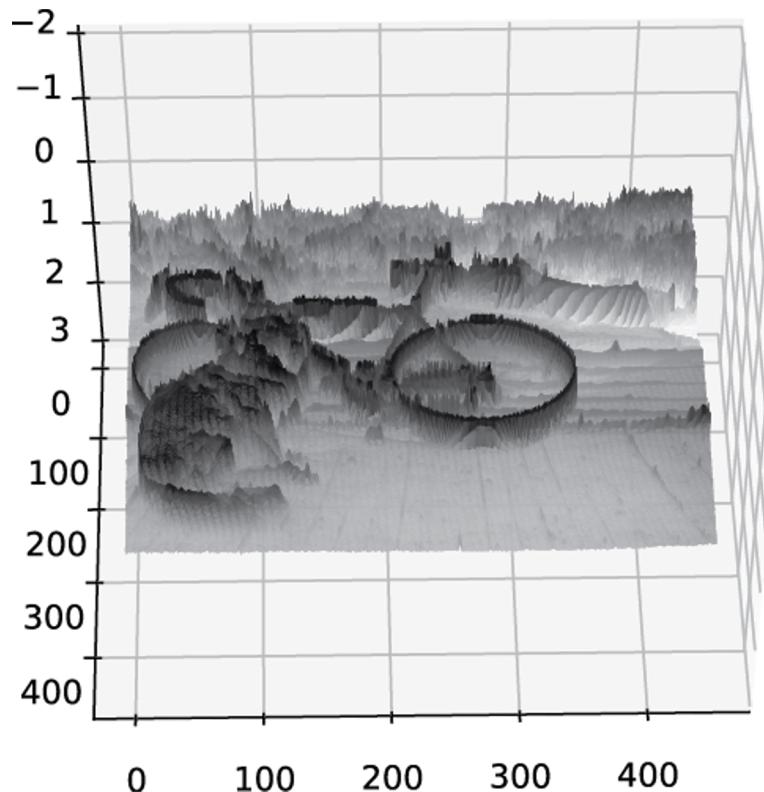
$$(\frac{\partial^2}{\partial x^2} h) \star f$$



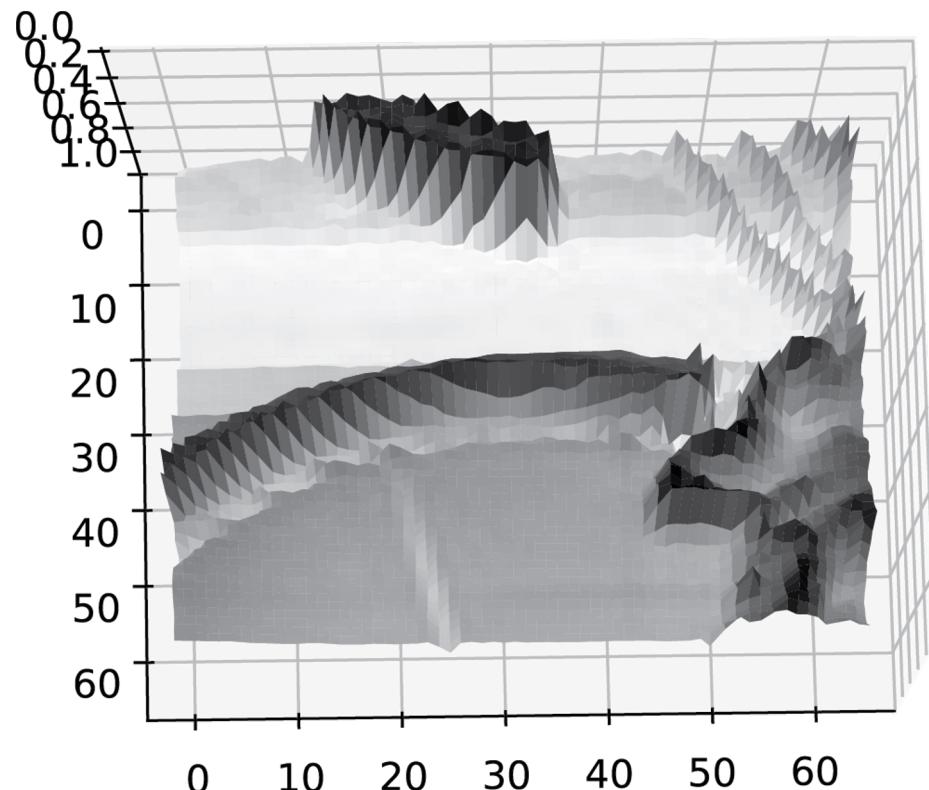
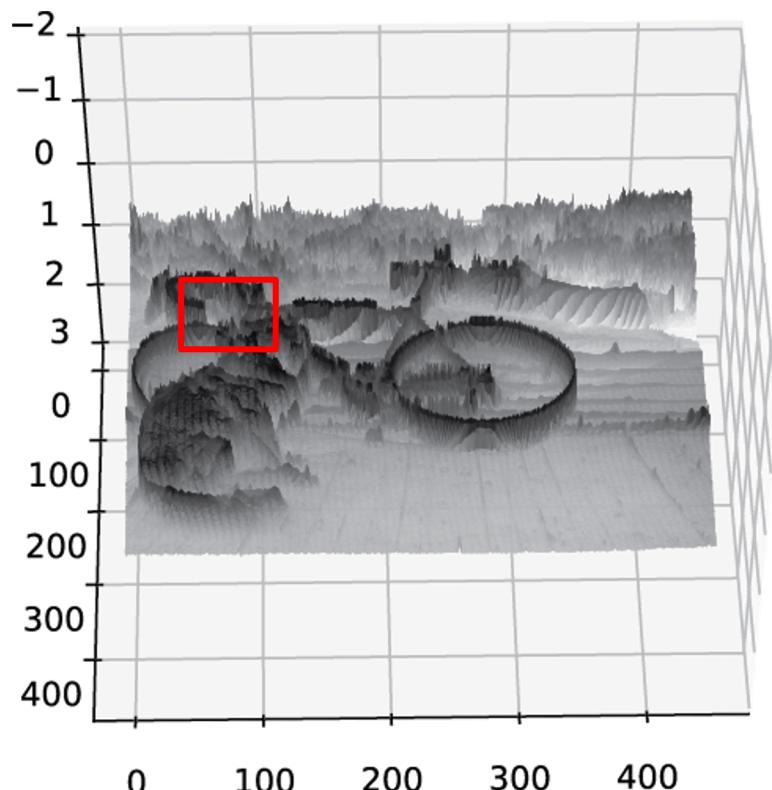
Where is the edge?

Zero-crossing of the function $(\frac{\partial^2}{\partial x^2} h) \star f$

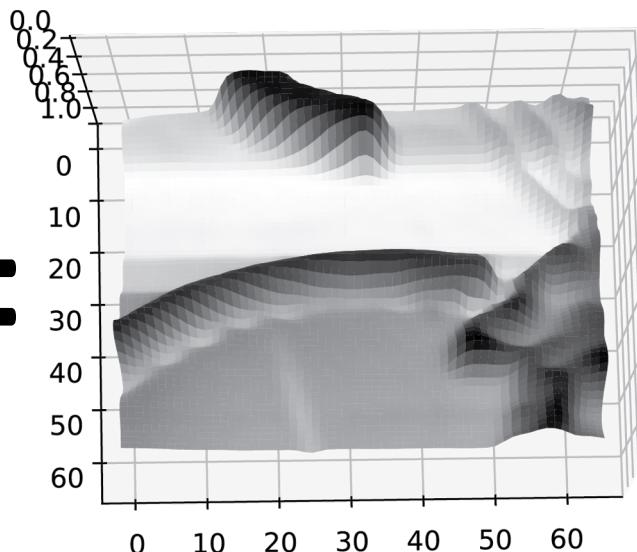
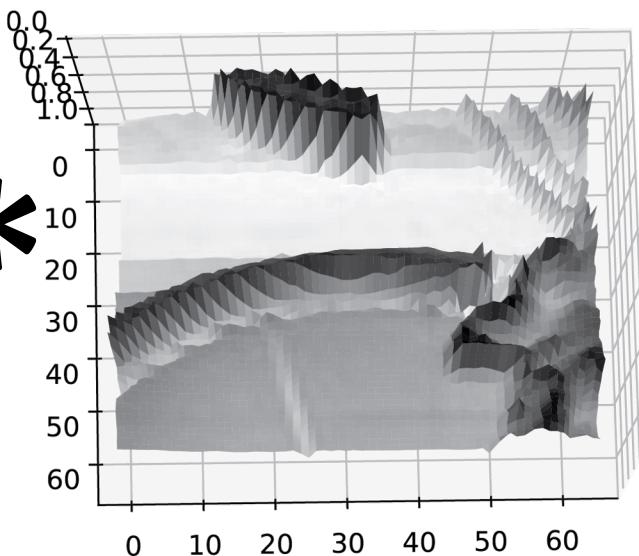
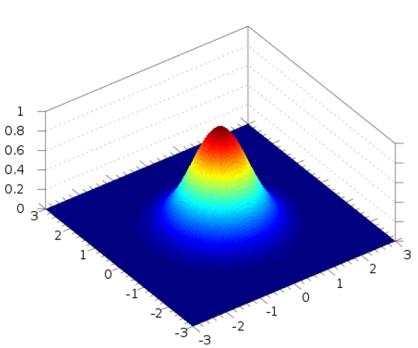
2D: Images are noisy



2D: Images are noisy



Solution: smooth first



Smooth first, then derivative

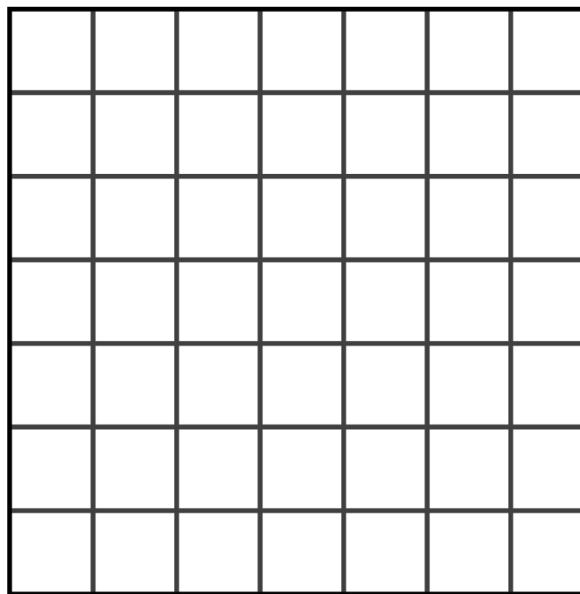
$$\frac{1}{2} \times (-1 \ 0 \ 1) \ast ($$

derivative filter

$$(\begin{array}{ccc} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{array})$$

Gaussian filter

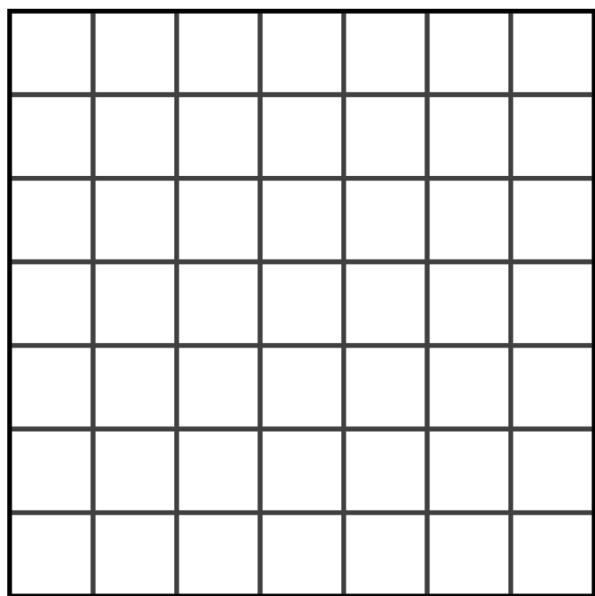
\ast



)

Smooth first, then derivative

$$\frac{1}{2} \times (-1 \quad 0 \quad 1) * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} *$$



Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$ \rightarrow $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$ \rightarrow $\frac{1}{2} \times \begin{bmatrix} 2 & & \\ & & \\ & & \end{bmatrix}$

The diagram illustrates a convolution operation. A 1x3 kernel $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$ is multiplied by a 3x3 input matrix $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$. The result is scaled by $\frac{1}{2}$. A pink arrow points from the kernel to the input matrix, indicating the receptive field of the first element in the kernel.

Smooth first, then derivative

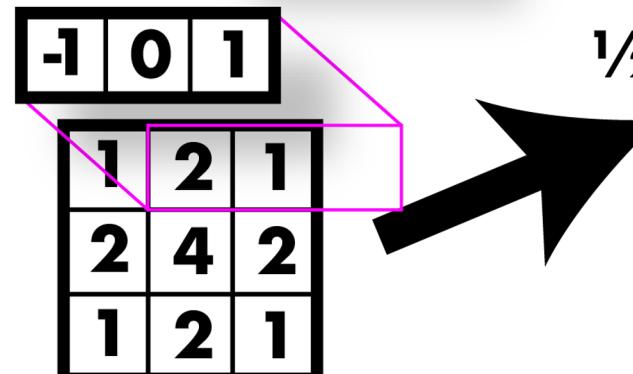
$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$ \rightarrow $\frac{1}{2} \times \begin{bmatrix} 2 & 0 & \\ & & \\ & & \end{bmatrix}$

The diagram illustrates a convolution operation. A 1x3 kernel $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$ is applied to a 3x3 input matrix. The result is scaled by $\frac{1}{2}$. The highlighted element in the input matrix is 2.

Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

$\frac{1}{2} \times$  $\begin{bmatrix} 2 & 0 & -2 \end{bmatrix}$

Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

The diagram illustrates a convolution operation. A 1x3 kernel $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$ is applied to a 3x3 input matrix $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$. The result is a 1x3 output matrix $\begin{bmatrix} 2 & 0 & -2 \end{bmatrix}$. The input matrix is highlighted with a pink box, and the output matrix is also highlighted with a pink box.

Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

The diagram illustrates a convolution operation. A 3x3 input matrix is multiplied by a 3x3 kernel matrix, with the result being scaled by $\frac{1}{2}$. The input matrix has values 1, 2, 1 in its top row, 2, 4, 2 in its middle row, and 1, 2, 1 in its bottom row. The kernel matrix has values -1, 0, 1 in its top row, 2, 4, 2 in its middle row, and 1, 2, 1 in its bottom row. The result is a 2x2 output matrix with values 2, 0, -2, 4, 0 in its top row and two empty cells in its bottom row. A pink arrow points from the input matrix to the kernel matrix, indicating the receptive field of the output unit at position (0,0).

$$\frac{1}{2} \times \begin{bmatrix} 2 & 0 & -2 \\ 4 & 0 & \end{bmatrix}$$

Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

The diagram illustrates a two-step convolution process. In the first step, a 3x3 input matrix is convolved with a 1x3 kernel. The result of this step is highlighted with a pink box. In the second step, this intermediate result is convolved with a 3x1 kernel. The final output is scaled by $\frac{1}{2}$.

Input Matrix:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

Kernel 1 (1x3):

$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$$

Intermediate Result (3x1 highlighted in pink):

$$\begin{bmatrix} 2 & 4 & 2 \end{bmatrix}$$

Kernel 2 (3x1):

$$\begin{bmatrix} 2 & 0 & -2 \\ 4 & 0 & -4 \\ \vdots & \vdots & \vdots \end{bmatrix}$$

Final Output (scaled by $\frac{1}{2}$):

$$\frac{1}{2} \times \begin{bmatrix} 2 & 0 & -2 \\ 4 & 0 & -4 \\ \vdots & \vdots & \vdots \end{bmatrix}$$

Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

The diagram illustrates a convolution operation. A 1x3 kernel $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$ is applied to a 3x3 input matrix $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$. The result is scaled by $\frac{1}{2}$. A pink arrow points from the input matrix to the output matrix $\begin{bmatrix} 2 & 0 & -2 \\ 4 & 0 & -4 \\ 2 & & \end{bmatrix}$.

Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

The diagram illustrates a convolution operation. A 3x3 input matrix (smoothed version) is multiplied by a 1x3 kernel (derivative operator). The result is a 3x1 output vector.

Input Matrix (Smoothed):

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

Kernel (Derivative):

$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$$

Output Vector:

$$\begin{bmatrix} 2 \\ 4 \\ 2 \end{bmatrix}$$

Smooth first, then derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$

The diagram illustrates a convolution operation. A 3x3 input matrix (smoothed image) is multiplied by a 1x3 kernel (derivative filter). The result is a 3x1 output vector (gradient map).

The input matrix is:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

The kernel is:

$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$$

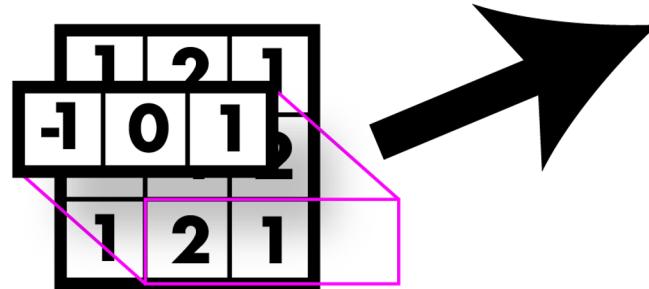
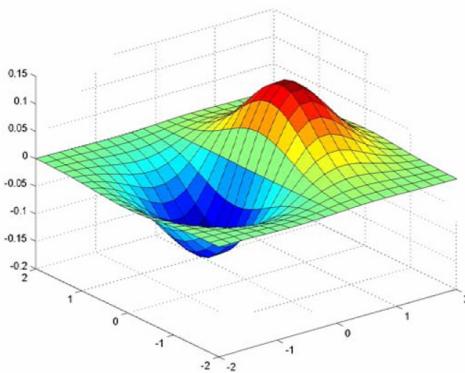
The output vector is:

$$\begin{bmatrix} 2 \\ 4 \\ 2 \end{bmatrix}$$

A pink arrow points from the input matrix to the output vector, indicating the flow of data.

Sobel filter: smooth & derivative

$$\frac{1}{2} \times \left(\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right)$$



$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Derivative of Gaussian filters - 2D

$$d_x \otimes (h \otimes I) = (d_x \otimes h) \otimes I$$

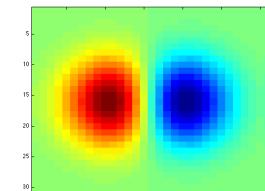
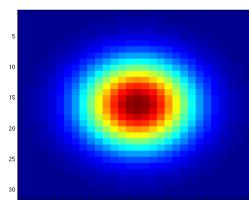
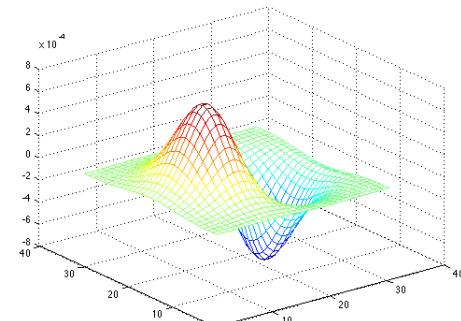
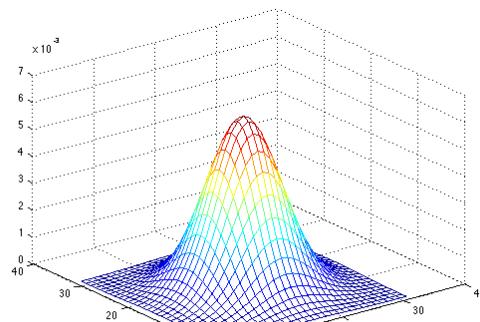
derivation filter Gaussian filter image

$$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$



$$\left[\begin{array}{ccccc} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{array} \right]$$

$$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$



Derivative of Gaussian filters - 2D

$$d_y \otimes (h \otimes I) = (d_y \otimes h) \otimes I$$

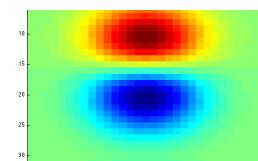
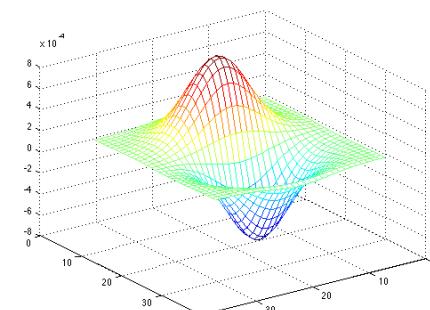
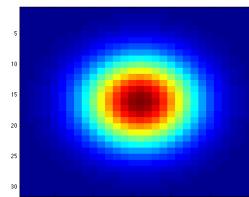
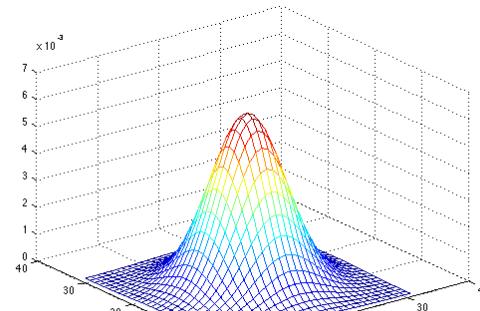
$$\begin{bmatrix} 0 & -1 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$



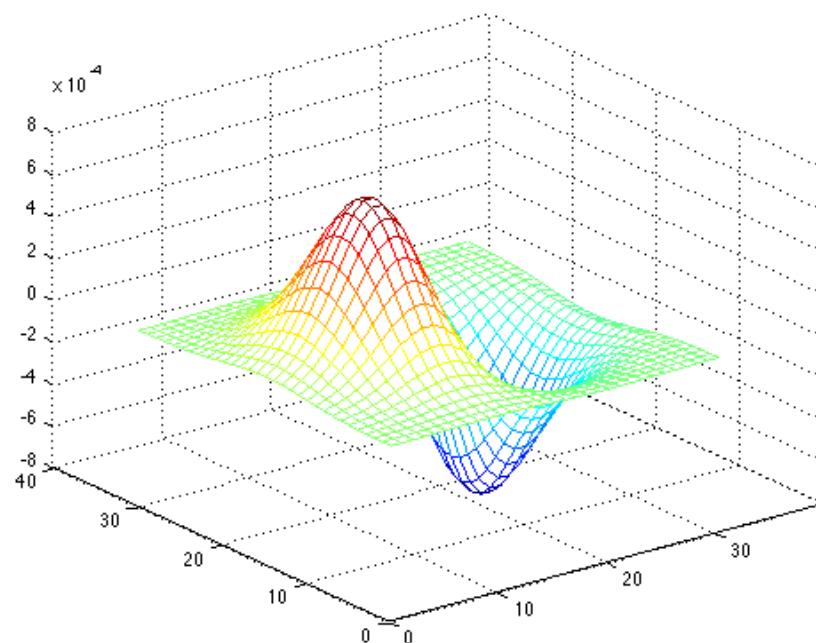
$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$



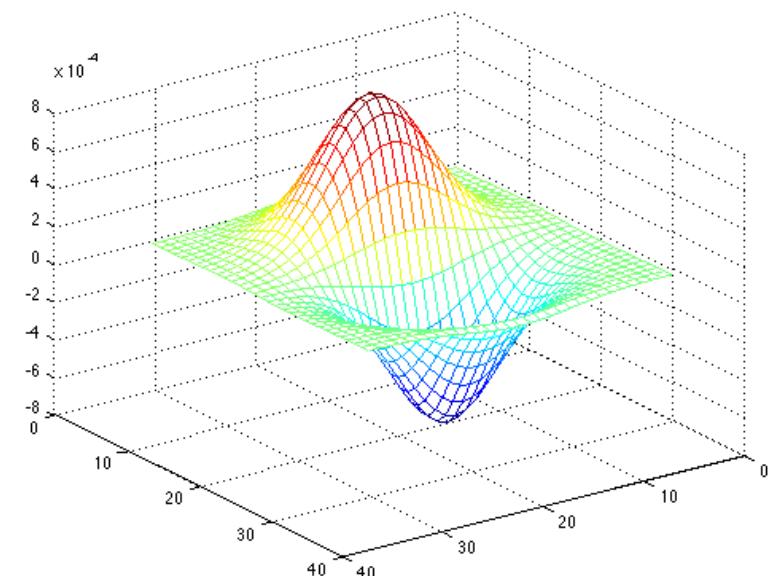
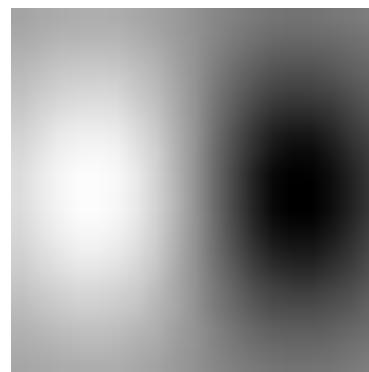
$$\begin{bmatrix} 0 & -1 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$



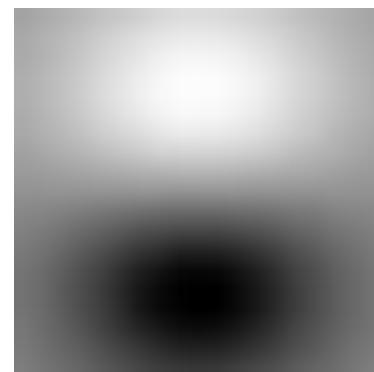
Derivative of Gaussian filters



x - direction

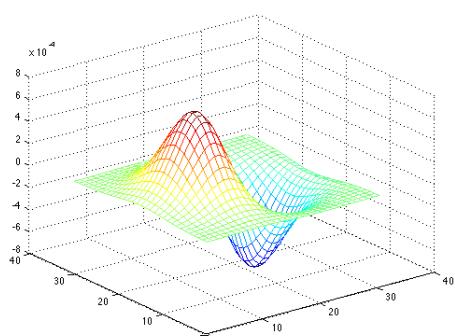


y - direction

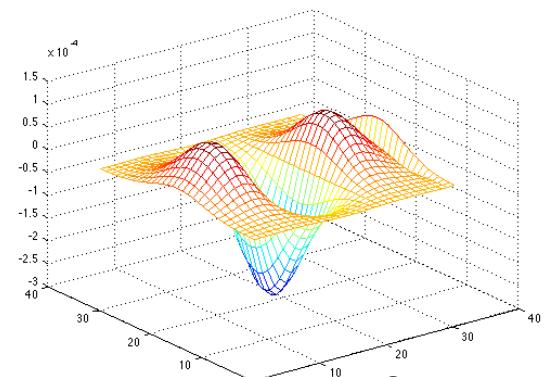


Derivative of Gaussian filters - 2D

$$\begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$



=

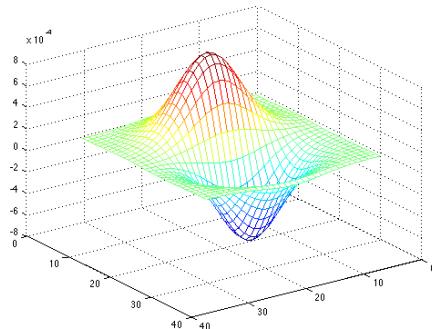


direction x

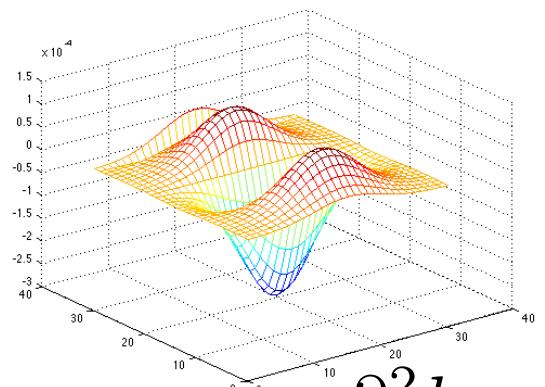
$$\frac{\partial h}{\partial x}$$

$$\frac{\partial^2 h}{\partial x^2}$$

$$\begin{bmatrix} 0 & -1 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$



=



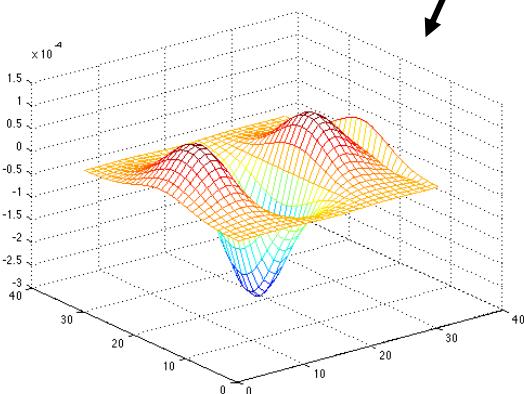
direction y

$$\frac{\partial h}{\partial y}$$

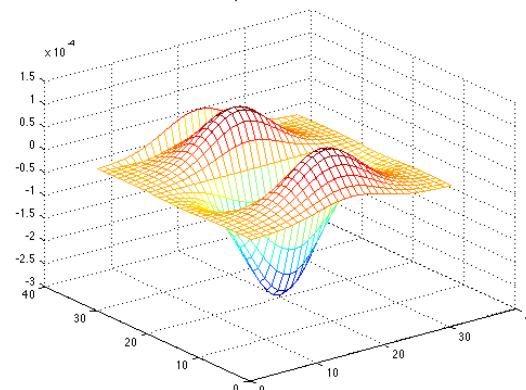
$$\frac{\partial^2 h}{\partial y^2}$$

Laplacian of Gaussian

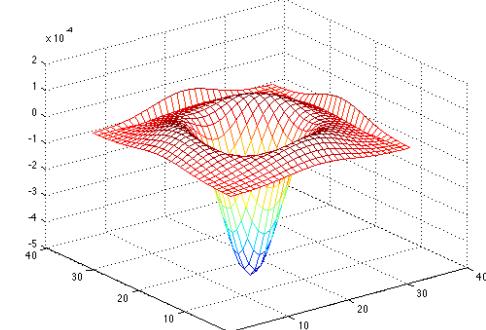
$$\nabla^2 h = \frac{\partial^2 h}{\partial x^2} + \frac{\partial^2 h}{\partial y^2}$$



+

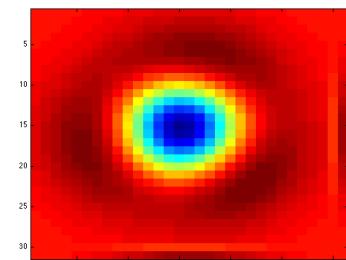


=



0	1	0
1	-4	1
0	1	0

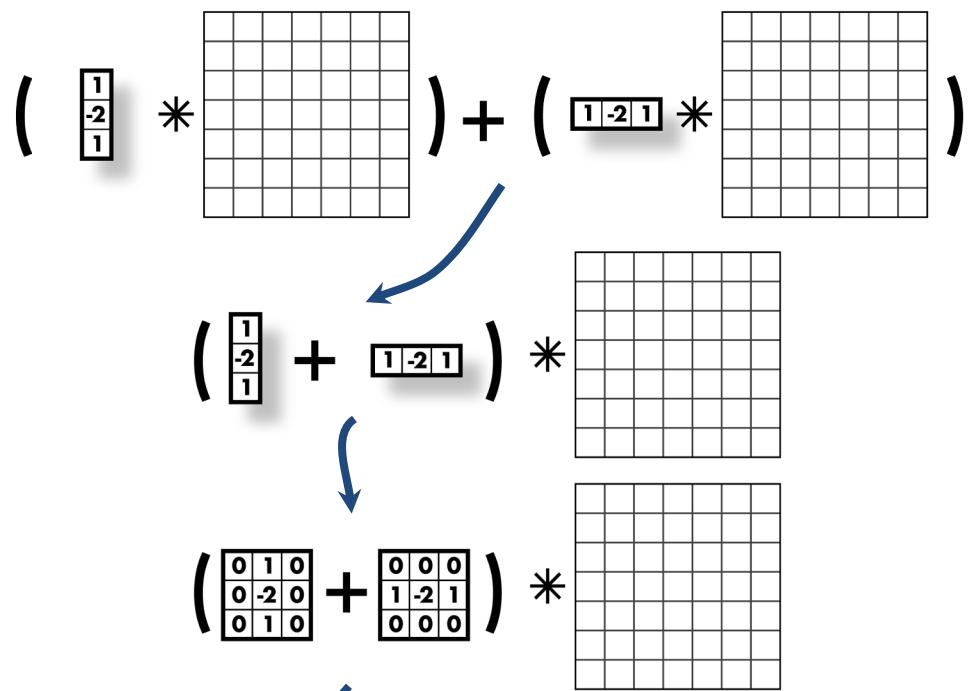
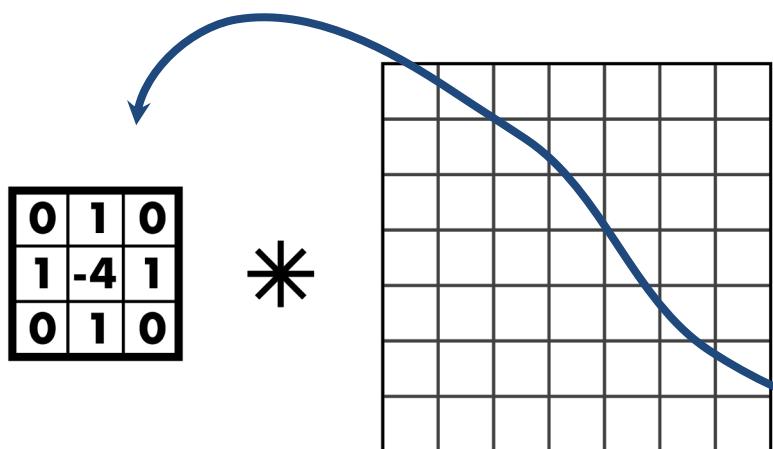
Example of a
Laplacian filter



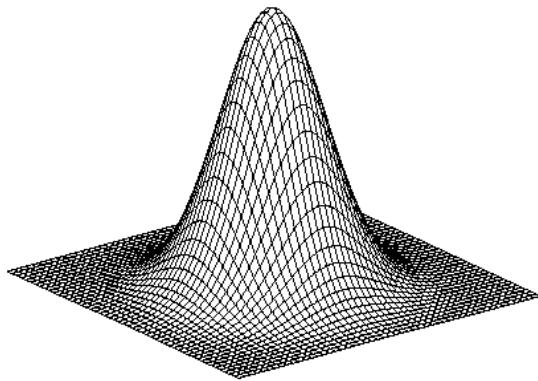
Laplacian of Gaussian

- ### - Laplacian:

$$\Delta f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

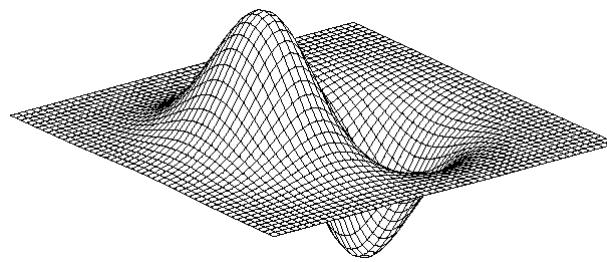


2D edge detection filters



Gaussian

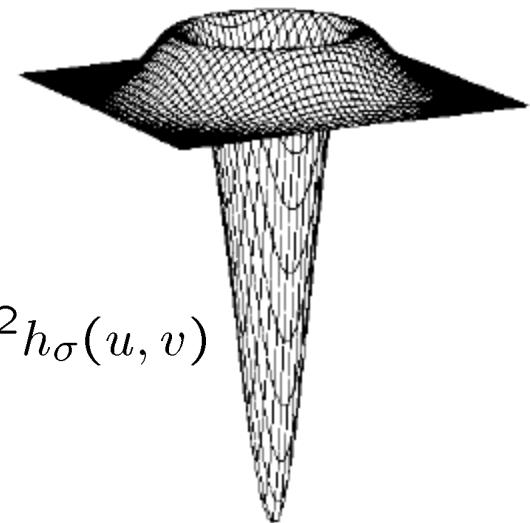
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



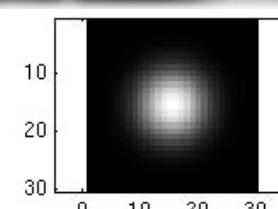
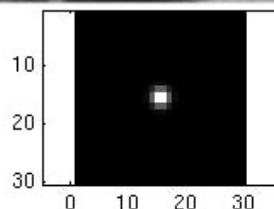
$$\nabla^2 h_\sigma(u, v)$$

- ∇^2 is the Laplacian operator:

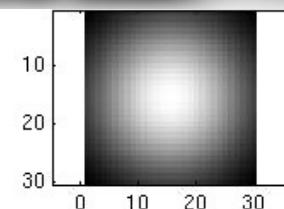
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Smoothing with a Gaussian

Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



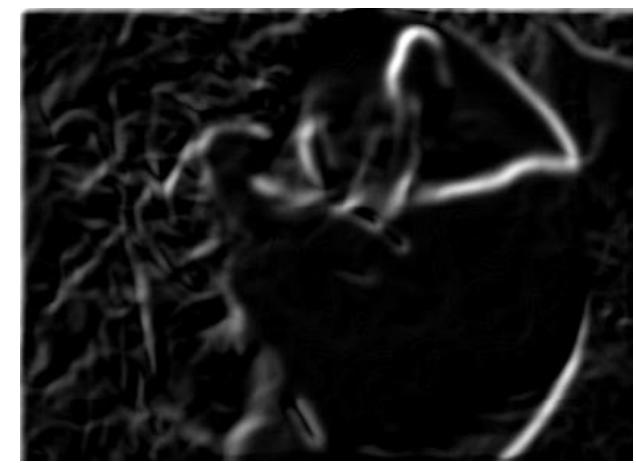
...



Effect of σ on derivatives



$\sigma = 1$ pixel



$\sigma = 3$ pixels

The apparent structures differ depending on Gaussian's scale parameter.

Larger values: larger scale edges detected
Smaller values: finer features detected

Mask properties

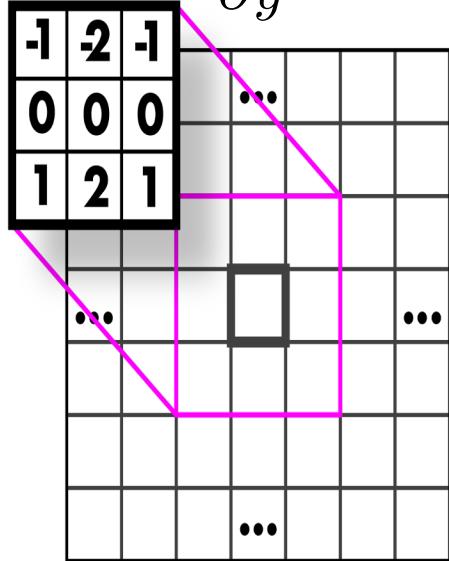
- Smoothing
 - Positive values
 - Sum to 1 → constant regions same as input
 - Amount of smoothing proportional to mask size
 - Remove “high-frequency” components; “low-pass” filter
- Derivatives
 - Opposite signs used to get high response in regions of high contrast
 - Sum to 0 → no response in constant regions
 - High absolute value at points of high contrast

Another approach: gradient magnitude

- Don't need 2nd derivatives
- Just use magnitude of gradient
- Are we done? No!

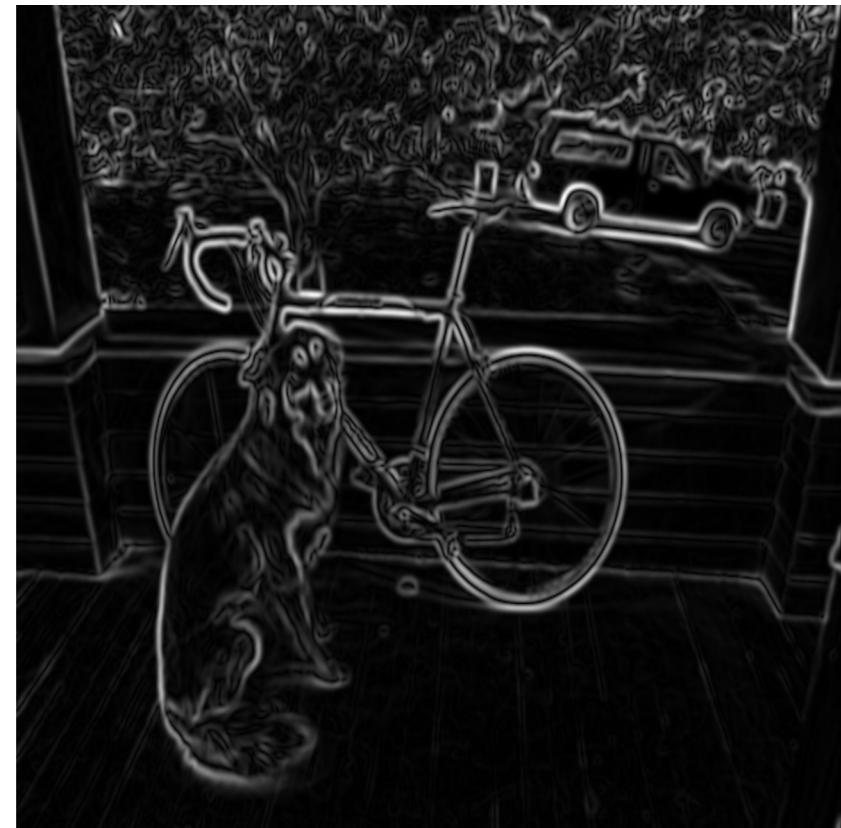
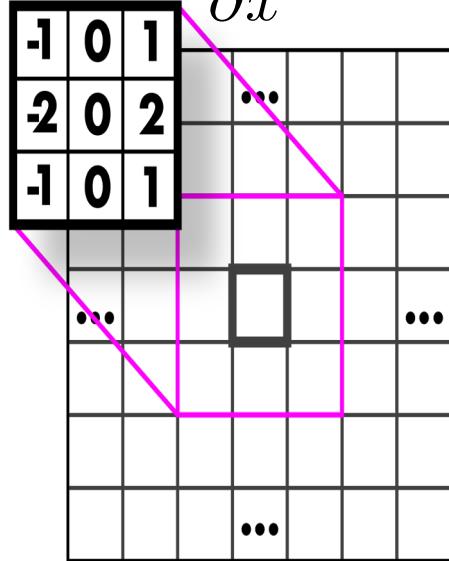
Horizontal Sobel filter
for computing the
partial derivative

$$\frac{\partial f(x, y)}{\partial y}$$



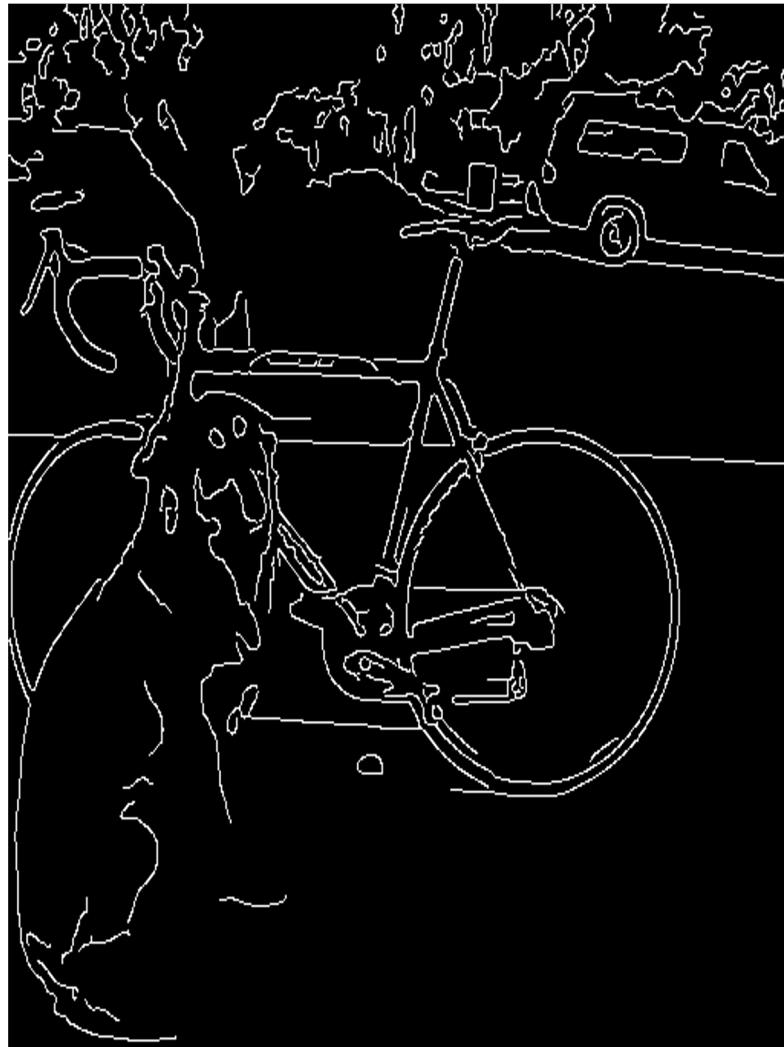
Vertical Sobel filter
for computing the
partial derivative

$$\frac{\partial f(x, y)}{\partial x}$$





What we really want: line drawing



Canny Edge Detection

Algorithm:

- Smooth image (only want “real” edges, not noise)
- Calculate gradient direction and magnitude
- Non-maximum suppression perpendicular to edge
- Threshold into strong, weak, no edge
- Connect together components

Smooth image

- Apply a Gaussian filter

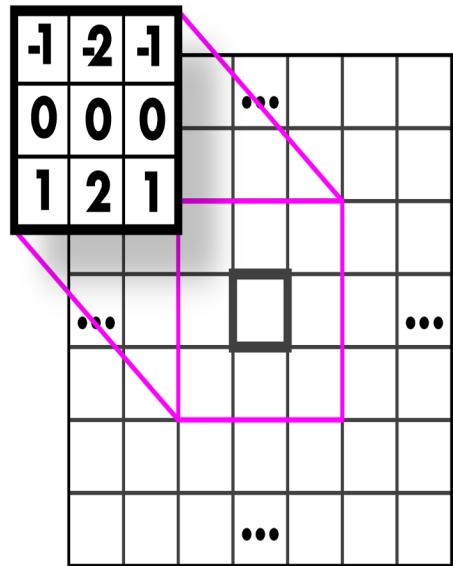


Gradient magnitude and direction

- Sobel filter

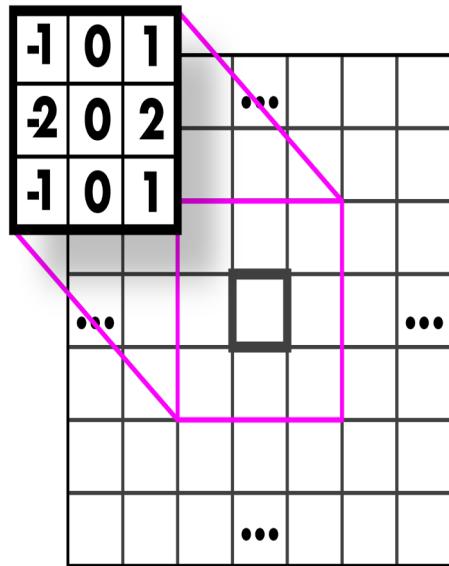
Horizontal Sobel
filter for computing
the partial derivative

$$\frac{\partial f(x, y)}{\partial y}$$



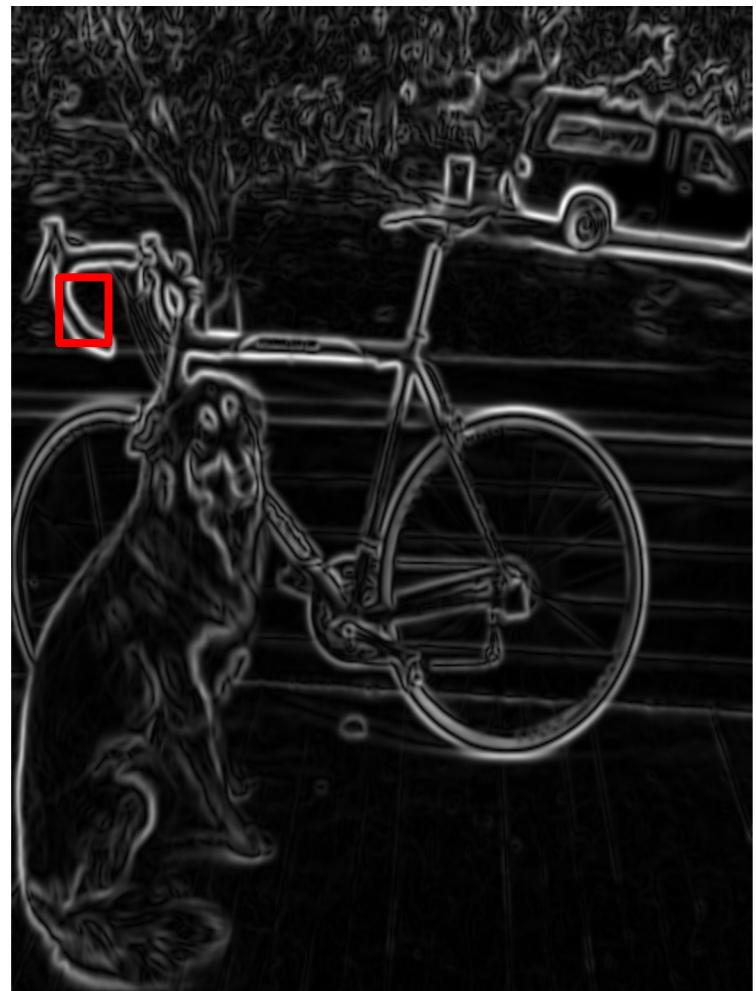
Vertical Sobel filter
for computing the
partial derivative

$$\frac{\partial f(x, y)}{\partial x}$$

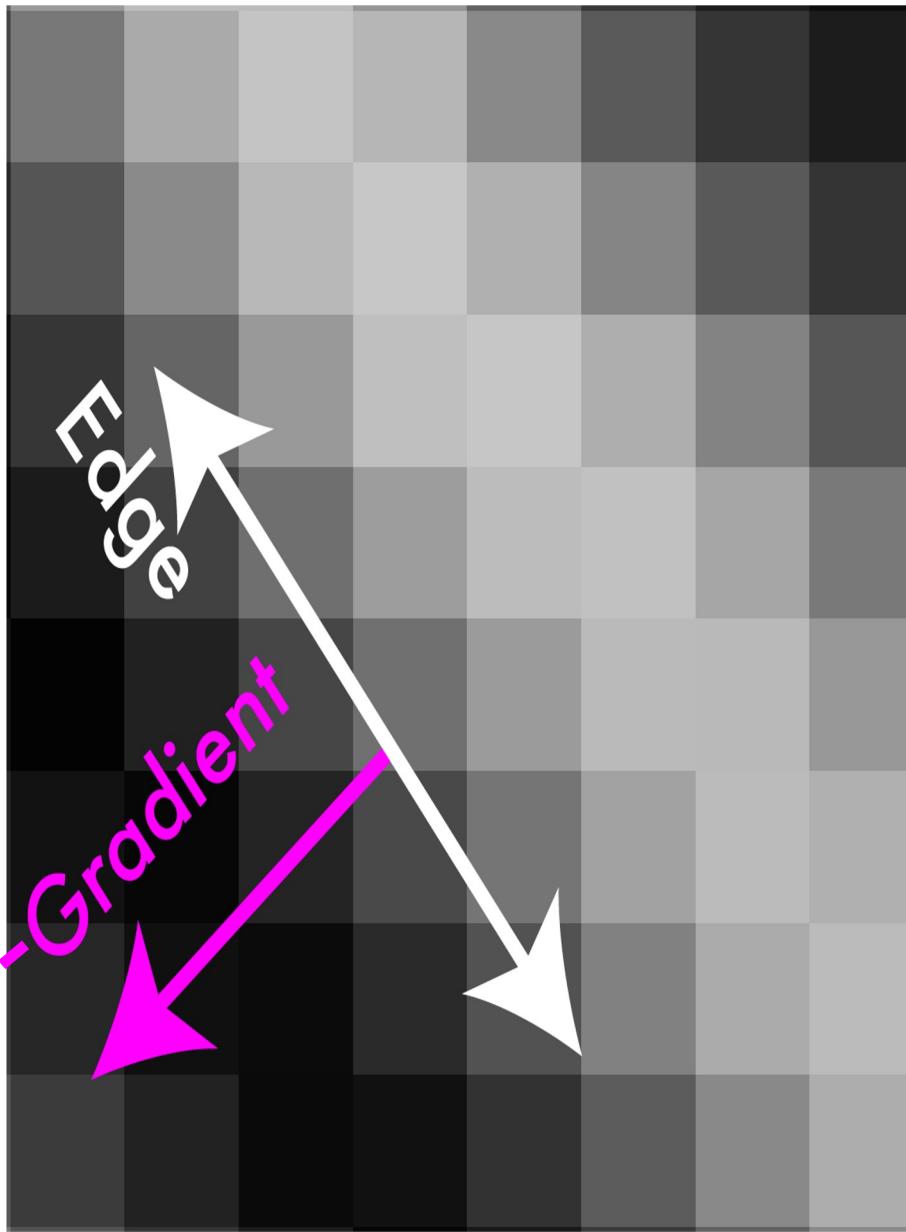


Non-maximum suppression

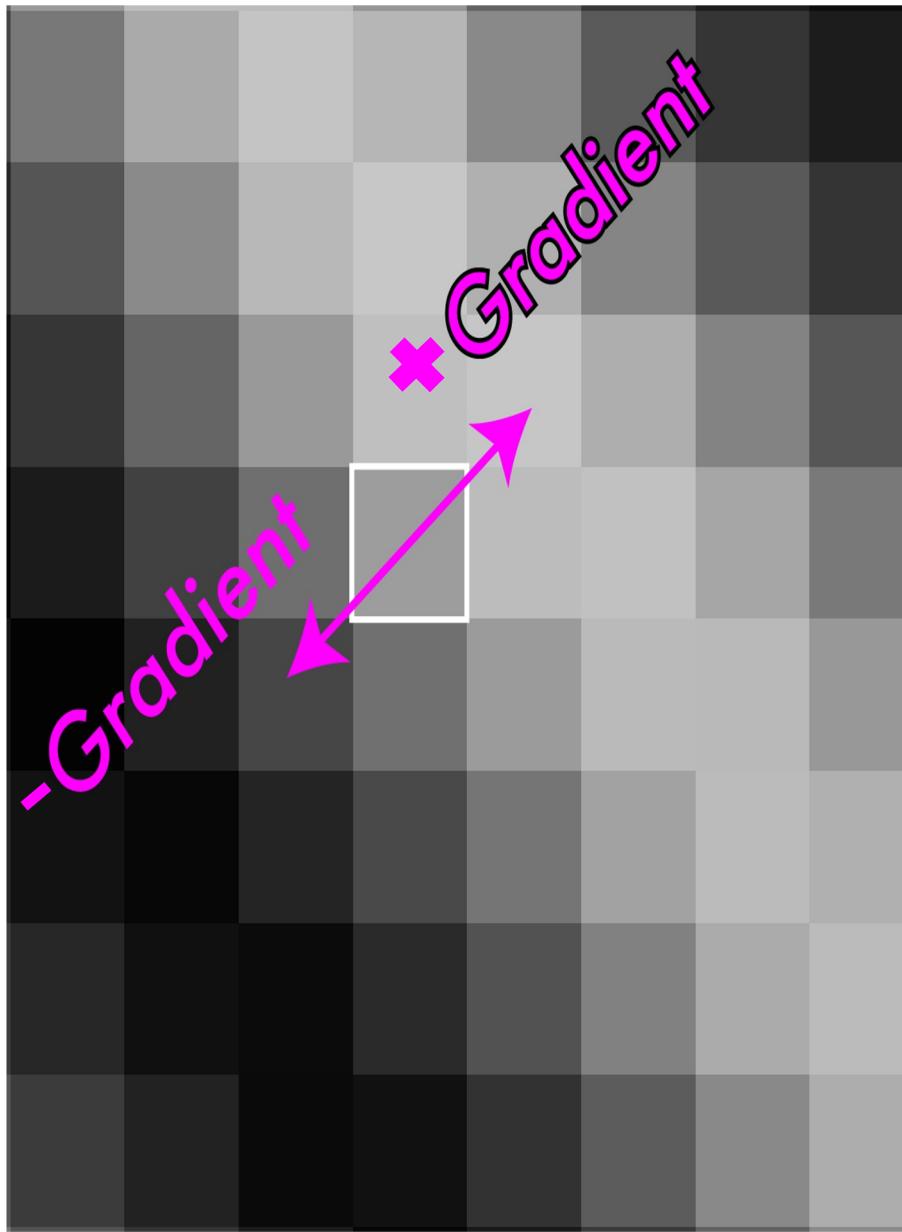
- Want single pixel edges, not thick blurry lines
- Need to check nearby pixels
- See if response is highest



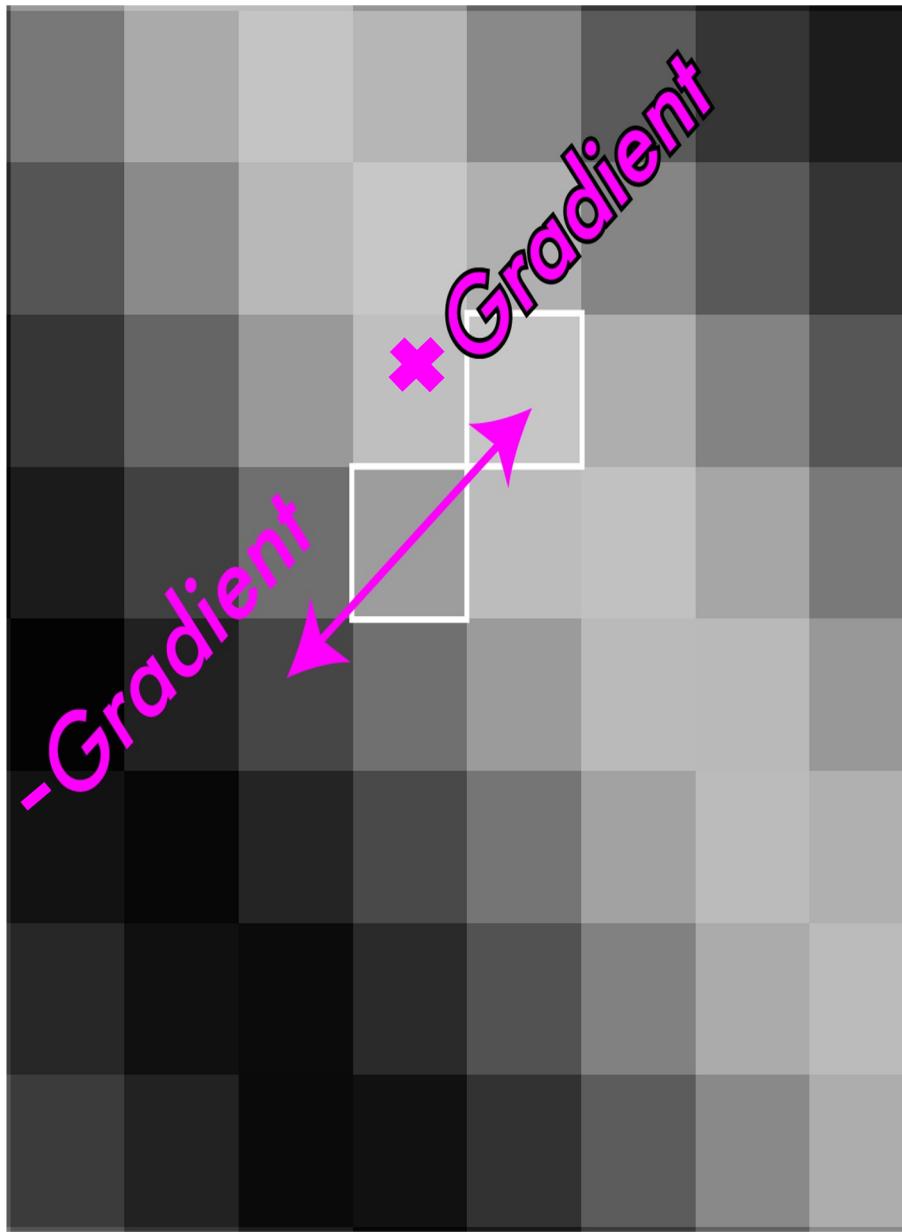
Non-maximum suppression



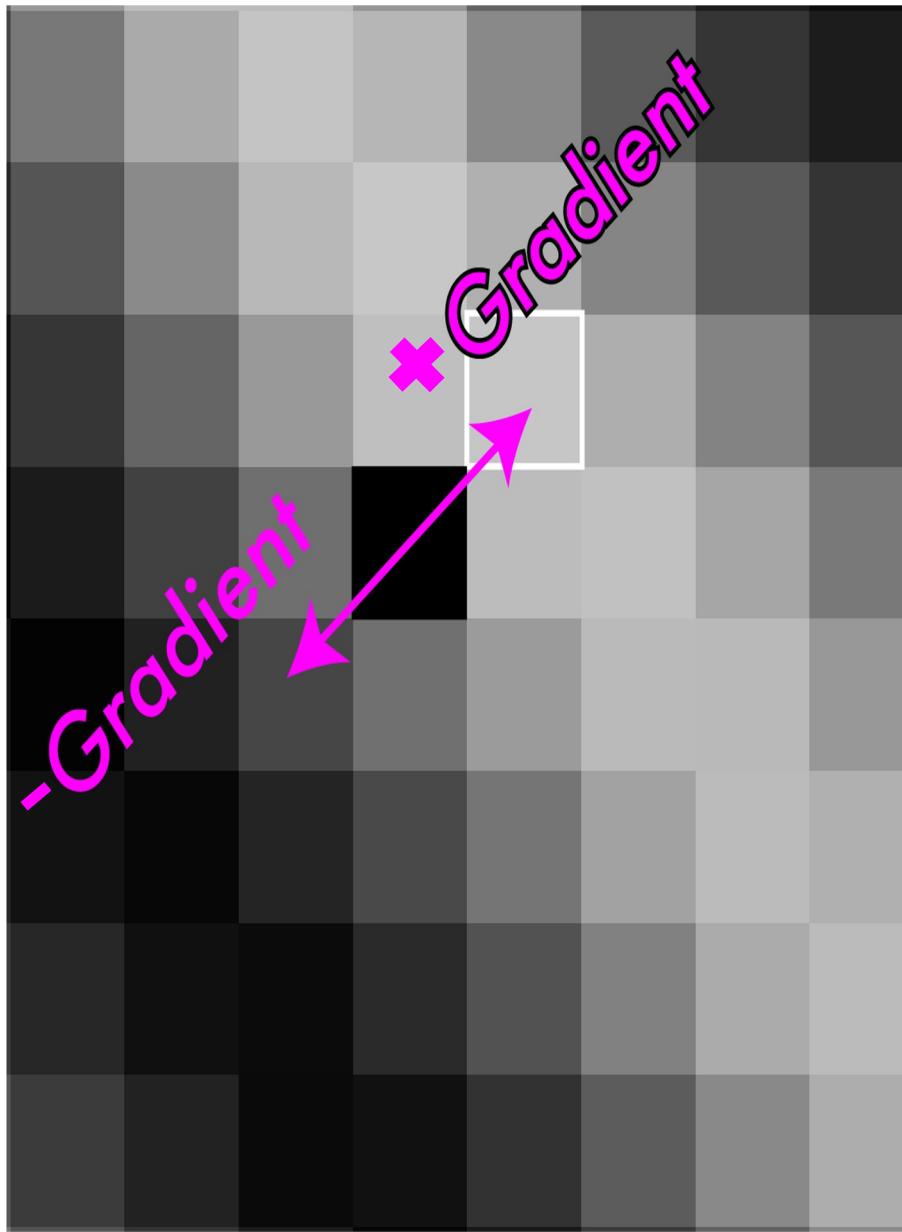
Non-maximum suppression



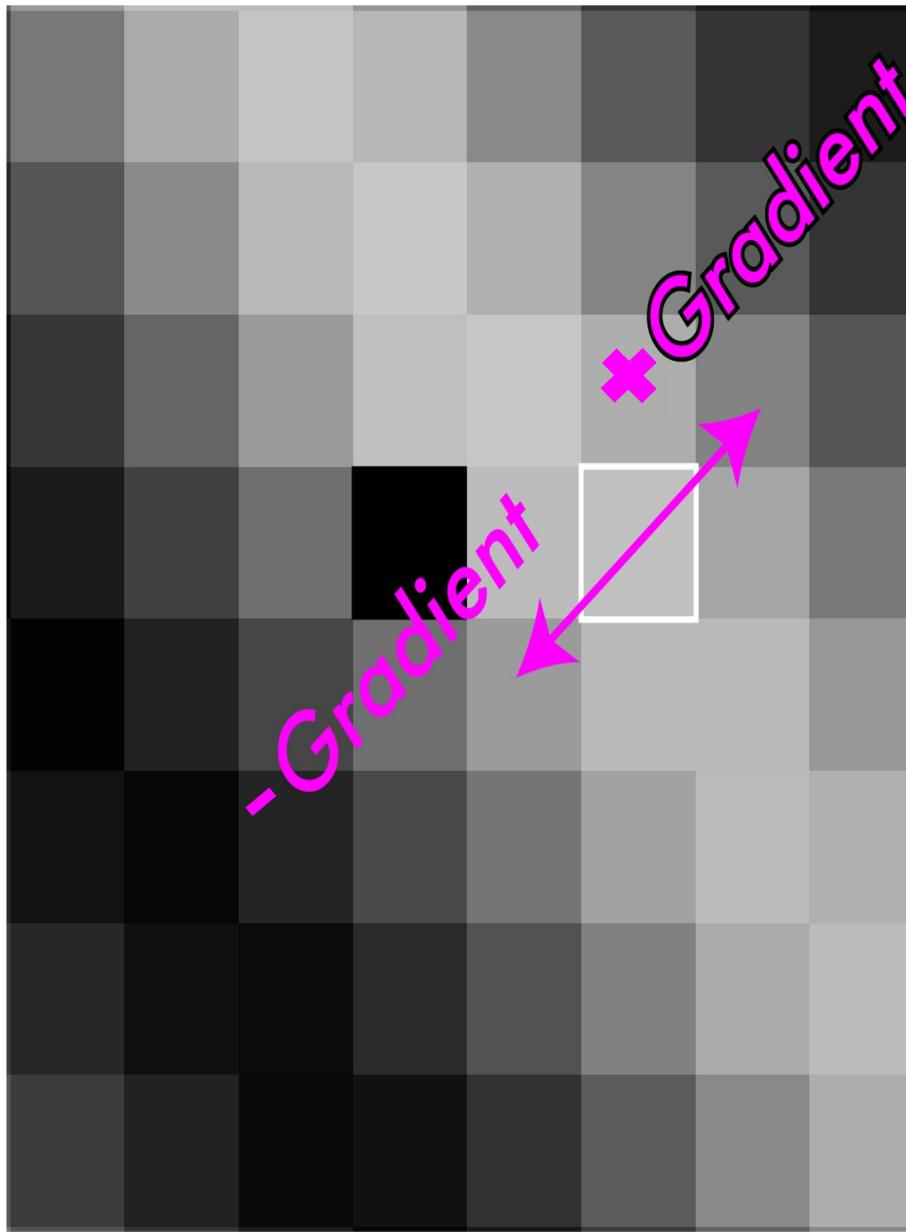
Non-maximum suppression



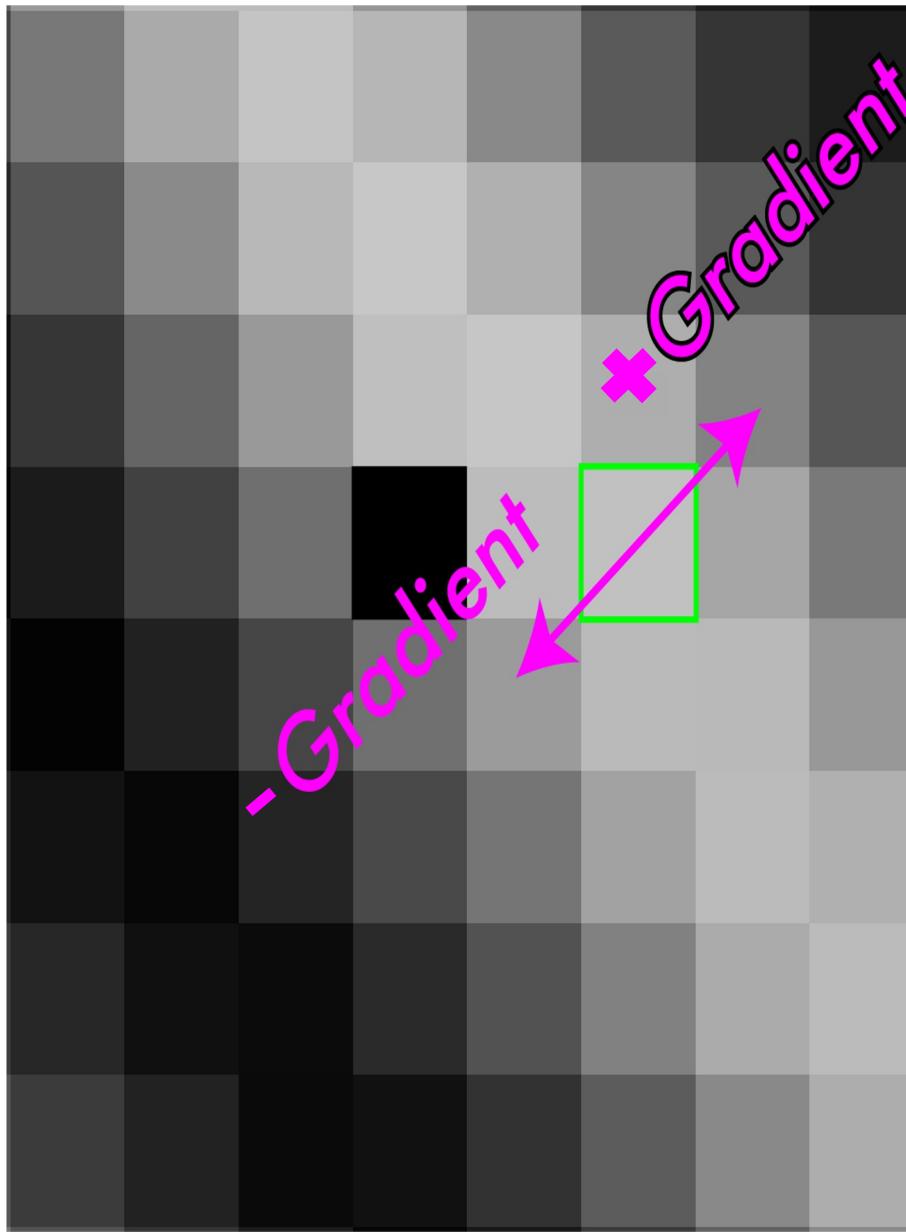
Non-maximum suppression



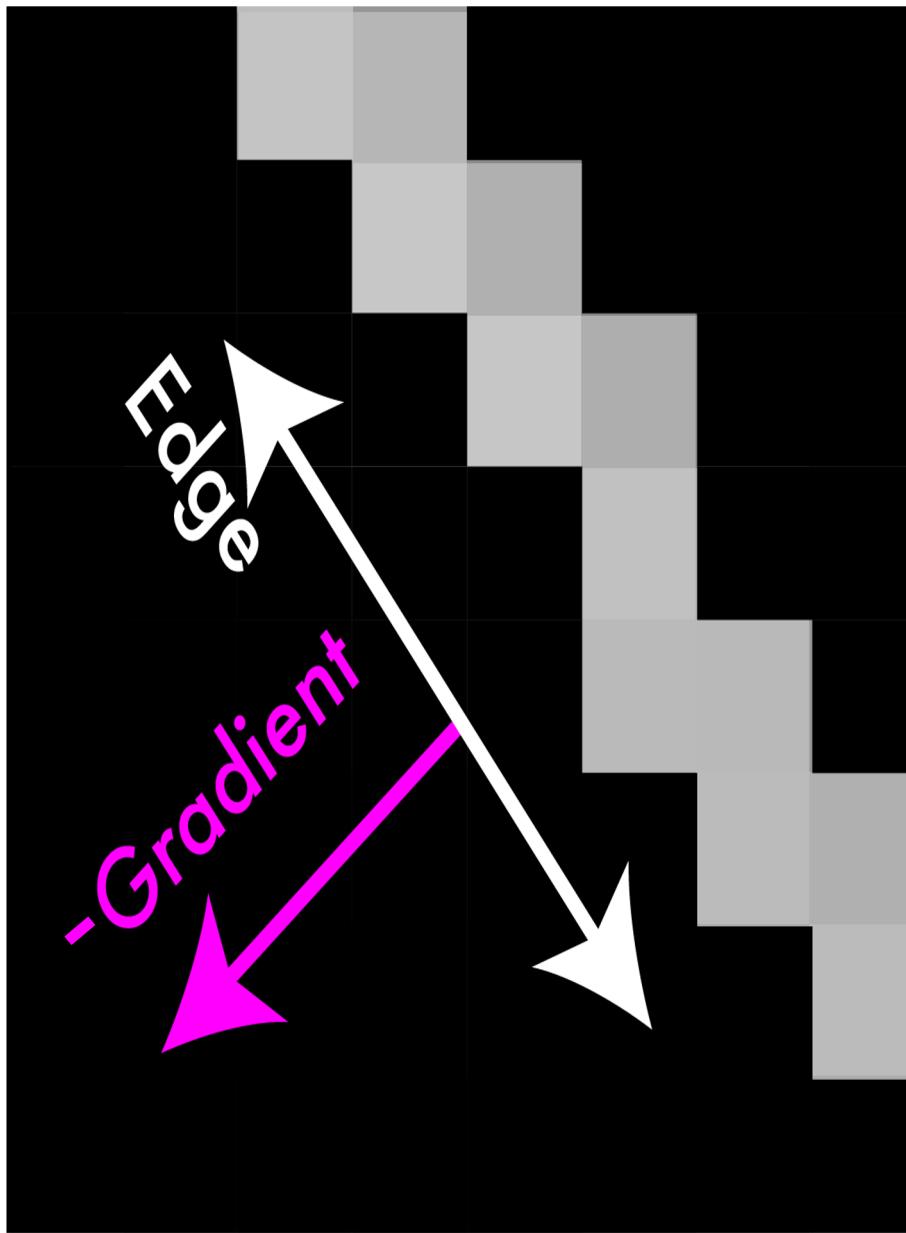
Non-maximum suppression



Non-maximum suppression



Non-maximum suppression

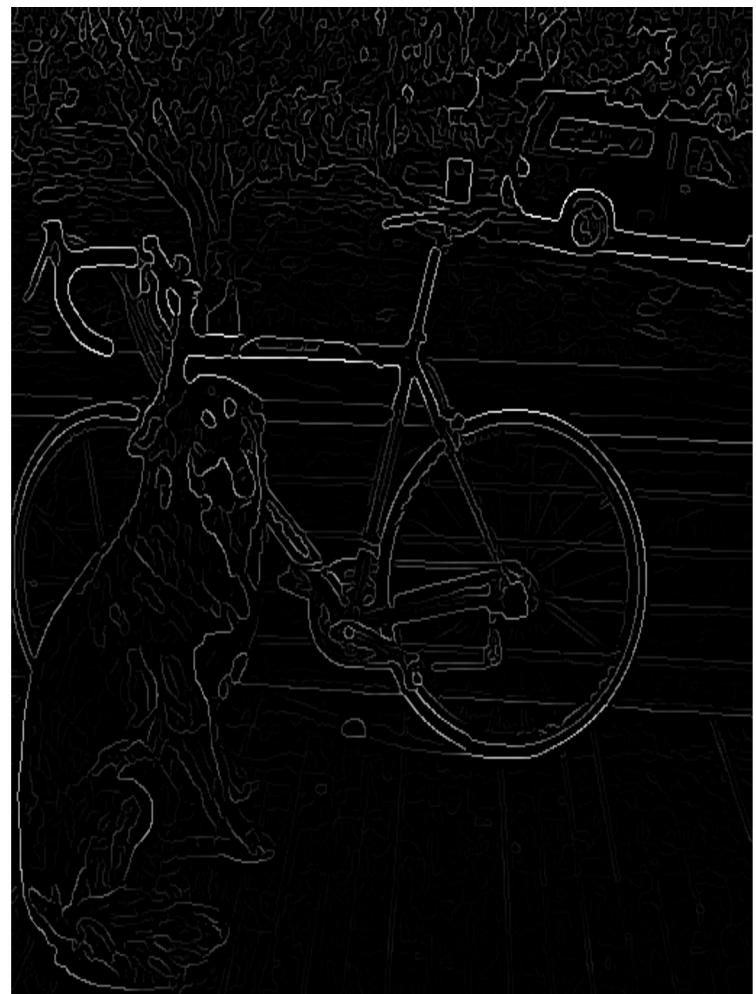


Non-maximum suppression



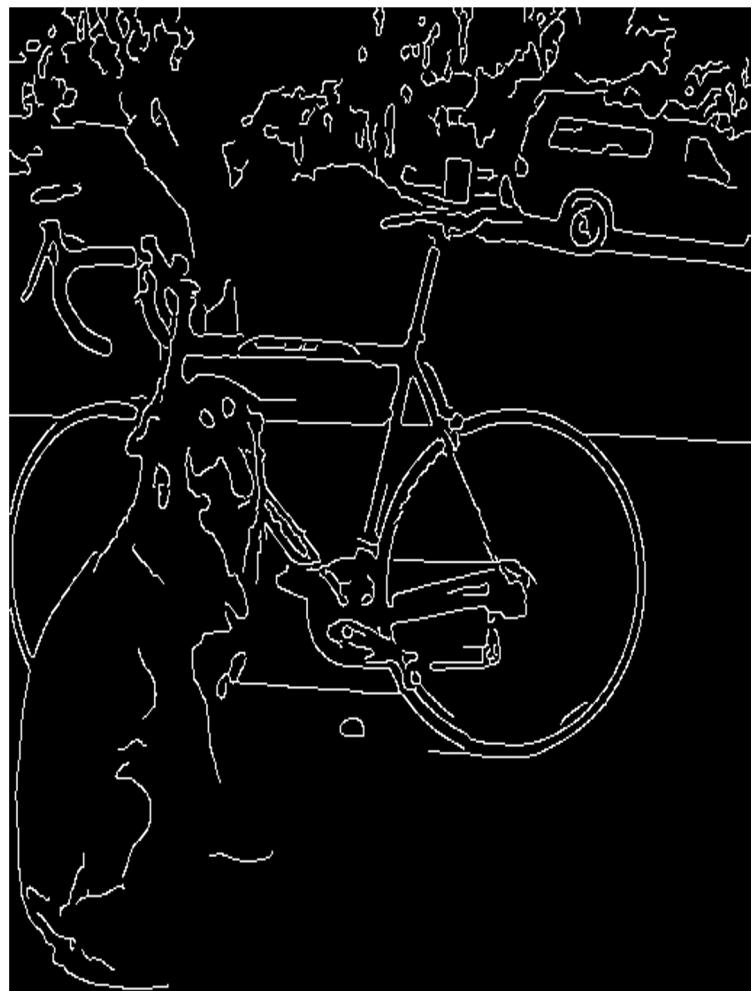
Threshold edges

- Still some noise
- Only want strong edges
- 2 thresholds, 3 cases
 - $R > T$: strong edge
 - $R < T$ but $R > t$: weak edge
 - $R < t$: no edge
- Why two thresholds?



Connect them up!

- Strong edges are edges!
- Weak edges are edges
iff they connect to strong edges
- Look in some neighborhood
(usually 8 closest)

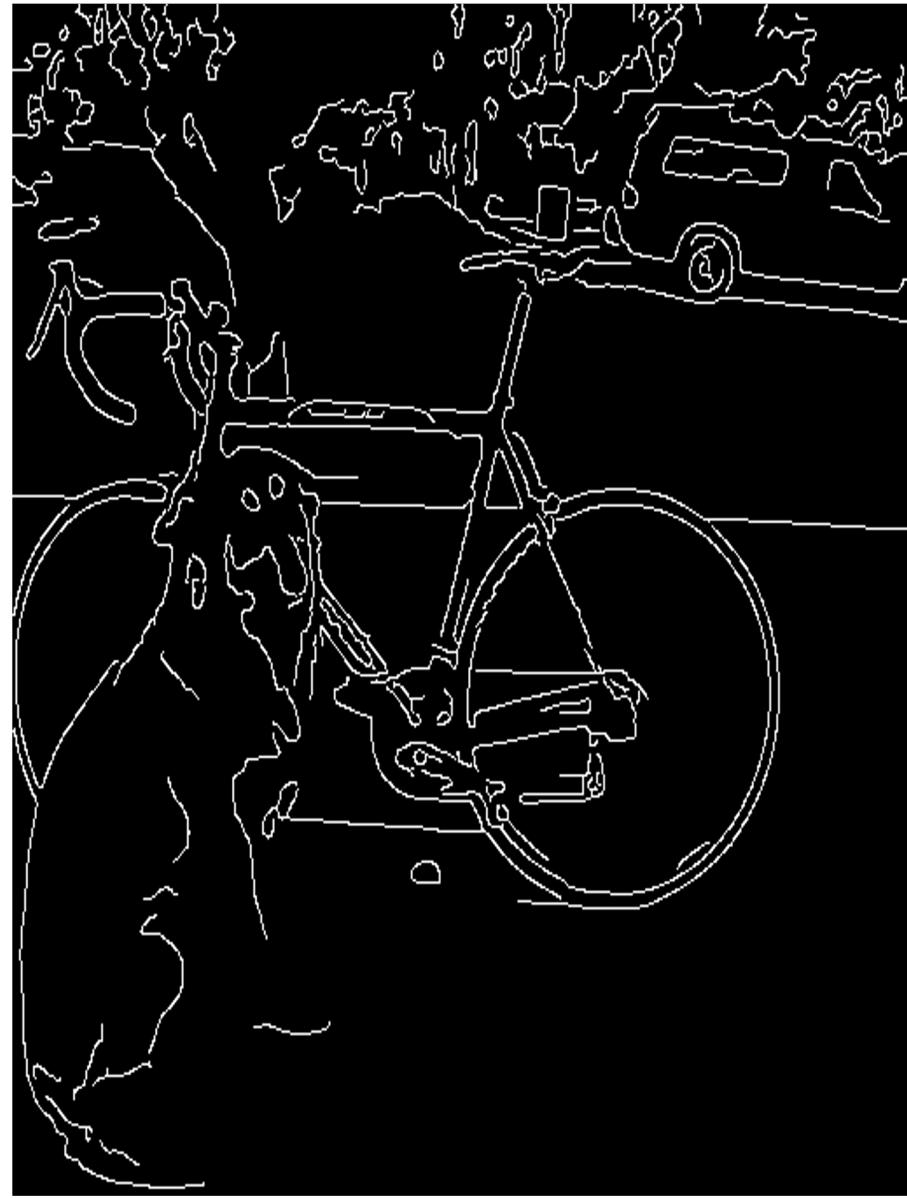


Canny Edge Detection

Algorithm:

- Smooth image (only want “real” edges, not noise)
- Calculate gradient direction and magnitude
- Non-maximum suppression perpendicular to edge
- Threshold into strong, weak, no edge
- Connect together components
- Tunable: Sigma, thresholds

Canny Edge Detection



Canny Edge Detection - demo

<http://bigwww.epfl.ch/demo/ip/demos/edgeDetector/>

Course structure

1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

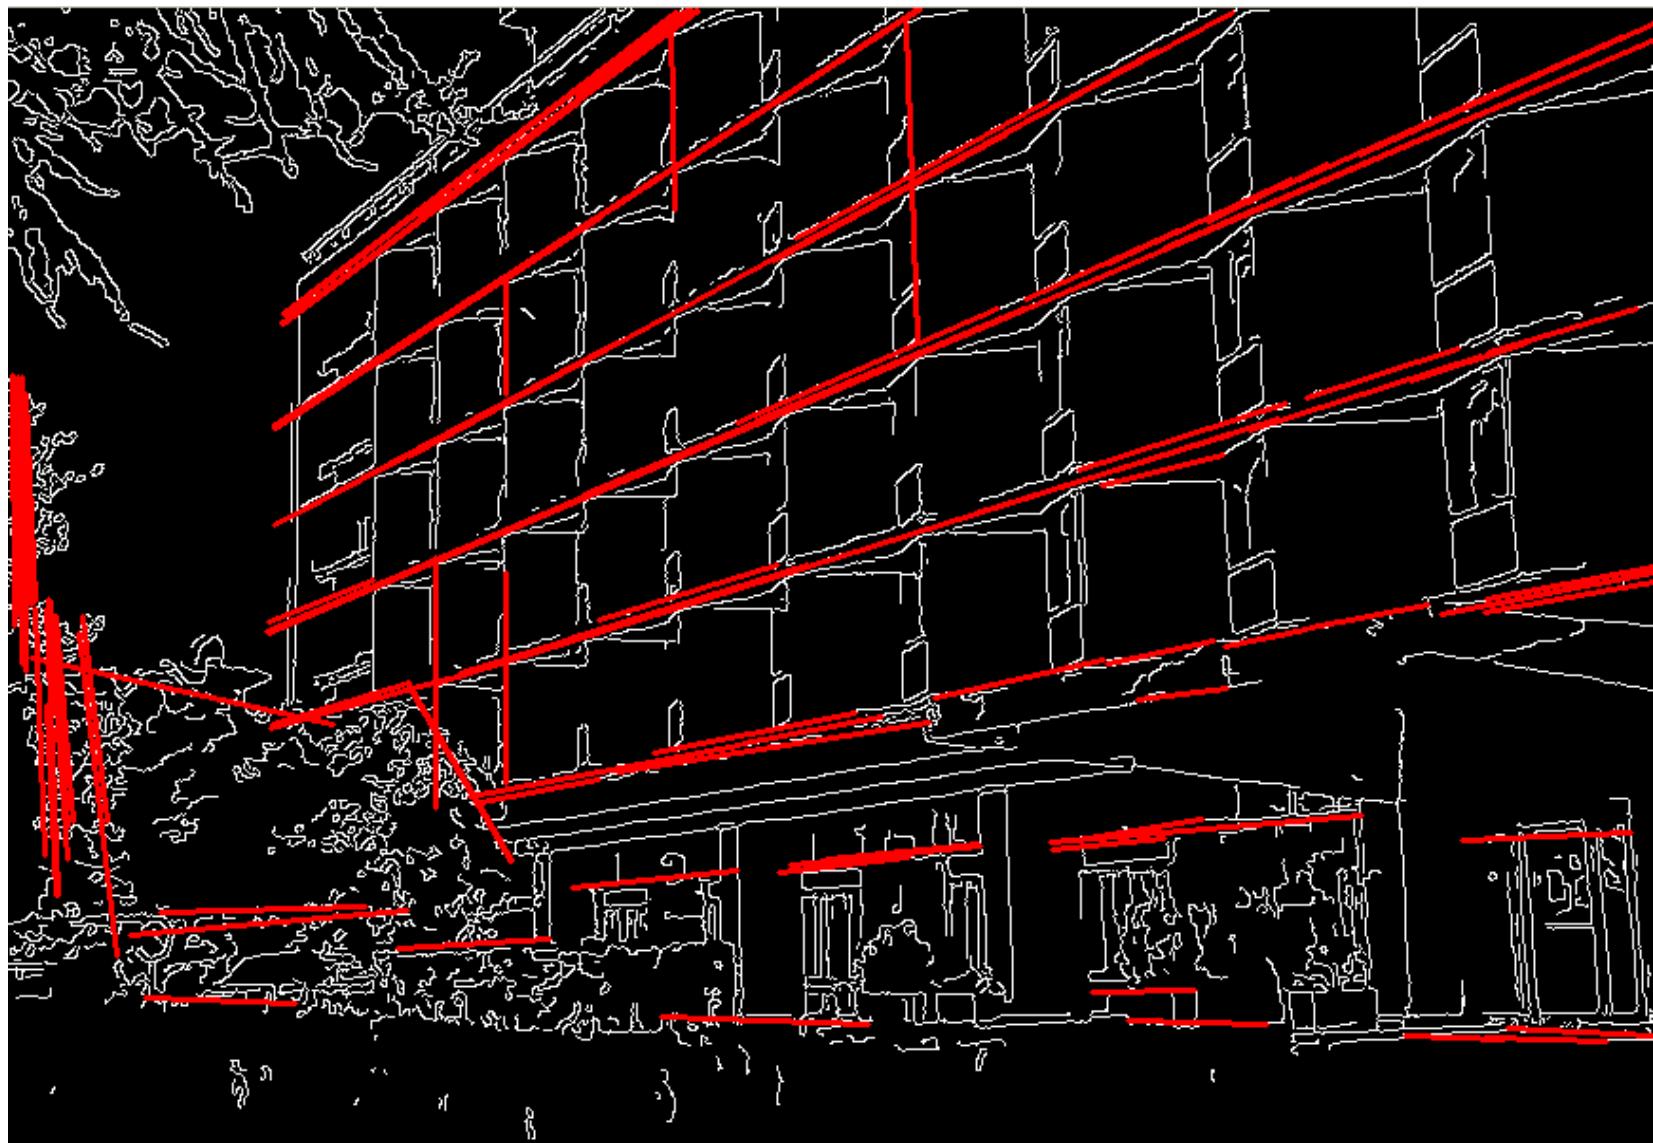
4. Object Recognition: high – level vision

Object classification, object detection, part based models, bovw models

5. Video understanding

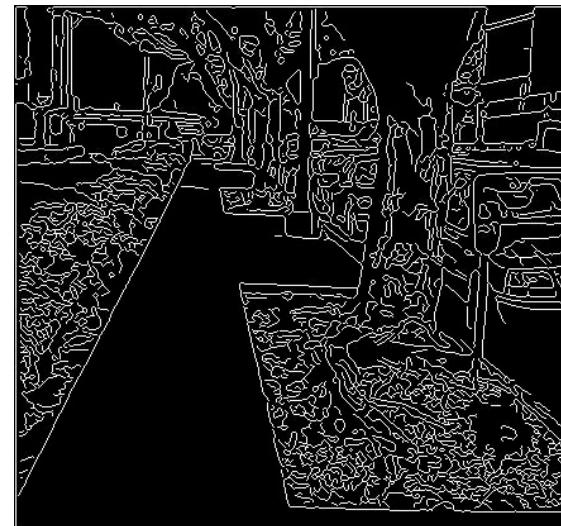
Object tracking, background subtraction, motion descriptors, optical flow

Fitting



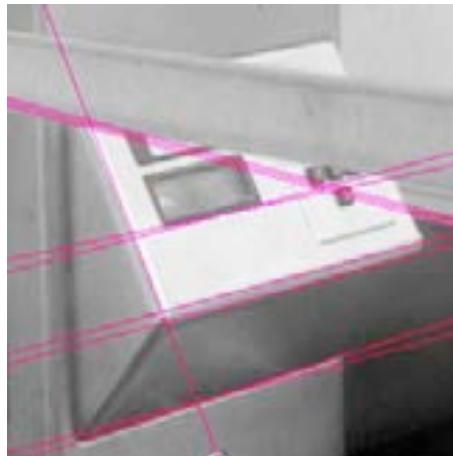
Fitting

- We have learned how to detect edges
- We would like to form a higher-level, more compact representation of the features in the image by grouping multiple features according to a simple model



Fitting

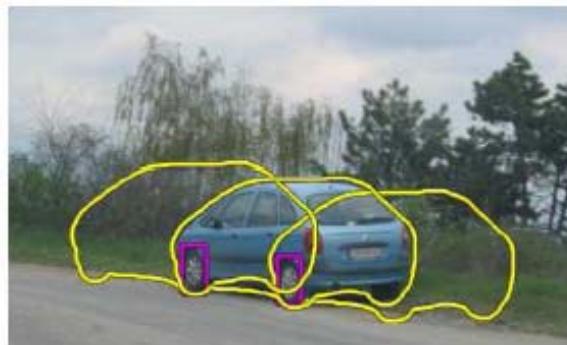
- Choose a *parametric model* to represent a set of features



simple model: lines



simple model: circles



complex model: car



Lab 2: Line detection - application

PROBĂ DE CONCURS

Domeniu... *CTI* **730**

Sesiunea... *IULIE 2018*

Nota pe lucrare *7,90 (patru 5,50)* Notă după contestație *7,90*

PDLX



TEST GRILĂ

INFORMATICĂ

MATEMATICĂ

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5		X		
6		X		
7		X		
8				X
9			X	
10	X			
11		X		
12		X		
13			X	
14			X	
15			X	

FIZICĂ

Număr întrebare	Răspuns			
	A	B	C	D
1	X			
2			X	
3			X	
4	X			
5		X		
6	X			
7	X			
8		X		
9			X	
10			X	
11		X		
12				X
13			X	
14	X			
15			X	

NOTĂ : Se bifează X în căsuța corespunzătoare răspunsului corect.

Lab 2: Line detection - application

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				
9				X
10	X			
11		X		
12		X		
13			X	
14				X
15			X	

	Număr întrebare	Răspuns			
		A	B	C	D
1			X		
2			X		
3				X	
4		X			
5			X		
6				X	
7		X			
8					X
9				X	
10		X			
11			X		
12			X		
13				X	
14				X	
15				X	

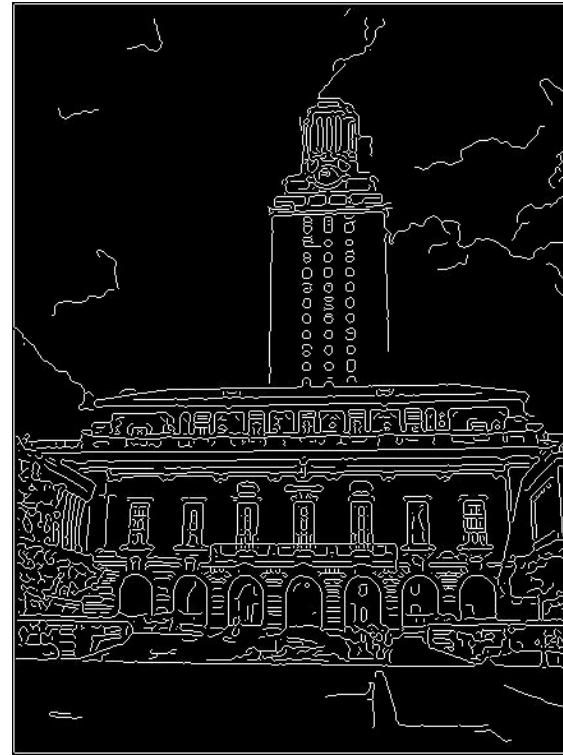
Lab 2: Line detection - application

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2	X			
3		X		
4	X			
5		X		
6		X	X	
7	X			
8				X
9			X	
10	X			X
11		X	X	
12	X			
13		X		
14			X	
15			X	

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3			X	
4	X			
5			X	
6			X	
7		X		
8				X
9			X	
10	X			
11		X		
12		X		
13			X	
14				X
15			X	

Fitting: challenges

Case study: Line detection

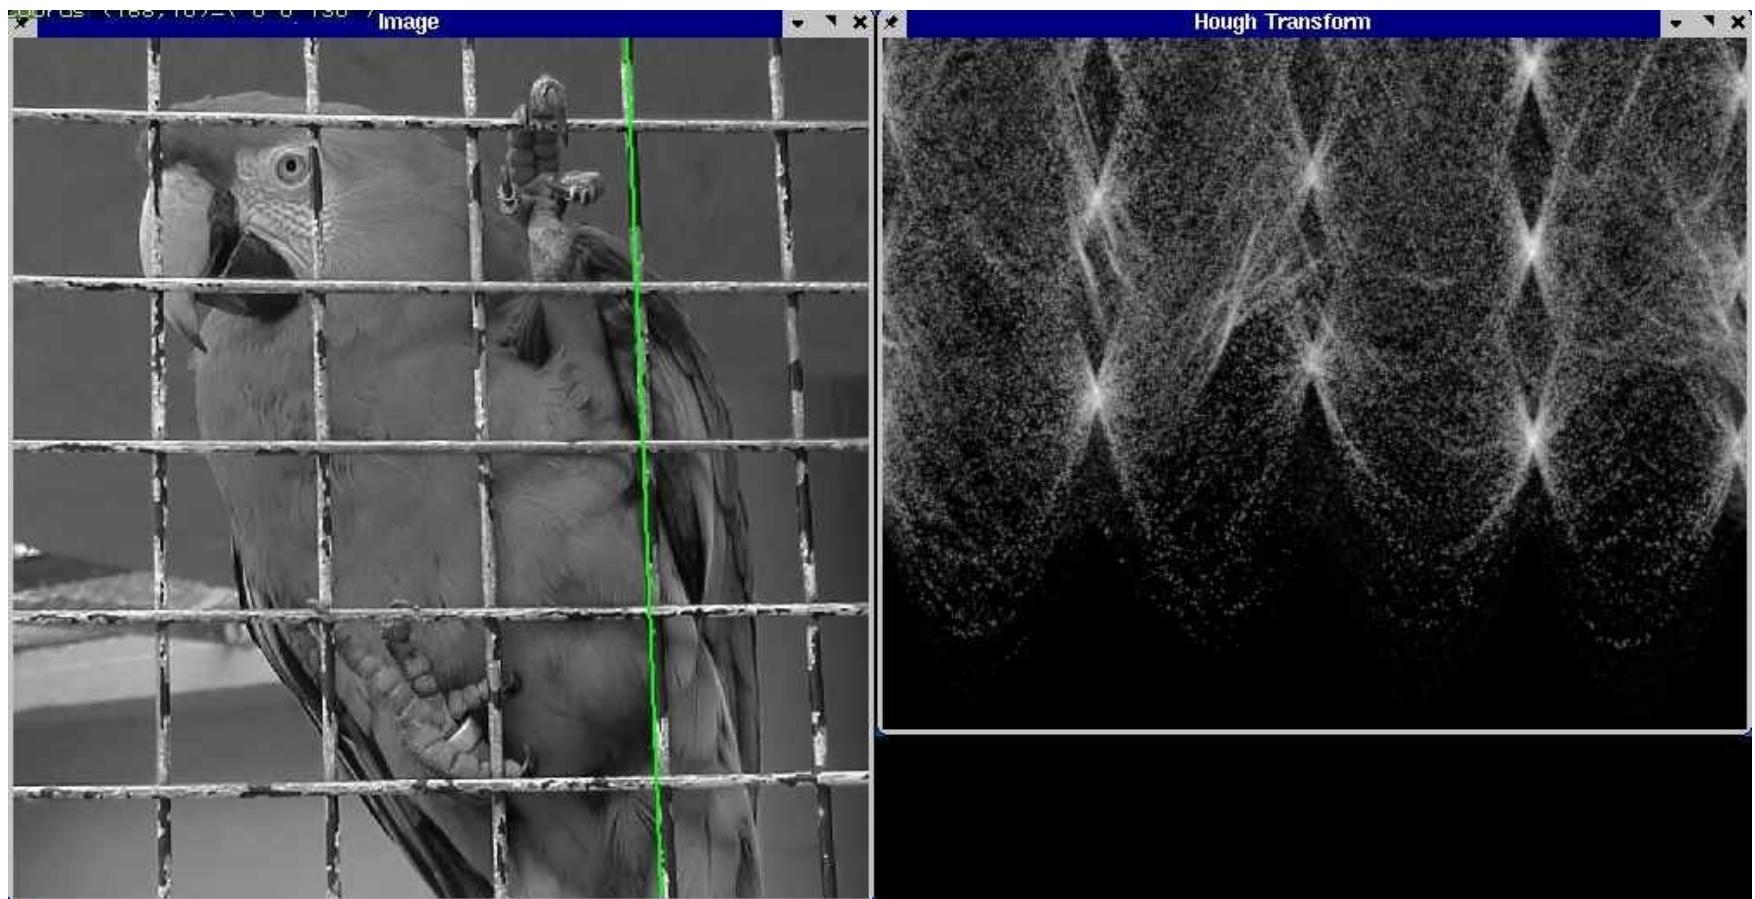


- **Noise** in the measured feature locations
- **Extraneous data:** clutter (outliers), multiple lines
- **Missing data:** occlusions

Fitting: challenges

- If we know which points belong to the line, how do we find the “optimal” line parameters?
 - Least squares
- What if there are outliers?
 - Robust fitting, RANSAC
- What if there are many lines?
 - Voting methods: RANSAC, Hough transform

Fitting: The Hough transform



Voting schemes

- Let each feature vote for all the models that are compatible with it
 - Hopefully the noise features will not vote consistently for any single model
 - Missing data doesn't matter as long as there are enough features remaining to agree on a good model

Hough transform

- An early type of voting scheme
 - Discretize *parameter space* into bins
 - For each feature point in the image, put a vote in every bin in the parameter space that could have generated this point
 - Find bins that have the most votes

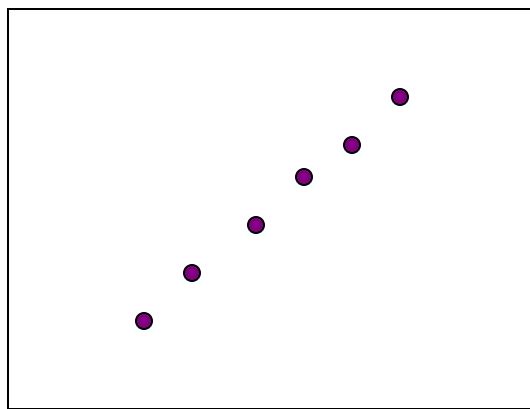
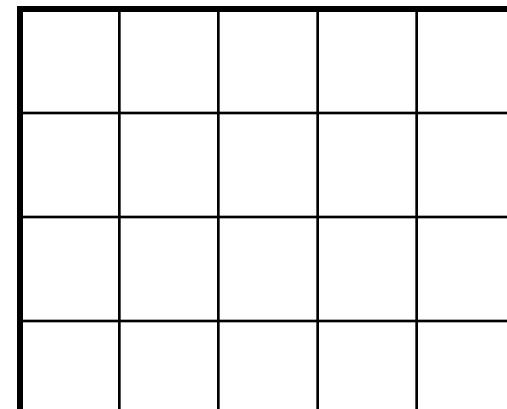
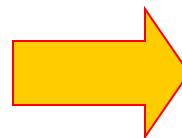


Image space

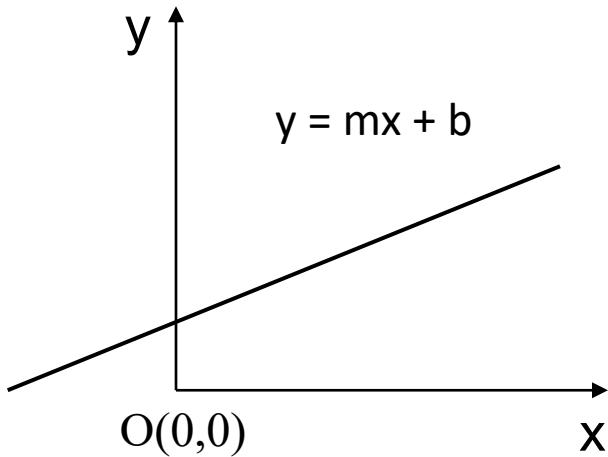


Hough parameter space

P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*,
Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959

Slide credit : S. Lazebnik

Parameterization of a line



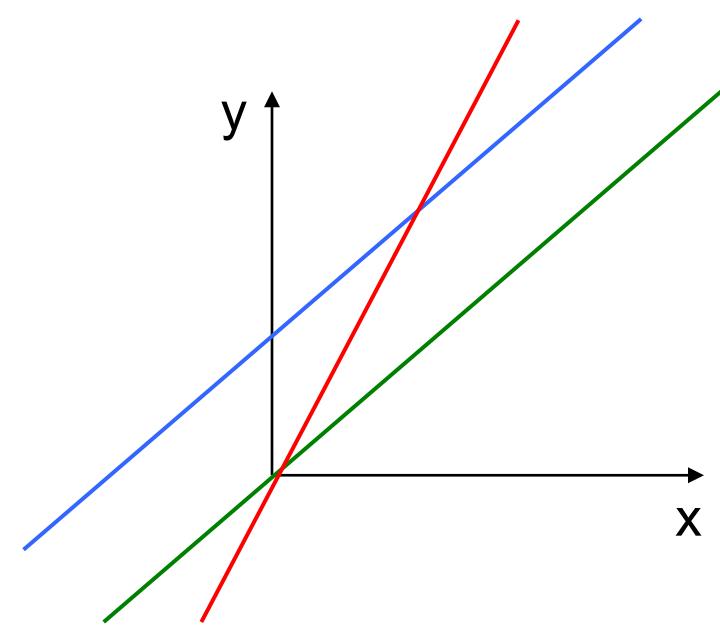
Parameterization of a line in image plane: $y = m * x + b$,

m – slope (defines the angle between the line and Ox axis),

b – intercept (bias to origin $O(0,0)$)

(m,b) are line parameters

Parameterization of a line - examples



$d_1: y = x$, so $m=1, b = 0$

$d_2: y = x + 2$, so $m=1, b = 2$

$d_3: y = 2x$, so $m=2, b = 0$

Parameterization of a line in image plane: $y = m * x + b$,

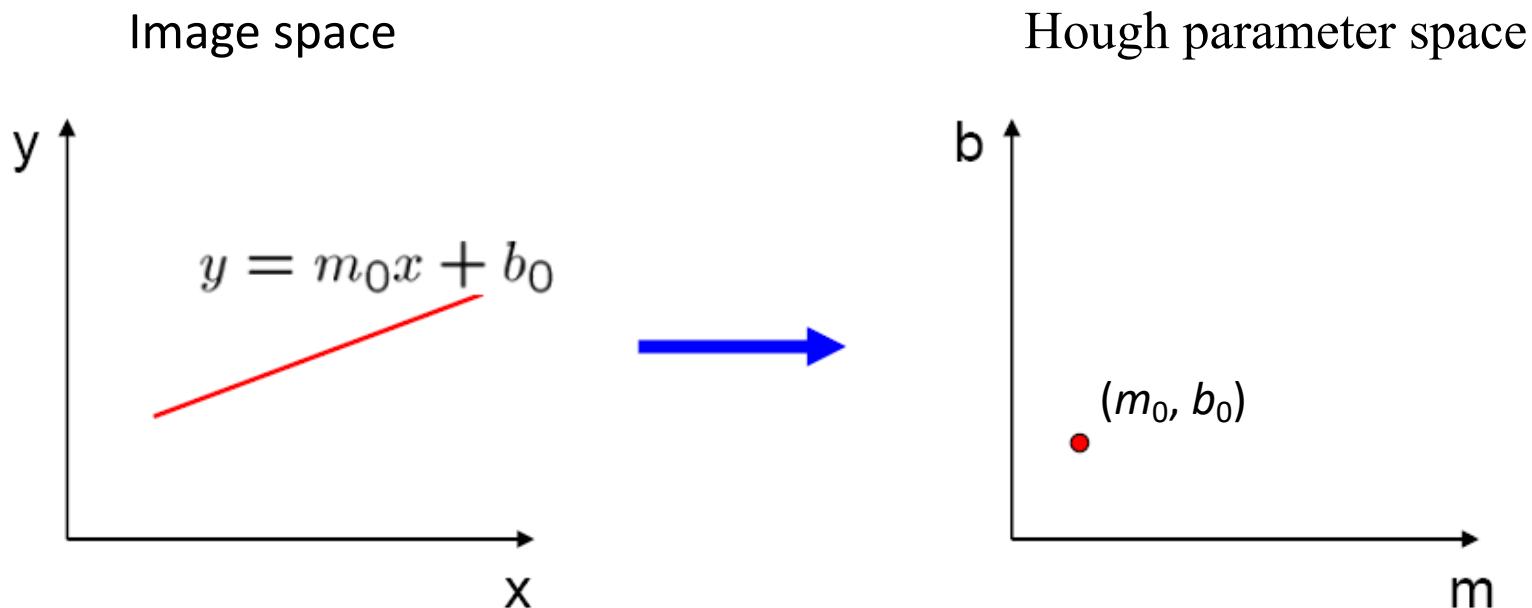
m – slope (defines the angle between the line and Ox axis),

b – intercept (bias to origin $O(0,0)$)

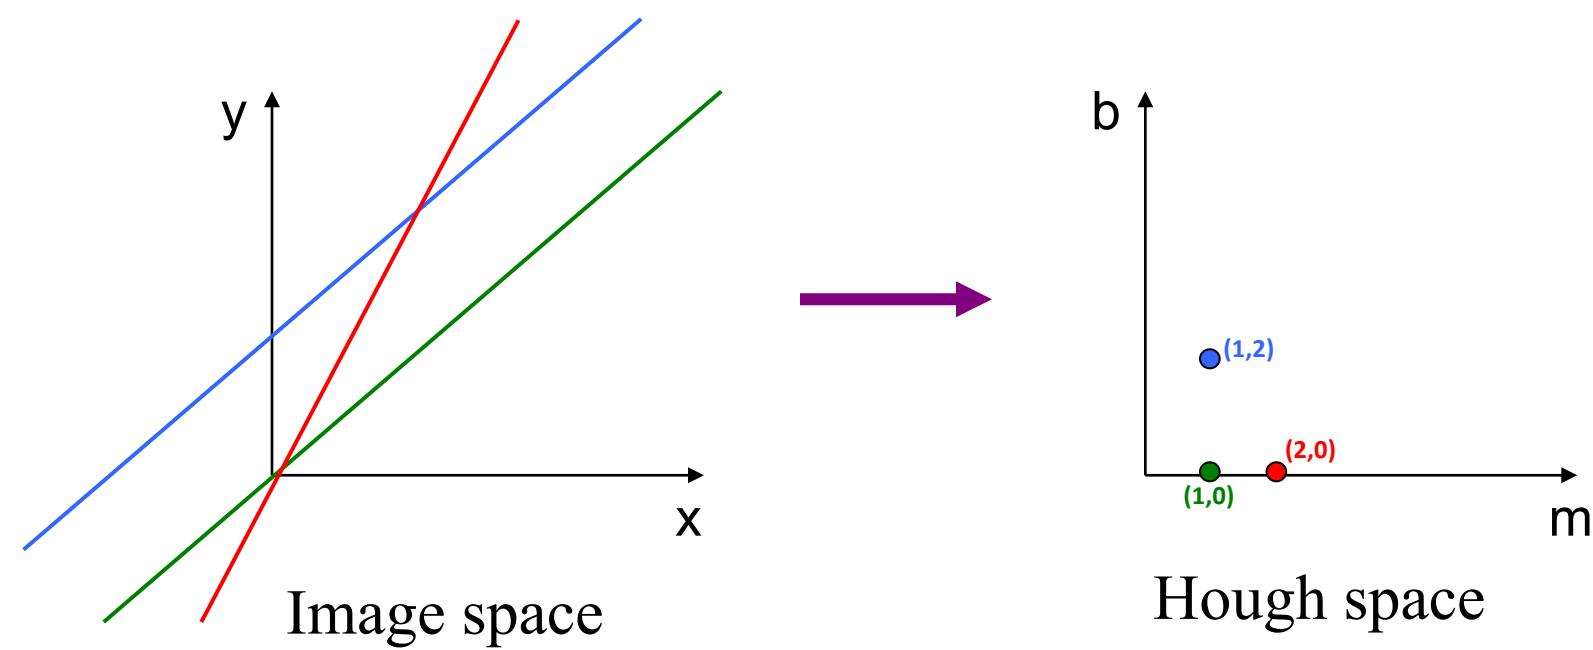
(m,b) are line parameters

Parameter space representation

- A **line** in the image corresponds to a **point** in Hough space



Hough parameter space



$d_1: y = x$, so $m=1, b = 0$

$d_2: y = x + 2$, so $m=1, b = 2$

$d_3: y = 2x$, so $m=2, b = 0$

$d_1: m=1, b = 0 \rightarrow (1, 0)$

$d_2: m=1, b = 2 \rightarrow (1, 2)$

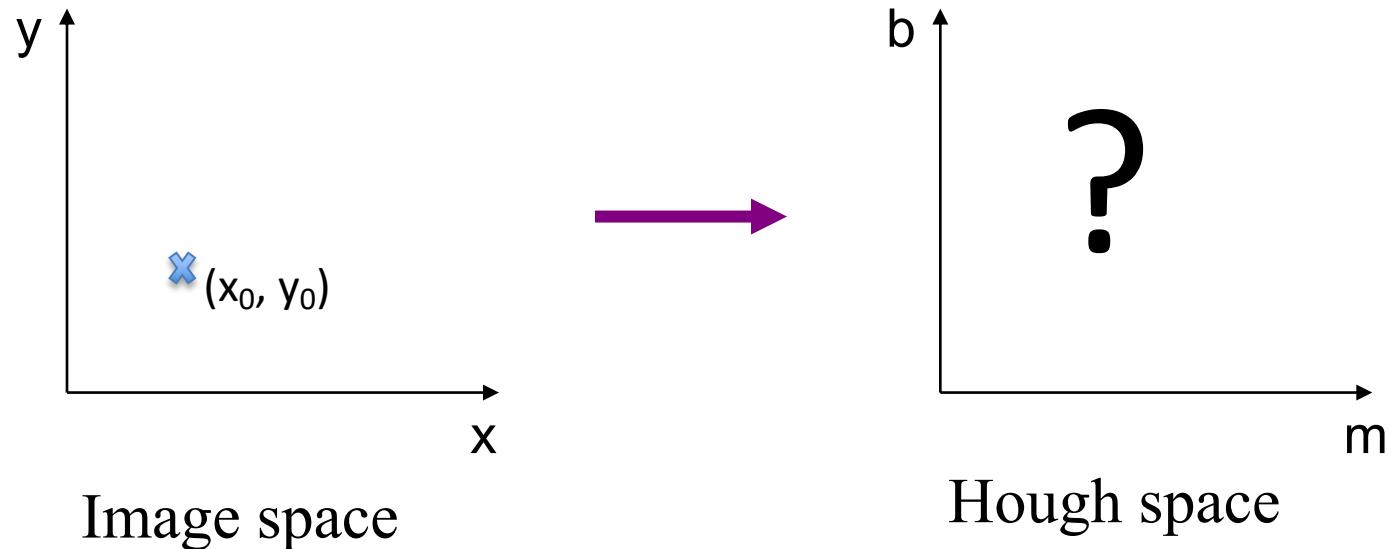
$d_3: m=2, b = 0 \rightarrow (2, 0)$

A line in the image space corresponds to a point in the Hough space.

Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.

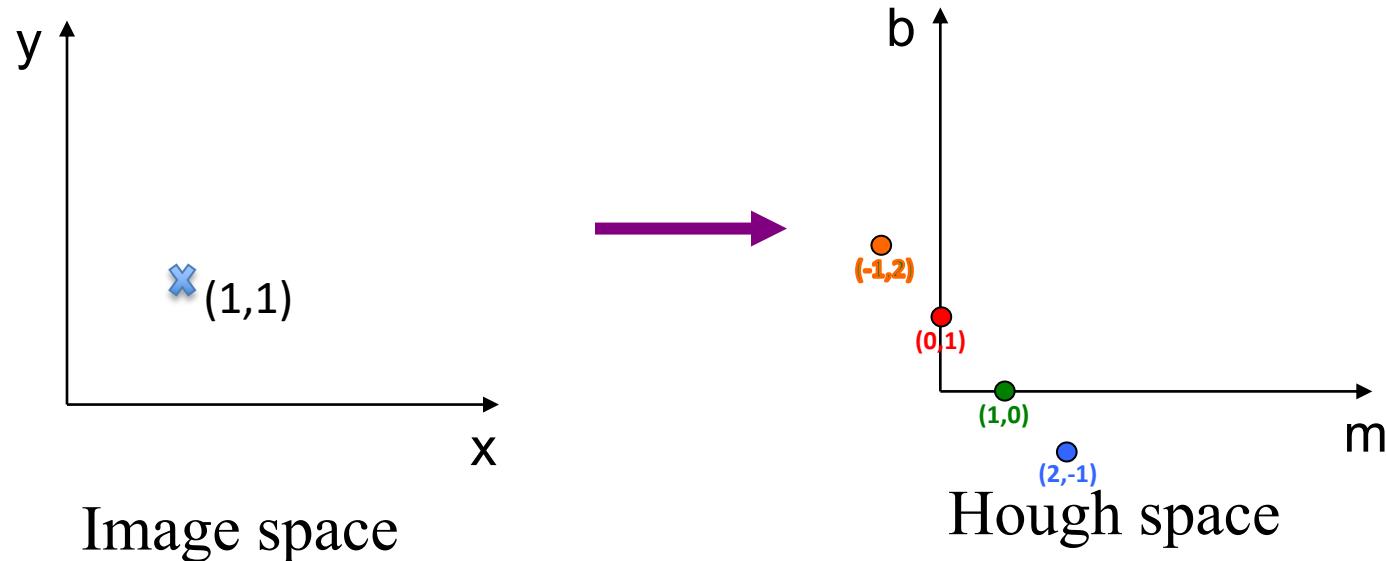
What does a **point** (x_0, y_0) in the image space map to in the Hough space?



Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.

What does a **point** (x_0, y_0) in the image space map to in the Hough space?



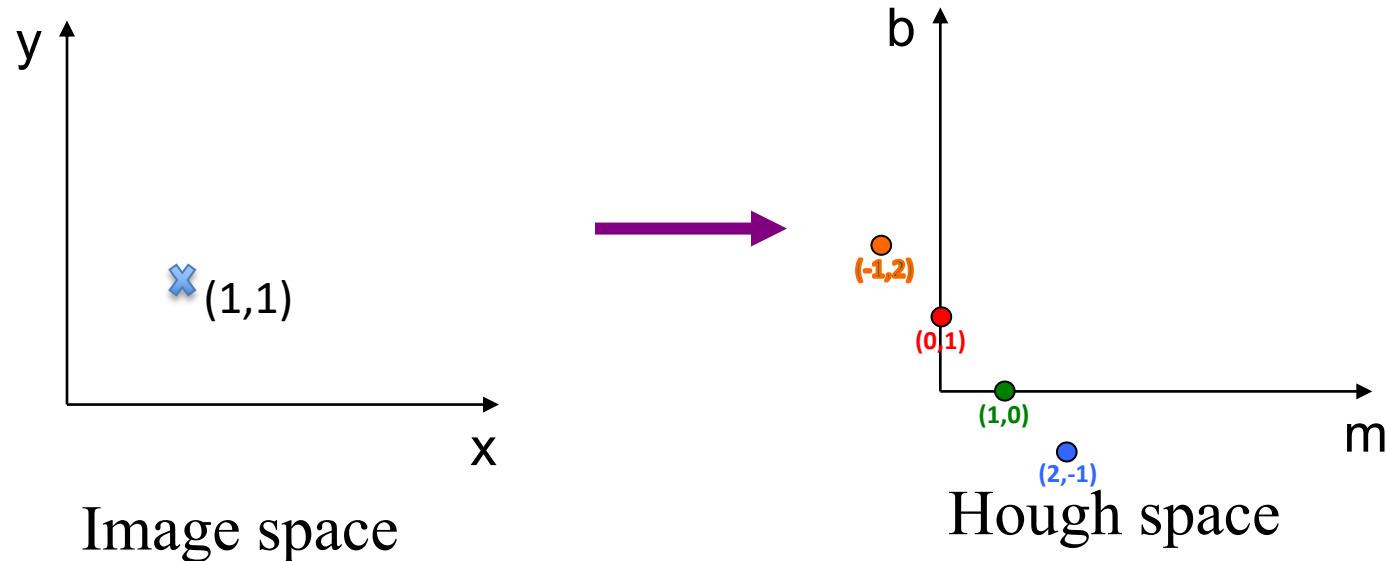
There is an infinity of lines of the form $y = mx + b$ that contain point $(1,1)$:

$$y = x \quad (m=1, b=0); \quad y = 2x-1 \quad (m=2, b=-1), \quad y = 1 \quad (m=0, b = 1), \quad y = -x + 2 \quad (m=-1, b = 2),$$

Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.

A **point** (x_0, y_0) in the image space maps to a line in the Hough space



There is an infinity of lines of the form $y = mx + b$ that contain point $(1,1)$:

$$y = x \quad (m=1, b=0); \quad y = 2x-1 \quad (m=2, b=-1), \quad y = 1 \quad (m=0, b = 1), \quad y = -x + 2 \quad (m=-1, b = 2),$$

All 4 points $(1,0)$, $(2,-1)$, $(0,-1)$, $(-1,2)$ are on the line $b = -m+1$

Correspondence between image and Hough spaces

A line in the image space corresponds to a point in the Hough space.

A **point** (x_0, y_0) in the image space maps to a line in the Hough space

Parameterization of a line in image plane: $y = m * x + b,$

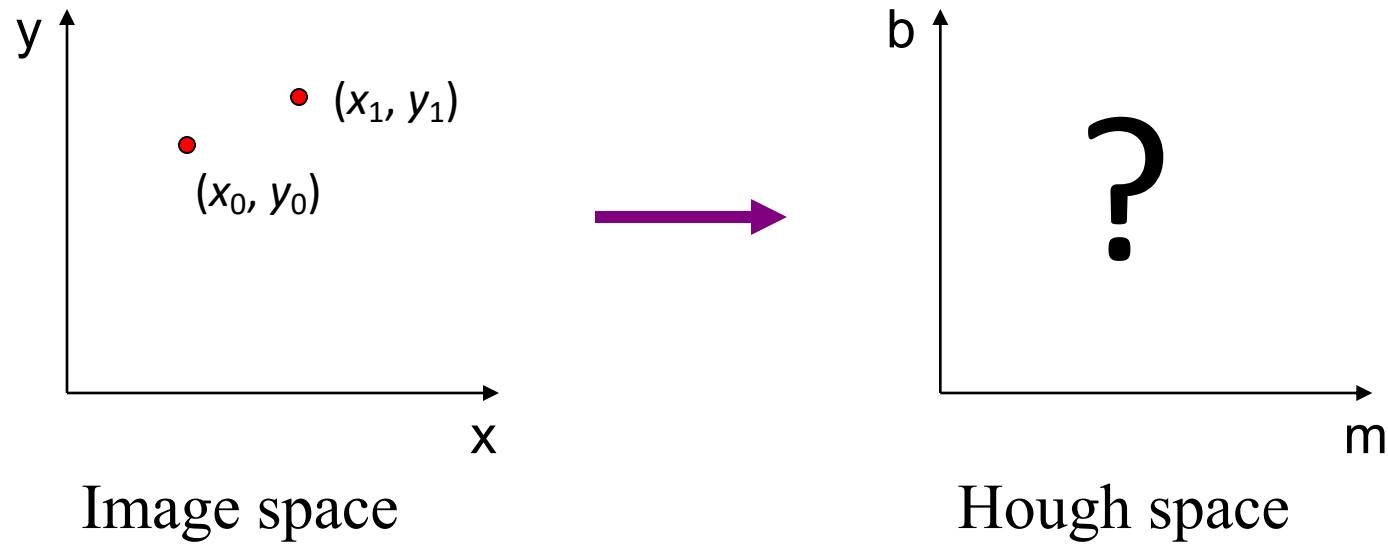
All lines that contain point (x_0, y_0) can have any slope m and any intercept b , the only condition that should be satisfied is

$$y_0 = m * x_0 + b.$$

$y_0 = m * x_0 + b \Leftrightarrow b = -x_0 * m + y_0$ (parameterization of the line in the Hough space with slope $-x_0$ and intercept y_0)

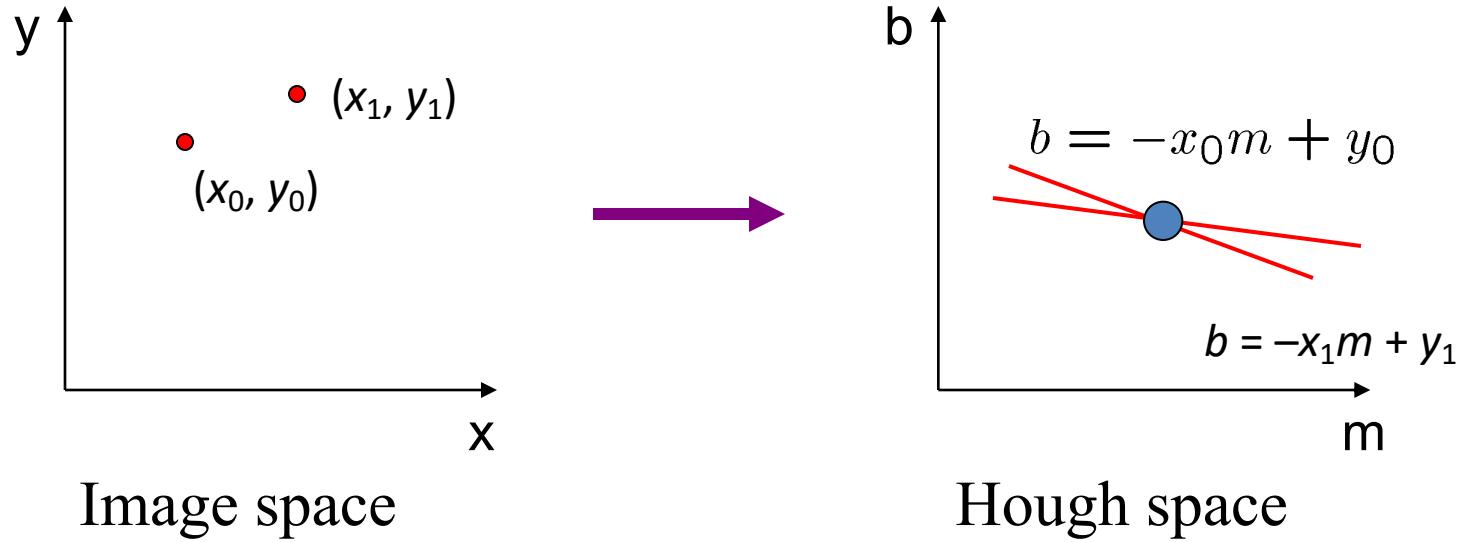
$$(x_0, y_0) = (1, 1) \Rightarrow b = -m + 1$$

Parameter space representation



Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?

Parameter space representation



Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?

Point (x_0, y_0) maps to the line in the Hough space: $b = -x_0m + y_0$

Point (x_1, y_1) maps to the line in the Hough space: $b = -x_1m + y_1$

The line that contains both (x_0, y_0) and (x_1, y_1) maps to a point in the Hough space. This point is the intersection of lines $b = -x_0m + y_0$ și $b = -x_1m + y_1$

Finding lines in an image with Hough

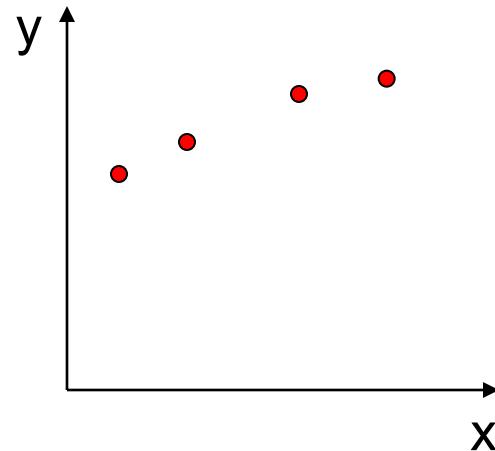
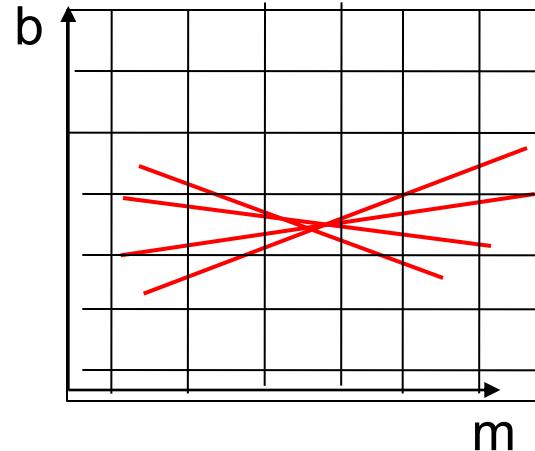


Image space



Hough space

How we can use the previous observations to find parameters (m, b) that define the most probable line in the image space?

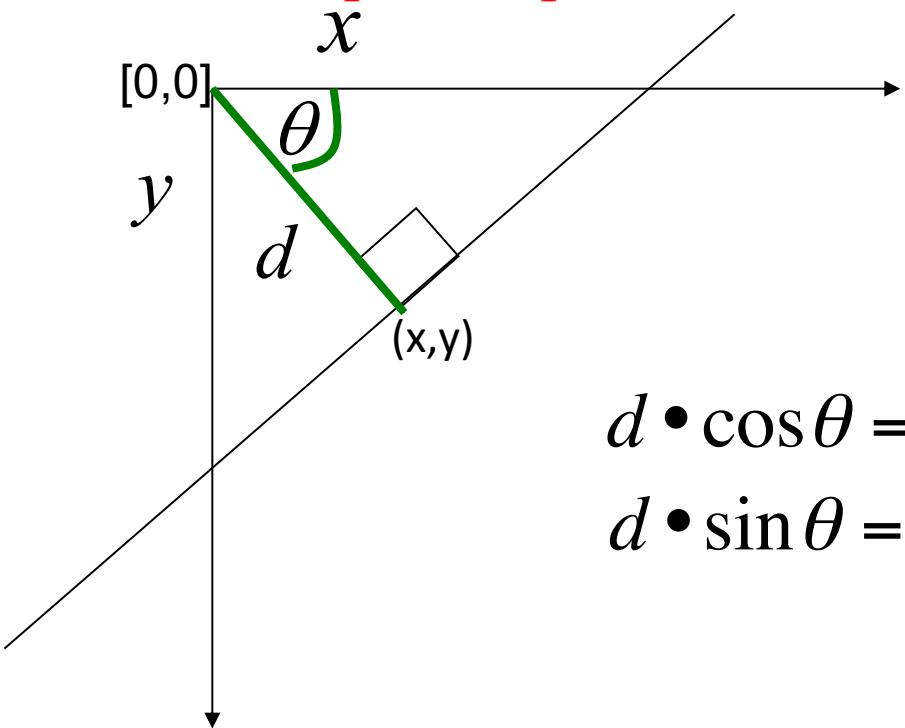
- each edgel point (part of an edge) from image space will vote (for an infinity of lines) in the Hough space (here a line is a point)
- discretize the Hough space in discrete intervals/cells; the cell with largest number of votes determines the parameters of the line

Polar representation

Problems with the (m,b) space:

- Unbounded parameter domains
- Vertical lines require infinite m

Alternative: **polar representation**



d : distance from origin to line
(length of the perpendicular)

θ : angle between the
perpendicular and Ox axis

$$d \cdot \cos \theta = x$$

$$d \cdot \sin \theta = y$$

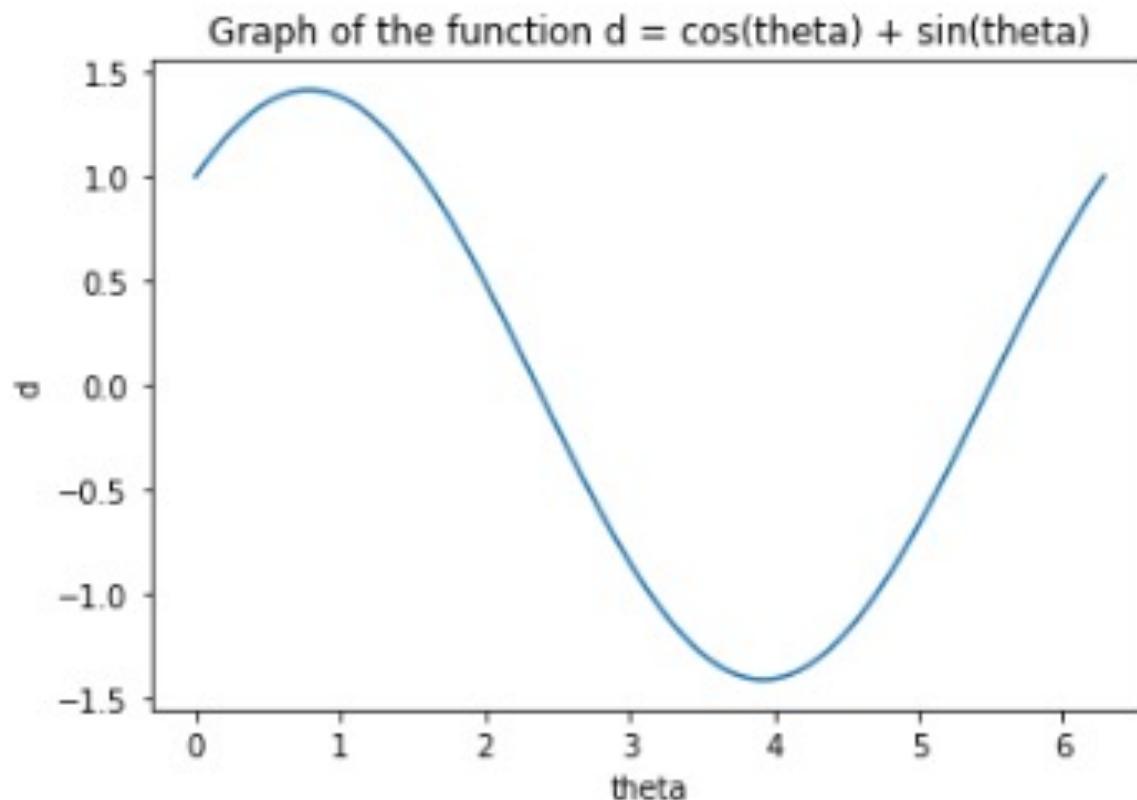
$$d \cdot (\cos \theta)^2 = x \cdot \cos \theta$$

$$d \cdot (\sin \theta)^2 = y \cdot \sin \theta$$

$$d = x \cdot \cos \theta + y \cdot \sin \theta$$

Point in the image space \rightarrow sinusoid curve in the Hough space

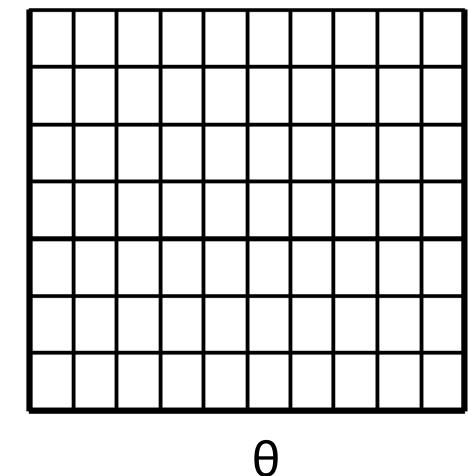
```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 x, y = 1,1
4 theta = np.linspace(0,2*np.pi,100)
5 d = x * np.cos(theta) + y * np.sin(theta)
6 plt.figure,
7 plt.plot(theta, d)
8 plt.xlabel('theta')
9 plt.ylabel('d')
10 plt.title('Graph of the function d = cos(theta) + sin(theta)')
11 plt.show()
```



Hough transform-algorithm outline

- Initialize accumulator H to all zeros
- For each feature point (x, y) in the image
 - For $\theta = 0$ to 180
 - $d = x \cos \theta + y \sin \theta$
 - $H(\theta, d) = H(\theta, d) + 1$
 - end
- end
- Find the value(s) of (θ, d) where $H(\theta, d)$ is a local maximum
 - The detected line in the image is given by
 $d = x \cos \theta + y \sin \theta$

H : accumulator array (votes)

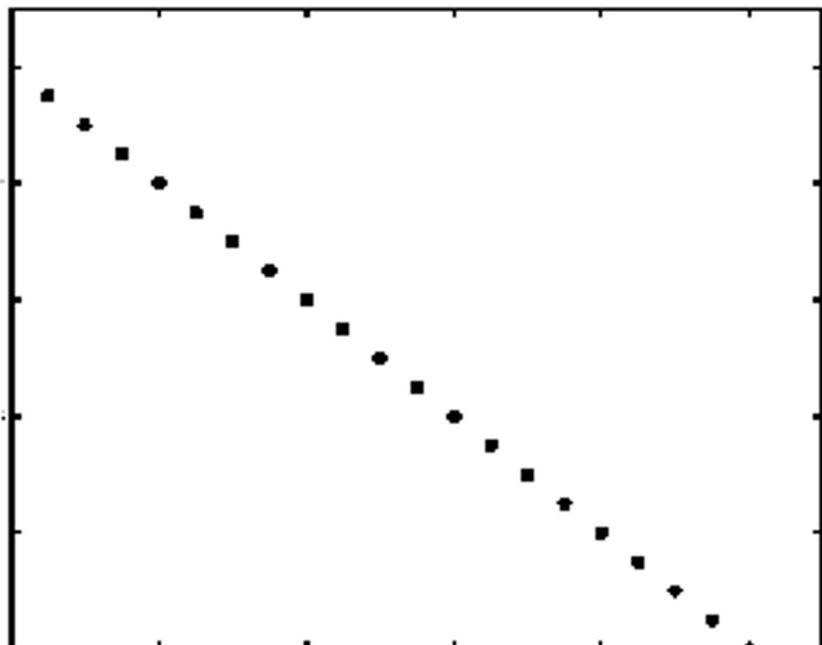


θ

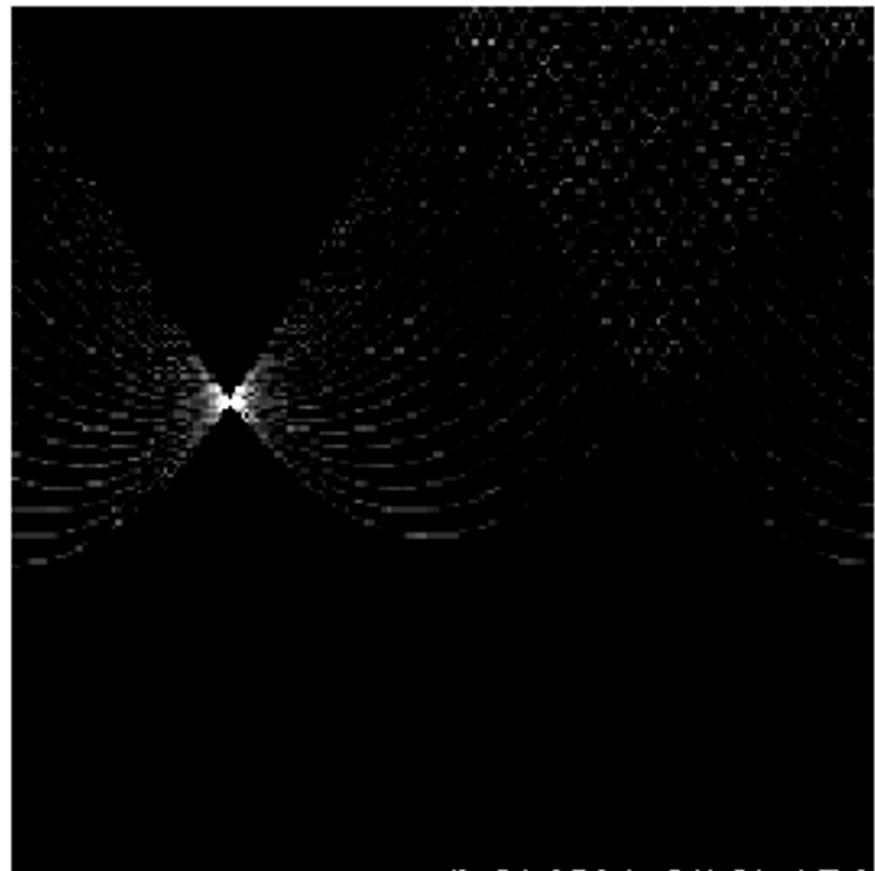
Hough transform - demo

<https://www.aber.ac.uk/~dcswww/Dept/Teaching/CourseNotes/current/CS34110/hough.html>

Basic illustration



features

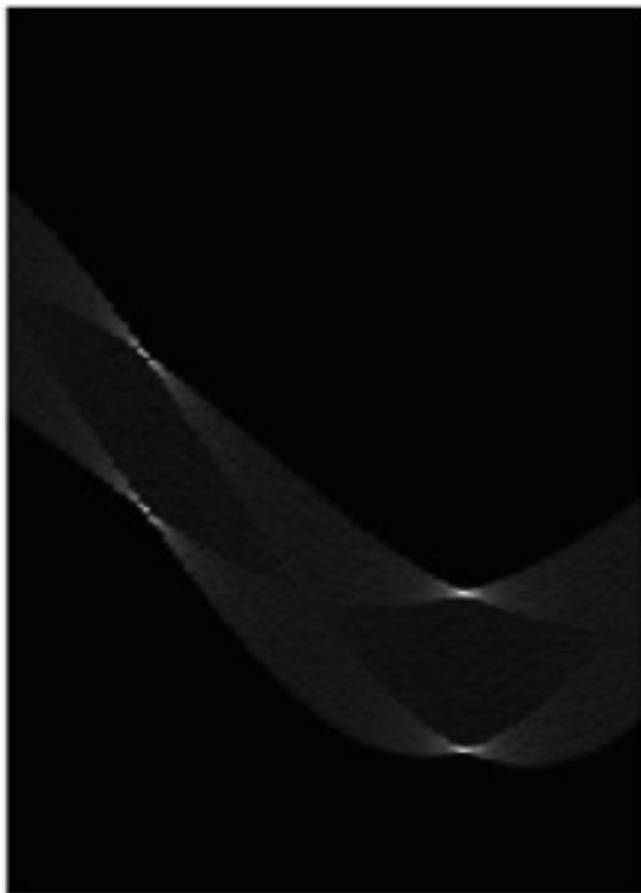


votes

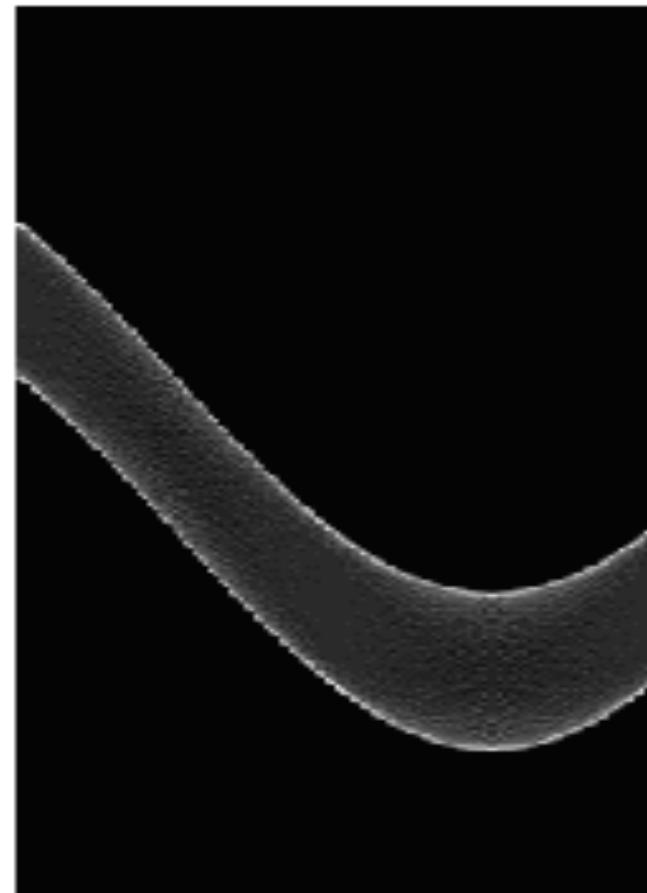
[Hough transform demo](#)

Other shapes

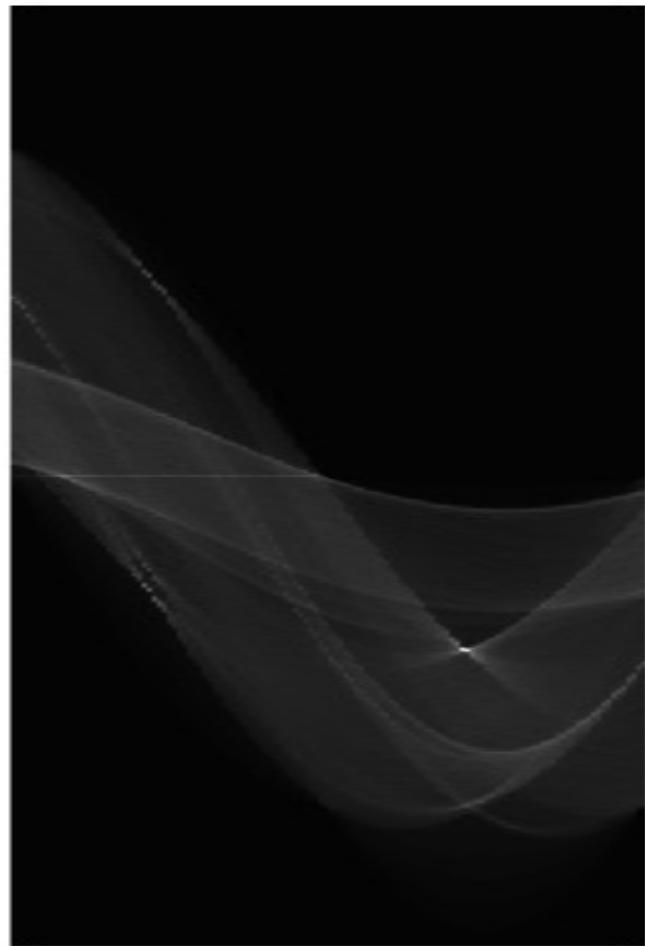
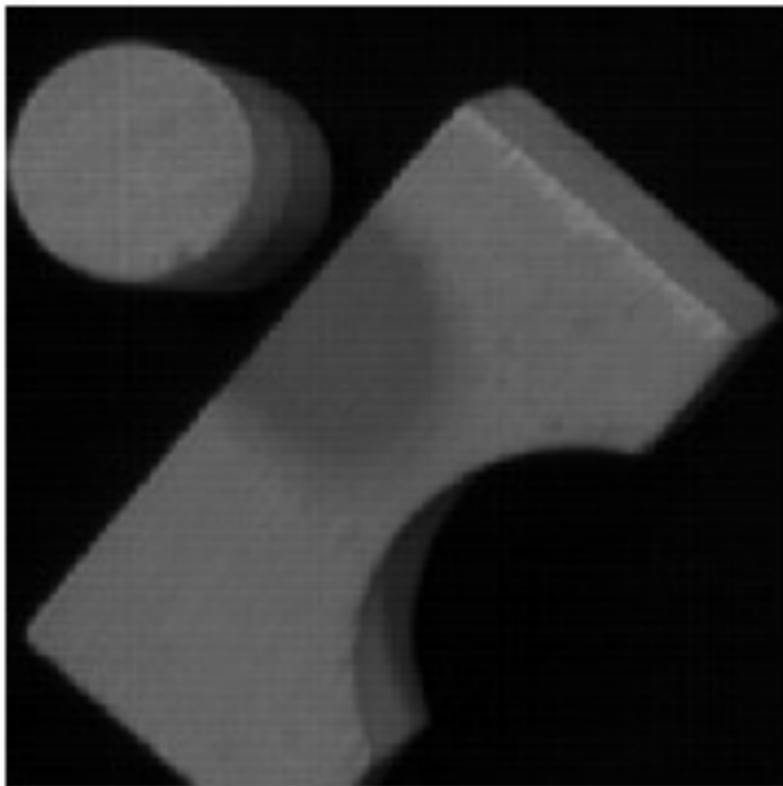
Square



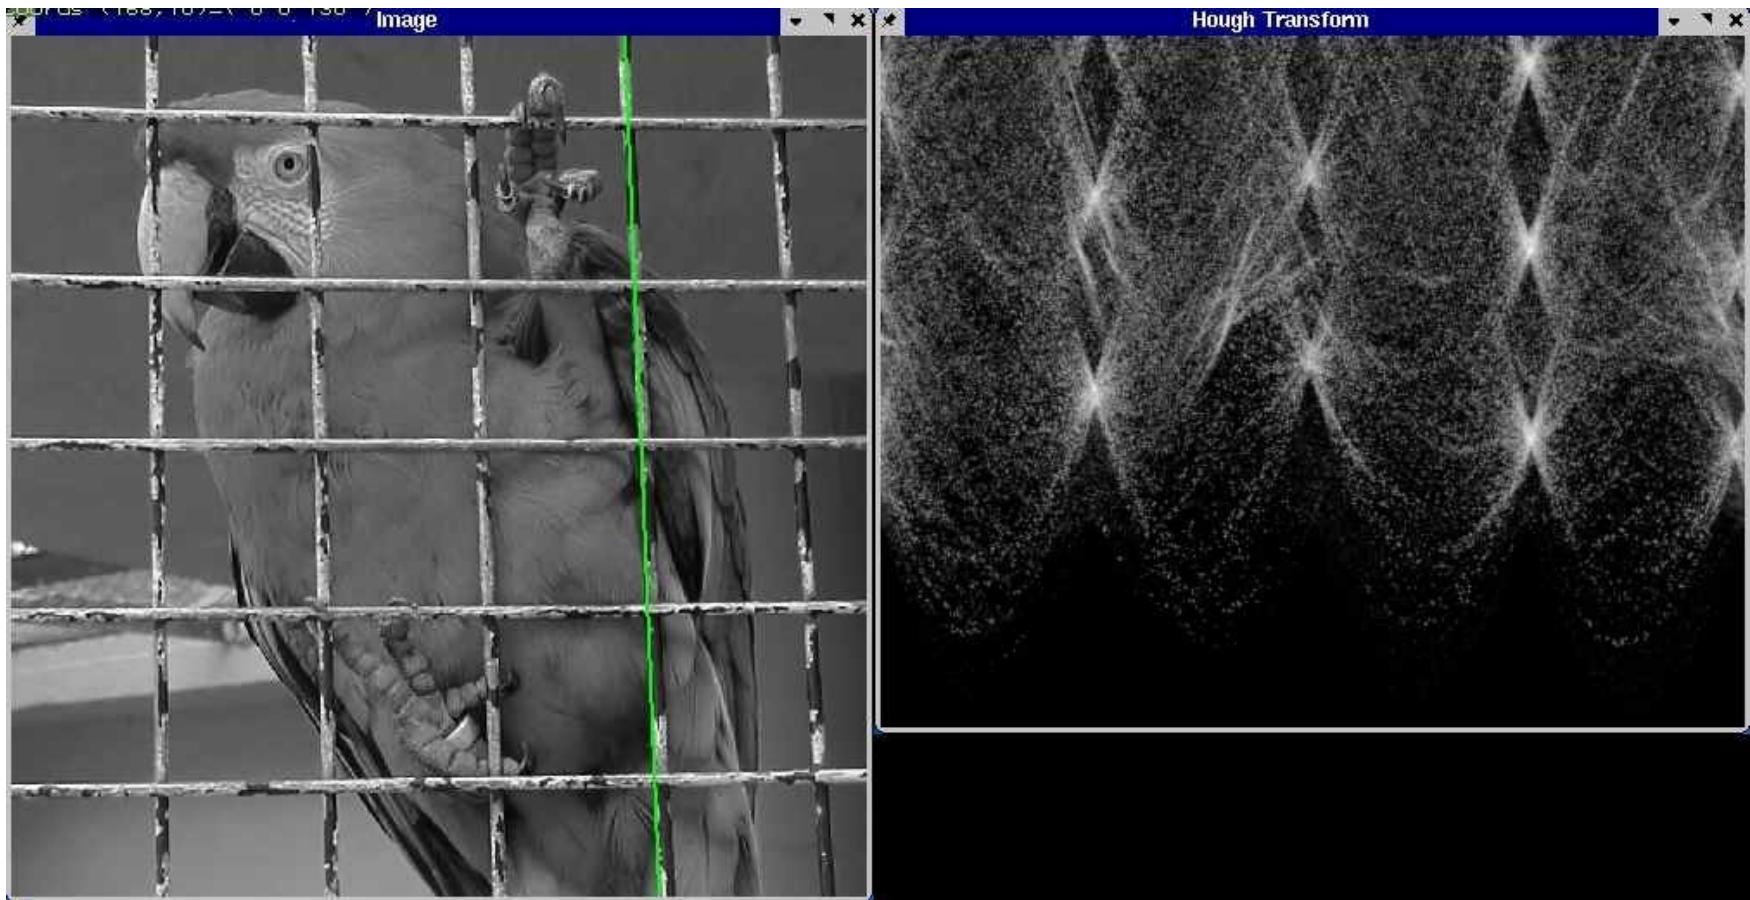
Circle



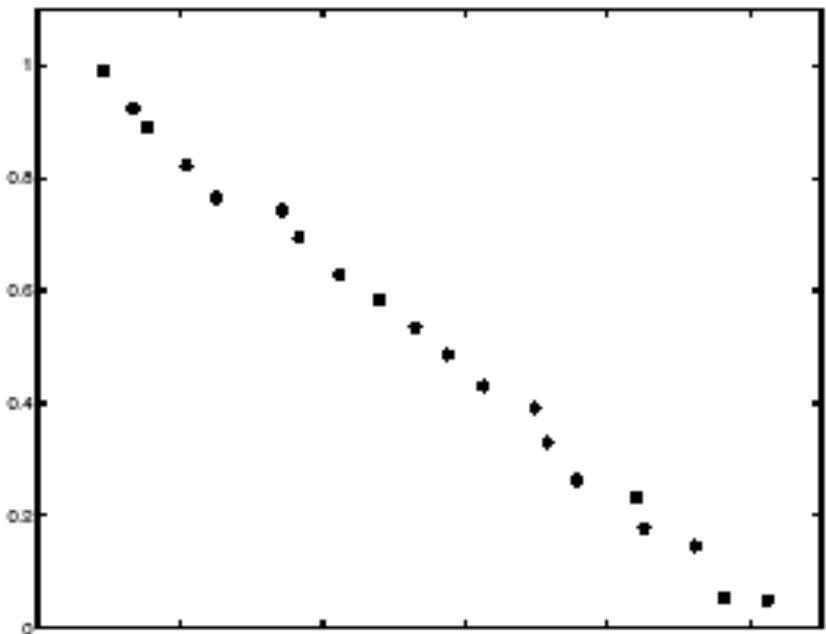
Several lines



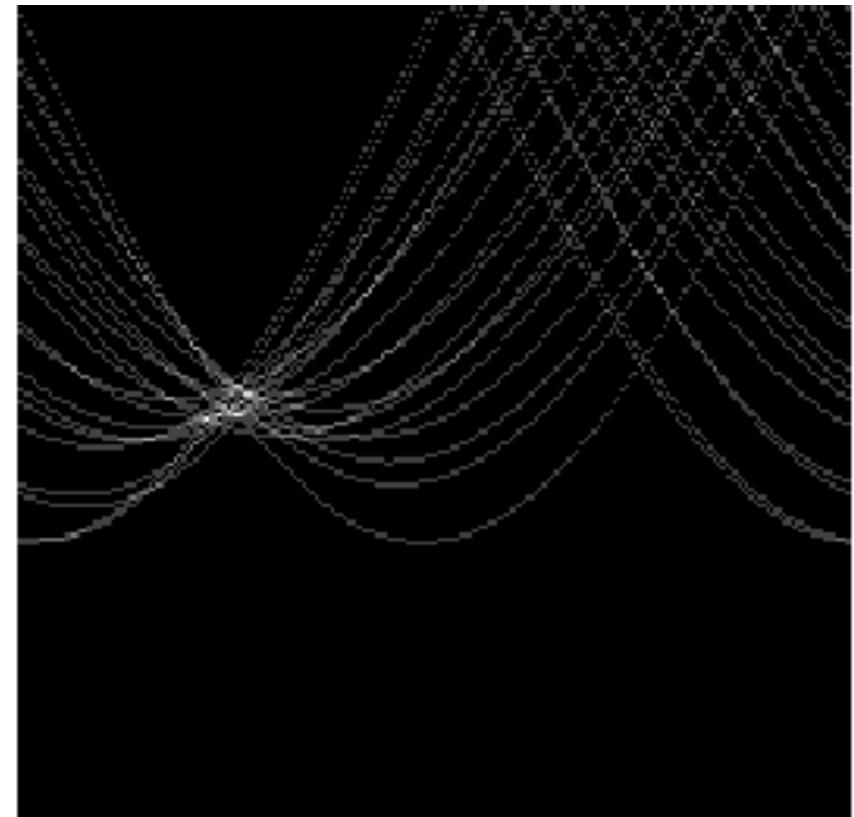
A more complex image



Effect of noise



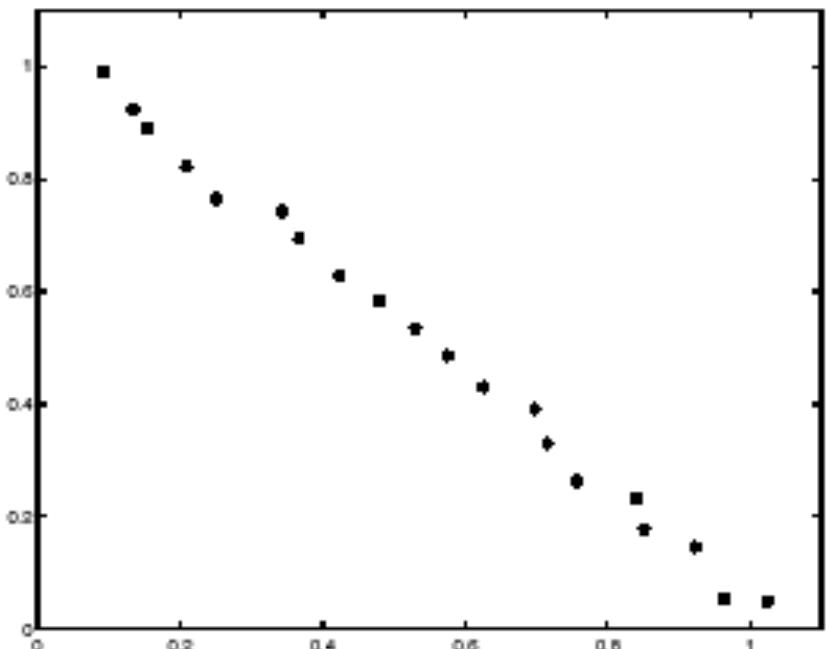
features



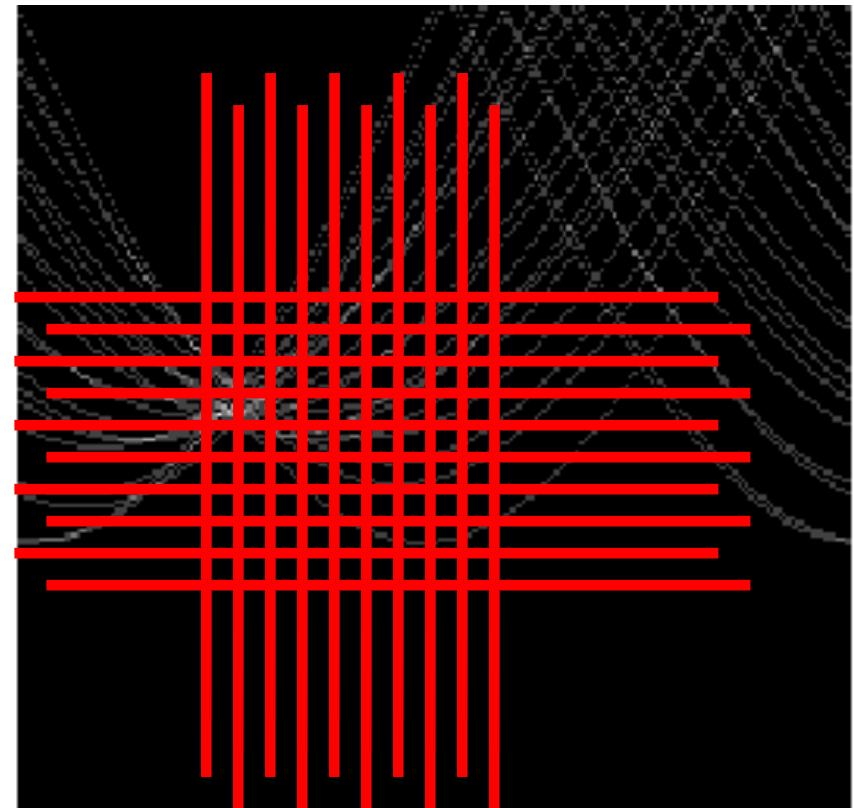
votes

- Peak gets fuzzy and hard to locate

Effect of noise



features

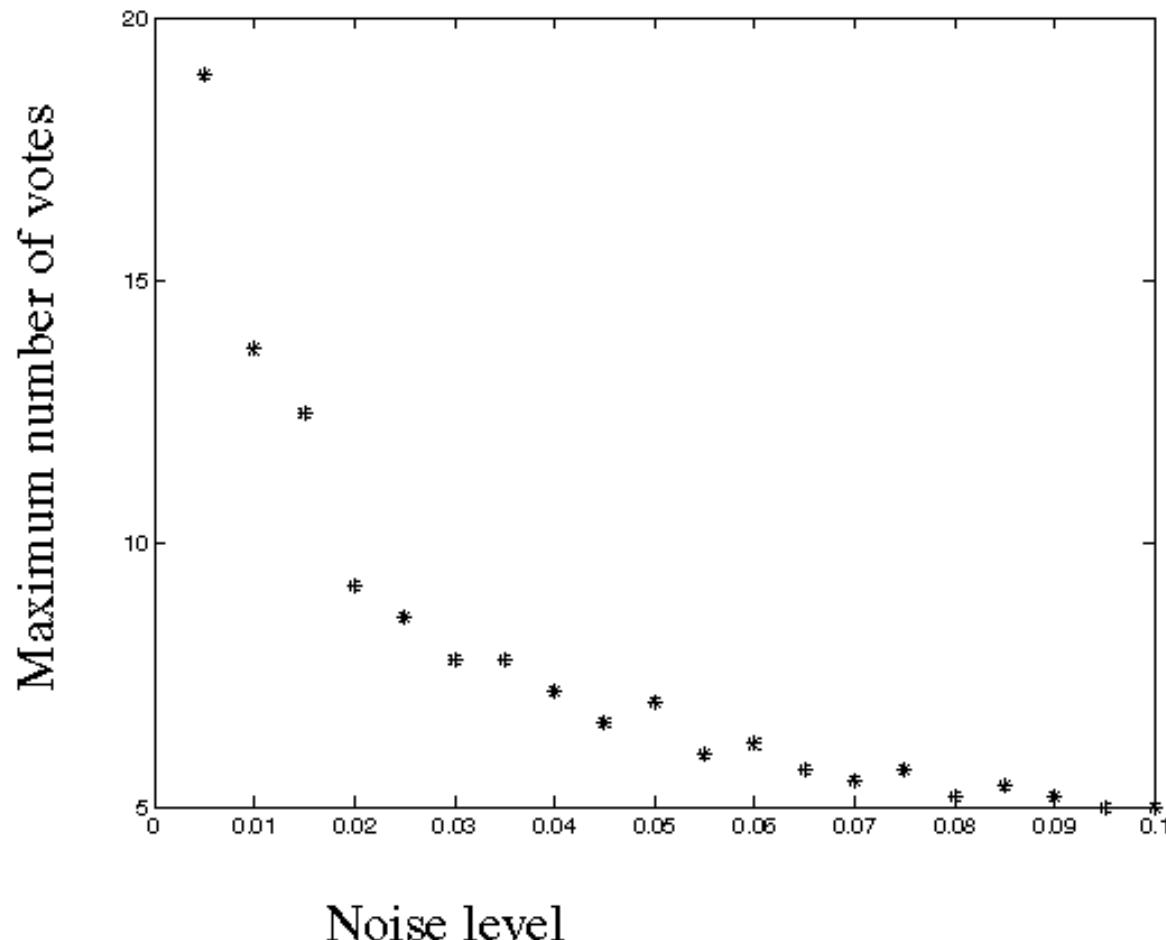


votes

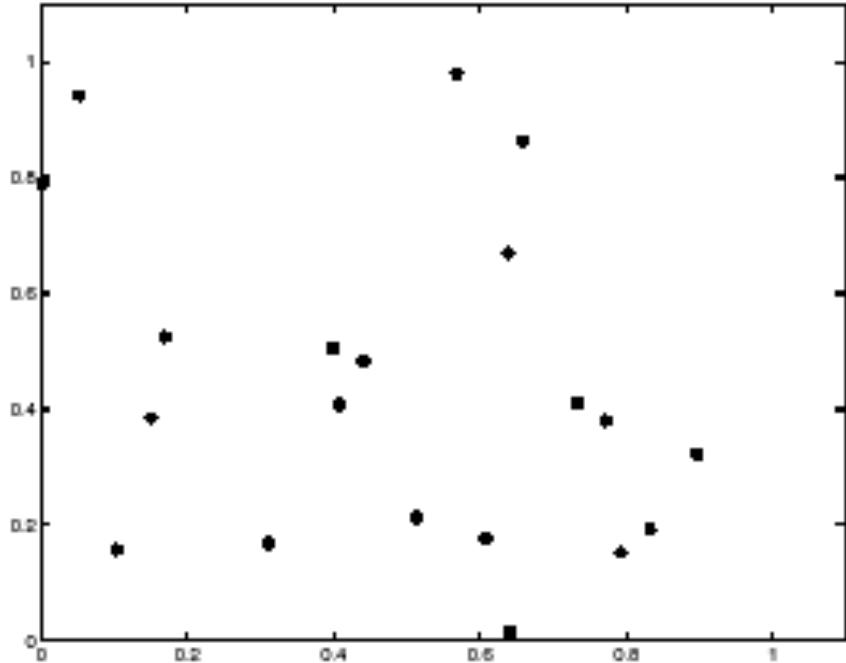
- Peak gets fuzzy and hard to locate

Effect of noise

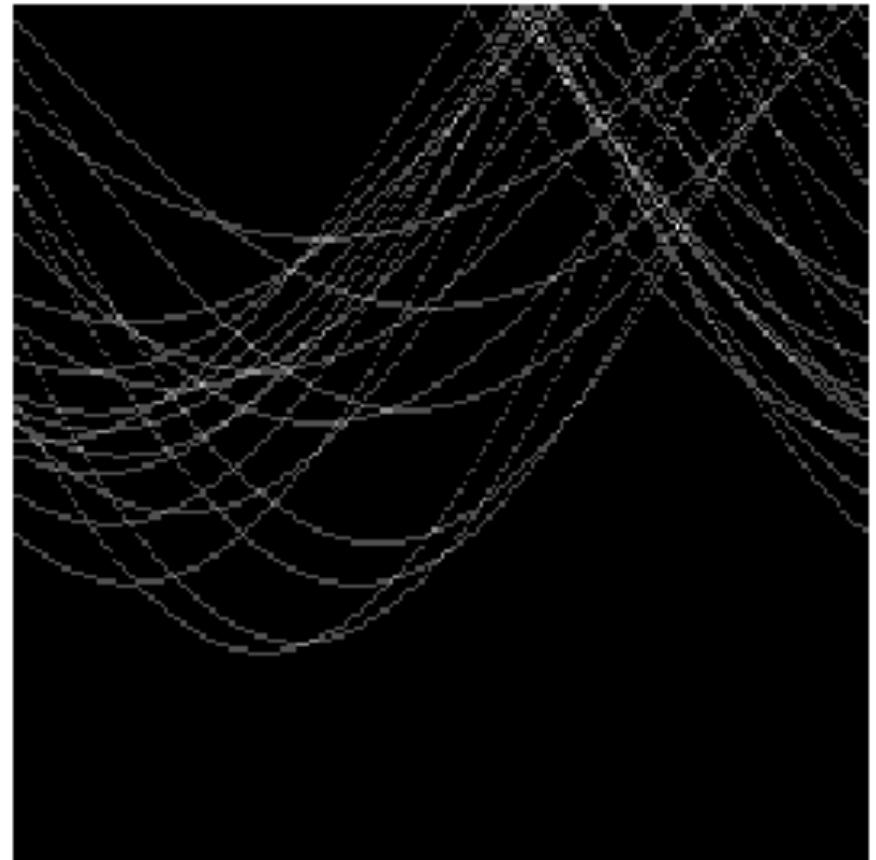
- Number of votes for a line of 20 points with increasing noise:



Random points



features

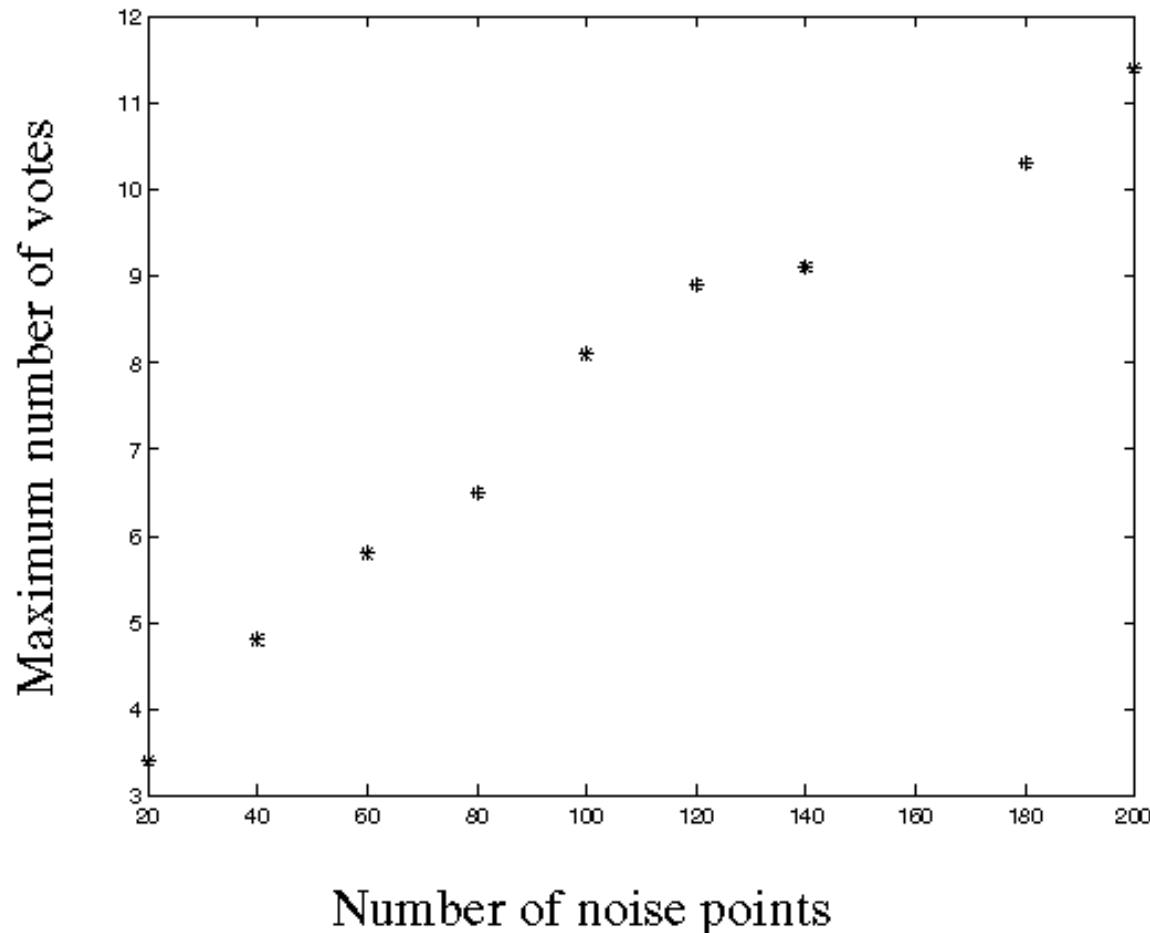


votes

- Uniform noise can lead to spurious peaks in the array

Random points

- As the level of uniform noise increases, the maximum number of votes increases too:

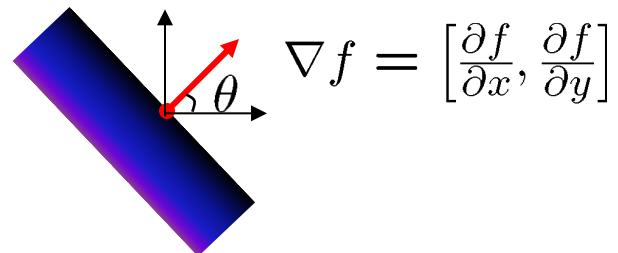


Dealing with noise

- Choose a good grid / discretization
 - **Too coarse:** large votes obtained when too many different lines correspond to a single bucket
 - **Too fine:** miss lines because some points that are not exactly collinear cast votes for different buckets
- Increment neighboring bins (smoothing in accumulator array)
- Try to get rid of irrelevant features
 - E.g., take only edge points with significant gradient magnitude

Incorporating image gradients

- When we detect an edge point, we also know its gradient orientation
- How does this constrain possible lines passing through the point?



- Modified Hough transform:
- For each edge point (x,y)

θ = gradient orientation at (x,y)

$d = x \cos \theta + y \sin \theta$

$H(\theta, d) = H(\theta, d) + 1$

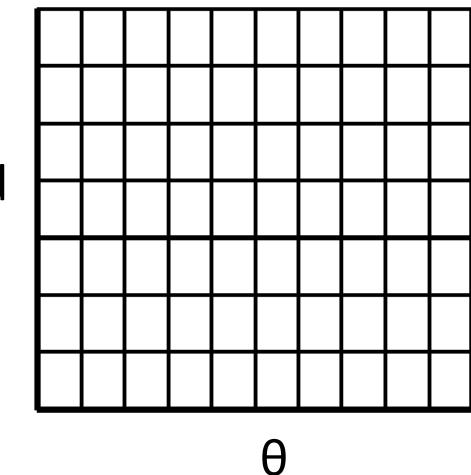
end

Algorithm outline with image gradients

- Initialize accumulator H to all zeros
- For each feature point (x,y) in the image
 - $\theta = \text{gradient orientation at } (x,y)$
 - $d = x \cos \theta + y \sin \theta$
 - $H(\theta, d) = H(\theta, d) + 1$

end

H : accumulator array (votes)

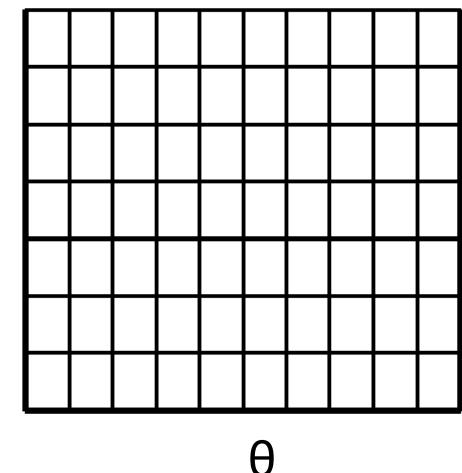


- Find the value(s) of (θ, d) where $H(\theta, d)$ is a local maximum
 - The detected line in the image is given by
 - $d = x \cos \theta + y \sin \theta$

Algorithm outline with image gradients

- Initialize accumulator H to all zeros
- For each feature point (x,y) in the image
 - $\theta = \text{gradient orientation at } (x,y)$
 - $\text{mag} = \text{gradient magnitude at } (x,y)$
 - $d = x \cos \theta + y \sin \theta$
 - $H(\theta, d) = H(\theta, d) + \text{mag}$
- end
- Find the value(s) of (θ, d) where $H(\theta, d)$ is a local maximum
 - The detected line in the image is given by $d = x \cos \theta + y \sin \theta$

H : accumulator array (votes)



Hough transform for circles

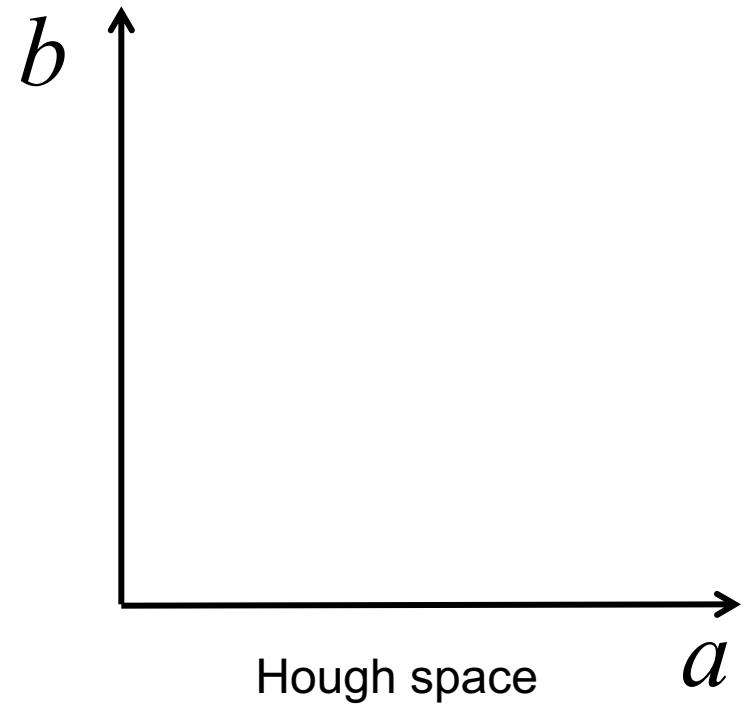
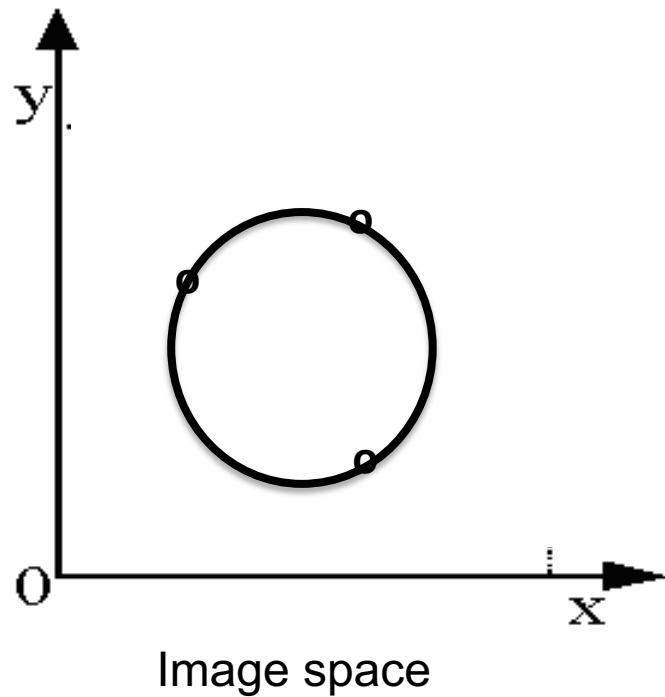
- How many dimensions will the parameter space have?
- Given an unoriented edge point, what are all possible bins that it can vote for?
- What about an *oriented* edge point?

Hough transform for circles

- Circle: center (a, b) and radius r . Equation?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For fixed radius r

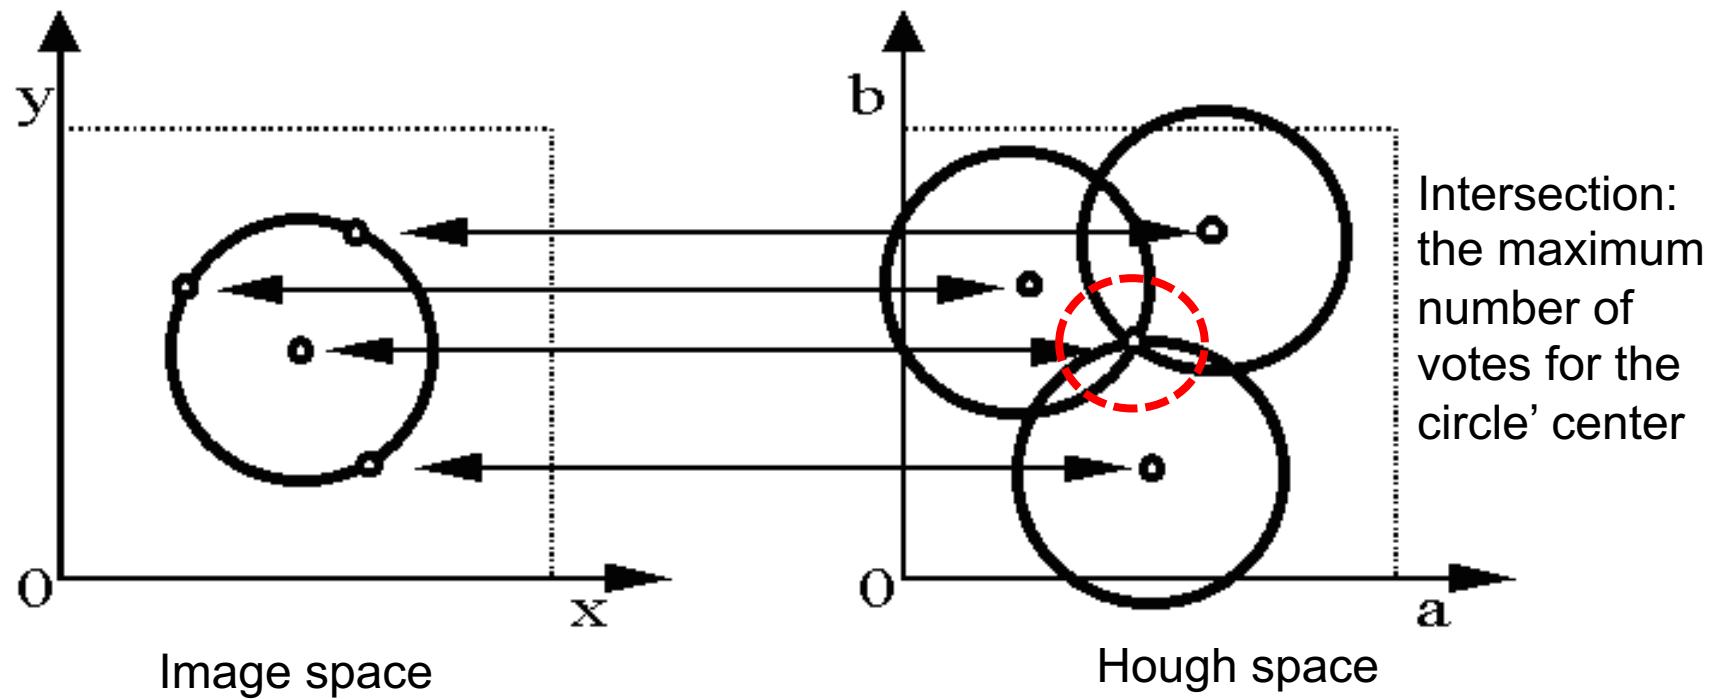


Hough transform for circles

- Circle: center (a, b) and radius r . Equation?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For fixed radius r

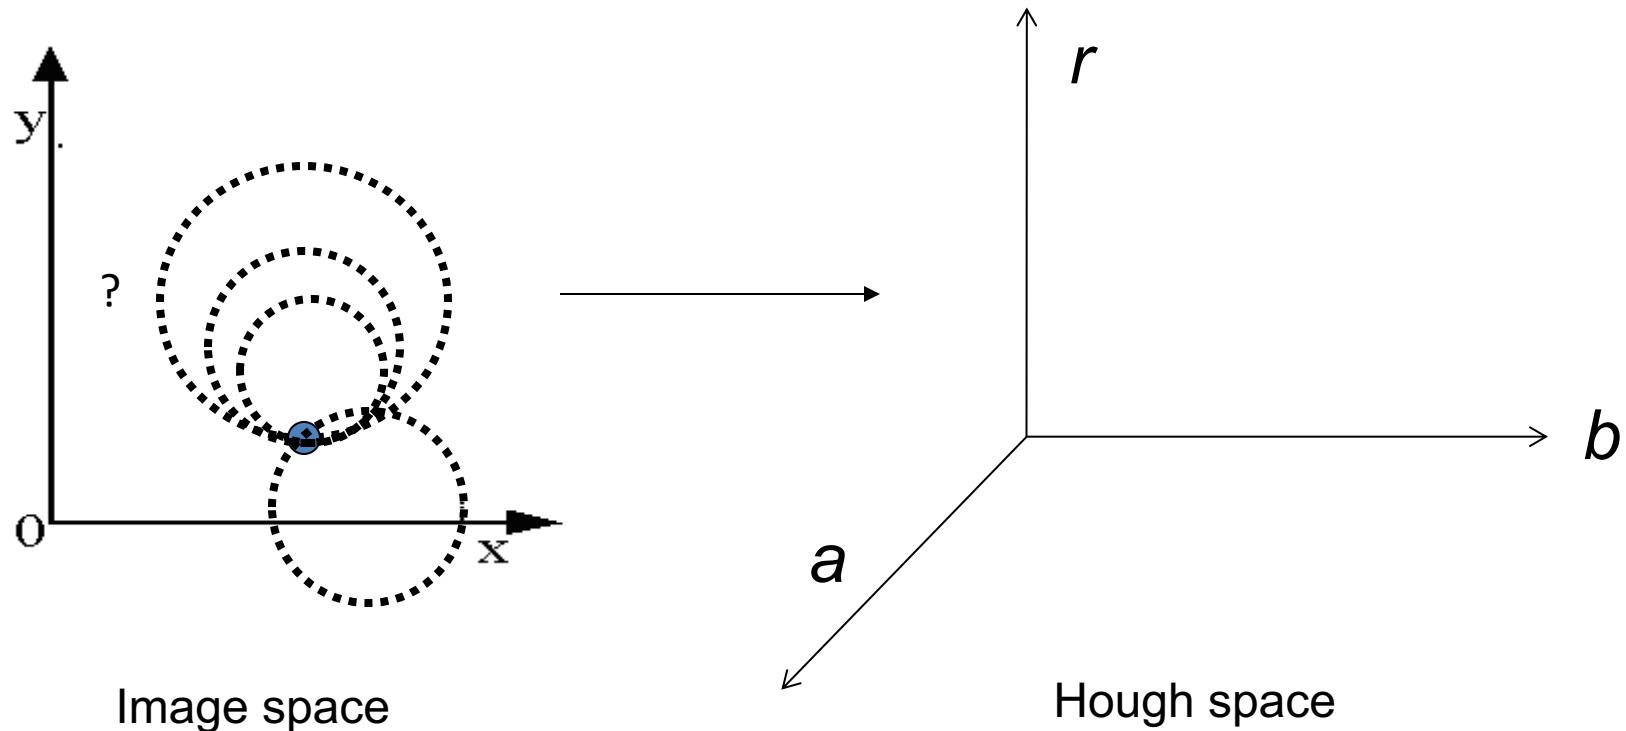


Hough transform for circles

- Circle: center (a, b) and radius r . Equation?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For unknown radius r



Hough transform for circles

- Circle: center (a, b) and radius r . Equation?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For unknown radius r

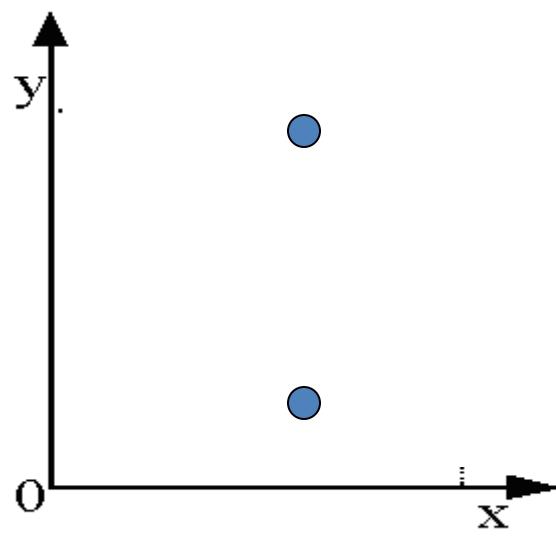
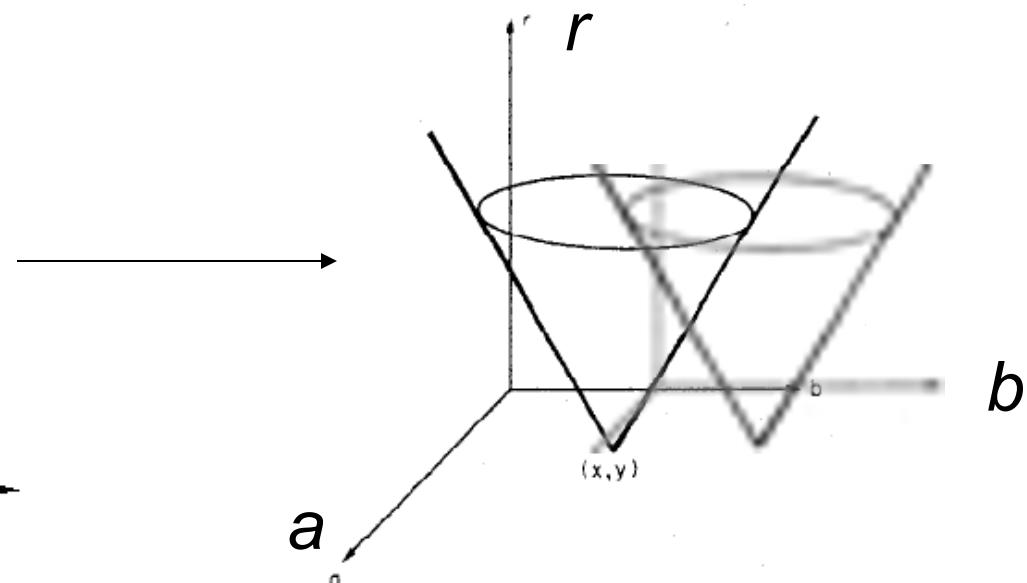
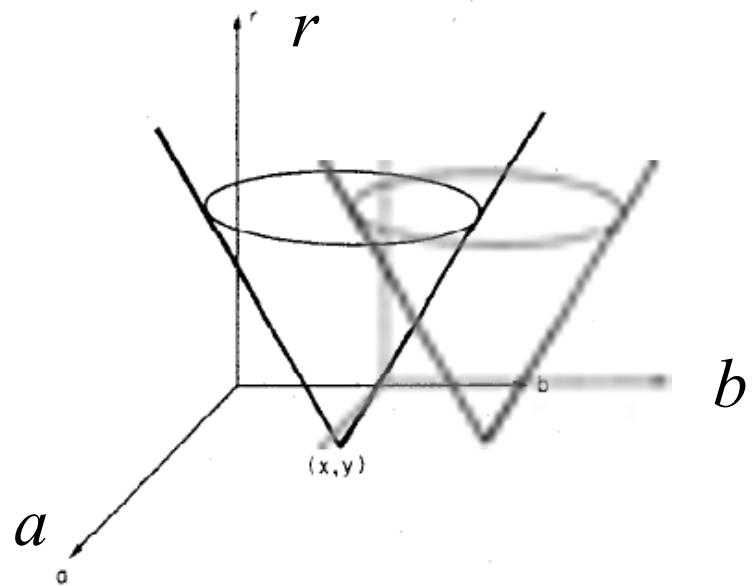
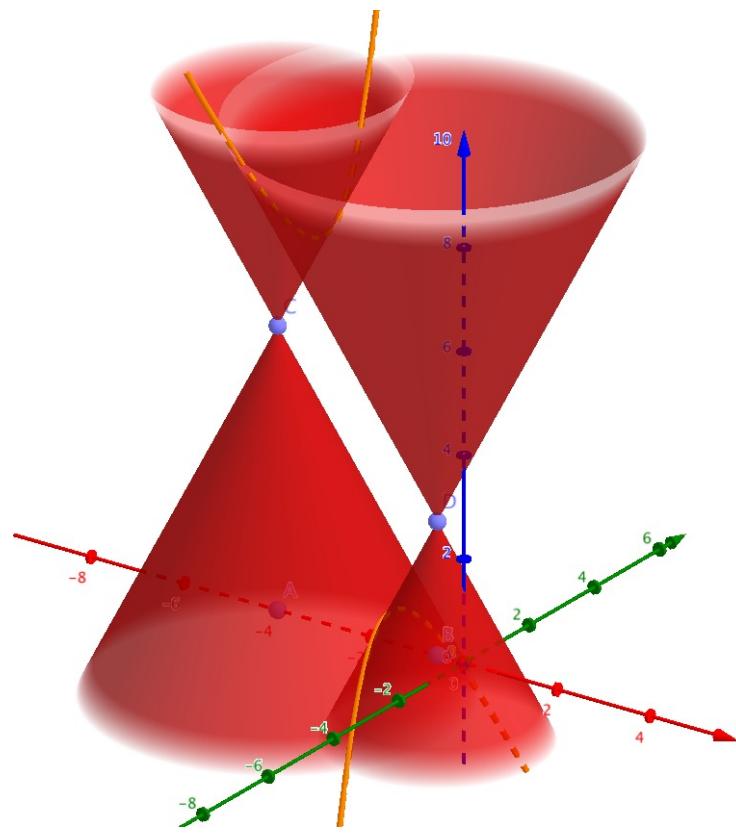


Image space



Hough space

Intersection of two cones



Hough space

Hough transform for circles

- Circle: center (a, b) and radius r . Equation?

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For unknown radius r , with known gradient direction

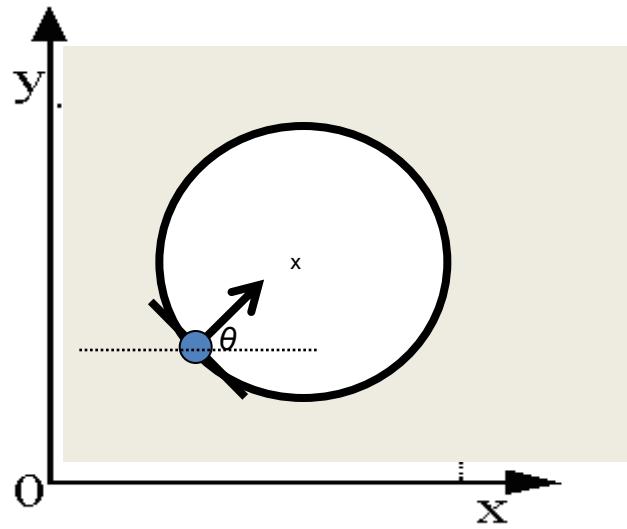
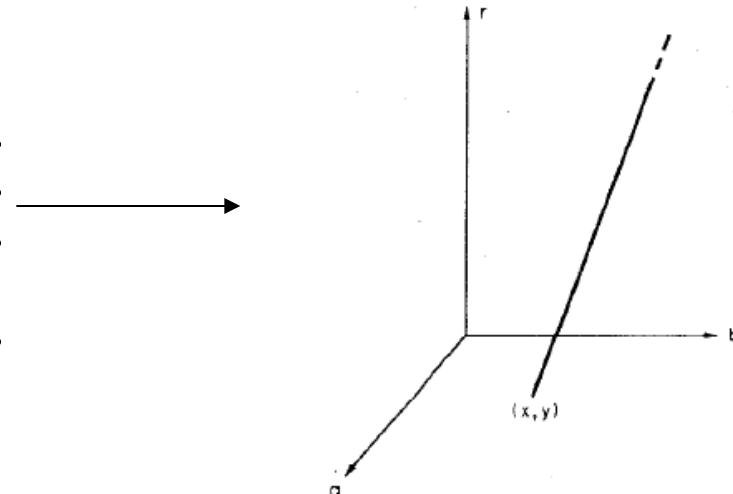


Image space



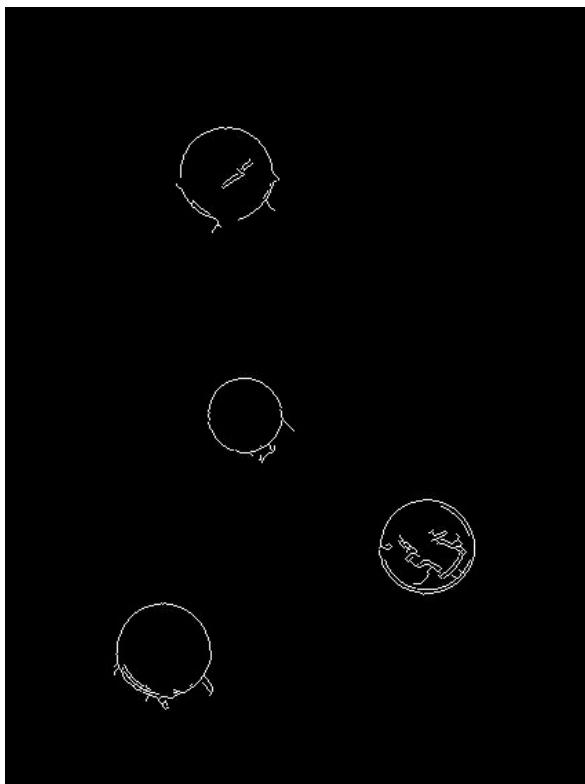
Hough space

Circle detection using Hough transform

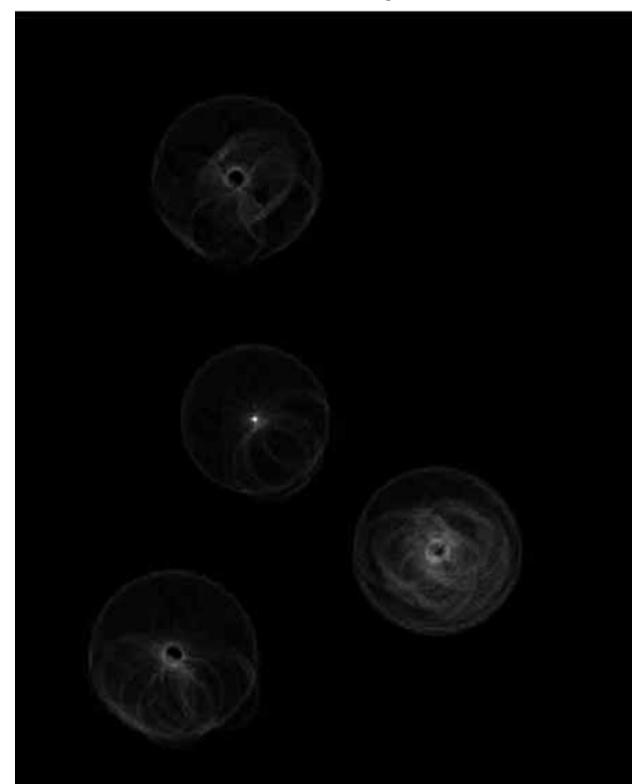
Original image



Edges



Votes ray 1



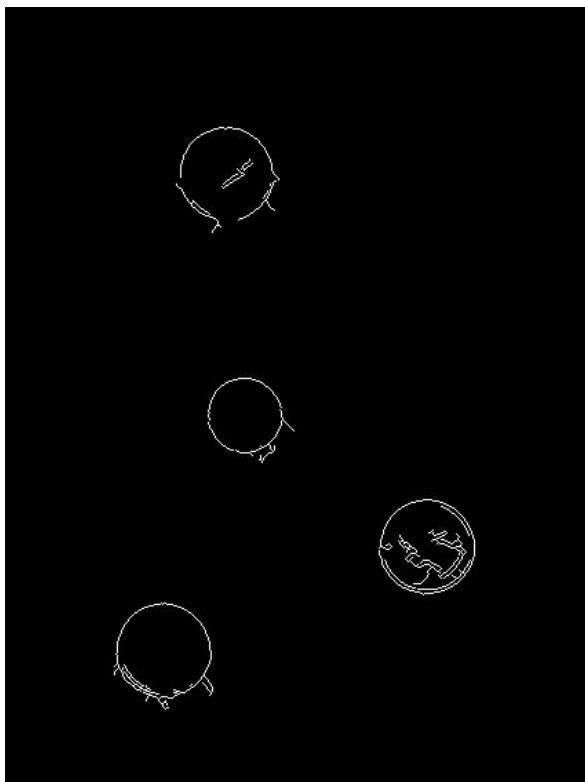
Use two Hough transforms, one for each ray (of known length)

Circle detection using Hough transform

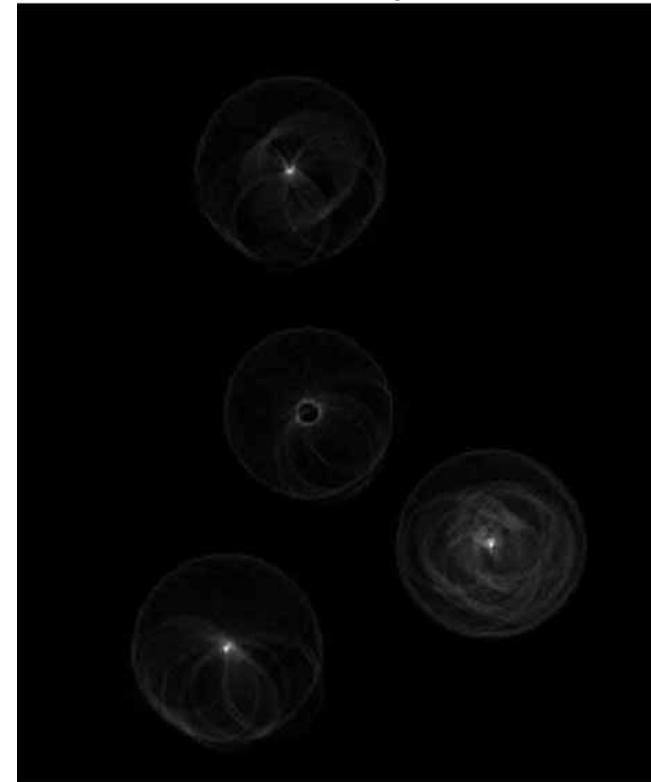
Original image



Edges



Votes ray 2



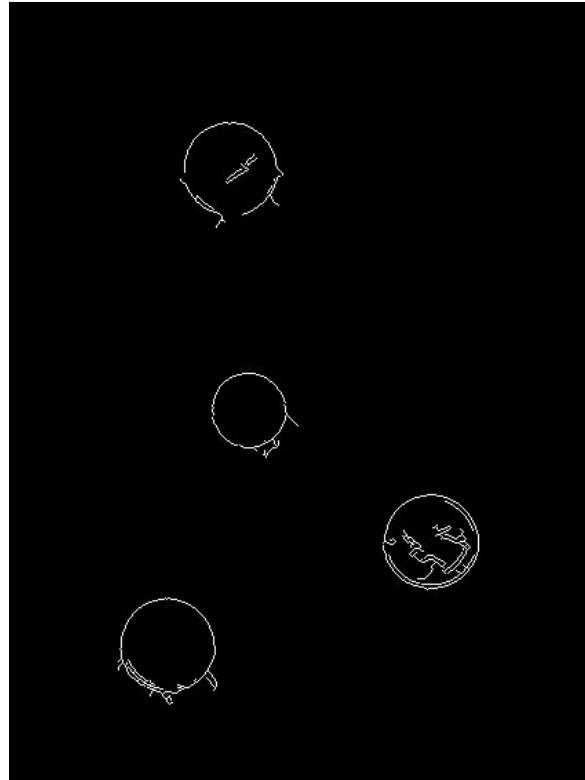
Use two Hough transforms, one for each ray (of known length)

Circle detection using Hough transform

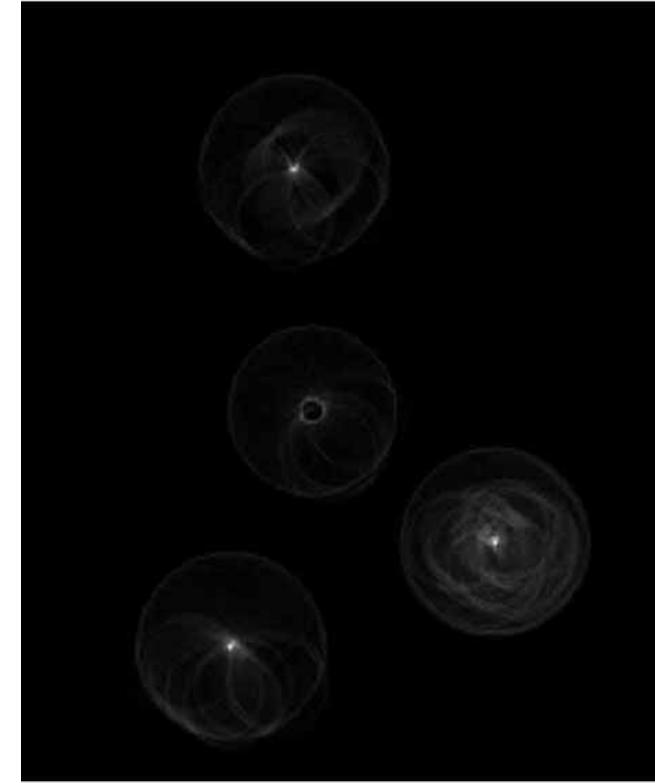
Combined detections



Edges



All votes



Hough transform: Pros and cons

- Pros
 - Can deal with non-locality and occlusion
 - Can detect multiple instances of a model
 - Some robustness to noise: noise points unlikely to contribute consistently to any single bin
- Cons
 - Complexity of search time increases exponentially with the number of model parameters
 - Non-target shapes can produce spurious peaks in parameter space
 - It's hard to pick a good grid size