

Computer Vision

Bogdan Alexe

bogdan.alexe@fmi.unibuc.ro

University of Bucharest, 2nd semester, 2023-2024

Course structure

1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

4. Object Recognition: high – level vision

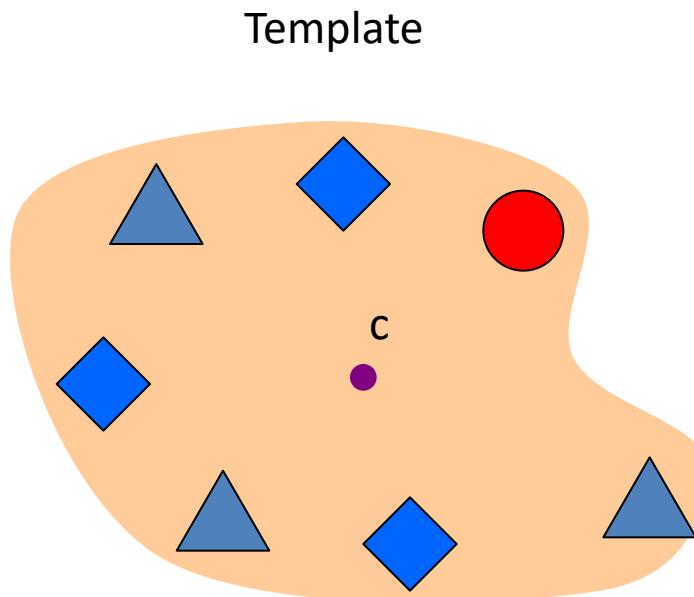
Object classification, object detection, part based models, bovw models

5. Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

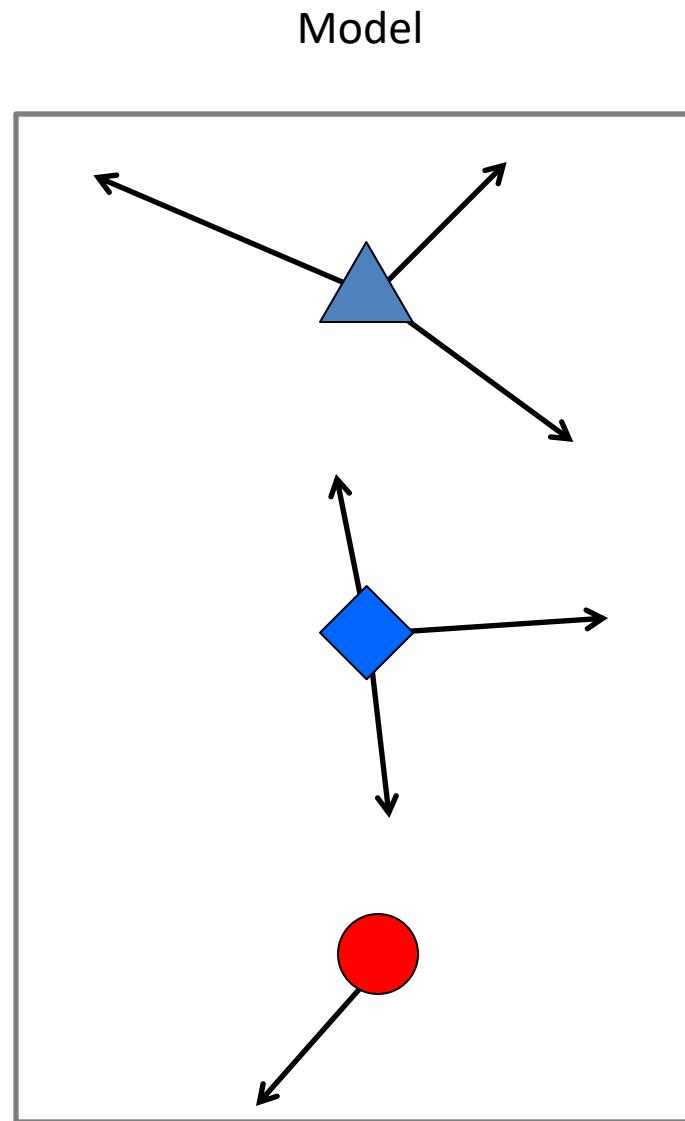
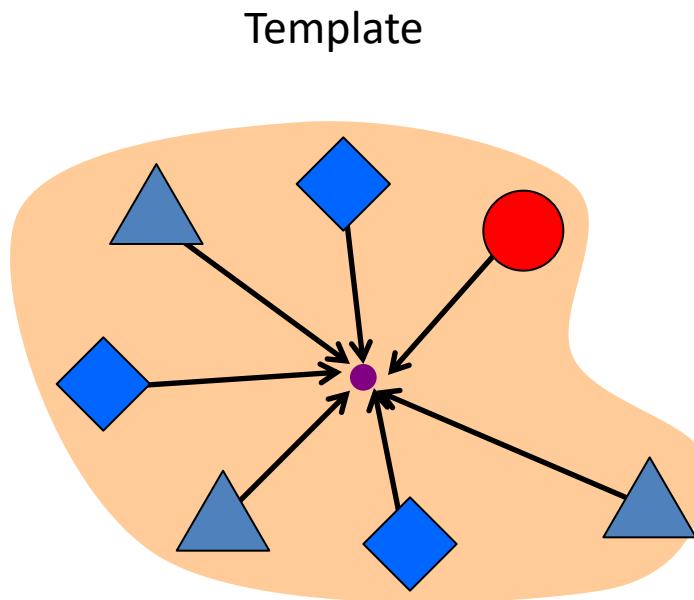
Generalized Hough transform

- We want to find a template defined by its reference point (center) and several distinct types of landmark points in stable spatial configuration



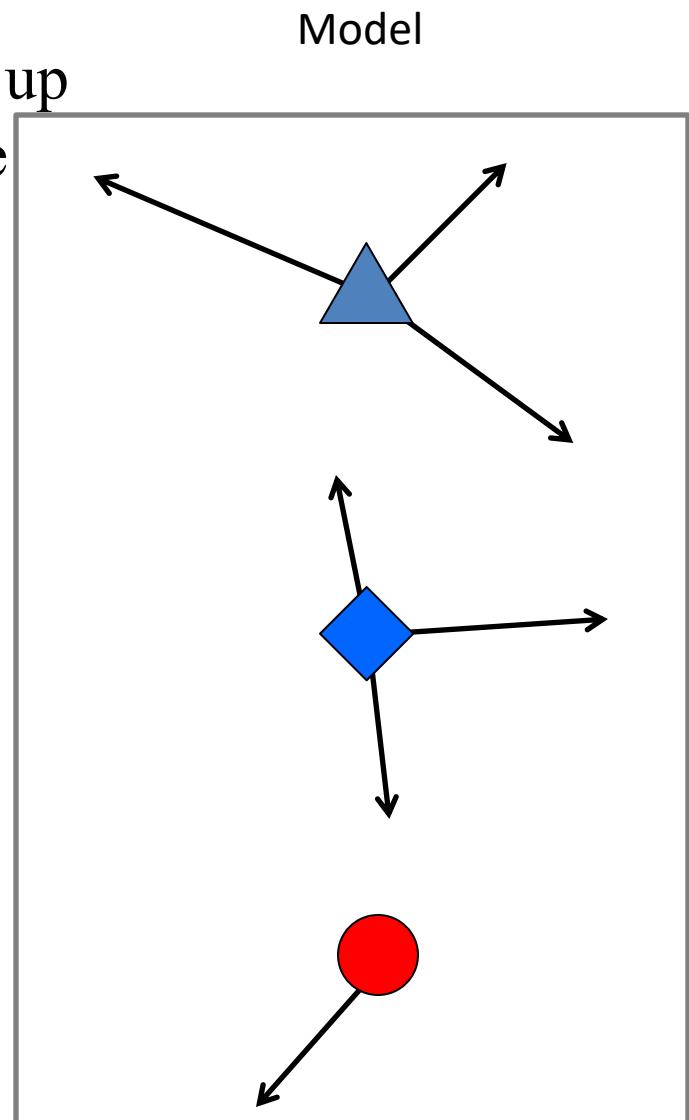
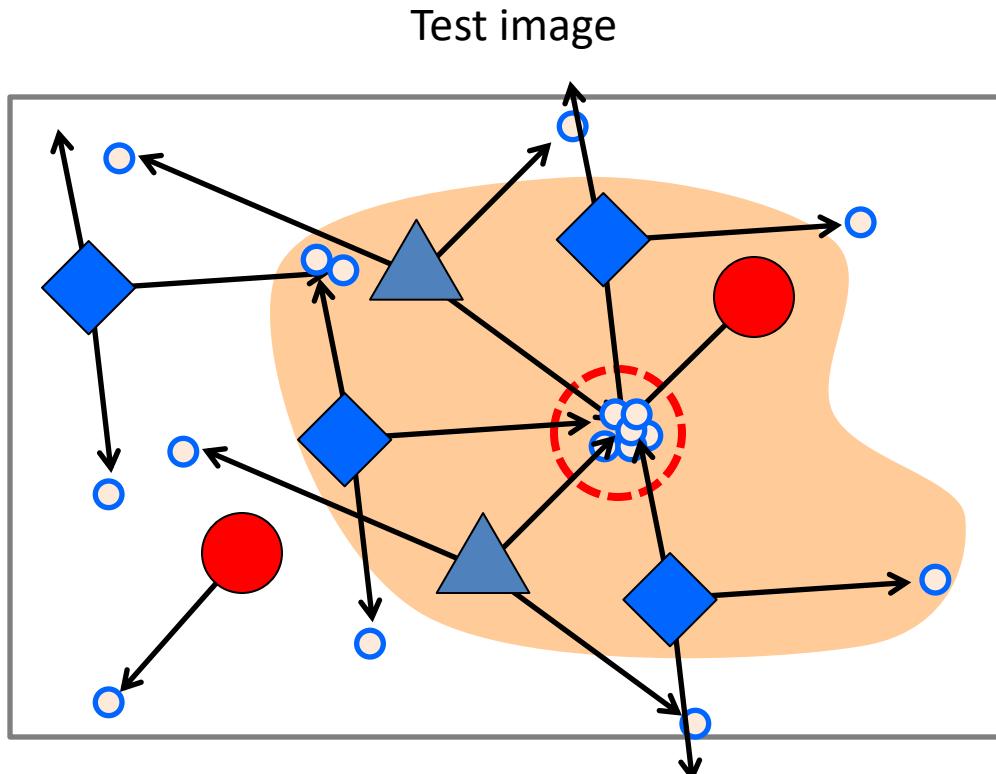
Generalized Hough transform

- Template representation: for each type of landmark point, store all possible displacement towards the center



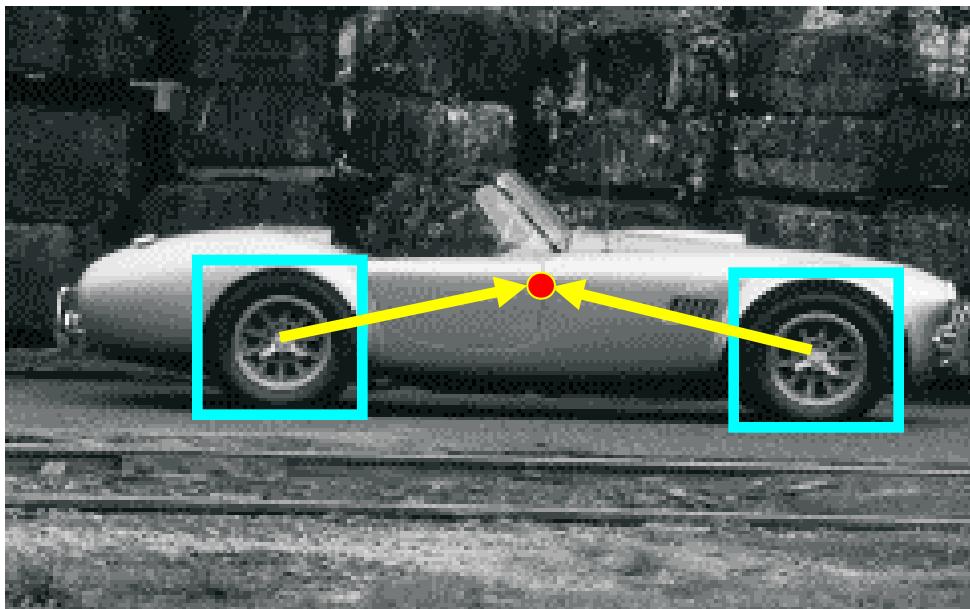
Generalized Hough transform

- Detecting the template:
 - For each feature in a new image, look up that feature type in the model and vote for the possible center locations associated with that type in the model

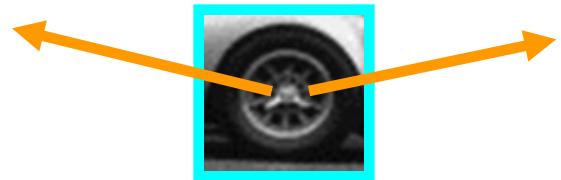


Application in recognition

- Index displacements by “visual codeword”



training image



visual codeword with
displacement vectors

Application in recognition

- Index displacements by “visual codeword”

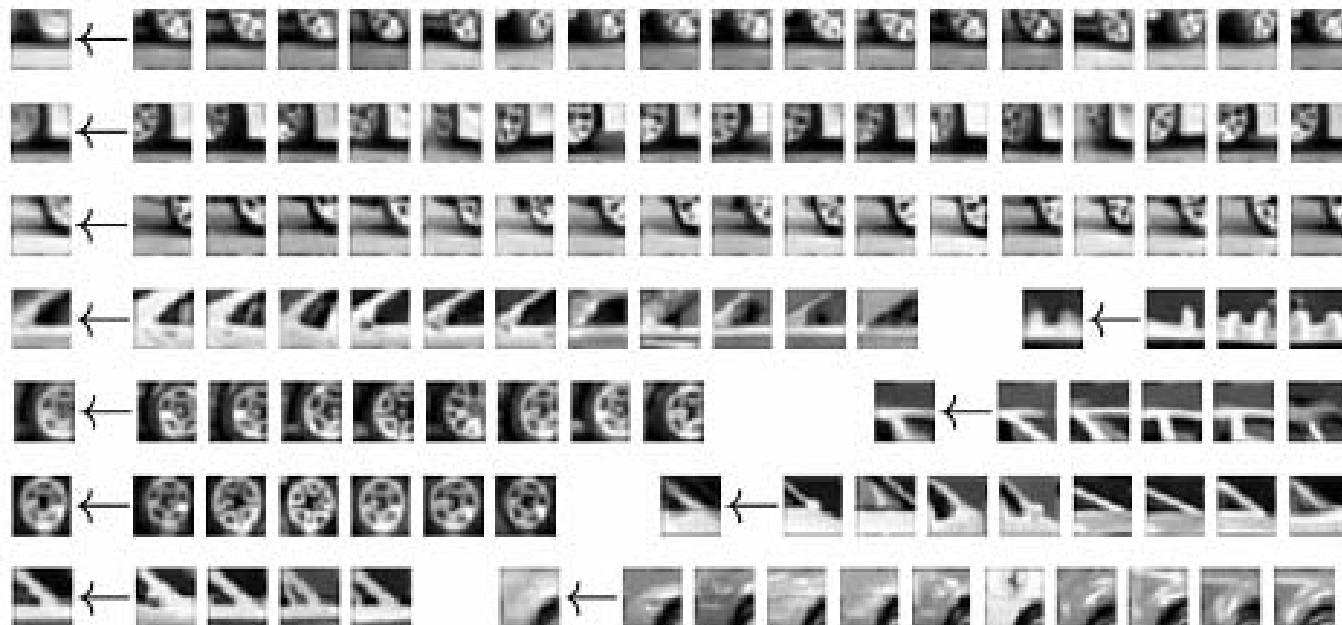


test image

B. Leibe, A. Leonardis, and B. Schiele, [Combined Object Categorization and Segmentation with an Implicit Shape Model](#), ECCV Workshop on Statistical Learning in Computer Vision 2004

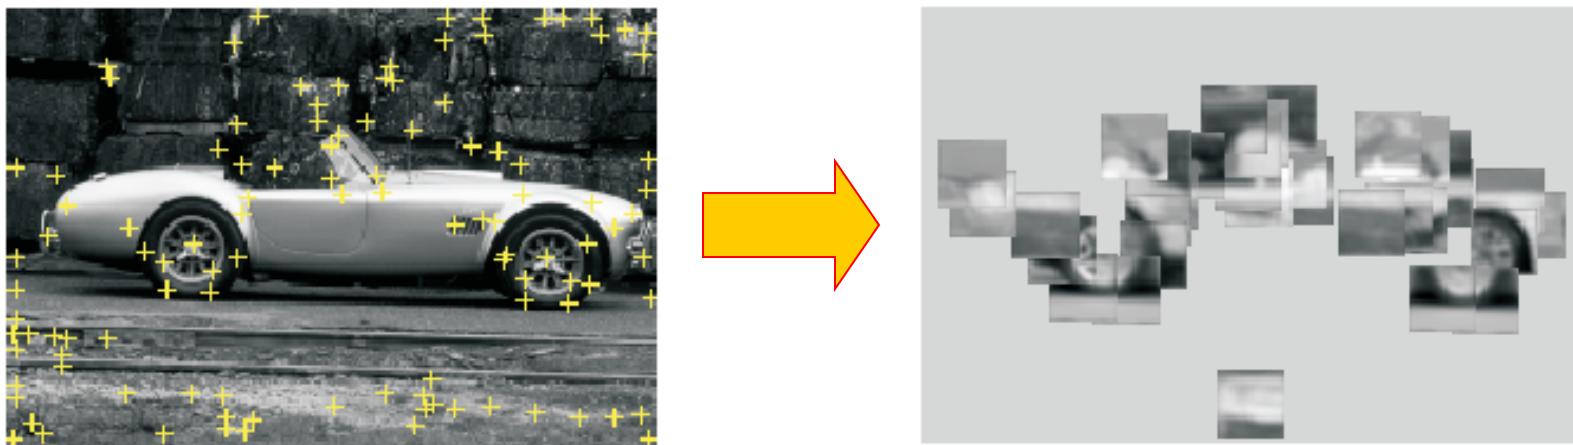
Implicit shape models: Training

1. Build *codebook* of patches around extracted interest points using clustering



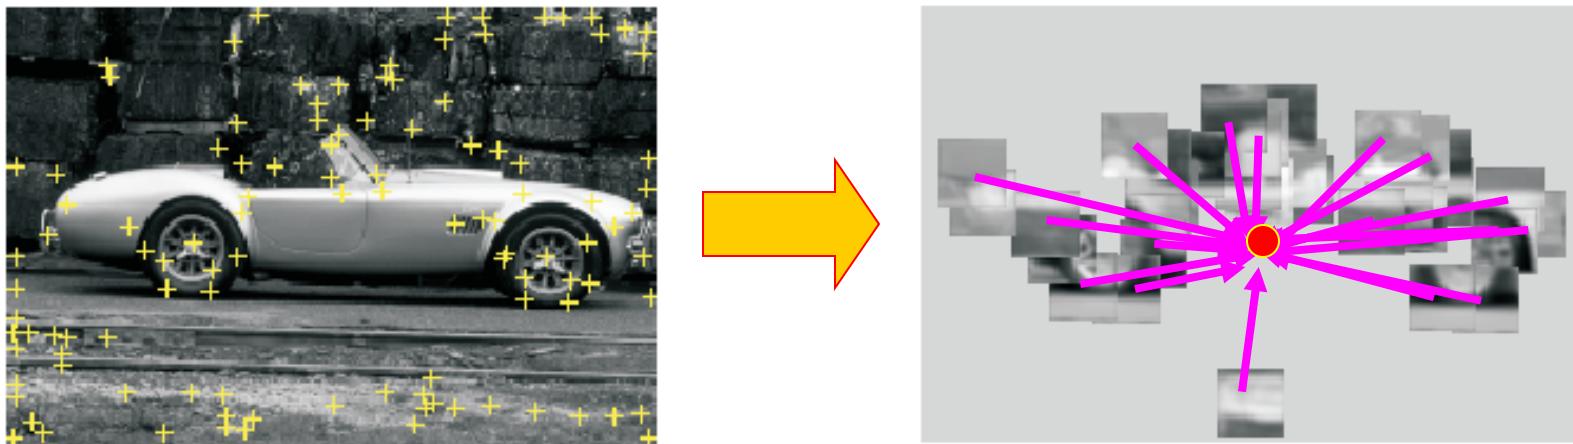
Implicit shape models: Training

1. Build *codebook* of patches around extracted interest points using clustering
2. Map the patch around each interest point to closest codebook entry



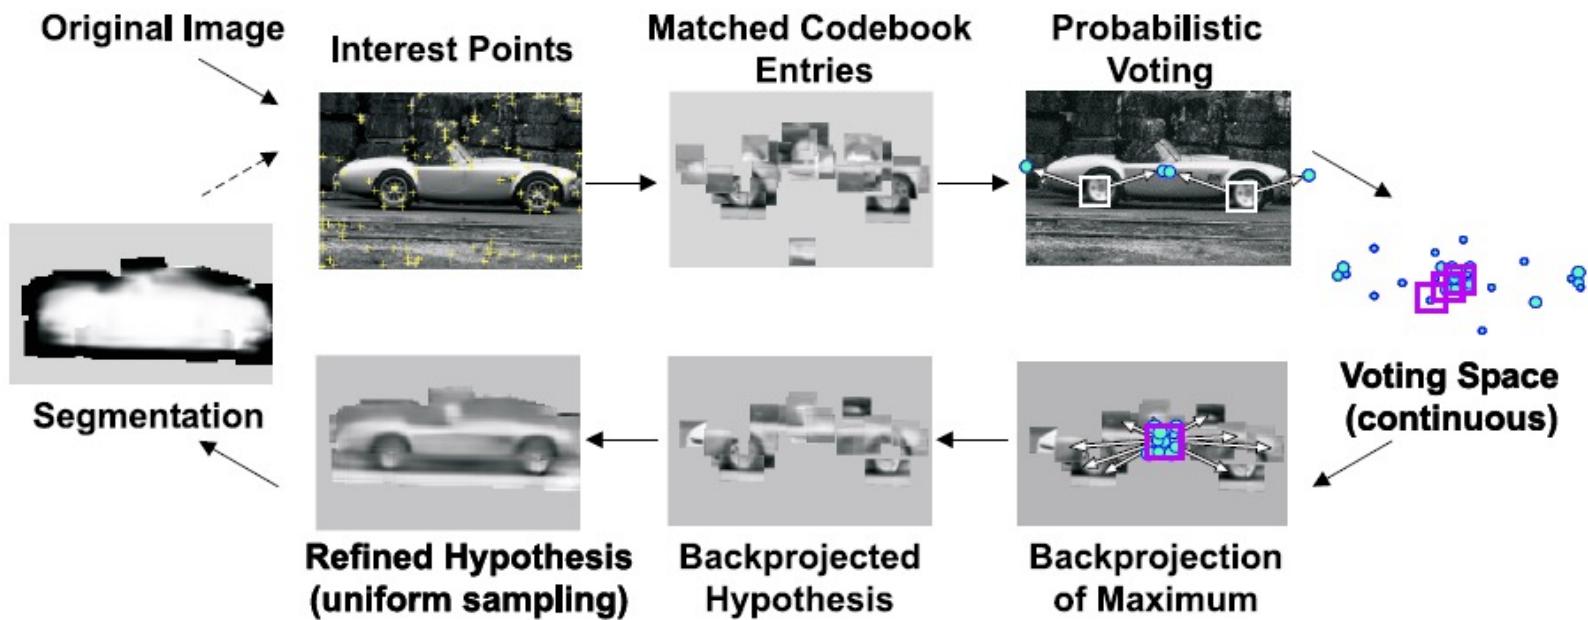
Implicit shape models: Training

1. Build *codebook* of patches around extracted interest points using clustering
2. Map the patch around each interest point to closest codebook entry
3. For each codebook entry, store all positions it was found, relative to object center

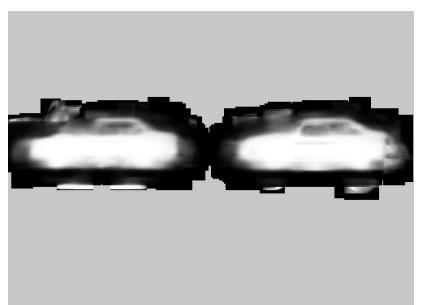
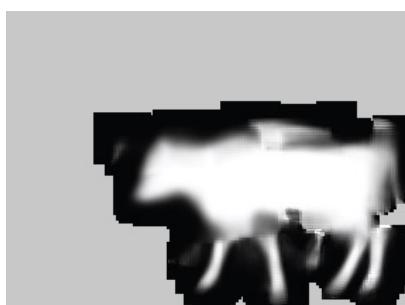
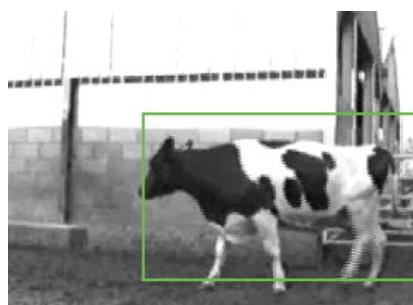
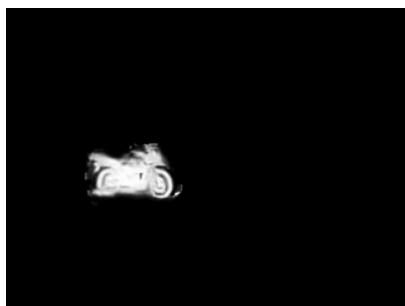
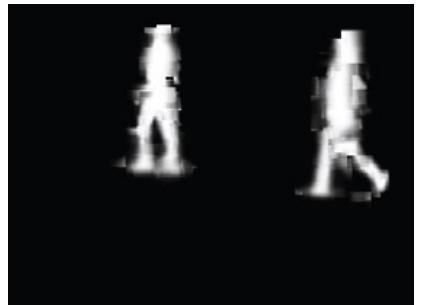
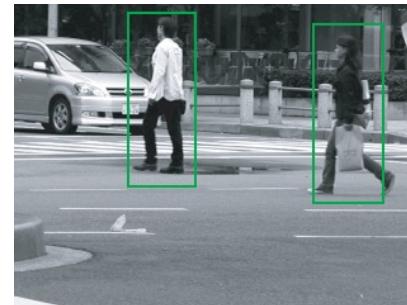
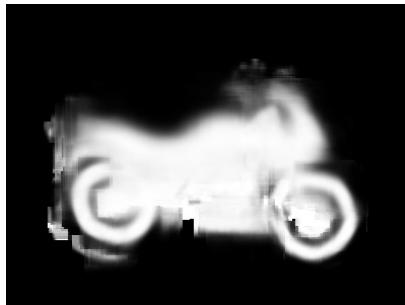


Implicit shape models: Testing

1. Given test image, extract patches, match to codebook entry
2. Cast votes for possible positions of object center
3. Search for maxima in voting space
4. Extract weighted segmentation mask based on stored masks for the codebook occurrences

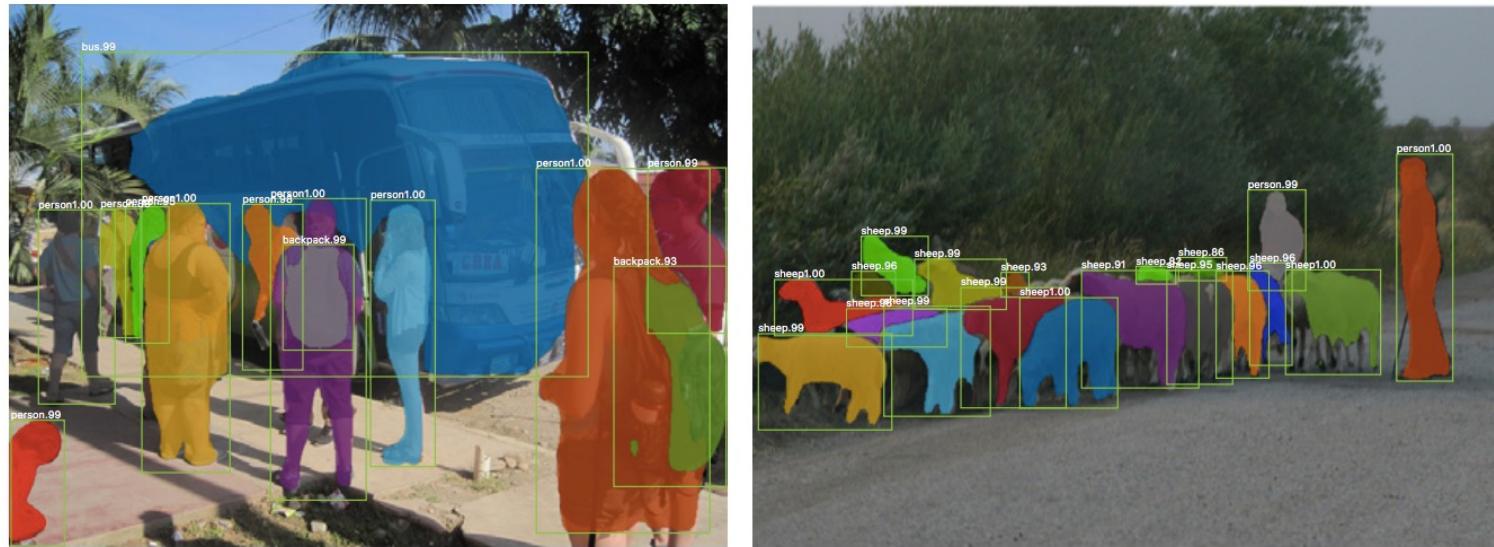
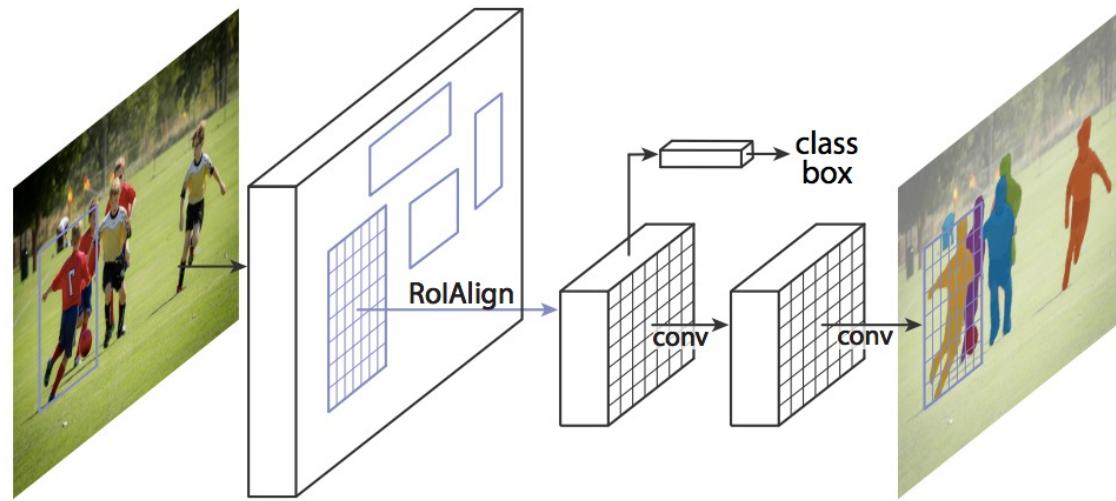


Additional examples



B. Leibe, A. Leonardis, and B. Schiele, [Robust Object Detection with Interleaved Categorization and Segmentation](#),
IJCV 77 (1-3), pp. 259-289, 2008.

These days: Mask R-CNN

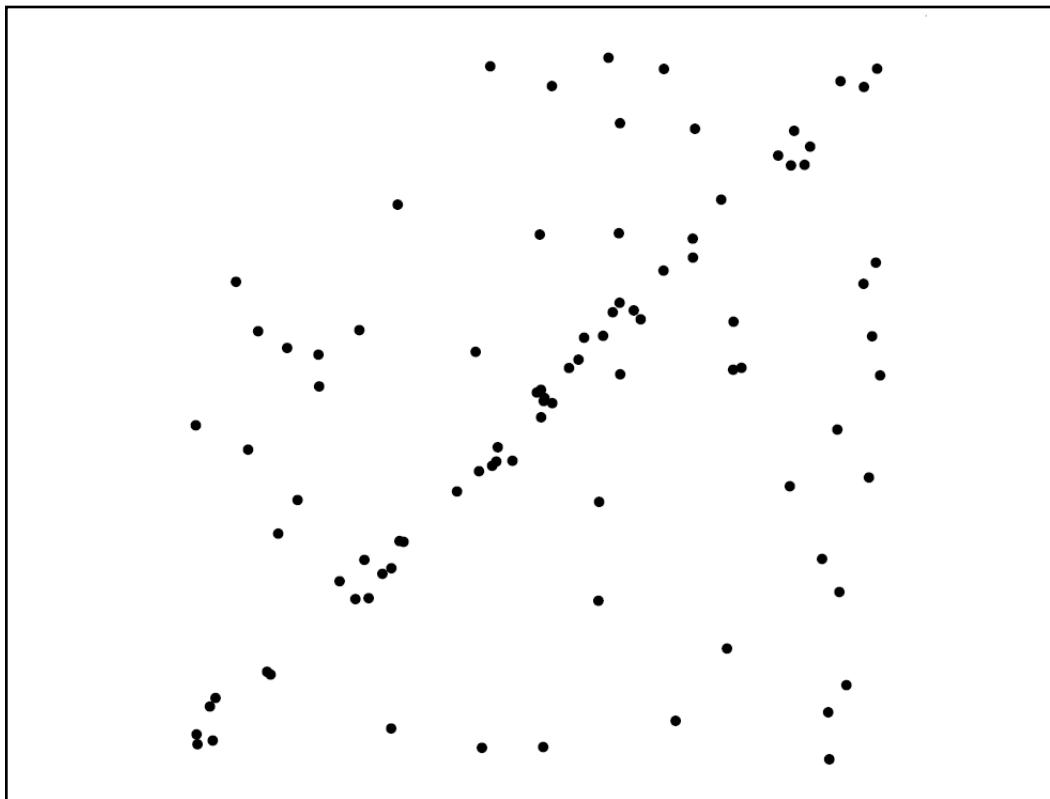


RANSAC

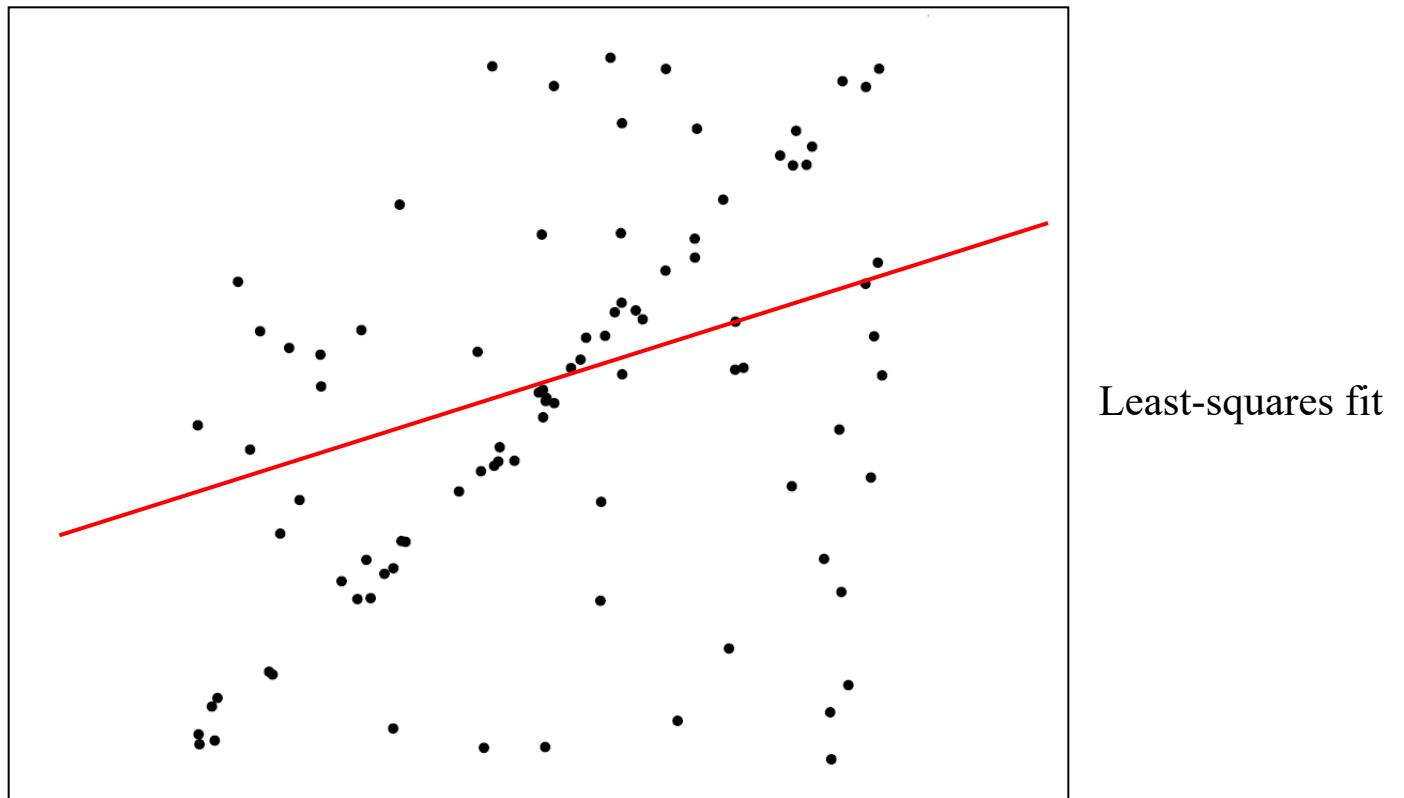
RANSAC

- Random sample consensus (RANSAC):
Very general framework for model fitting in
the presence of outliers
- Outline
 - Choose a small subset of points uniformly at random
 - Fit a model to that subset
 - Find all remaining points that are “close” to the model and
reject the rest as outliers
 - Do this many times and choose the best model

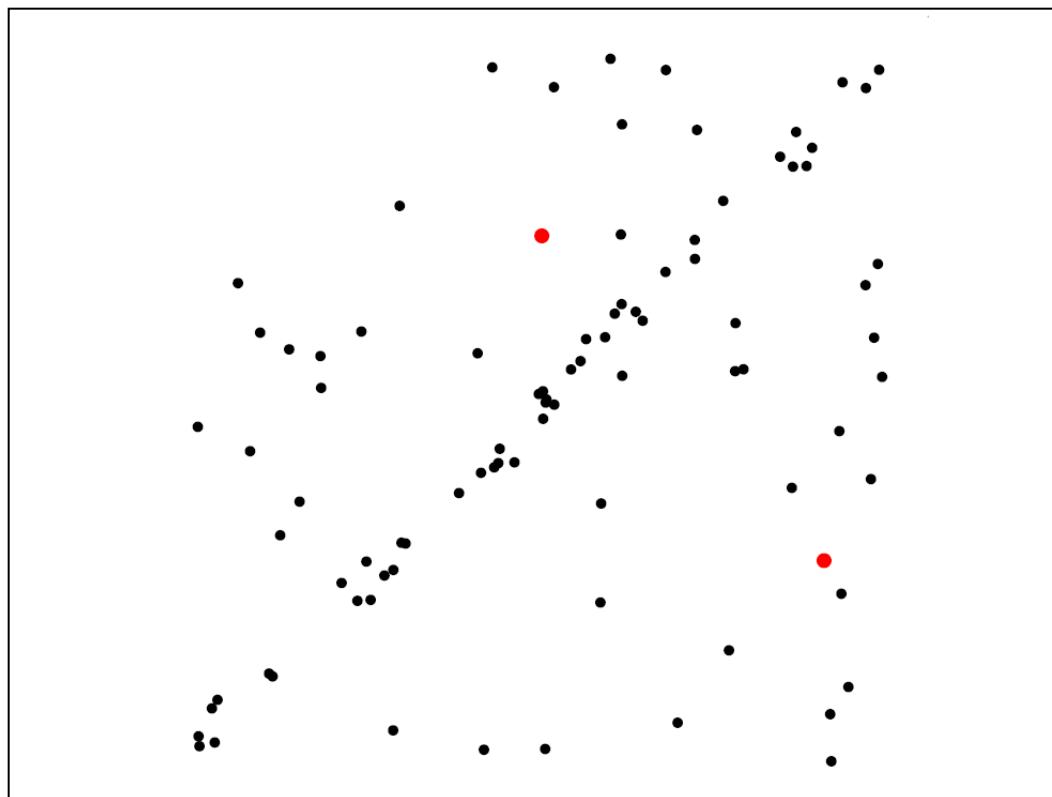
RANSAC for line fitting example



RANSAC for line fitting example

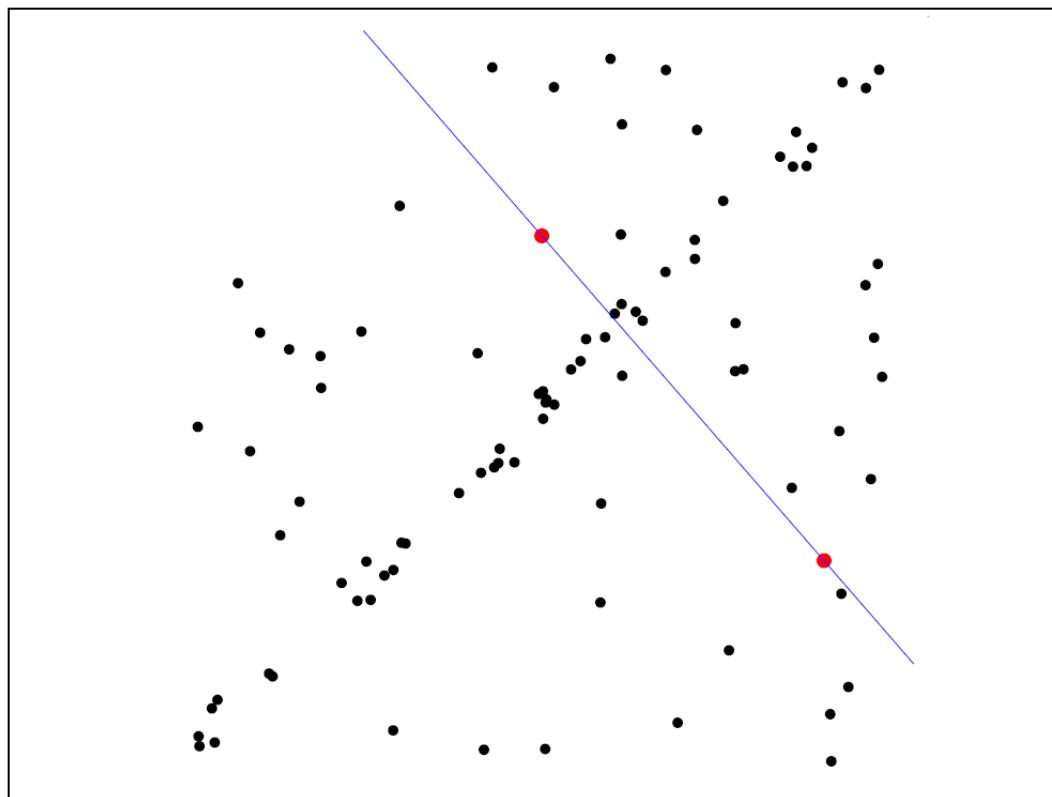


RANSAC for line fitting example



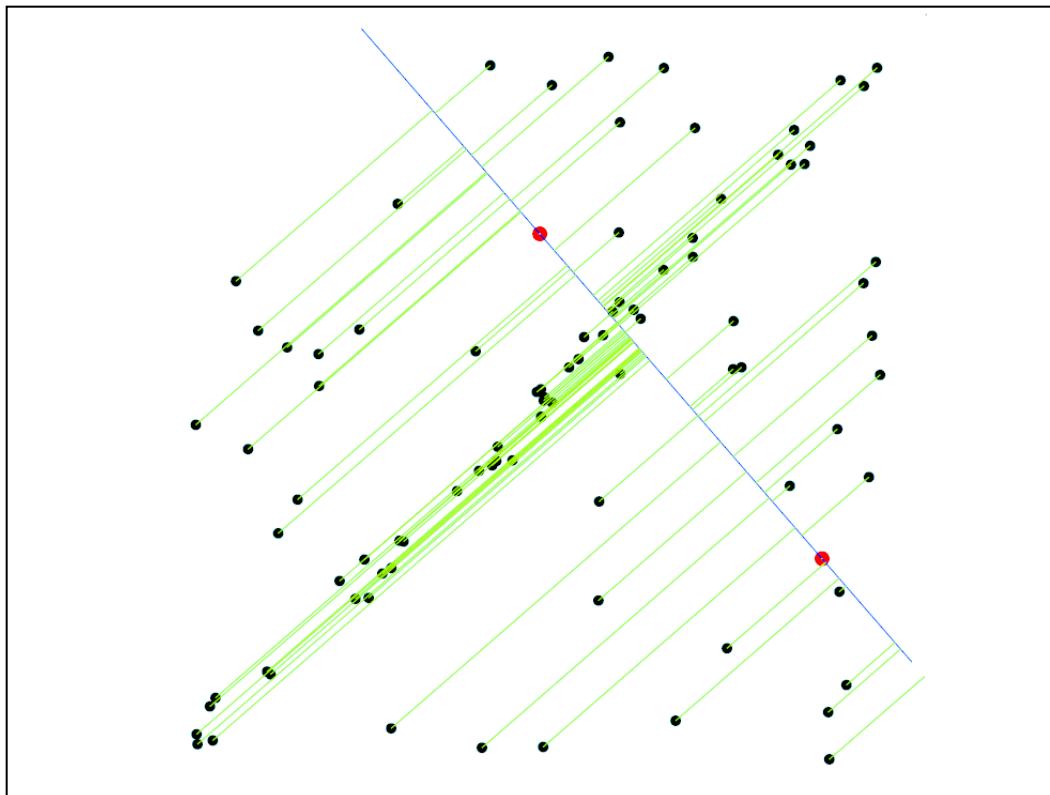
1. Randomly select minimal subset of points

RANSAC for line fitting example



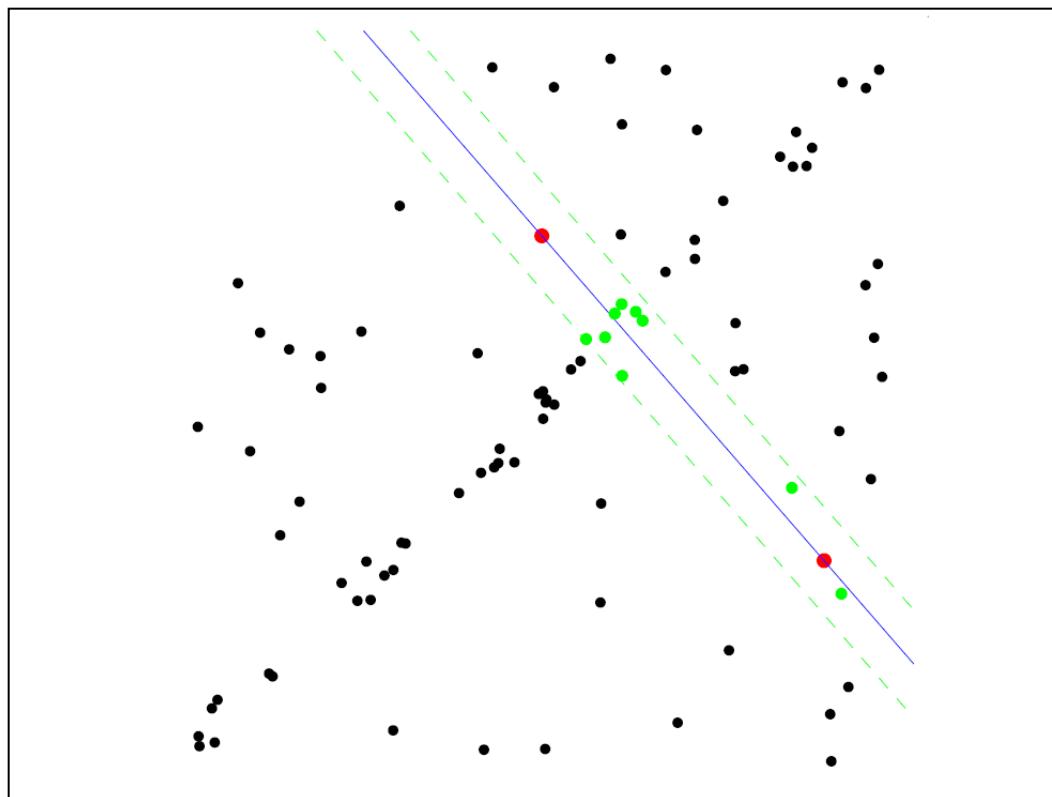
1. Randomly select minimal subset of points
2. Hypothesize a model

RANSAC for line fitting example



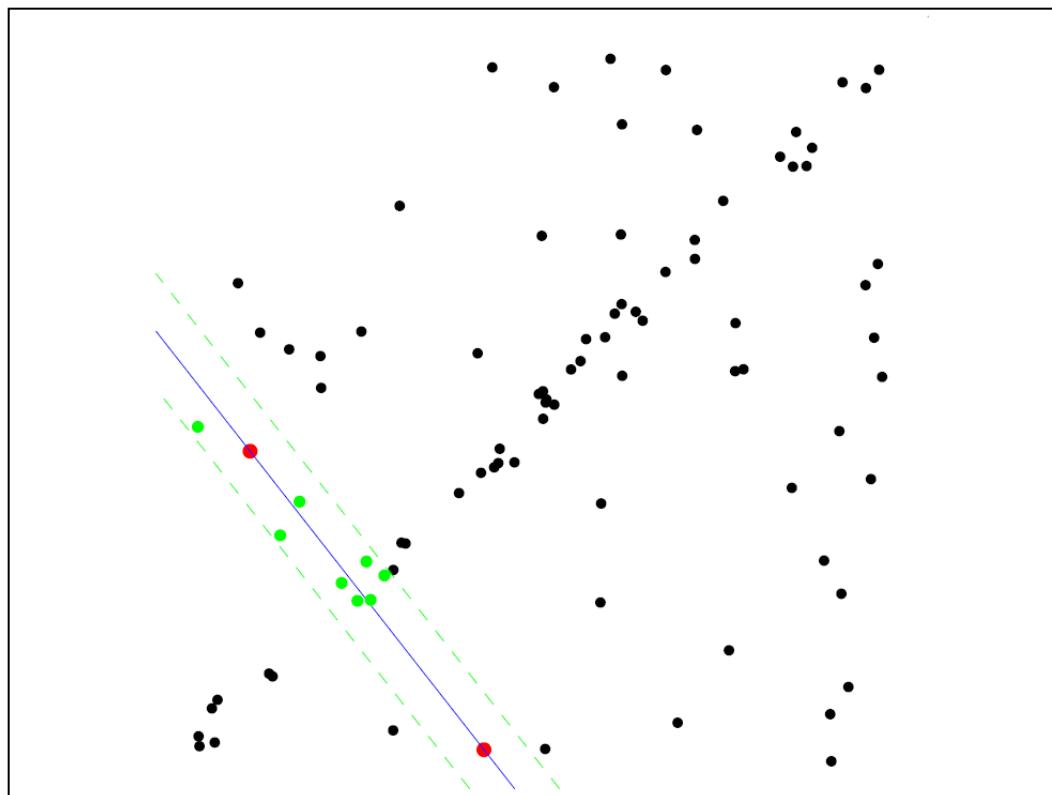
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function

RANSAC for line fitting example



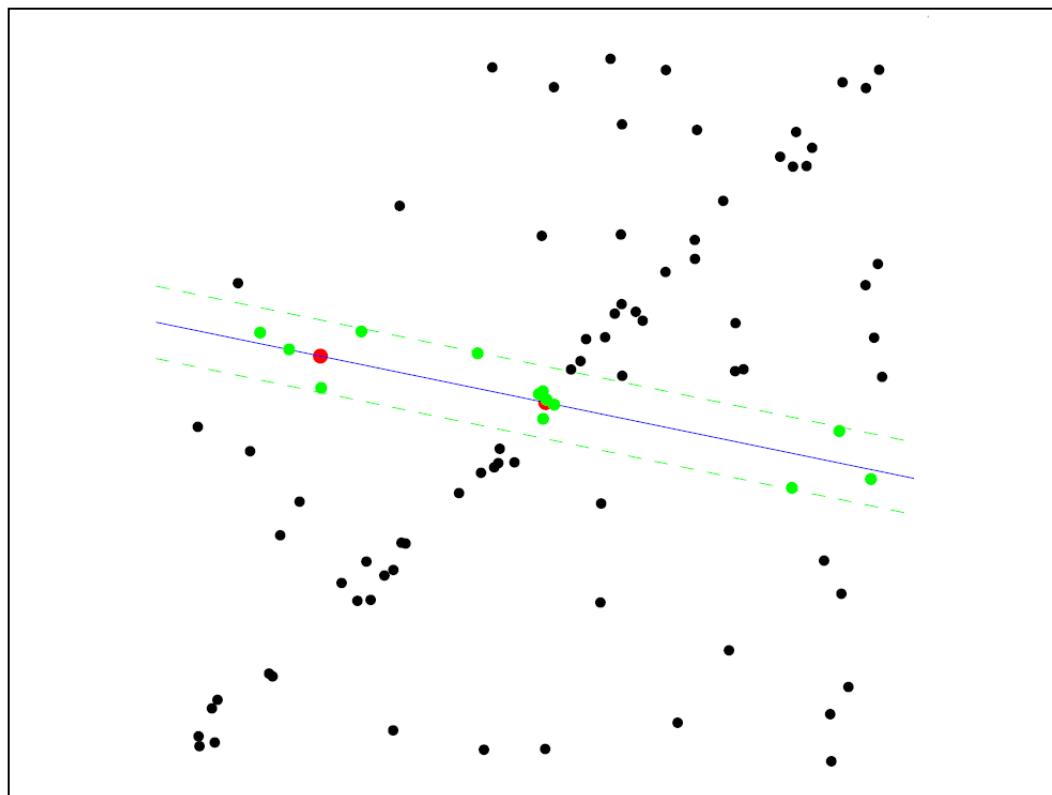
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model

RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

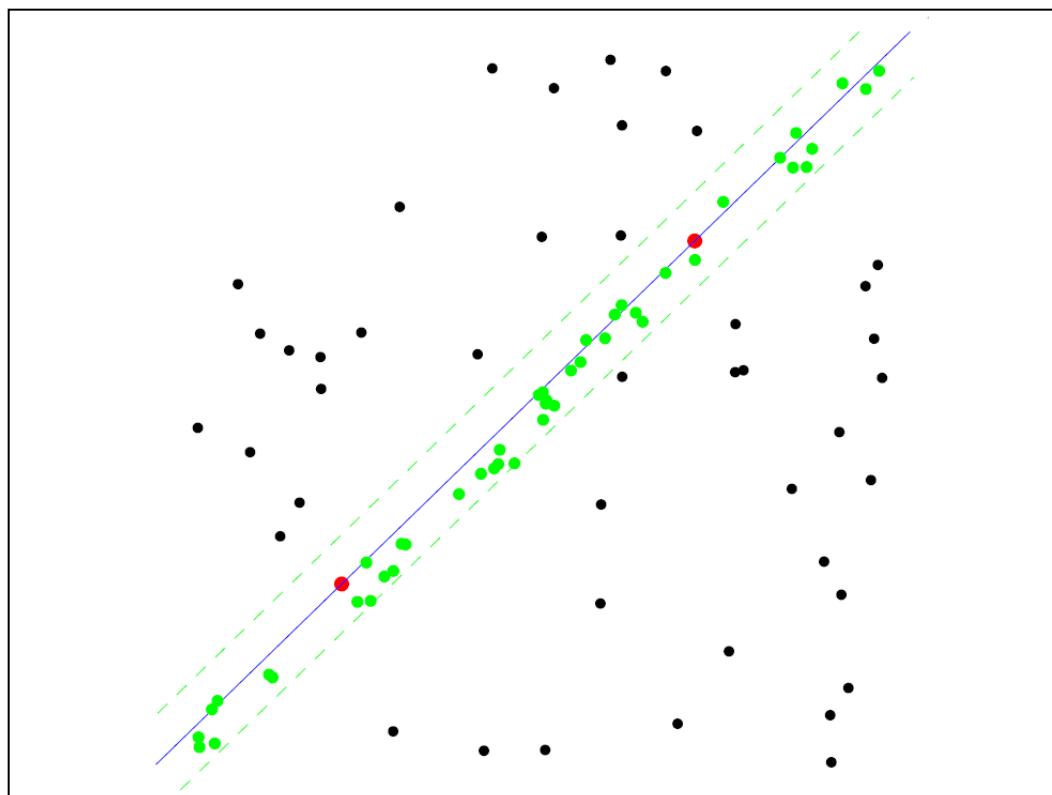
RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

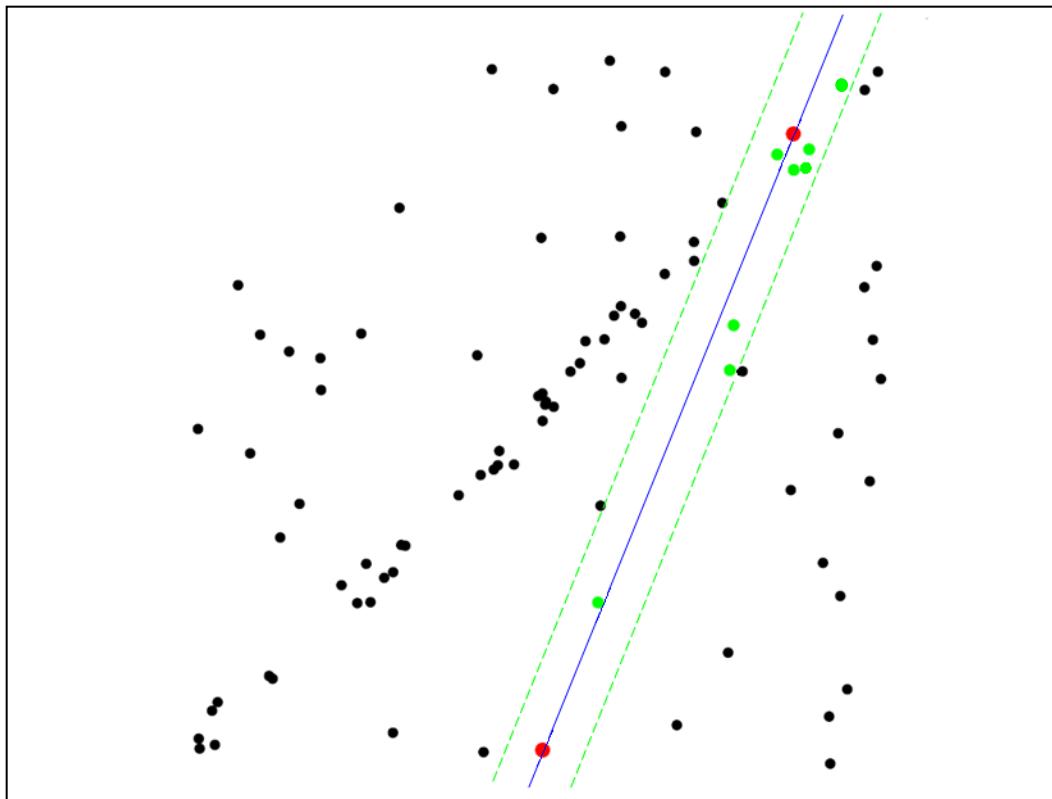
RANSAC for line fitting example

Uncontaminated sample



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

RANSAC for line fitting

- Repeat N times:
 - Draw s points uniformly at random
 - Fit a line (model) to these s points
 - Find *inliers* to this line among the remaining points (i.e., points whose distance from the line is less than t)
 - If there are d or more inliers, accept the line and refit using all inliers

Choosing the parameters

- Initial number of points s
 - Typically minimum number needed to fit the model
- Distance threshold t
 - Choose t so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2=3.84\sigma^2$, $t=1.9\sigma$
- Number of samples N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)

Choosing the parameters

- Initial number of points s
 - Typically minimum number needed to fit the model
- Distance threshold t
 - Choose t so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2=3.84\sigma^2$
- Number of samples N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$

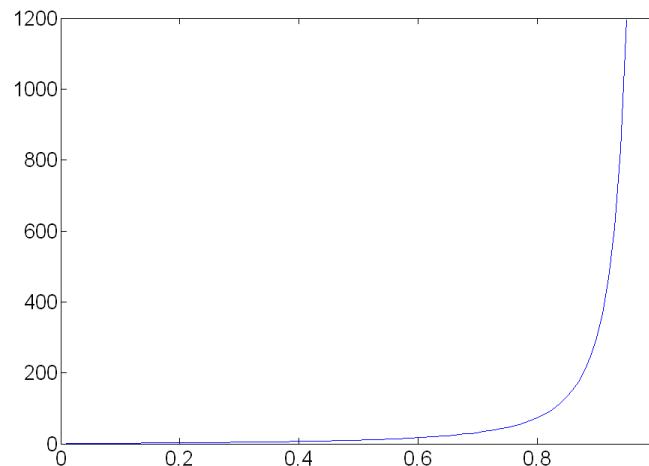
s	proportion of outliers e							
	5%	10%	20%	25%	30%	40%	50%	
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

Choosing the parameters

- Initial number of points s
 - Typically minimum number needed to fit the model
- Distance threshold t
 - Choose t so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2=3.84\sigma^2$
- Number of samples N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$



Choosing the parameters

- Initial number of points s
 - Typically minimum number needed to fit the model
- Distance threshold t
 - Choose t so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2=3.84\sigma^2$
- Number of samples N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)
- Consensus set size d
 - Should match expected inlier ratio

Adaptively determining the number of samples

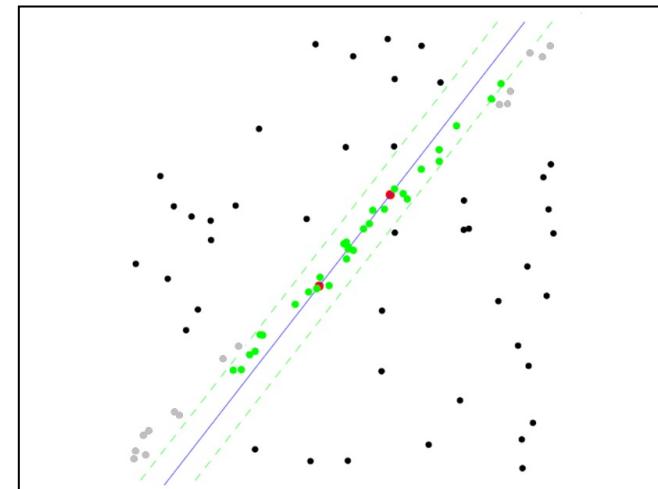
- Outlier ratio e is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield $e=0.2$
- Adaptive procedure:
 - $N = \infty$, $sample_count = 0$
 - While $N > sample_count$
 - Choose a sample and count the number of inliers
 - If inlier ratio is highest of any found so far, set $e = 1 - (\text{number of inliers})/(\text{total number of points})$
 - Recompute N from e :

$$N = \log(1 - p) / \log\left(1 - (1 - e)^s\right)$$

- Increment the $sample_count$ by 1

RANSAC pros and cons

- Pros
 - Simple and general
 - Applicable to many different problems
 - Often works well in practice
- Cons
 - Lots of parameters to tune
 - Doesn't work well for low inlier ratios (too many iterations, or can fail completely)
 - Can't always get a good initialization of the model based on the minimum number of samples



Course structure

1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

4. Object Recognition: high – level vision

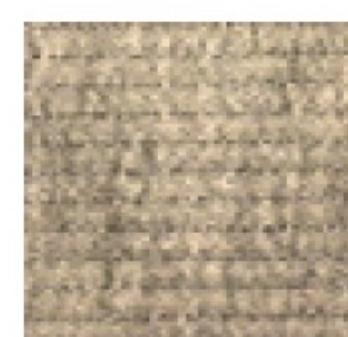
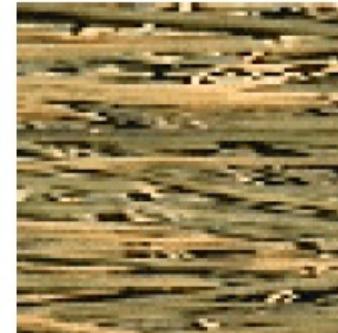
Object classification, object detection, part based models, bovw models

5. Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

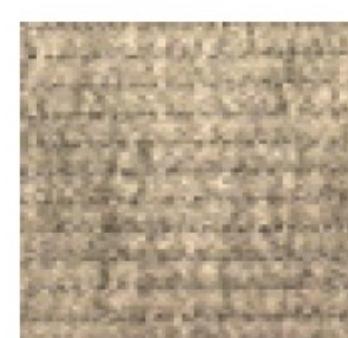
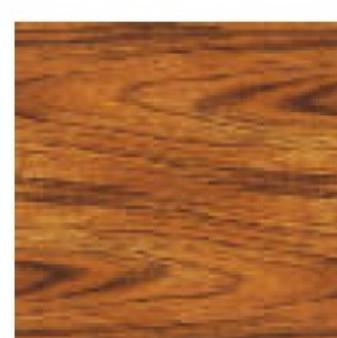
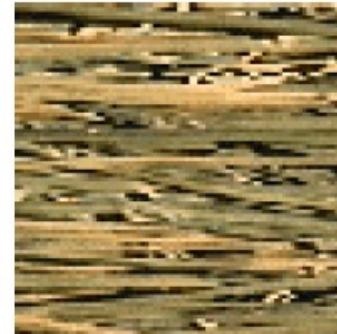
Texture

Texture



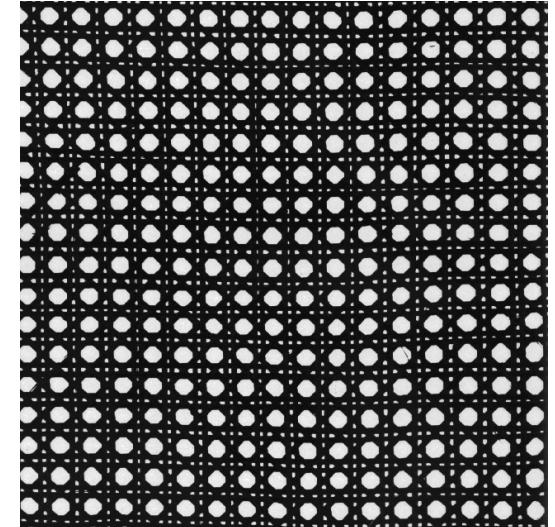
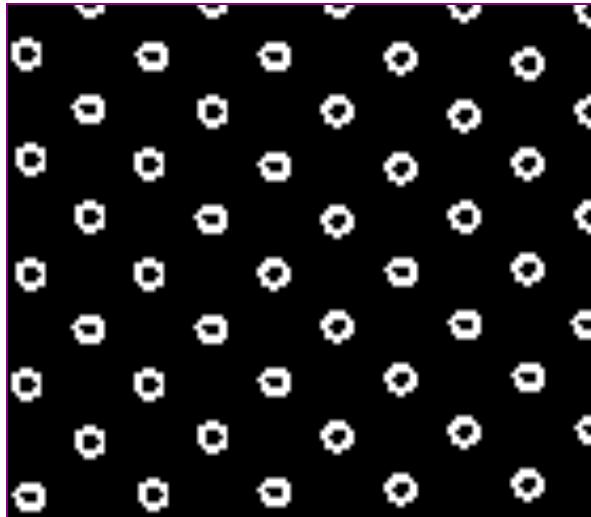
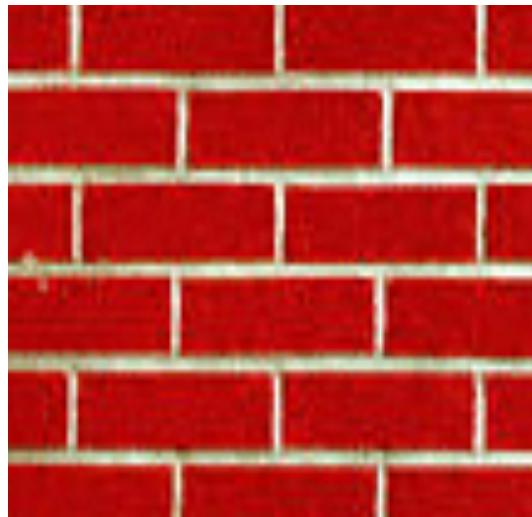
What defines a texture?

Texture



Similar structures of pixels (patterns) that repeat themselves
Similar aspect at a fixed scale
Follow some statistics properties

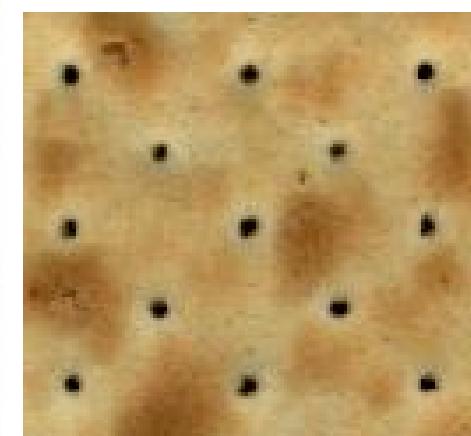
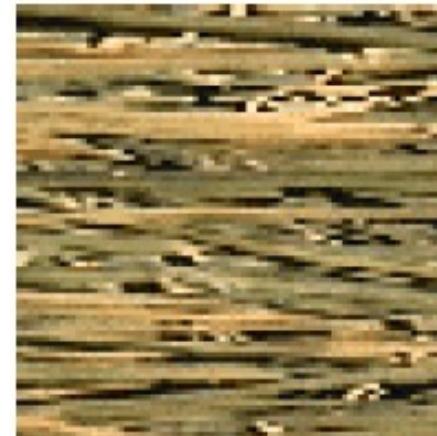
Includes: more regular patterns



Texture = set of *texels* in some regular pattern

Texels = *texture elements*

Includes: more random patterns



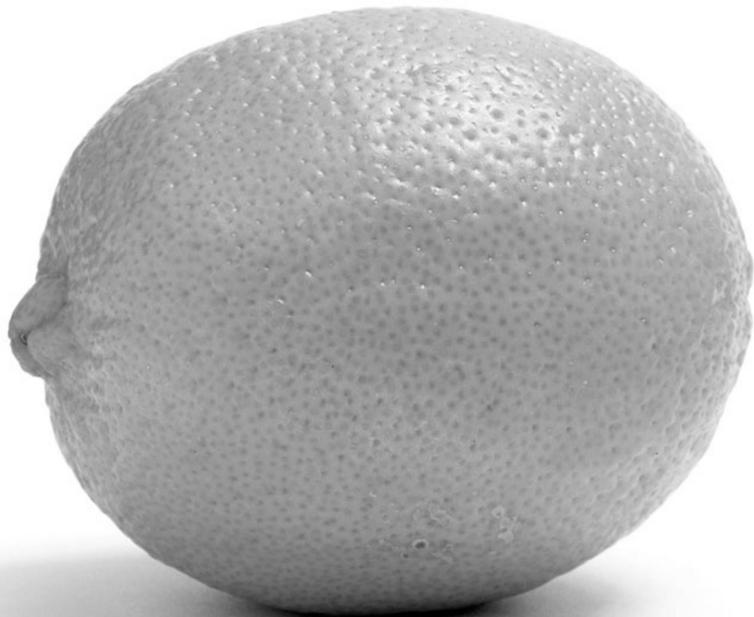
Texture = set of *texels* in some regular or repeated pattern

Texels = *texture elements*

Why analyze texture?

Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture



Slide credit: Kristen Grauman



Why analyze texture?

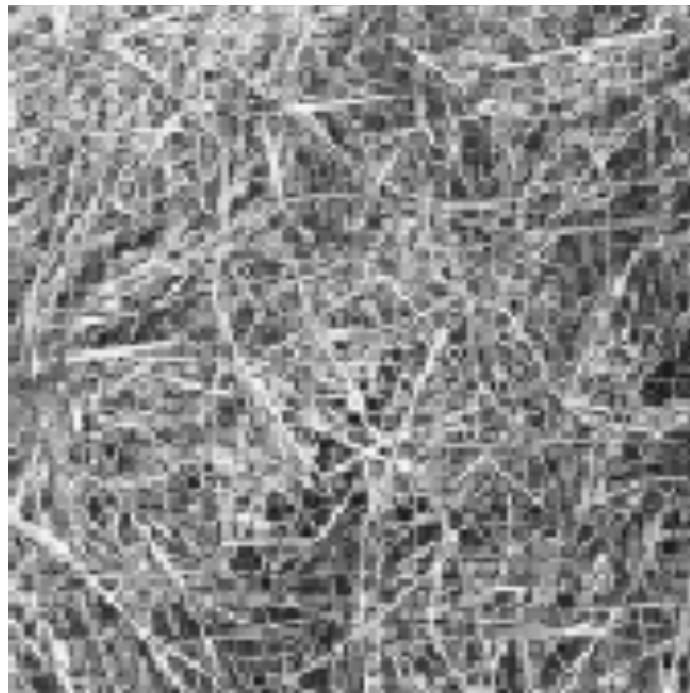
Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture

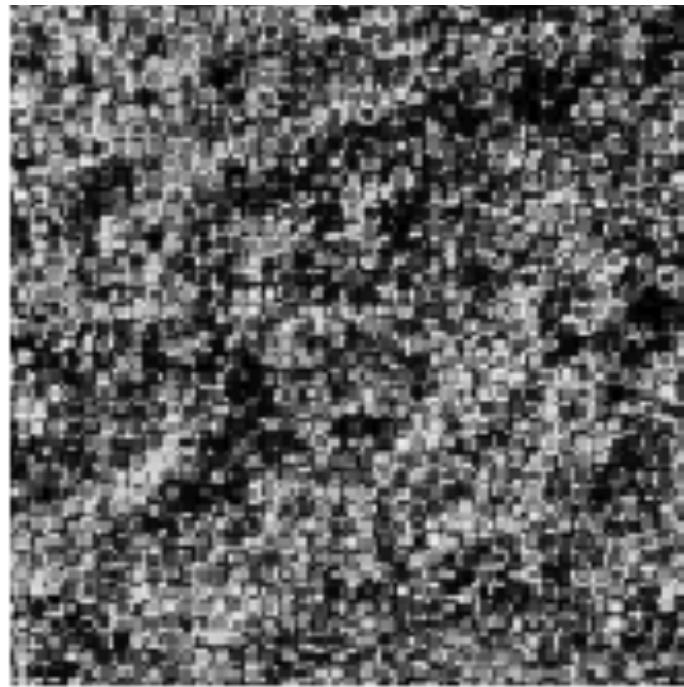
Technically:

- Representation-wise, we want a feature one step above “building blocks” of filters, edges.

Natural Textures



grass



leaves

What/Where are the texels?

The Case for Statistical Texture

- Segmenting out texels is difficult or impossible in real images.
- Numeric quantities or statistics that describe a texture can be computed from the grayscale (or colors) images alone.
- This approach is less intuitive, but is computationally efficient.
- It can be used for both classification and segmentation.

Some Simple Statistical Texture Measures

Edge Density and Direction

- Use an edge detector as the first step in texture analysis.
- The number of edge pixels in a fixed-size region tells us how busy that region is.
- The directions of the edges also help characterize the texture

Edge-based Texture Measures

1. edgeness per unit area

$$\text{Edgeness} = |\{ p \mid \text{gradient_magnitude}(p) \geq \text{threshold}\}| / N$$

where N is the size of the unit area

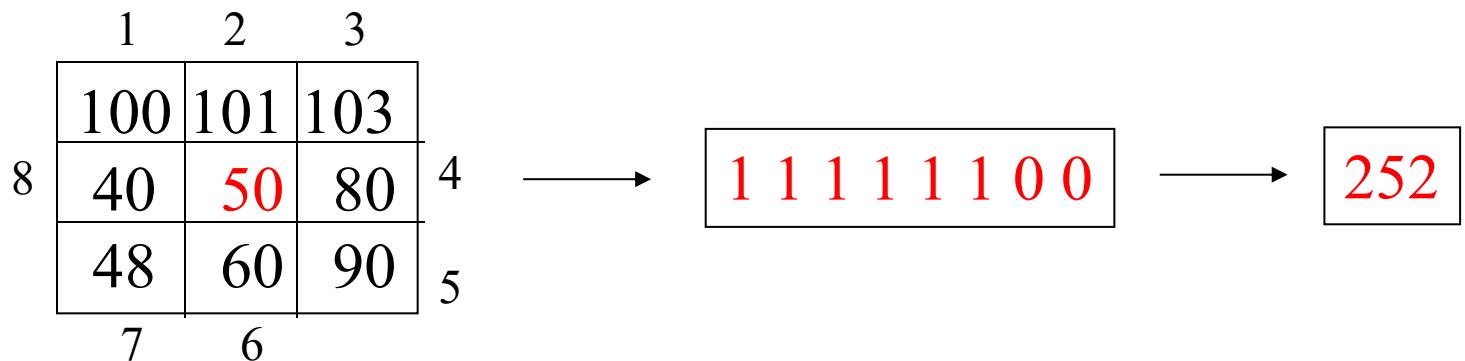
2. edge magnitude and direction histograms

$$F_{magdir} = (H_{magnitude}, H_{direction})$$

where these are the normalized histograms of gradient magnitudes and gradient directions, respectively.

Local Binary Pattern Measure

- For each pixel p , create an 8-bit number $b_1 b_2 b_3 b_4 b_5 b_6 b_7 b_8$, where $b_i = 0$ if neighbor i has value less than or equal to p 's value and 1 otherwise.
- Represent the texture in the image (or a region) by the histogram of these numbers.



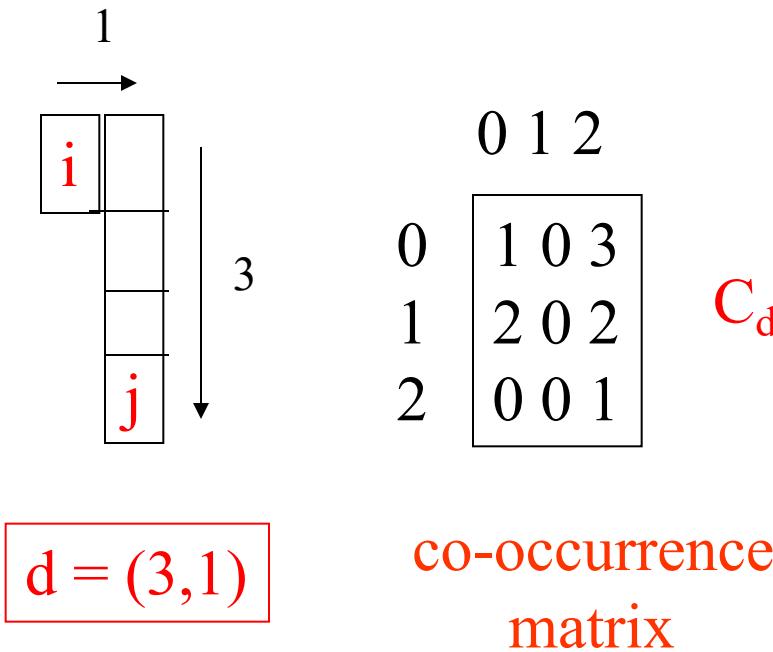
Co-occurrence Matrix Features

A co-occurrence matrix is a 2D array C in which

- Both the rows and columns represent a set of possible image values.
- $C_d(i,j)$ indicates how many times value i co-occurs with value j in a particular spatial relationship d .
- The spatial relationship is specified by a vector $d = (dr, dc)$.

Co-occurrence Example

1	1	0	0
1	1	0	0
0	0	2	2
0	0	2	2
0	0	2	2
0	0	2	2



grayscale
image

From C_d we can compute N_d , the normalized co-occurrence matrix, where each value is divided by the sum of all the values.

Co-occurrence Features

- extract numerical features from the co-occurrence matrix which can be used for representing and comparing textures

$$Energy = \sum_i \sum_j N_d^2(i, j)$$

$$Entropy = - \sum_i \sum_j N_d(i, j) \log_2 N_d(i, j)$$

$$Contrast = \sum_i \sum_j (i - j)^2 N_d(i, j)$$

$$Homogeneity = \sum_i \sum_j \frac{N_d(i, j)}{1 + |i - j|}$$

$$Correlation = \frac{\sum_i \sum_j (i - \mu_i)(j - \mu_j) N_d(i, j)}{\sigma_i \sigma_j}$$

where μ_i, μ_j are the means and σ_i, σ_j are the standard deviations of the row and column

But how do you choose d?

- This is actually a critical question with **all** the statistical texture methods.
- Are the “texels” tiny, medium, large, all three ...?
- Not really a solved problem.

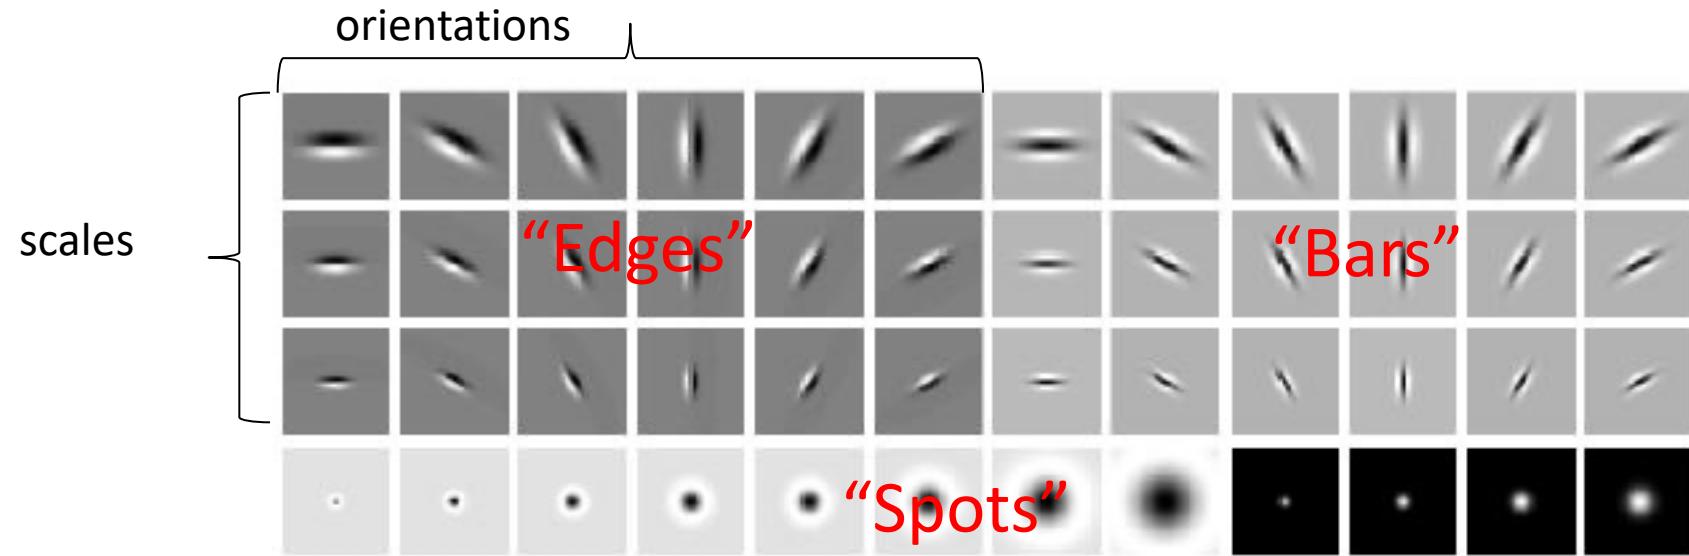
Alternative approach

Signal-processing-based algorithms use texture filters applied to the image to create filtered images from which texture features are computed.

Textures are made up of repeated local patterns, so:

- Find the patterns
 - Use filters that look like patterns (spots, bars, raw patches...)
 - Consider magnitude of response
- Describe their statistics within each local window, e.g.,
 - Mean, standard deviation
 - Histogram
 - Histogram of “prototypical” feature occurrences

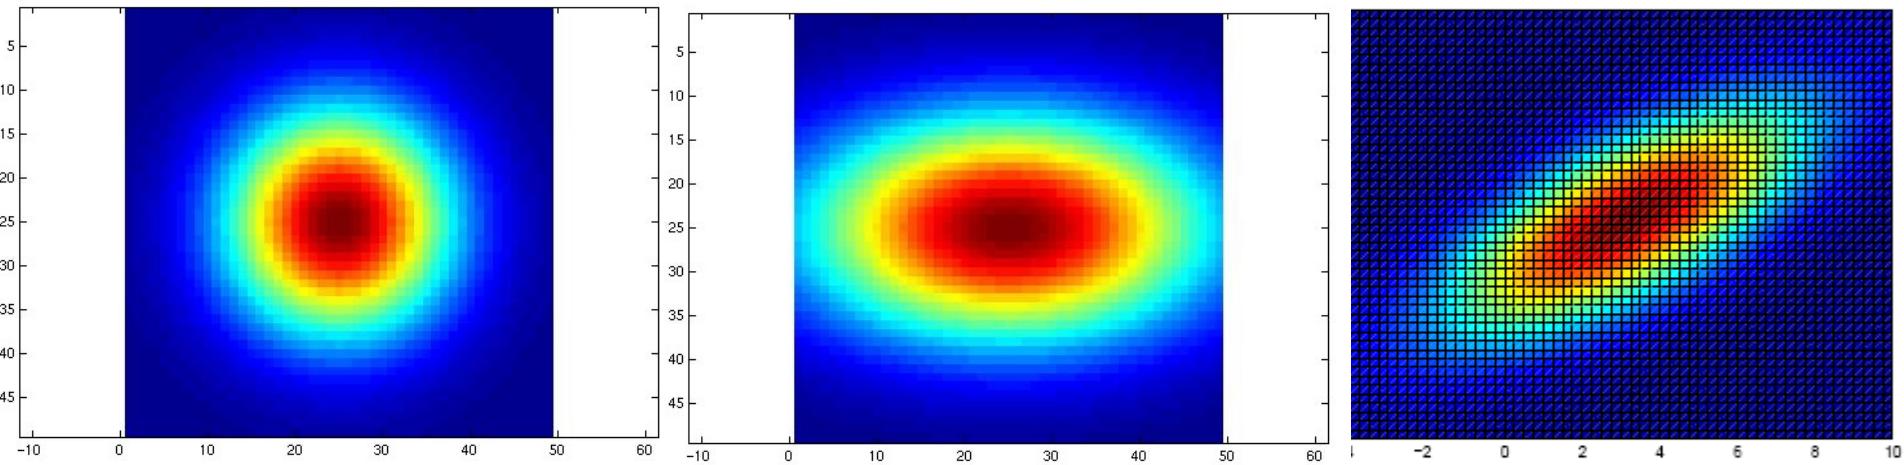
Filter banks



- What filters to put in the bank?
 - Typically we want a combination of scales and orientations, different types of patterns.

Multivariate Gaussian

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right).$$

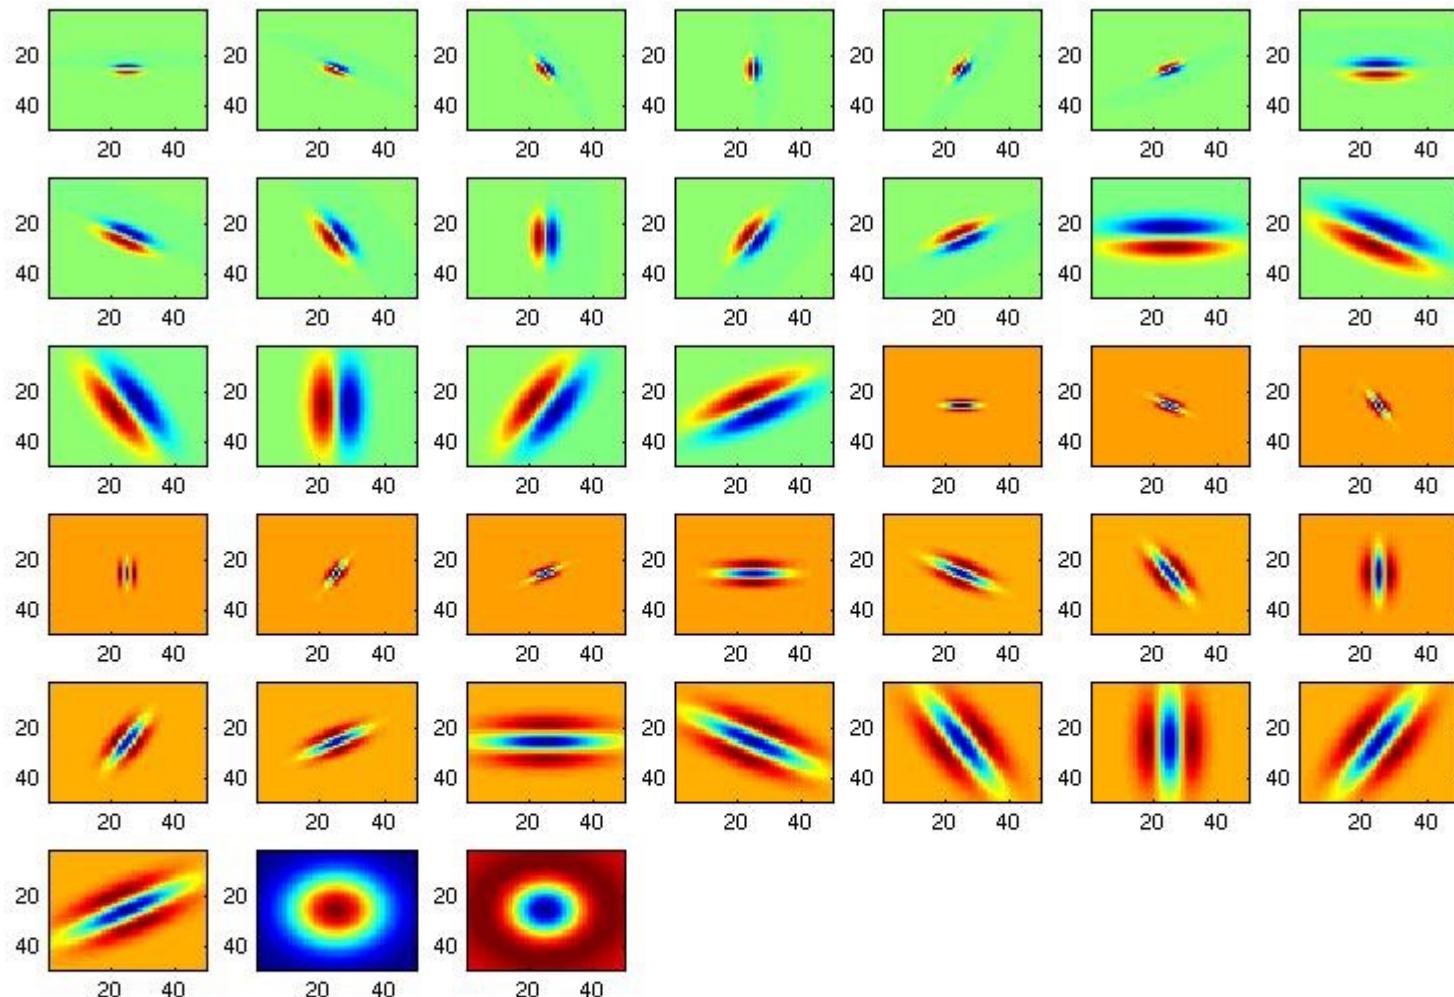


$$\Sigma = \begin{bmatrix} 9 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 16 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 10 & 5 \\ 5 & 5 \end{bmatrix}$$

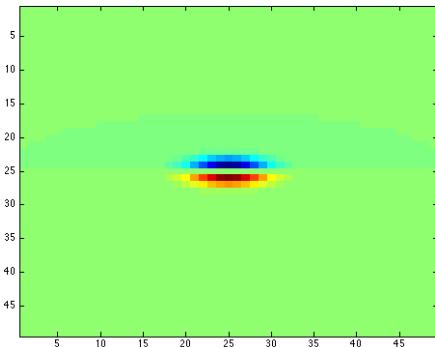
Filter bank



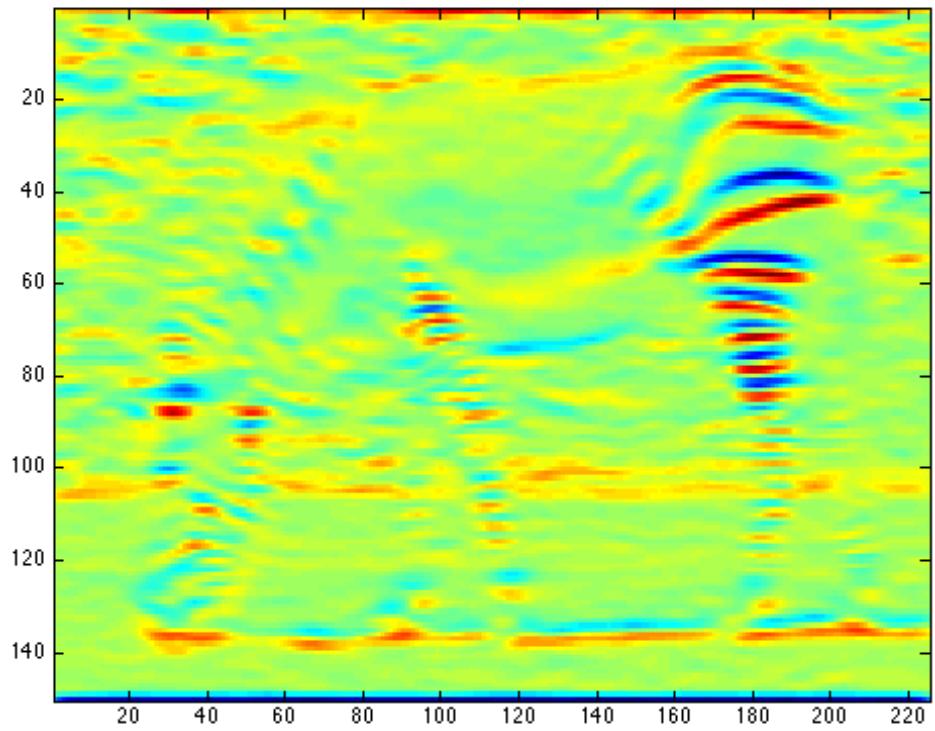
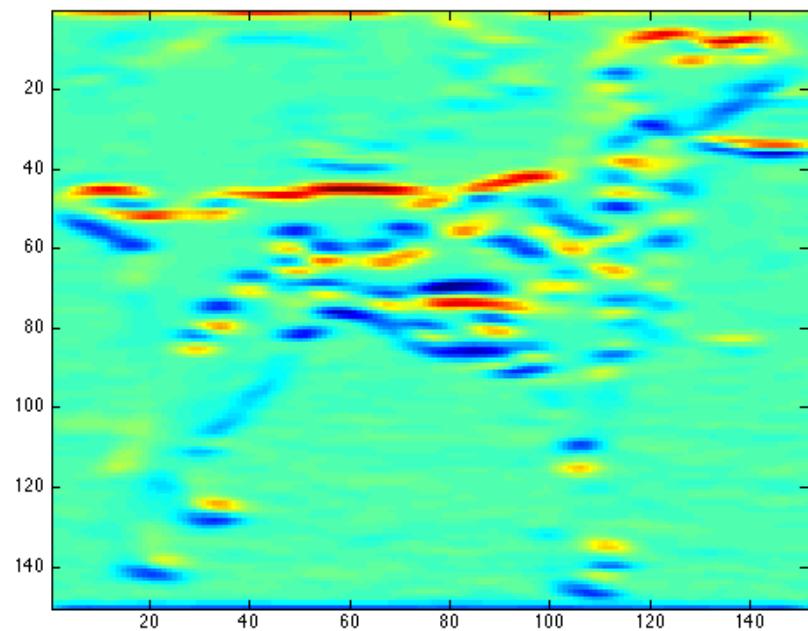
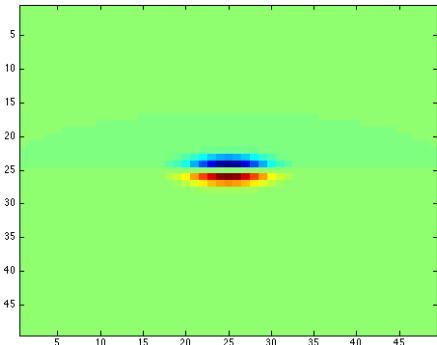
Initial images



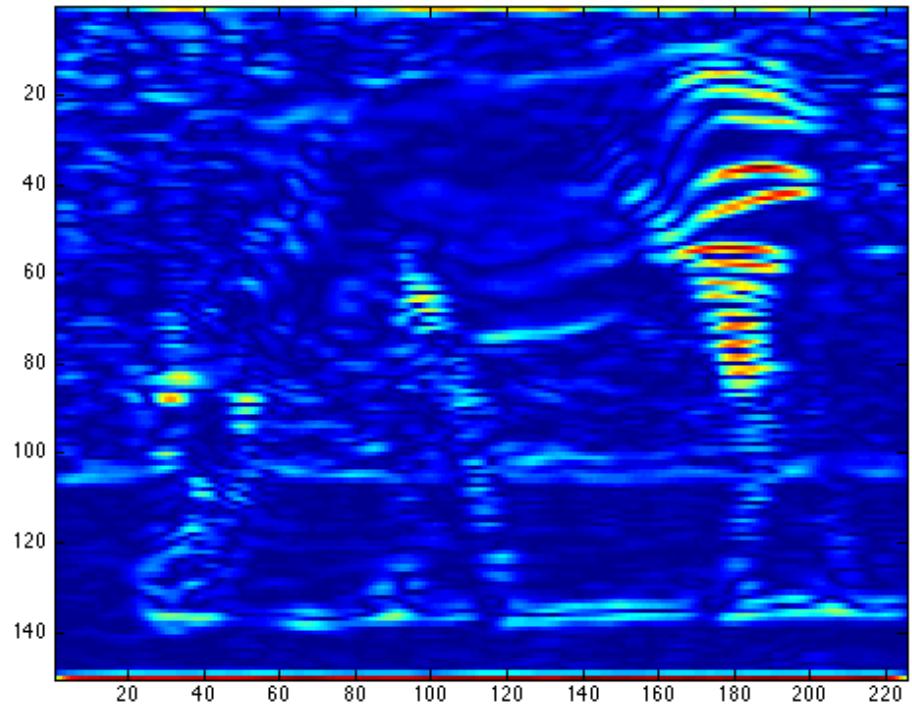
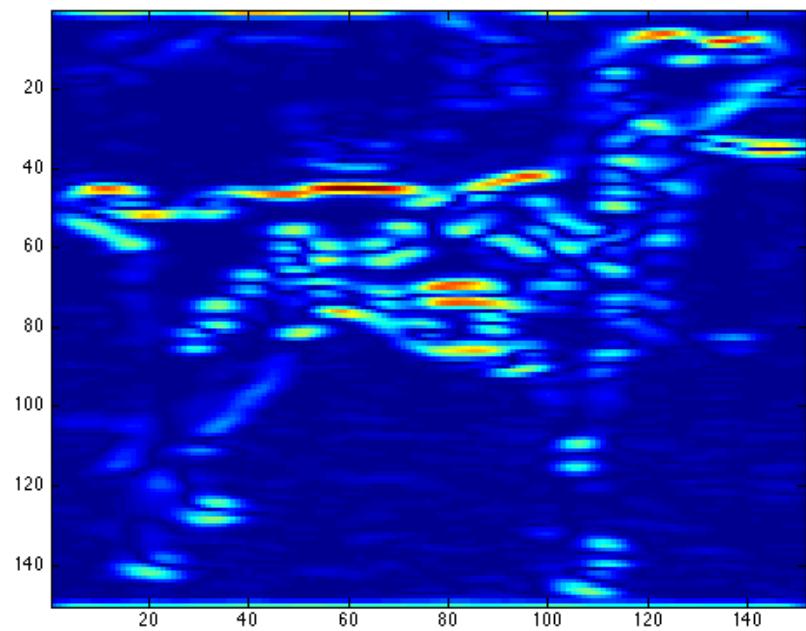
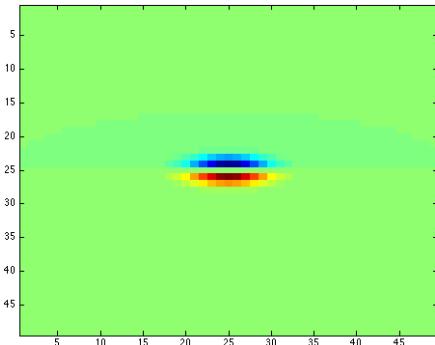
Initial images + filter



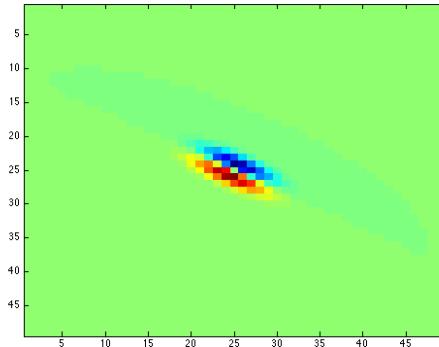
Filtered images



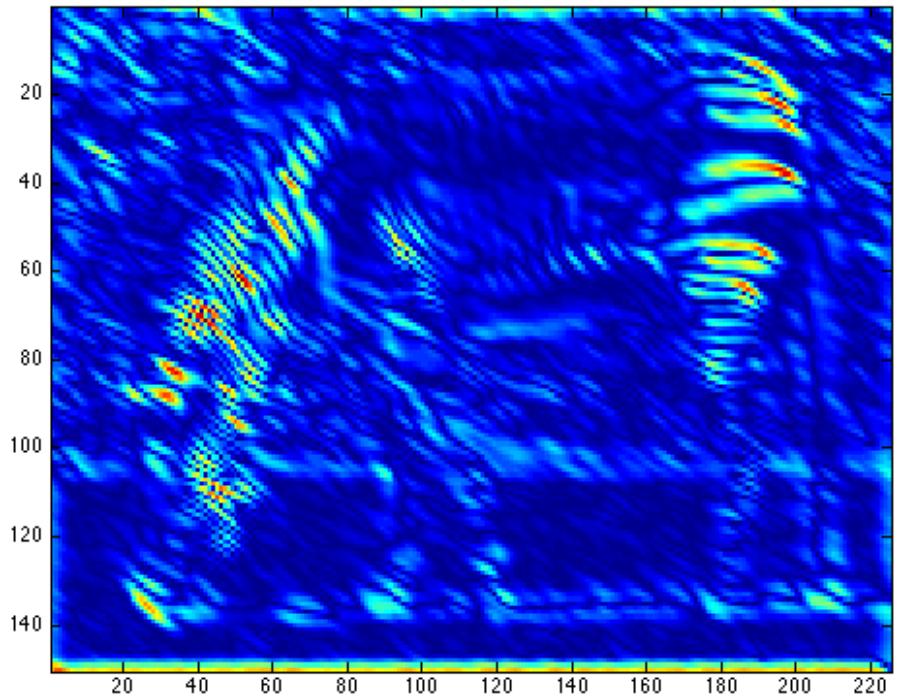
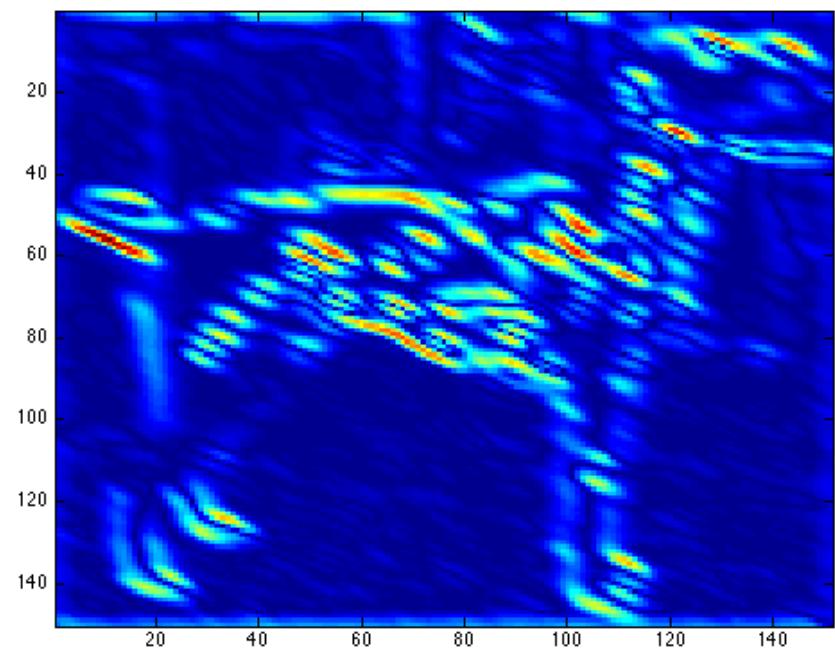
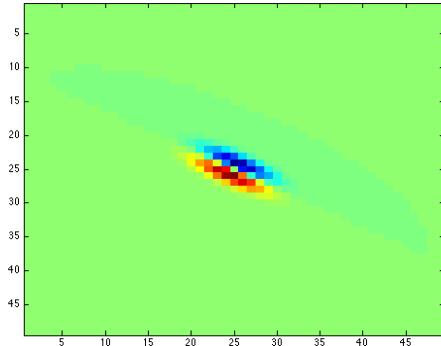
Filtered images with absolute response



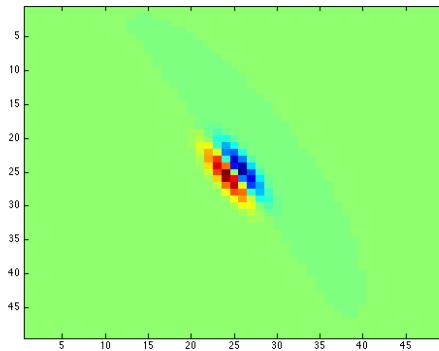
Initial images + filter



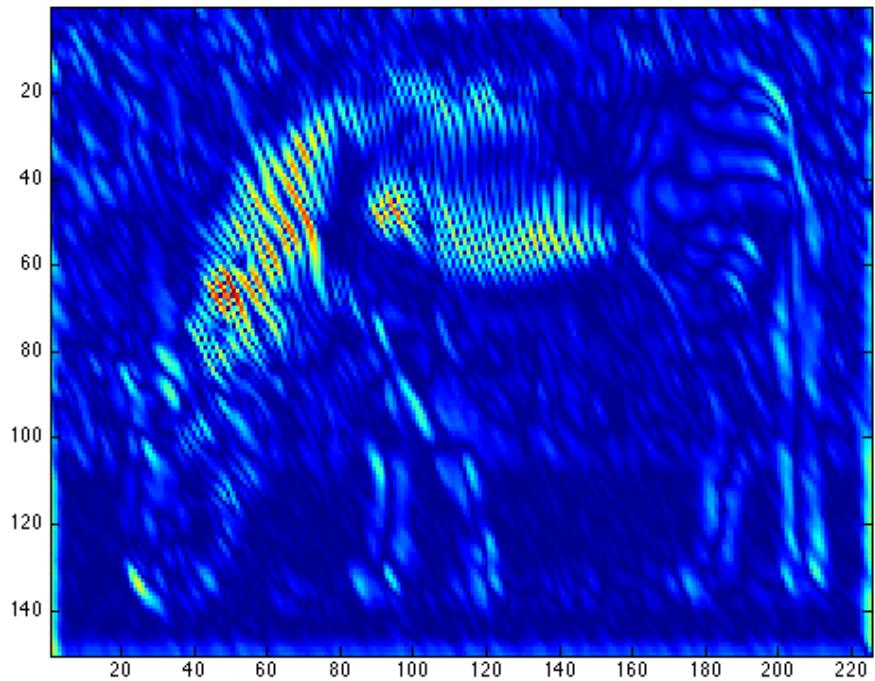
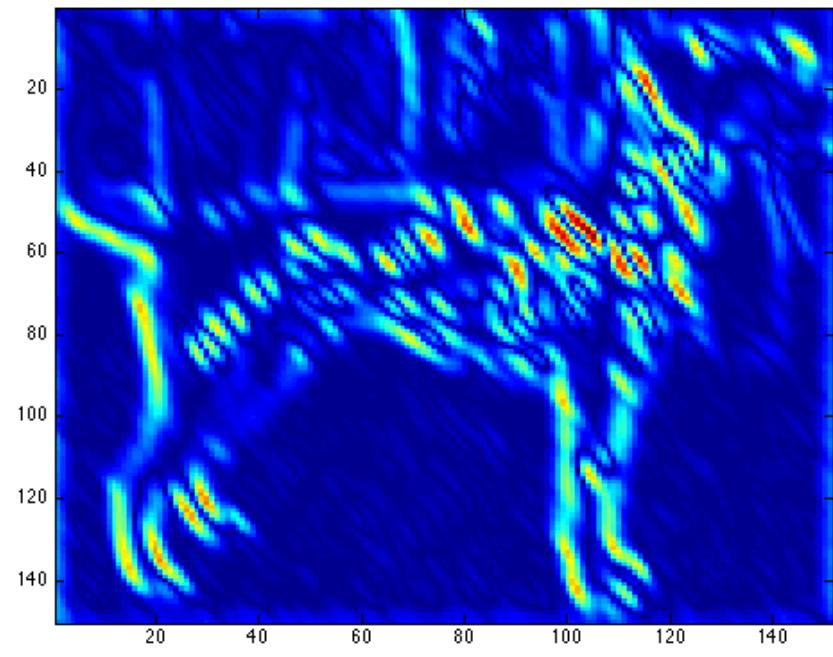
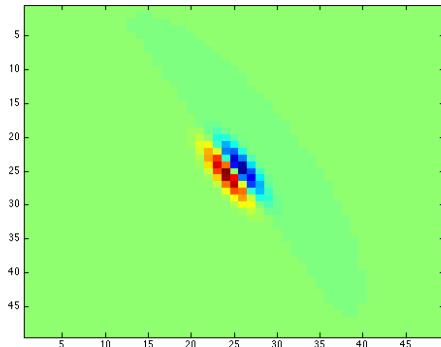
Filtered images with absolute response



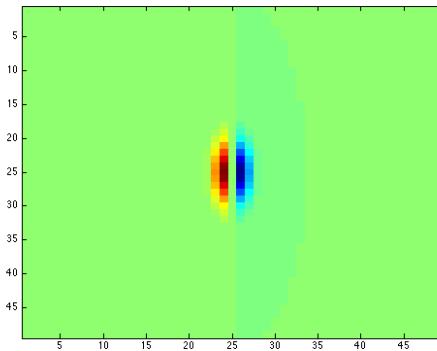
Initial images + filter



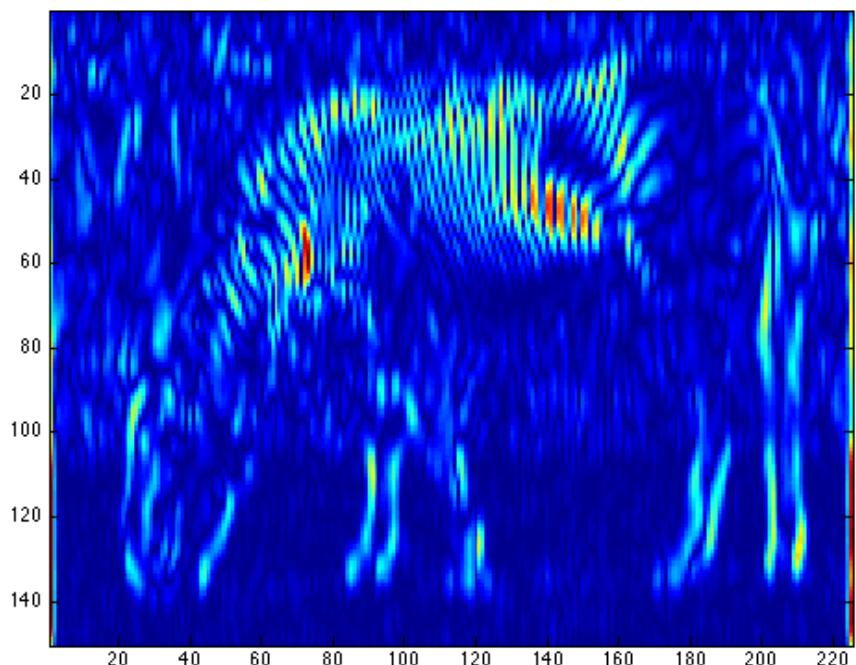
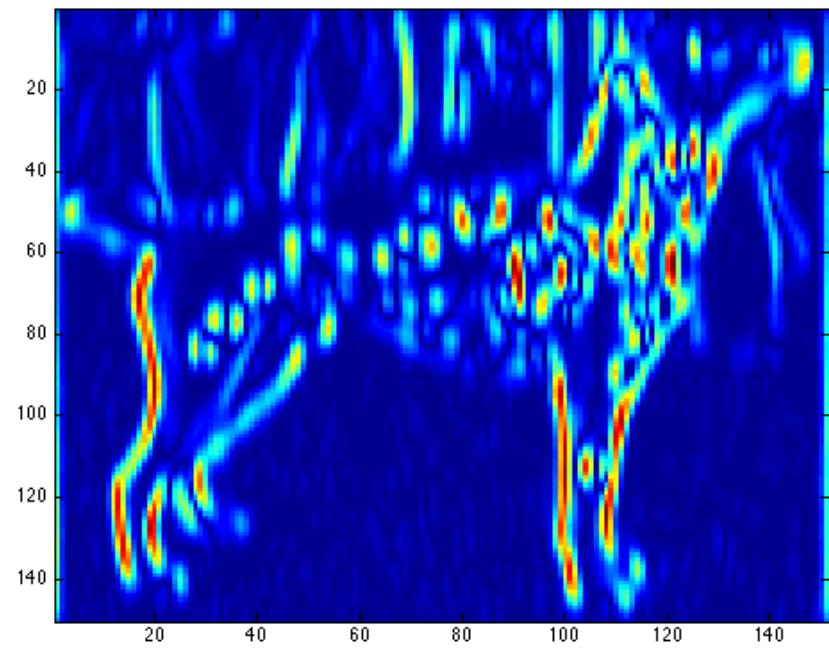
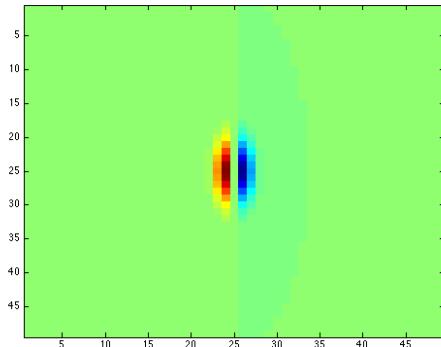
Filtered images with absolute response



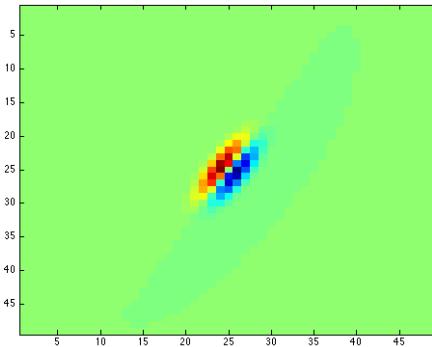
Initial images + filter



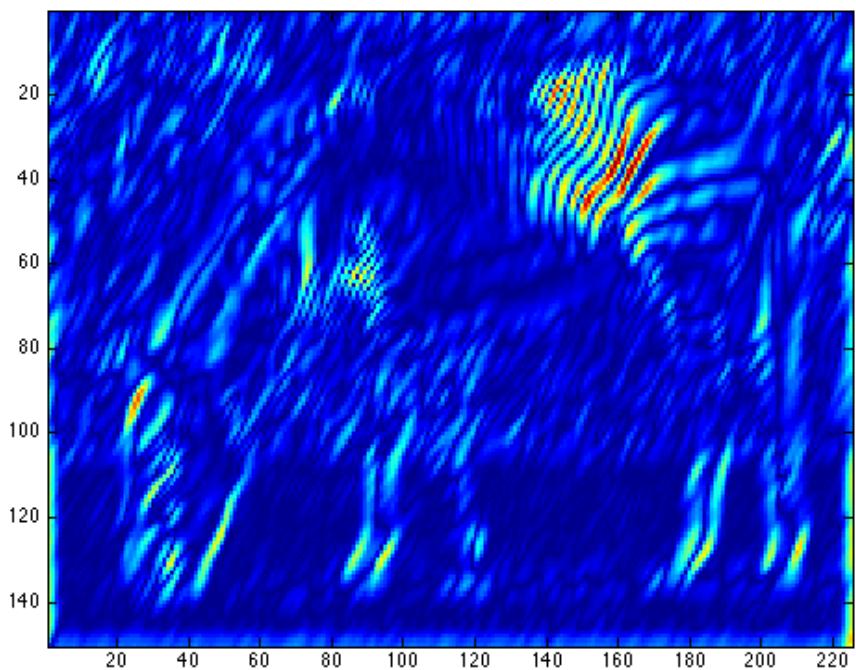
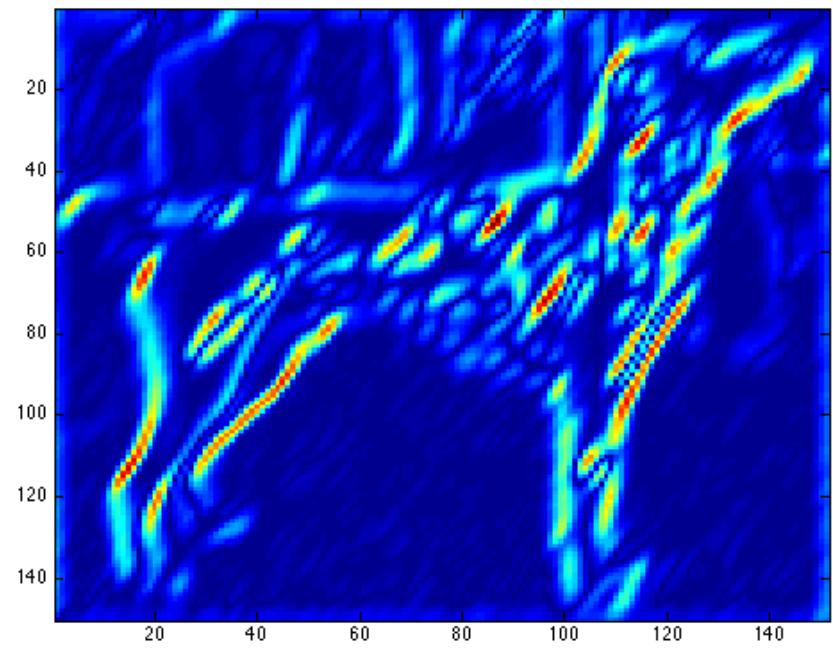
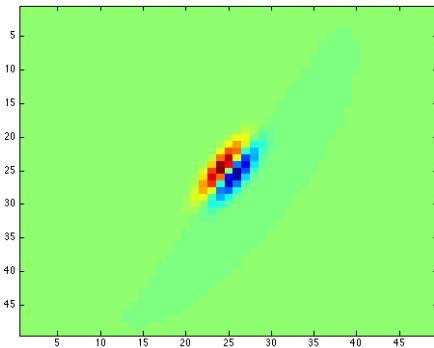
Filtered images with absolute response



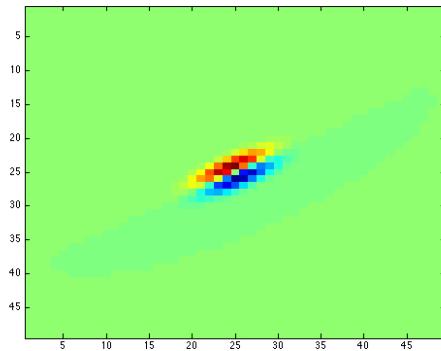
Initial images + filter



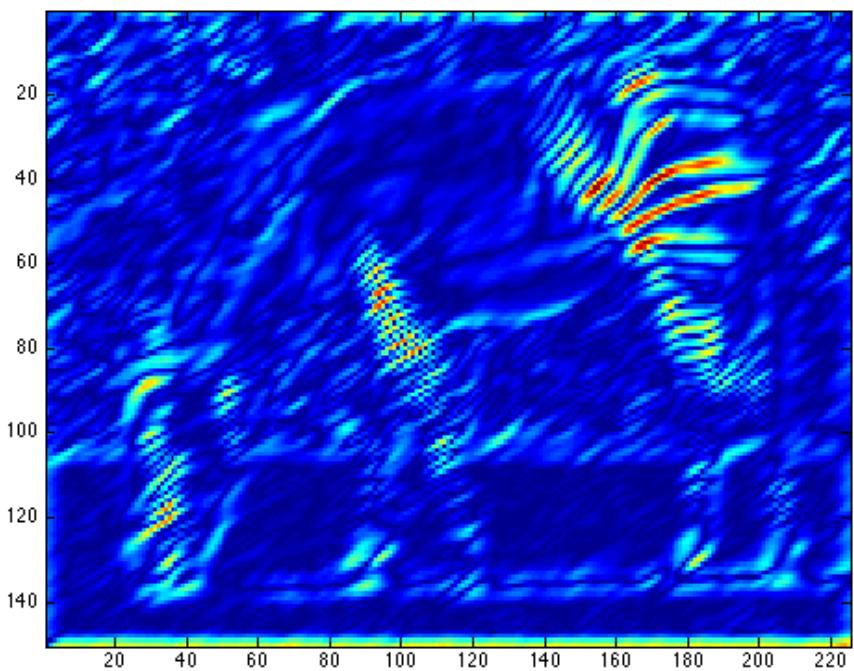
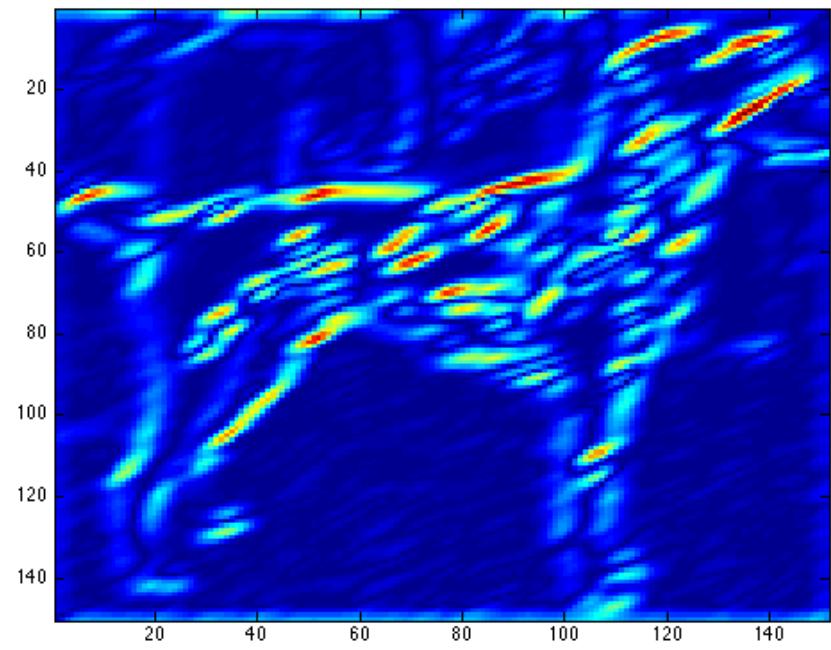
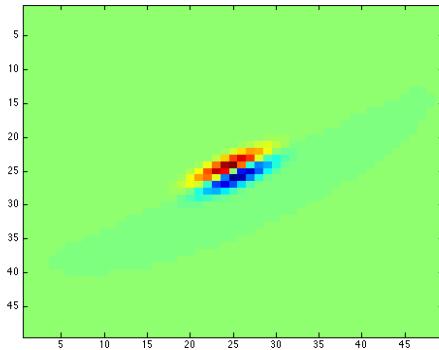
Filtered images with absolute response



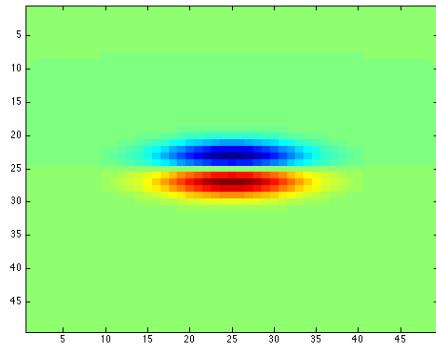
Initial images + filter



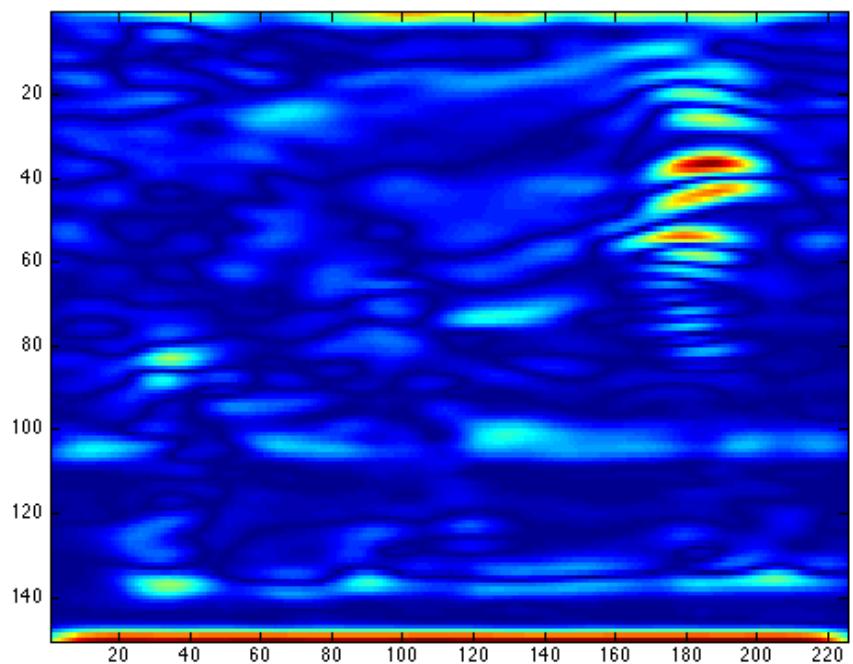
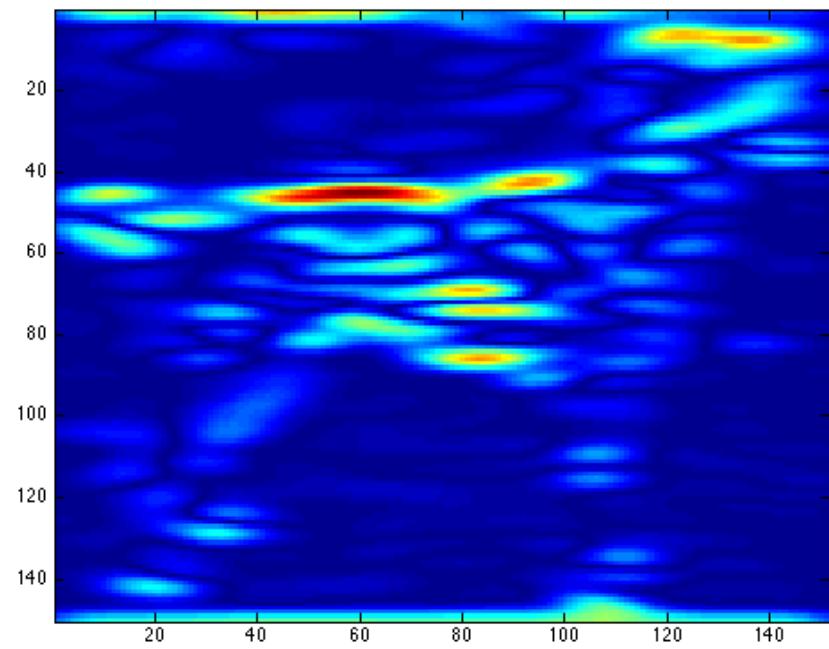
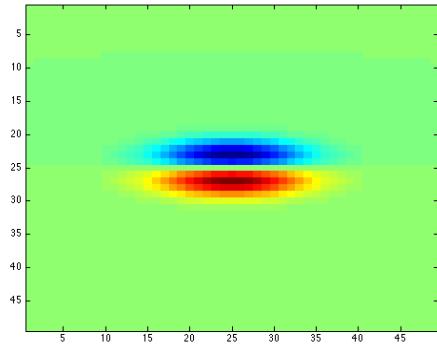
Filtered images with absolute response



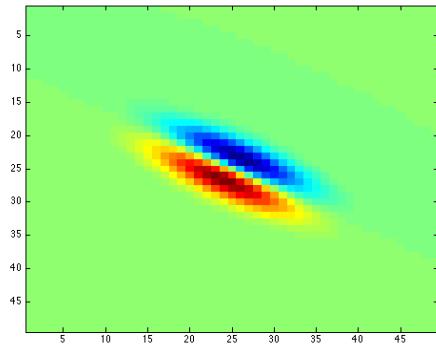
Initial images + filter



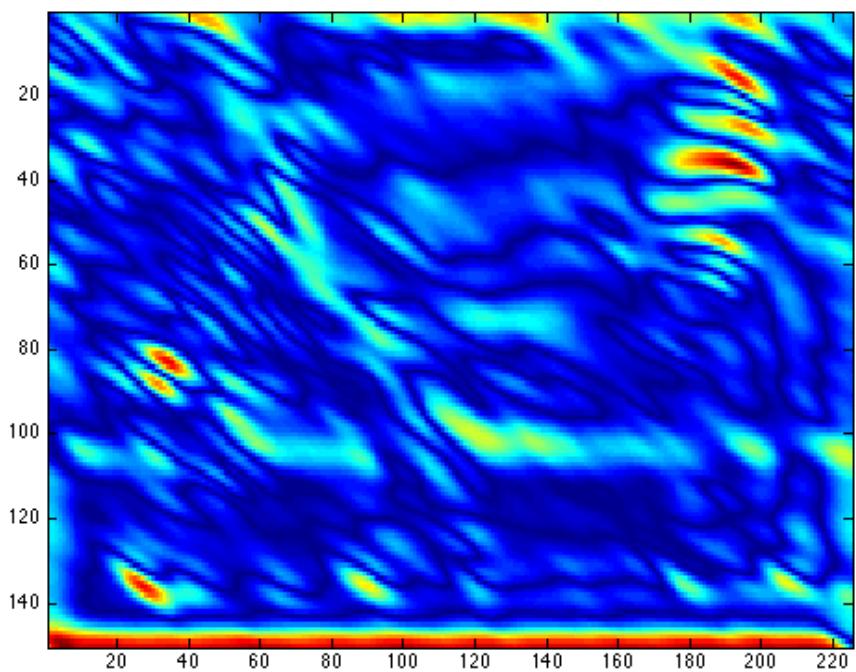
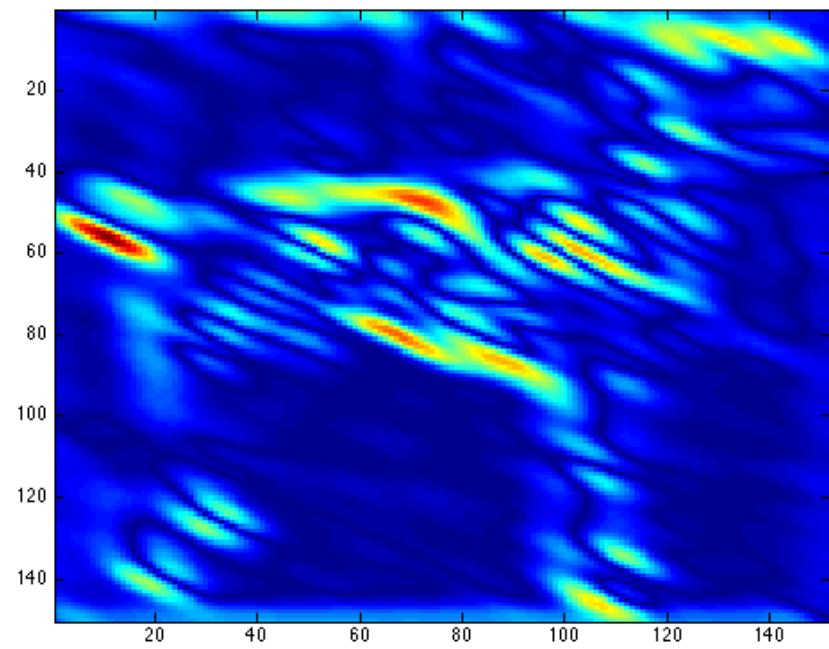
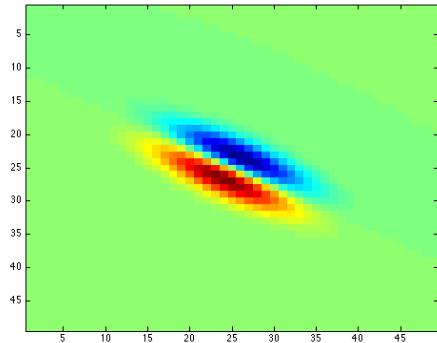
Filtered images with absolute response



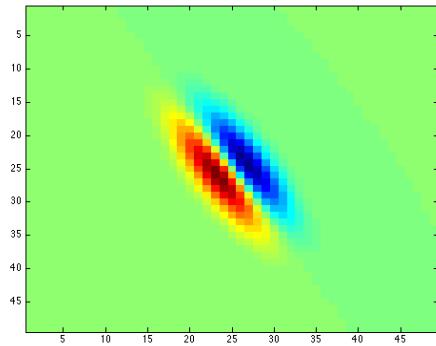
Initial images + filter



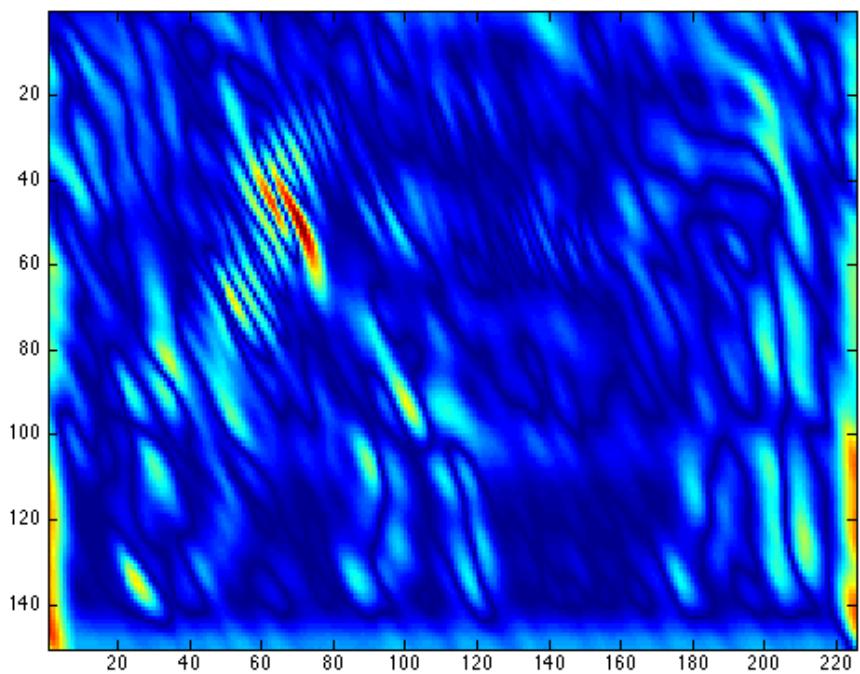
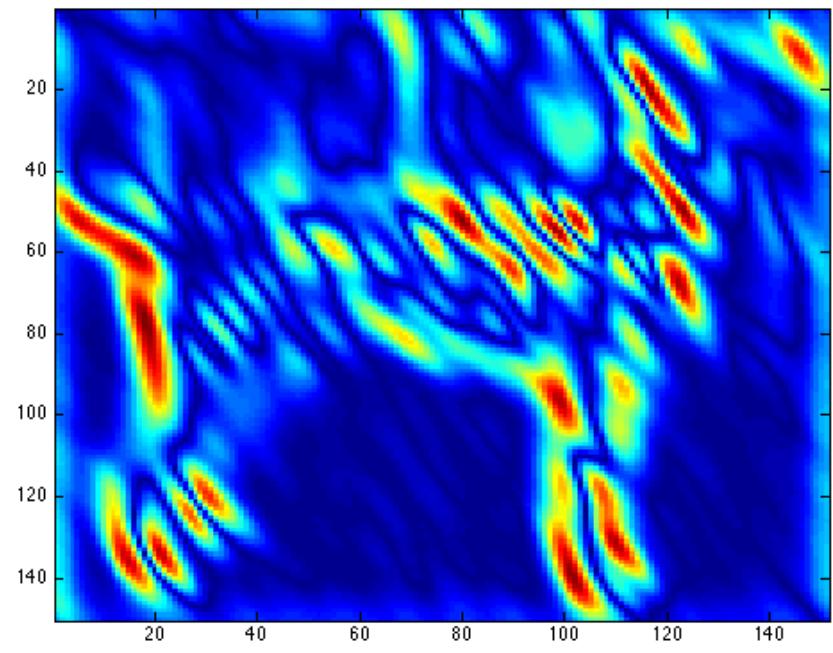
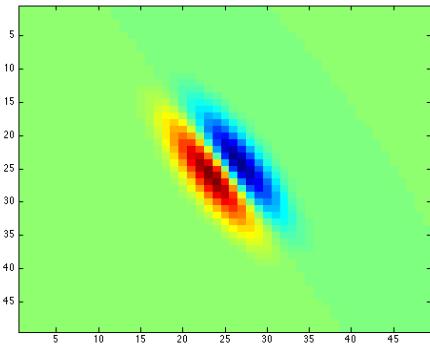
Filtered images with absolute response



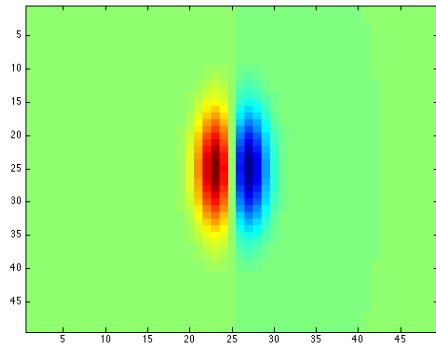
Initial images + filter



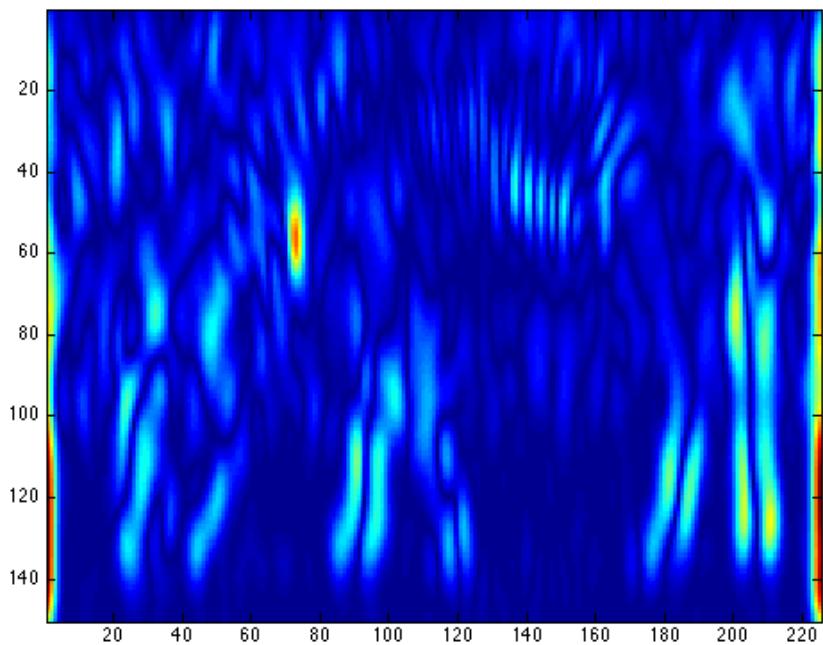
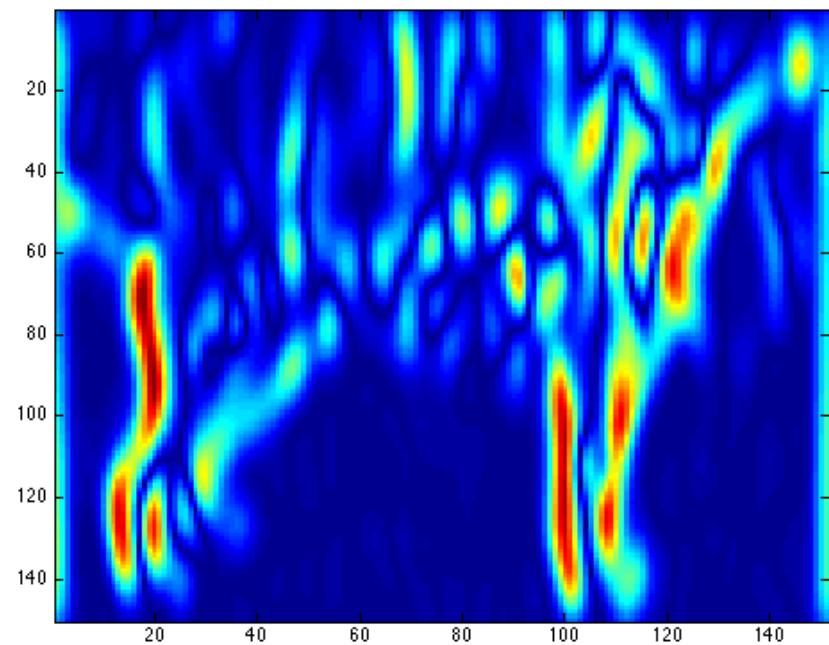
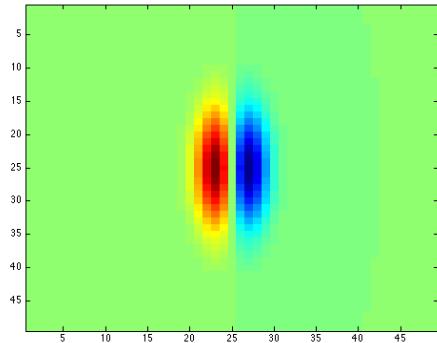
Filtered images with absolute response



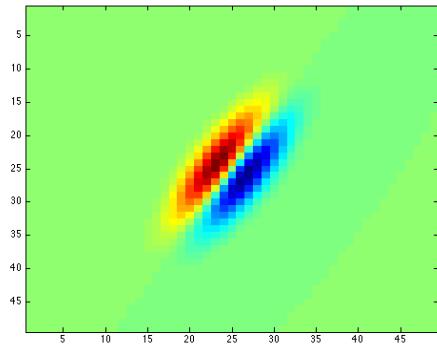
Initial images + filter



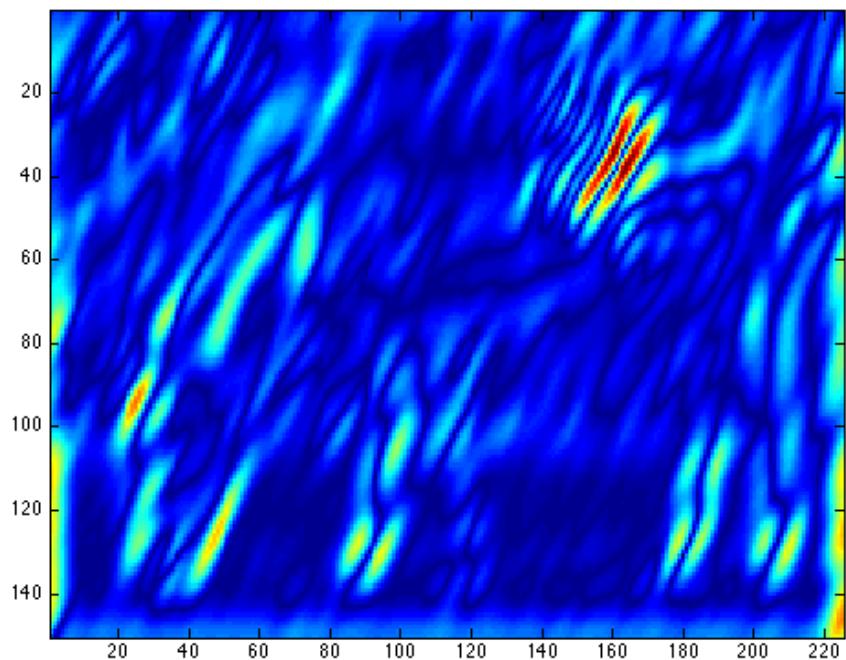
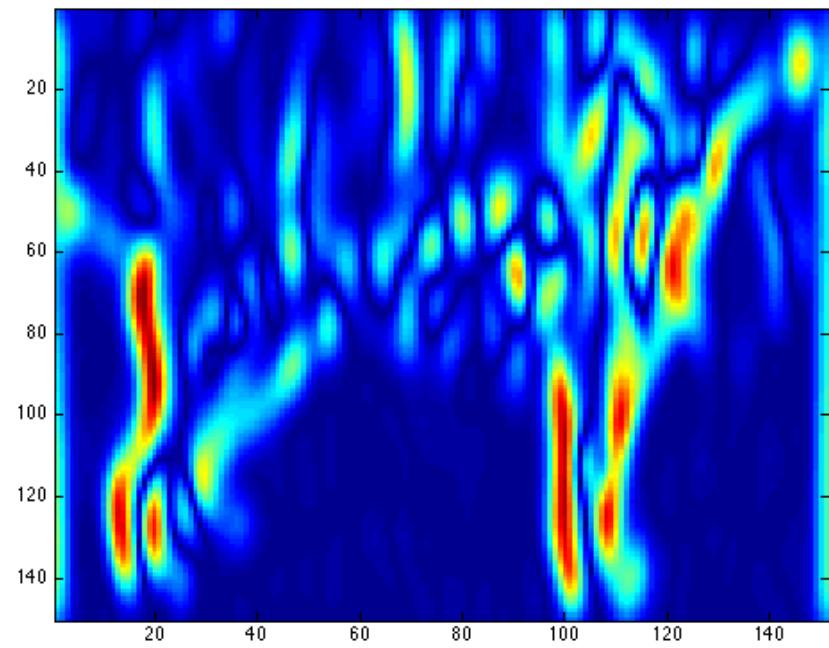
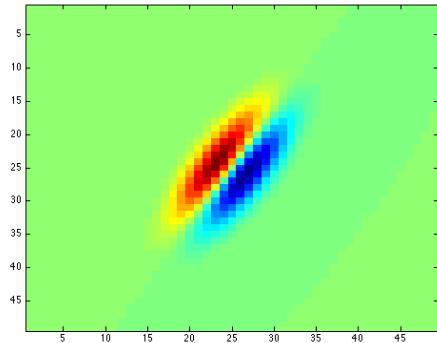
Filtered images with absolute response



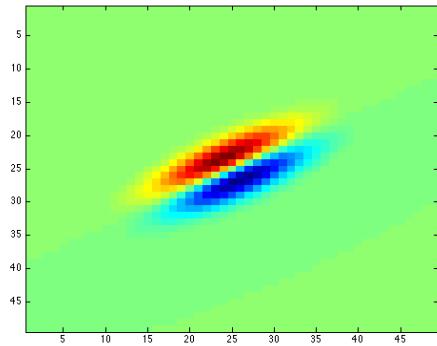
Initial images + filter



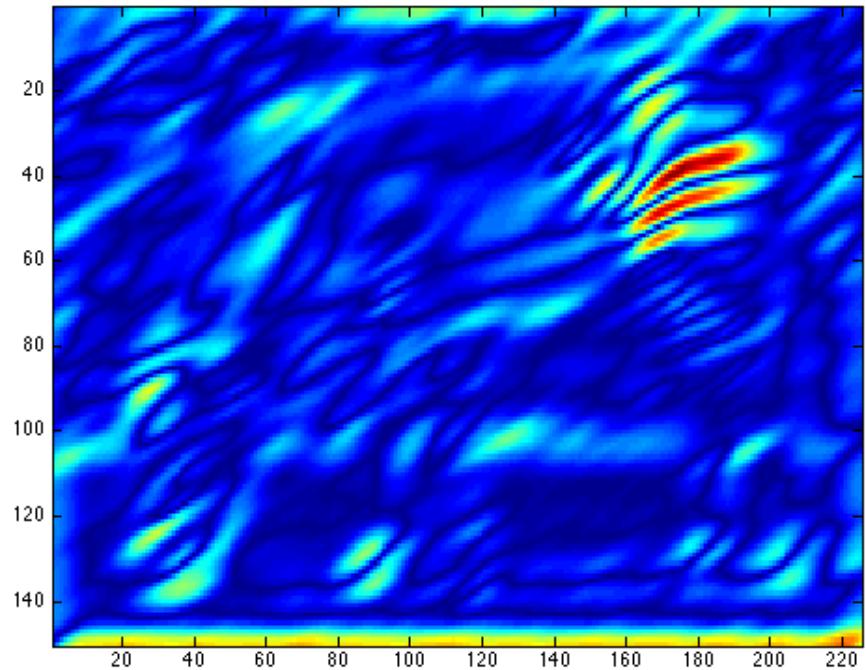
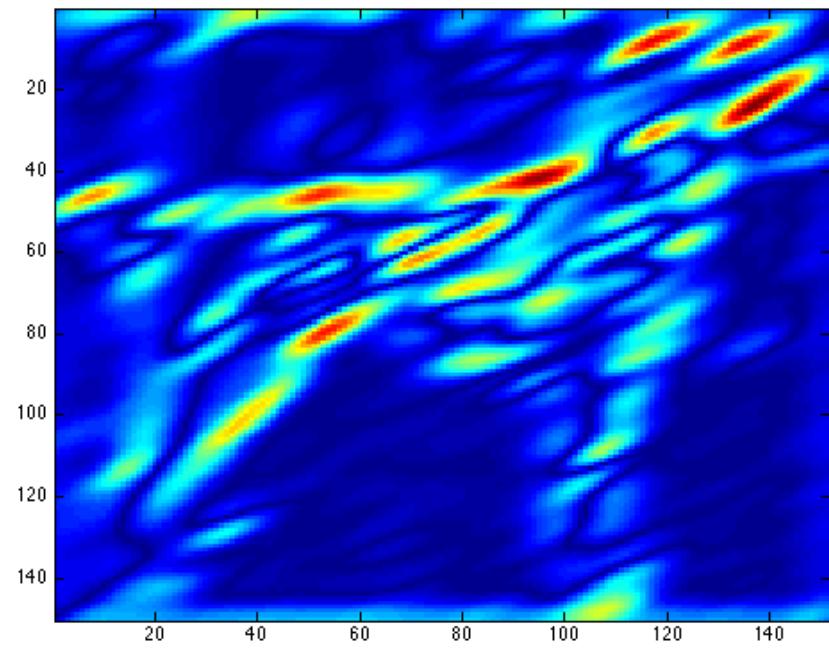
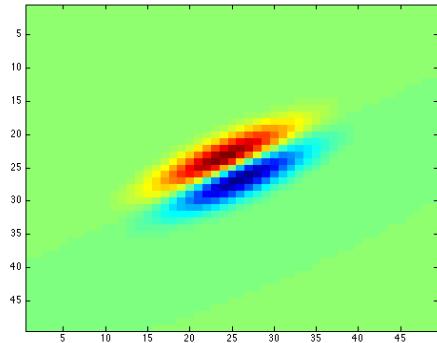
Filtered images with absolute response



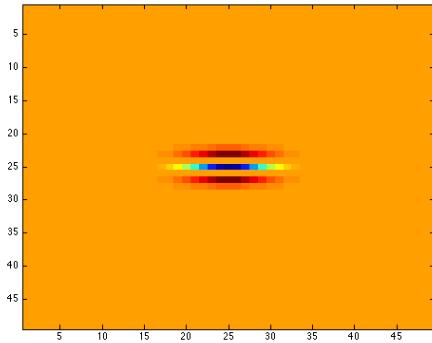
Initial images + filter



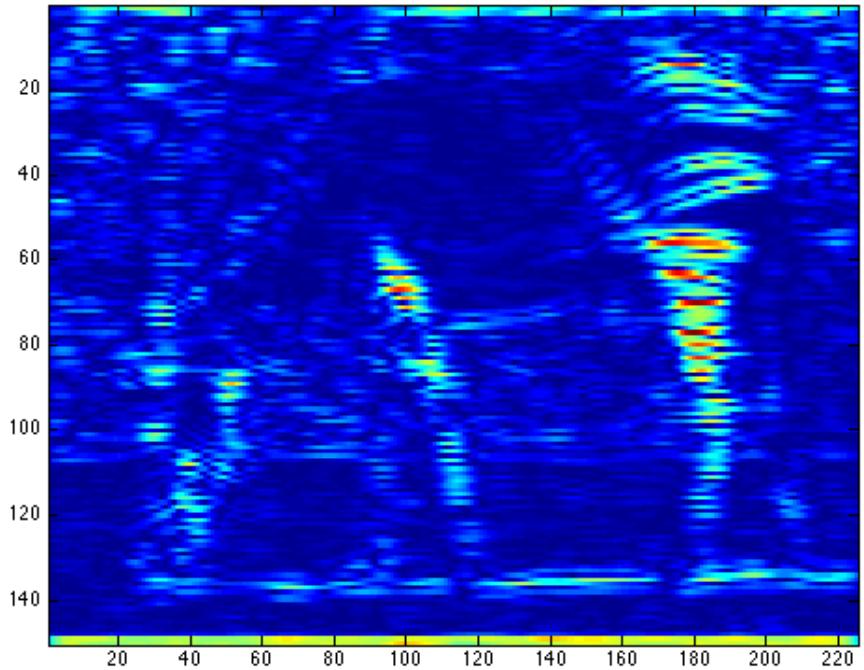
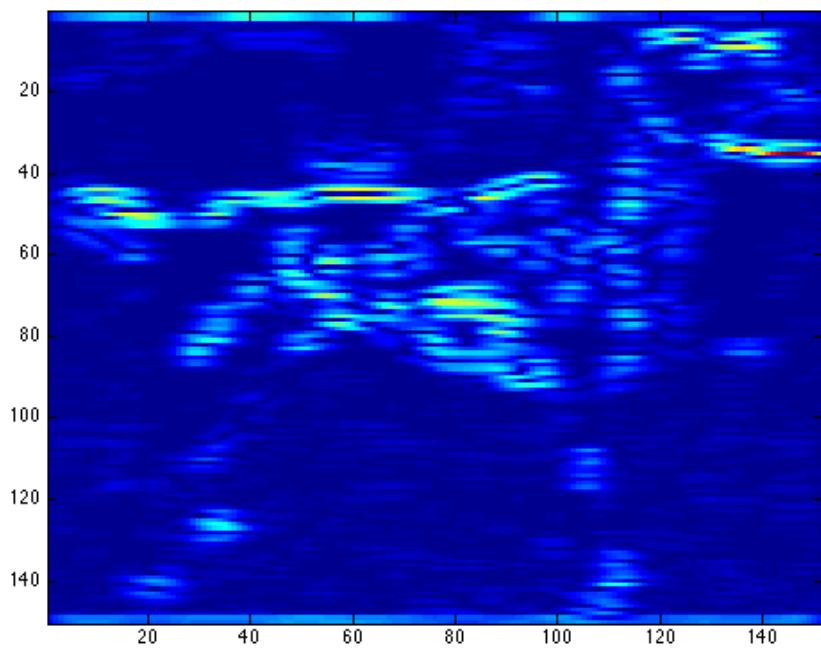
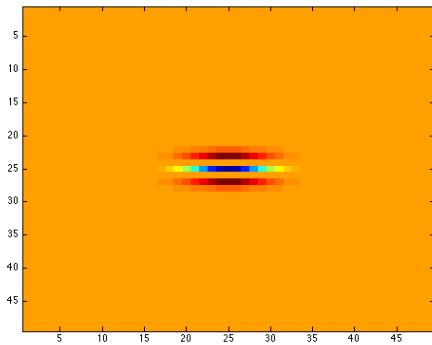
Filtered images with absolute response



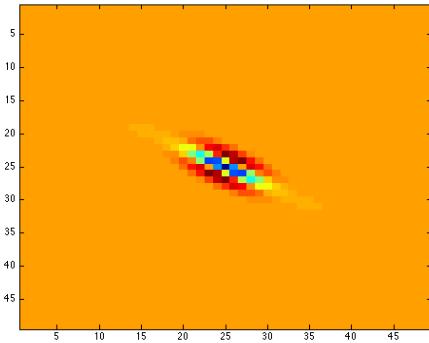
Initial images + filter



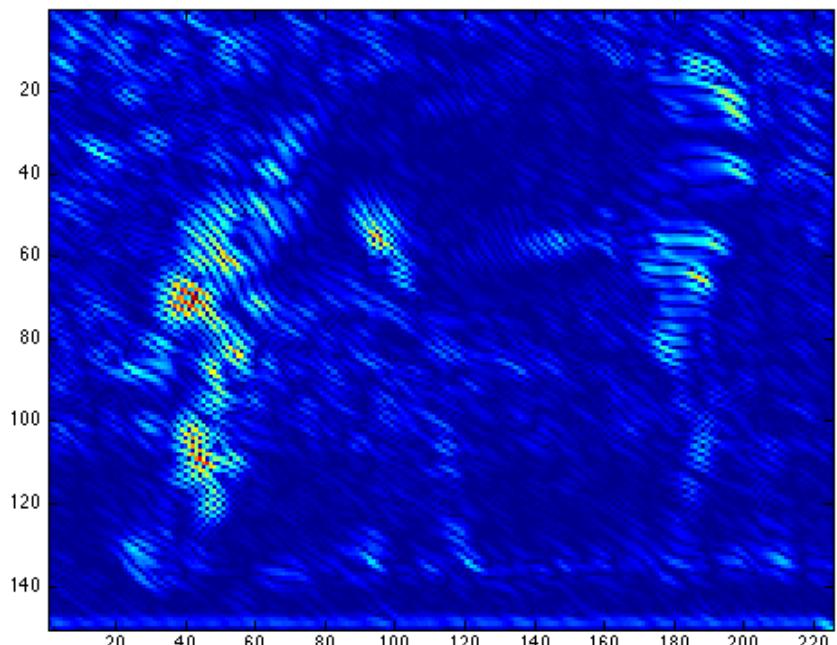
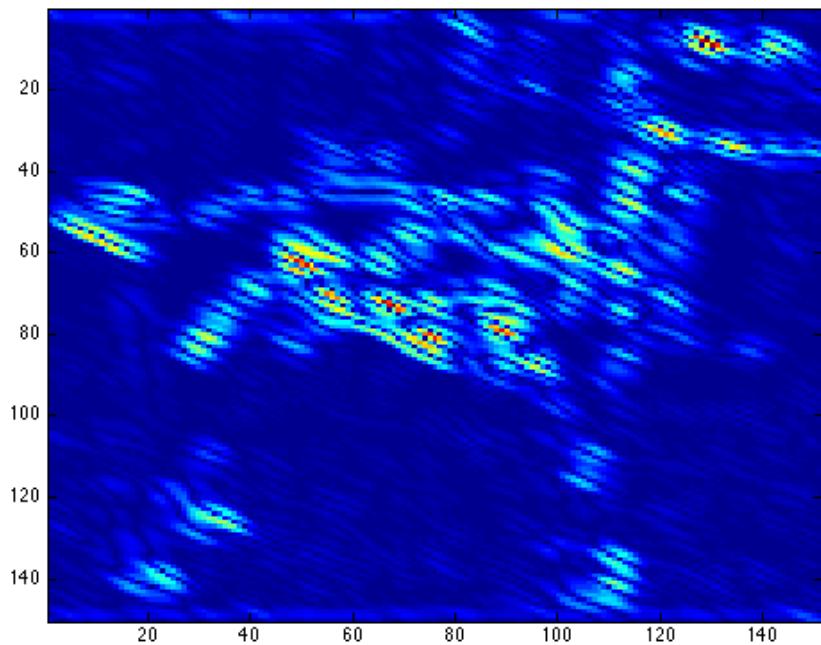
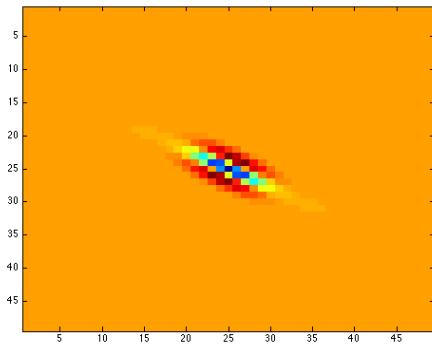
Filtered images with absolute response



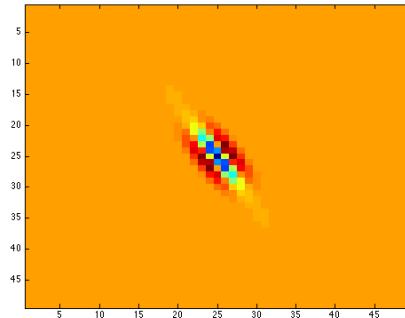
Initial images + filter



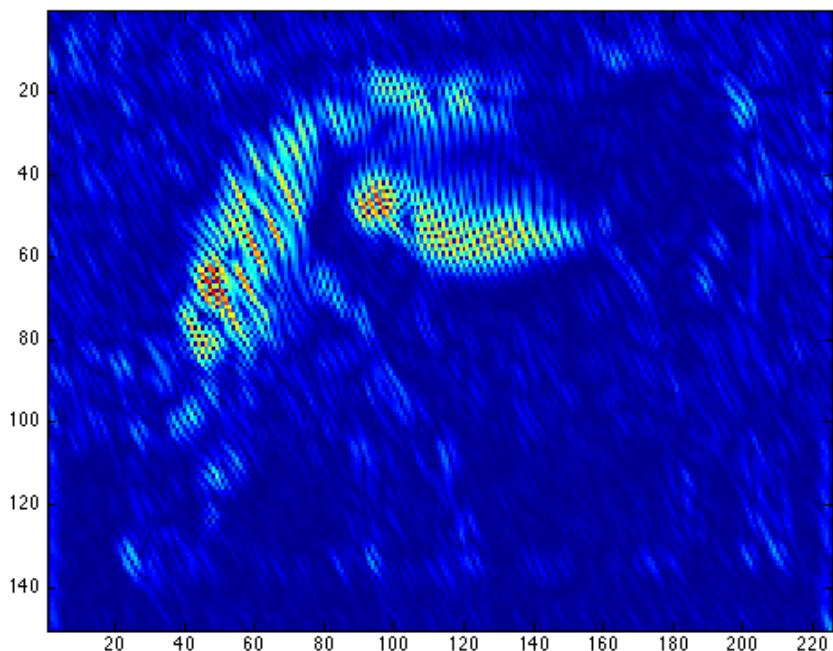
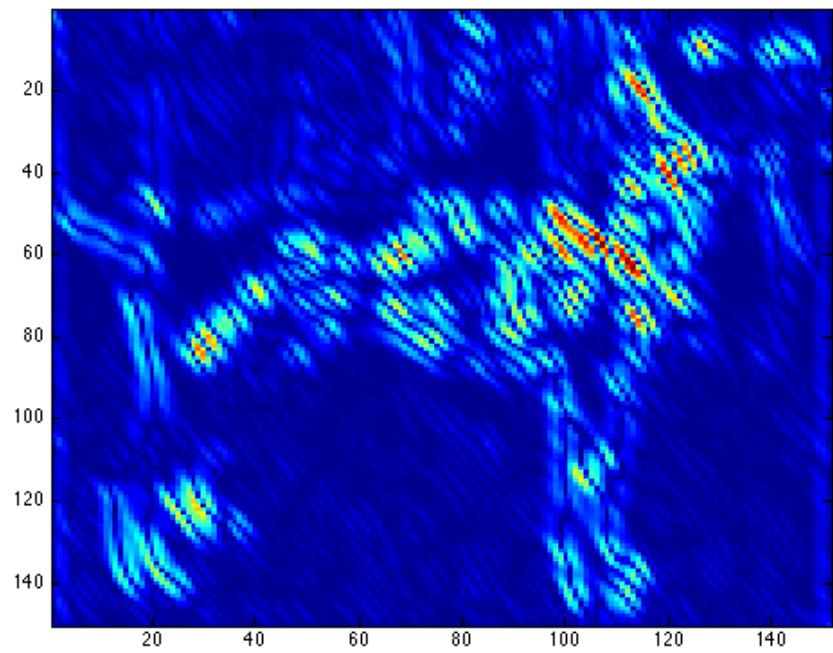
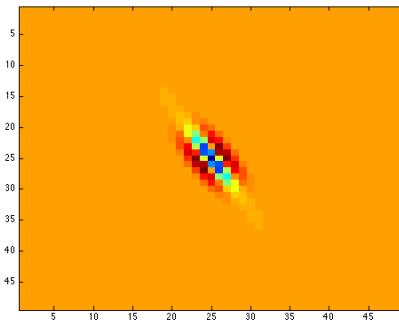
Filtered images with absolute response



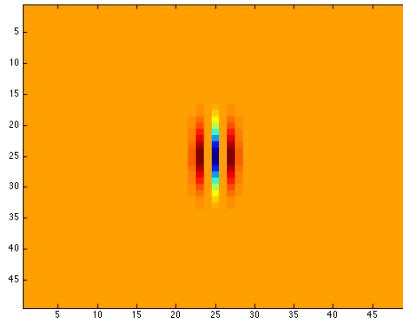
Initial images + filter



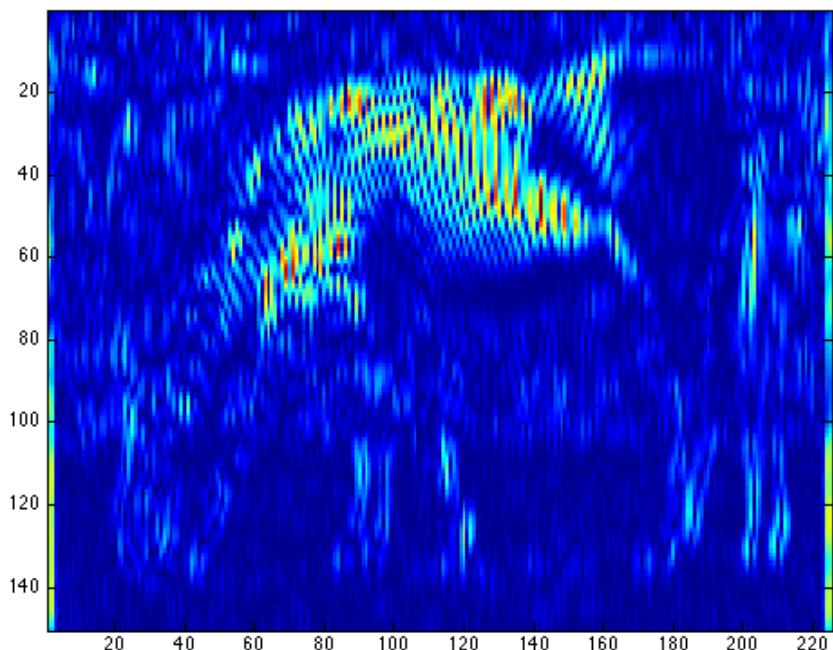
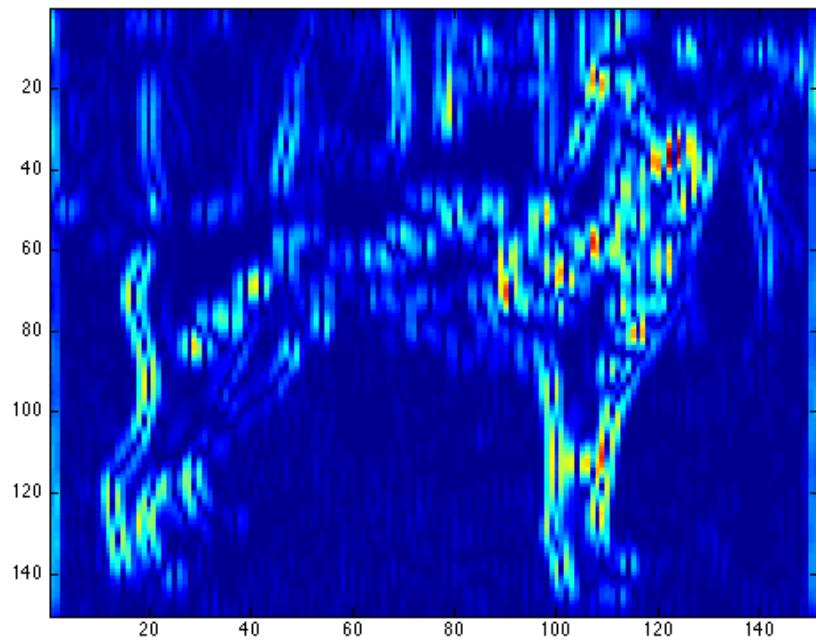
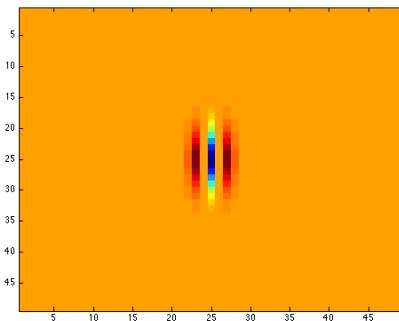
Filtered images with absolute response



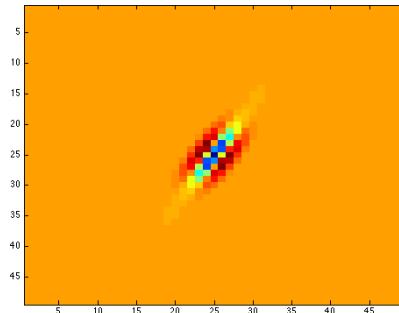
Initial images + filter



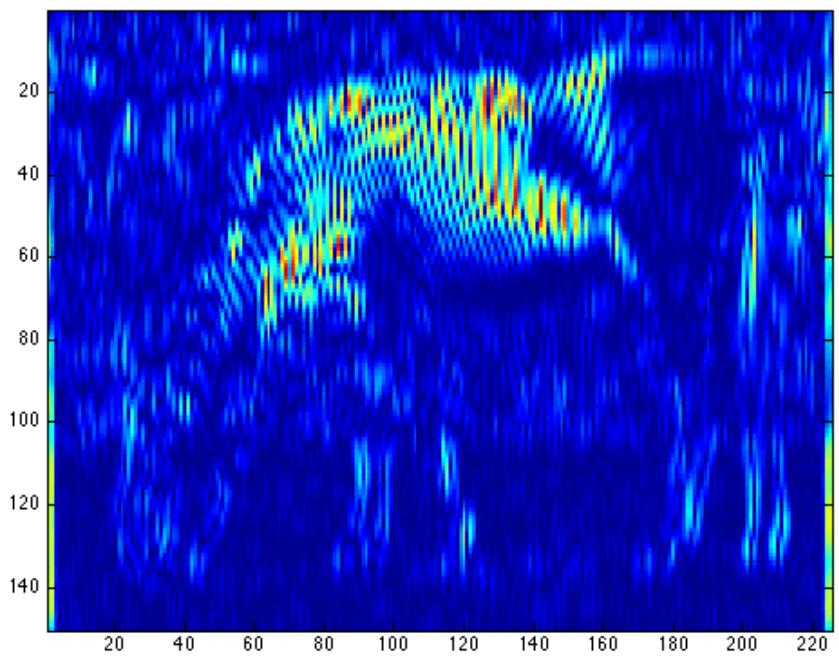
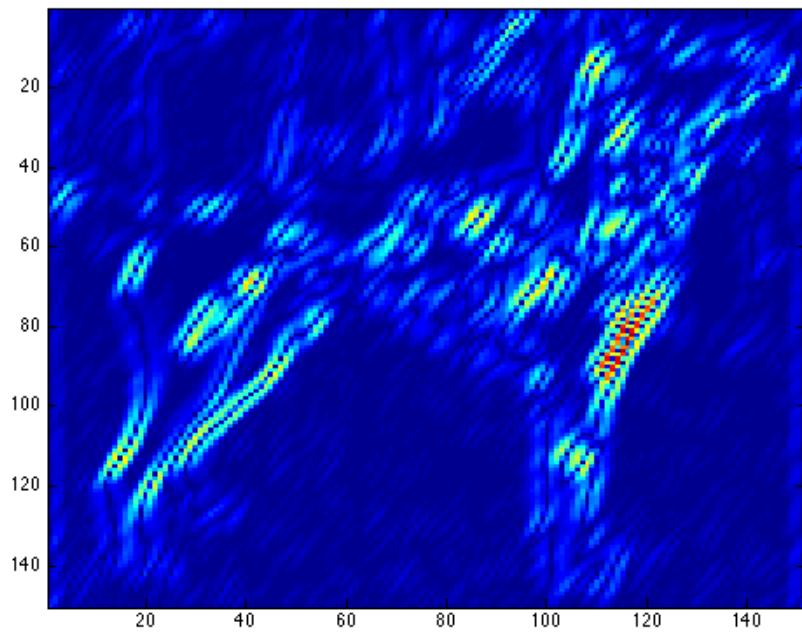
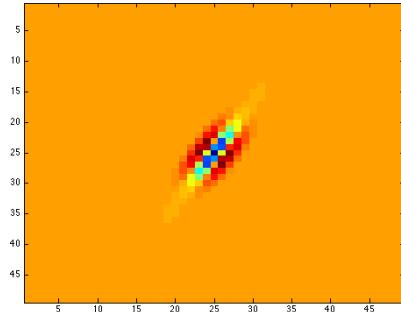
Filtered images with absolute response



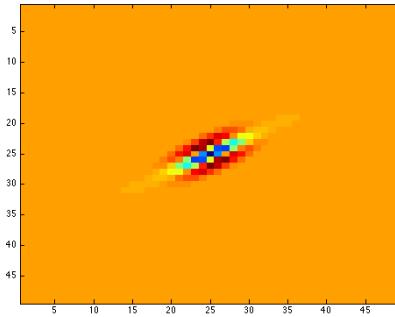
Initial images + filter



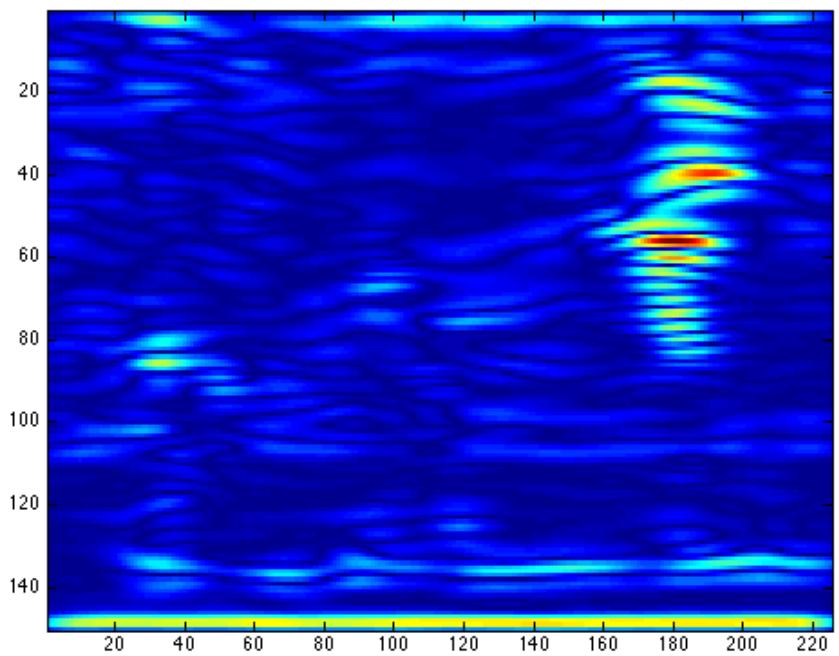
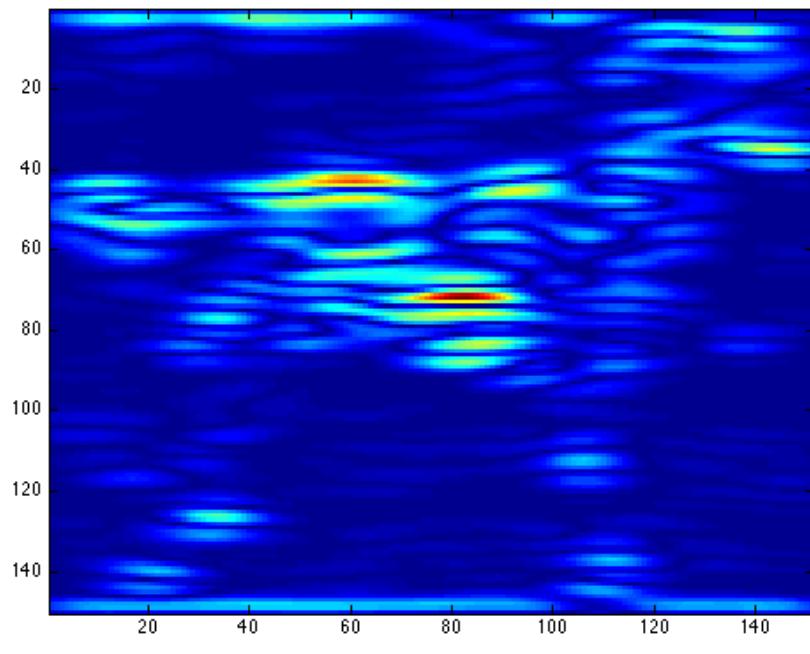
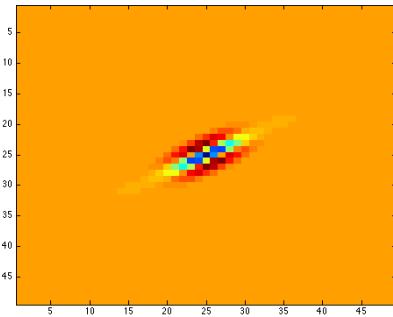
Filtered images with absolute response



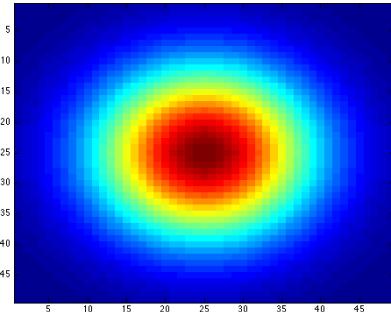
Initial images + filter



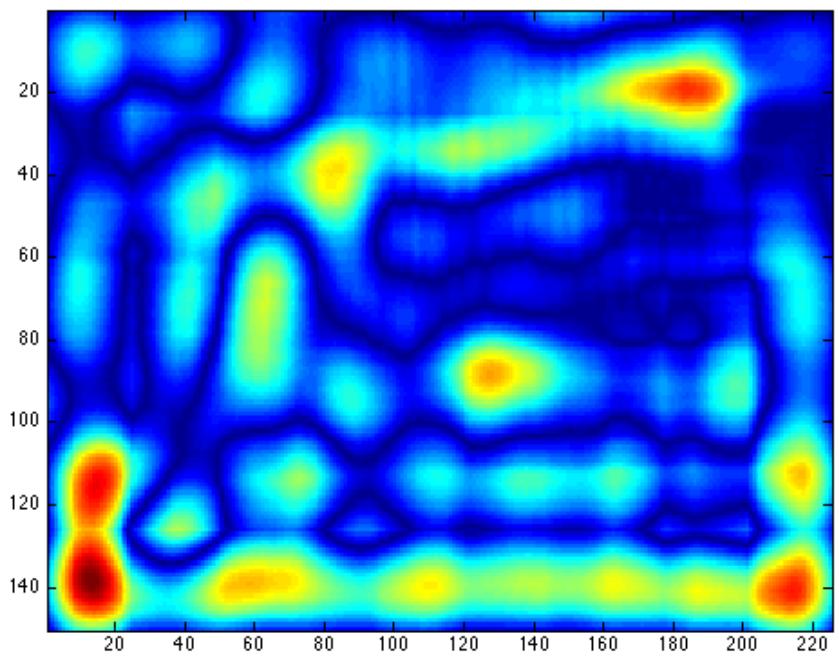
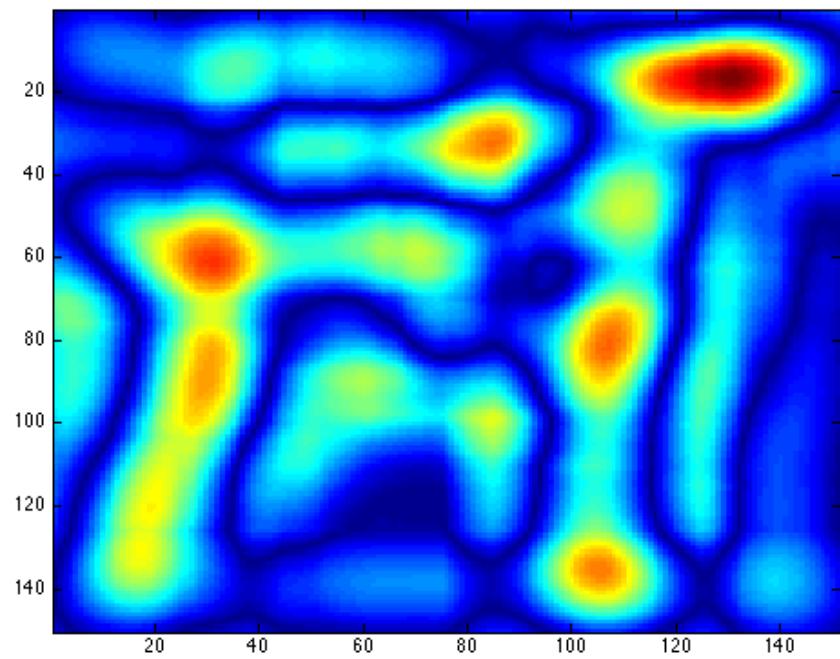
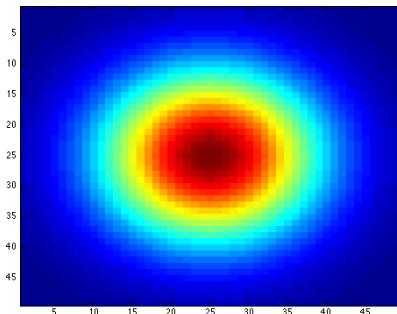
Filtered images with absolute response



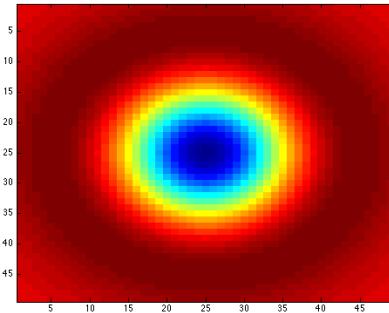
Initial images + filter



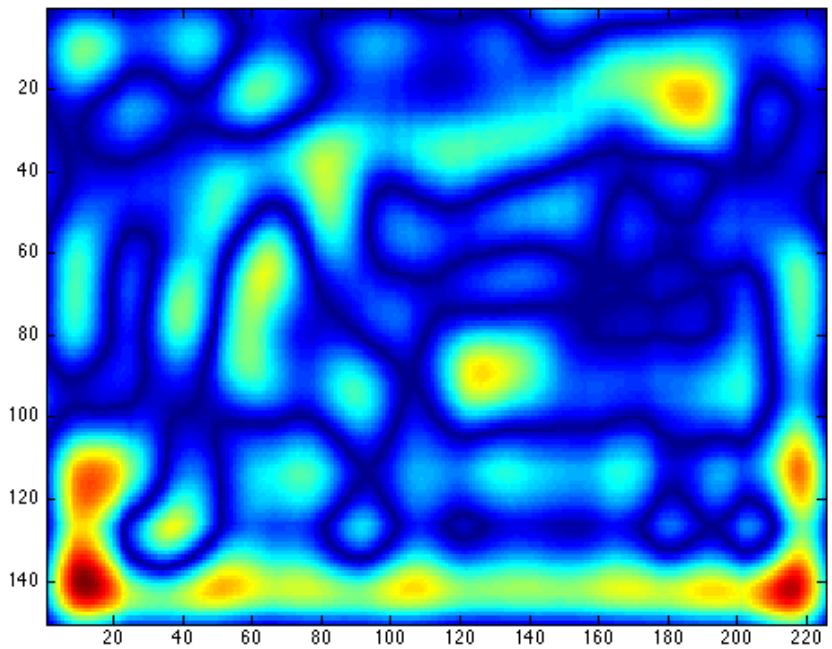
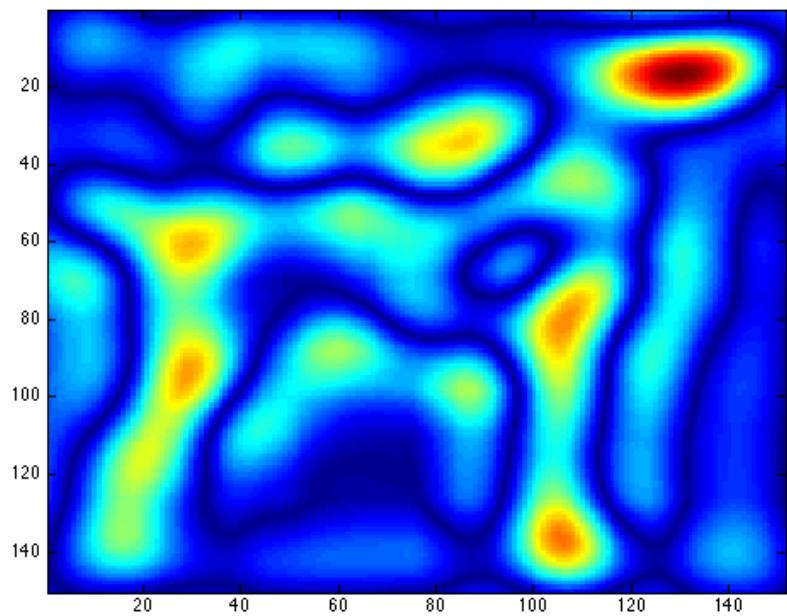
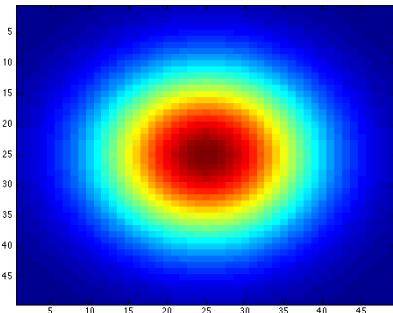
Filtered images with absolute response

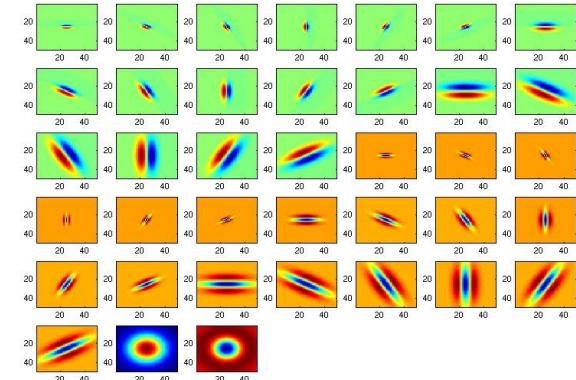
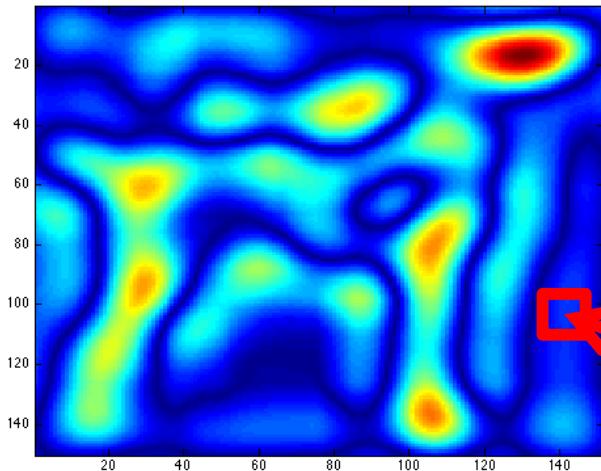


Initial images + filter

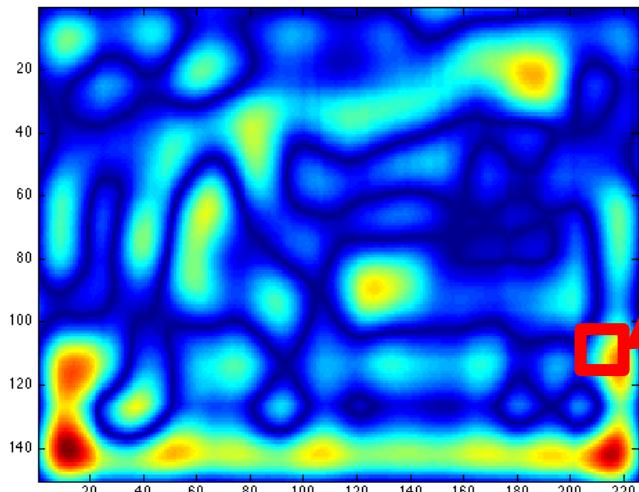


Filtered images with absolute response



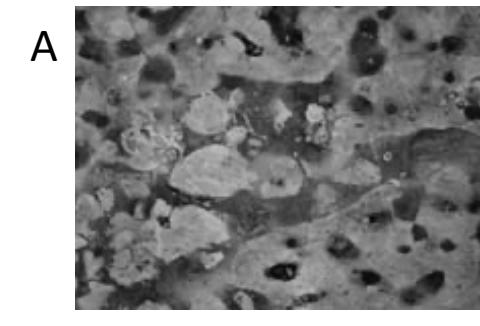
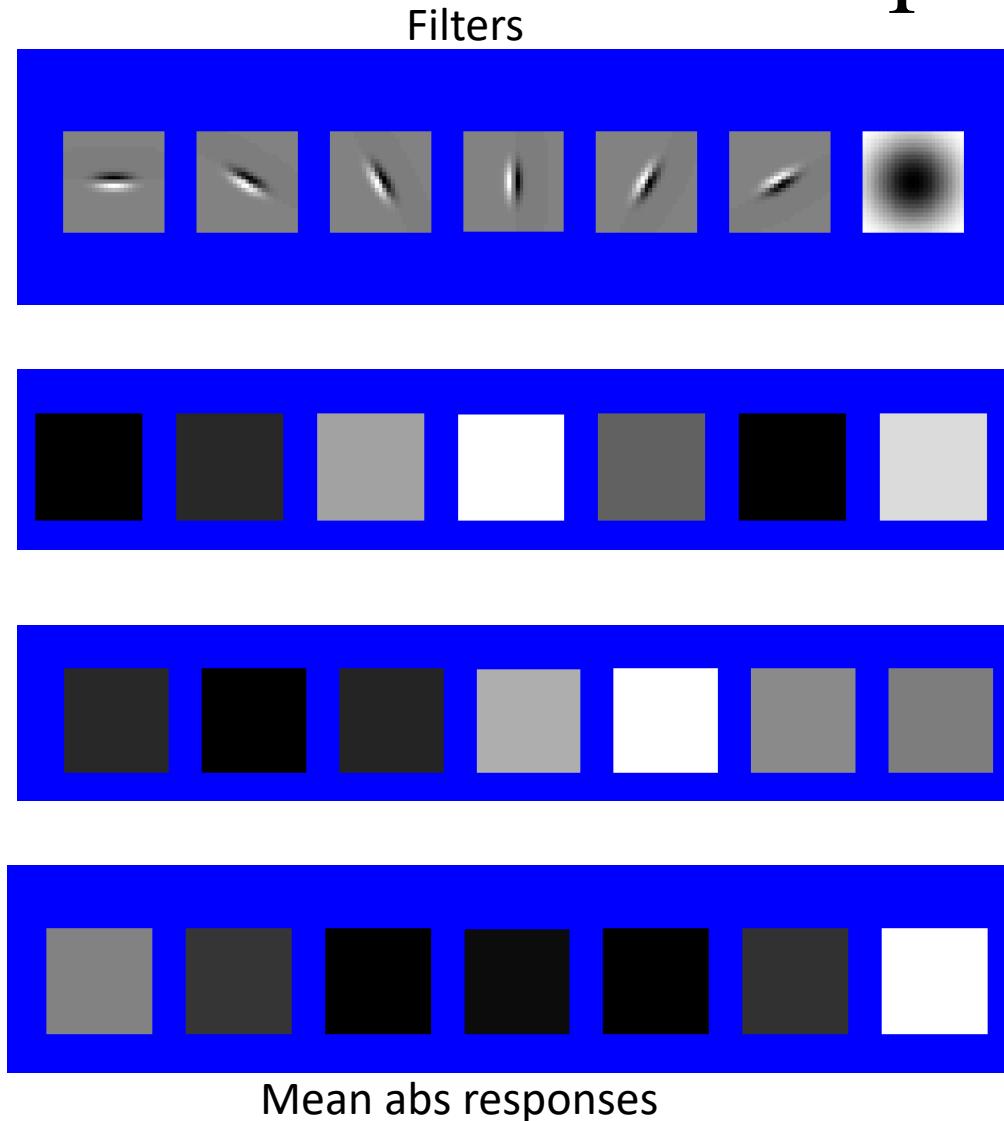


$[r_1, r_2, \dots, r_{38}]$

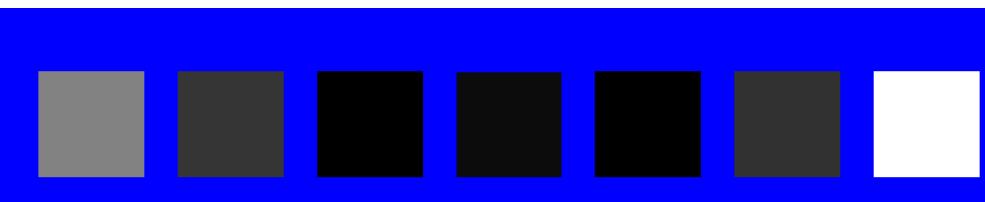
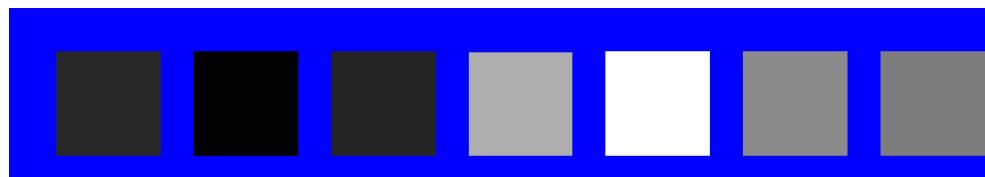
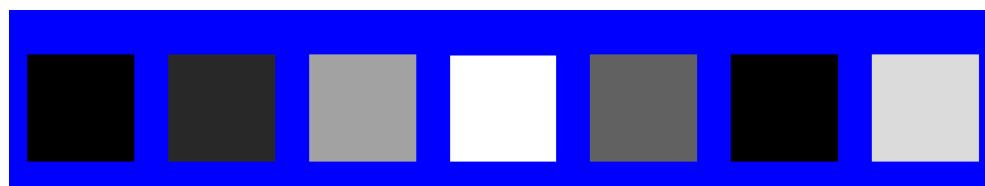
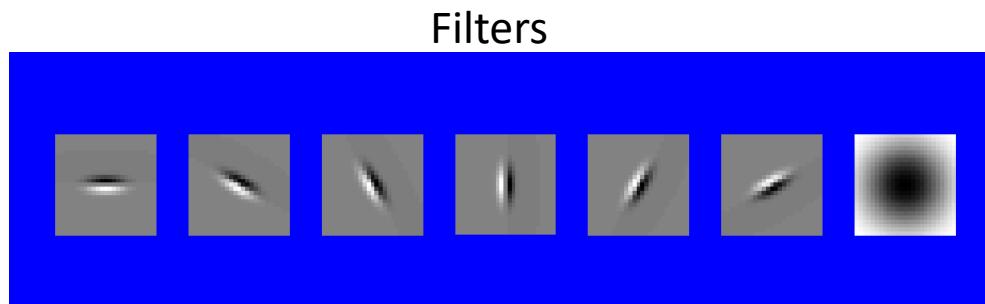
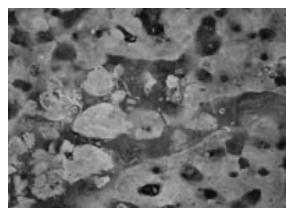
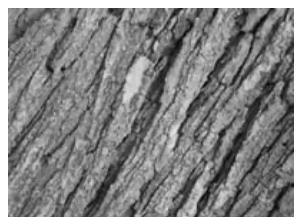


We can form a feature vector from the list of responses at each pixel.

You try: Can you match the texture to the response?



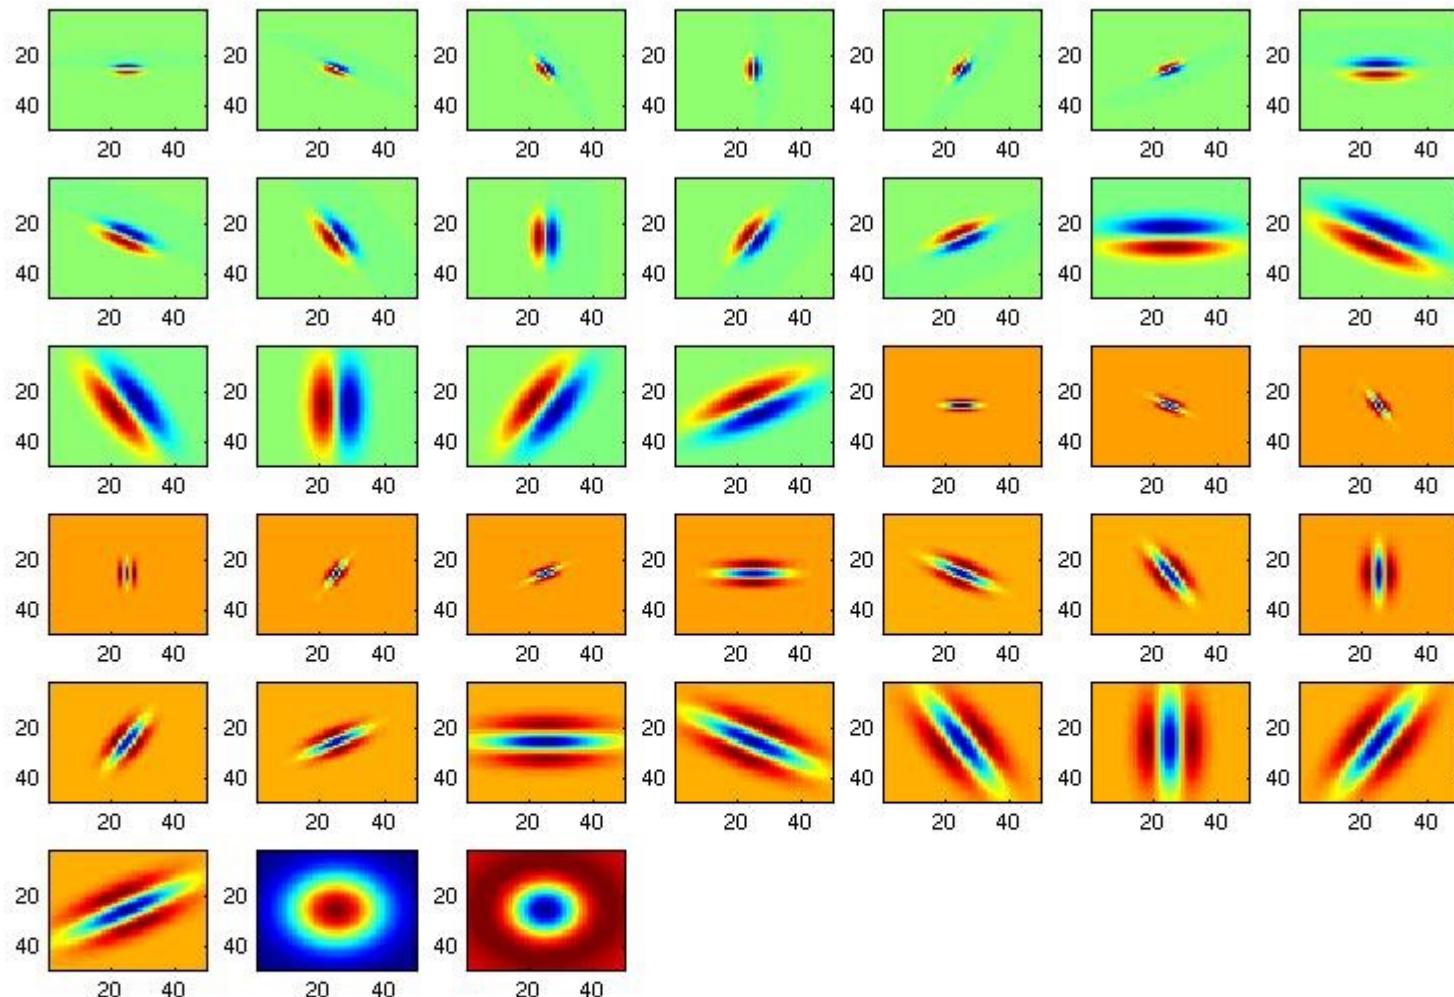
Representing texture by mean abs response



Mean abs responses

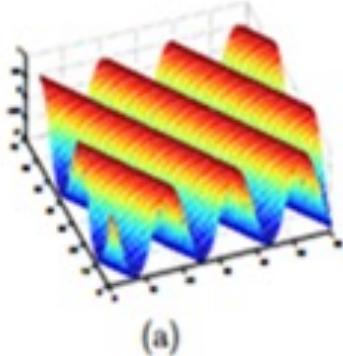
Slide credit Derek Hoiem

Filter bank

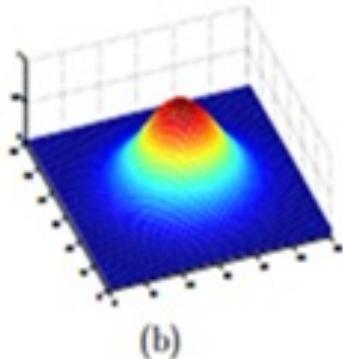


Gabor Filters

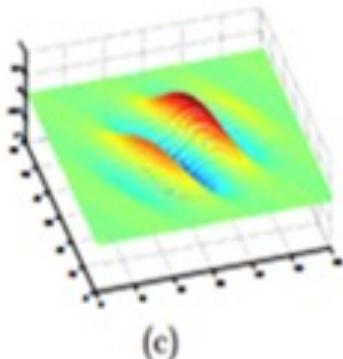
A Gabor filter is a combination of a Gaussian filter and a sinusoidal term.



A Sinusoid oriented 30° with X-axis



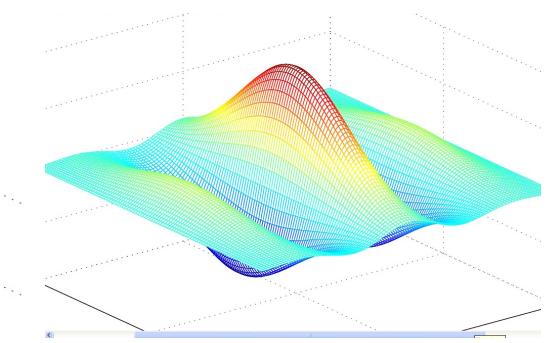
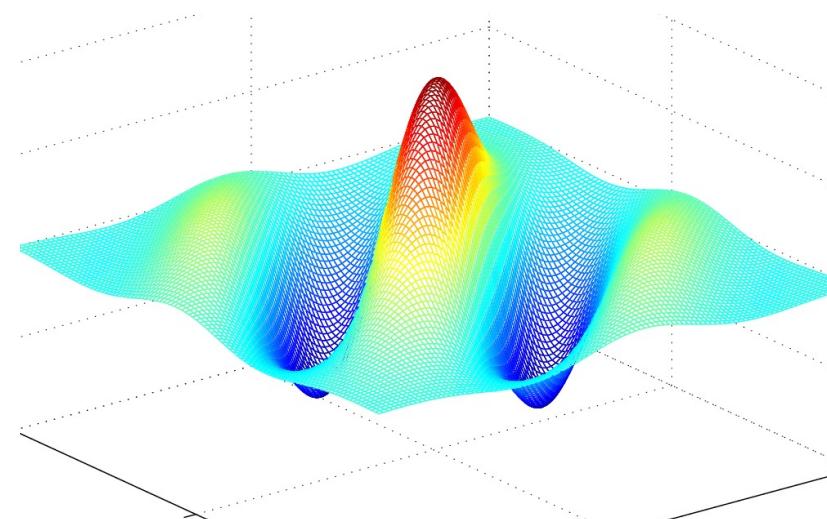
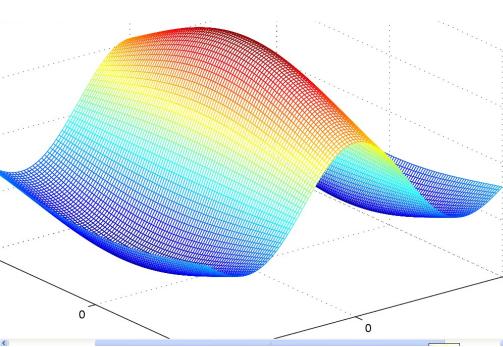
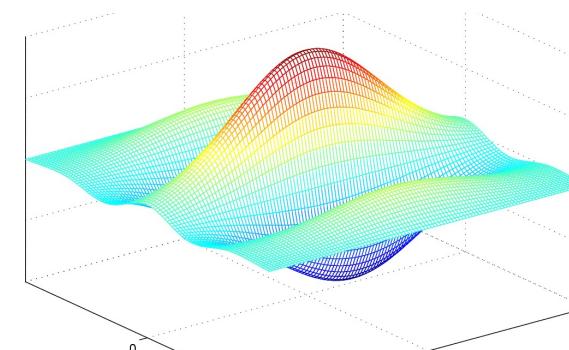
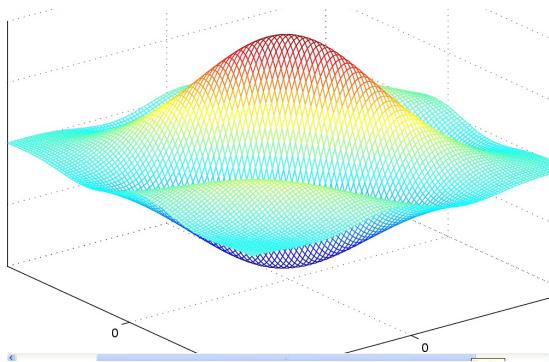
A 2-D Gaussian



The corresponding 2-D Gabor filter

Gabor Filters

- Wavelets at different frequencies and different orientations



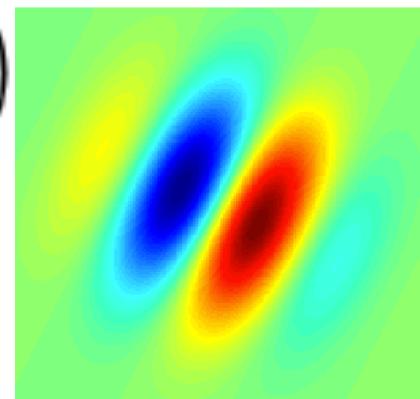
Gabor Filters

$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \cos\left(2\pi \frac{x'}{\lambda} + \psi\right)$$

where $x' = x \cos \theta + y \sin \theta$ and $y' = -x \sin \theta + y \cos \theta$.

The parameters λ, θ, ψ are parameters for the sinusoidal part (or factor).

λ and ψ are basic parameters of a sine wave.



λ - controls the Wavelength (<https://en.wikipedia.org/wiki/Wavelength>) of this sinusoid. A higher λ would "widen" the ellipses you see in the image above, while a lower λ would have the opposite.

θ - controls the rotation of the ellipses you see in the image above. These ellipses are the "parallel stripes". A value of $\theta = 0$ indicates no rotation and would make the ellipses vertical.

ψ - is the phase shift of the sinusoid i.e., how much the ellipses need to be shifted with respect to the center. It's a parameter that has to be specified.

σ - is the standard deviation (i.e., the spread) of the Gaussian part.

γ - controls the aspect ratio of the ellipses (roughly, how elongated the ellipses need to be).

Gabor filter values visualization by Juergen Mueller.

variable naming conventions as used by http://en.wikipedia.org/wiki/Gabor_filter.

Red: positive values. Blue: negative values

1. Wavelength is changed.

1. wavelength (lambda) : 5.00

2. orientation (theta) : 0.79

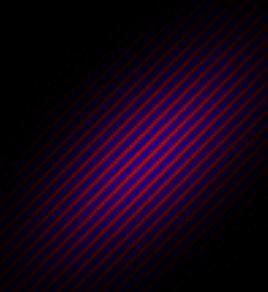
3. gaussvar (sigma) : 20.00

4. phaseoffset (phi) : 0.00

5. aspectratio (gamma) : 0.50

min gabor val : -0.91

max gabor val : 1.00



Color

Important cue for recognition



Important cue for segmentation

“GrabCut” — Interactive Foreground Extraction using Iterated Graph Cuts

Carsten Rother*

Vladimir Kolmogorov†
Microsoft Research Cambridge, UK

Andrew Blake‡

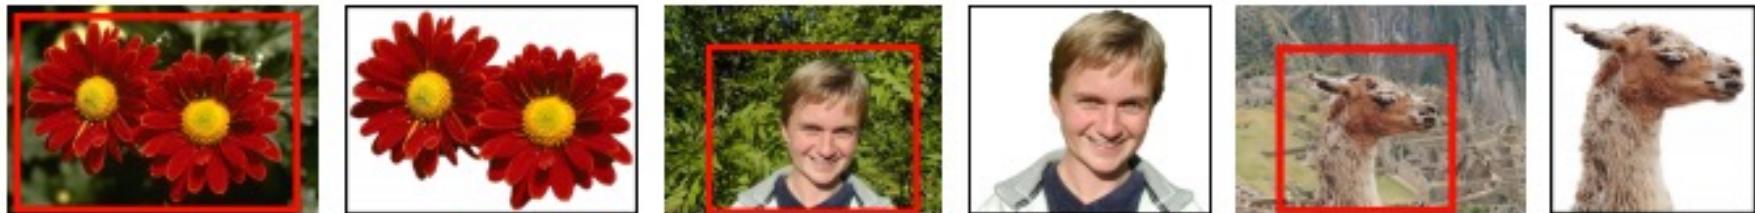


Figure 1: Three examples of GrabCut. The user drags a rectangle loosely around an object. The object is then extracted automatically.

Exercise lab 3

Grabcat - Foreground/background segmentation of cats in images

Input images



Ground-truth binary masks



Exercise lab 3

Grabcat - Foreground/background segmentation of cats in images



Exercise lab 3

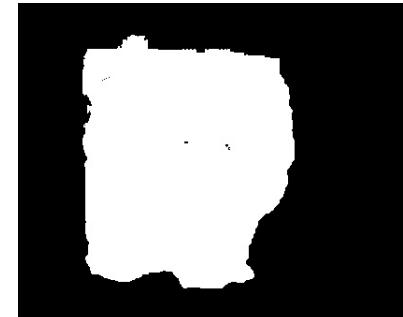
Grabcat - Foreground/background segmentation of cats in images



Input image



**Ground-truth
binary mask**



Grabcut



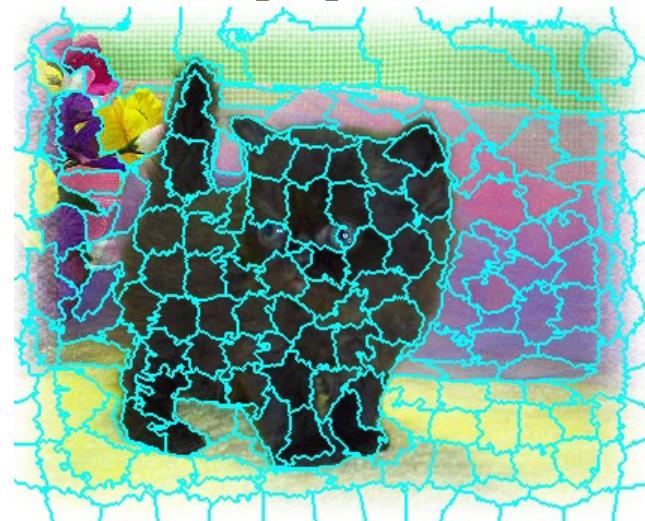
Grabcat

Exercise lab 3 - Grabcat

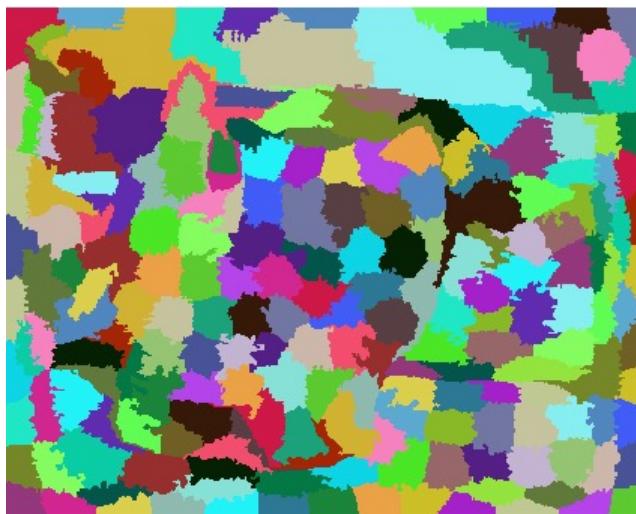
Input image



Superpixels



Superpixels color labeling



Best achievable mask using superpixels

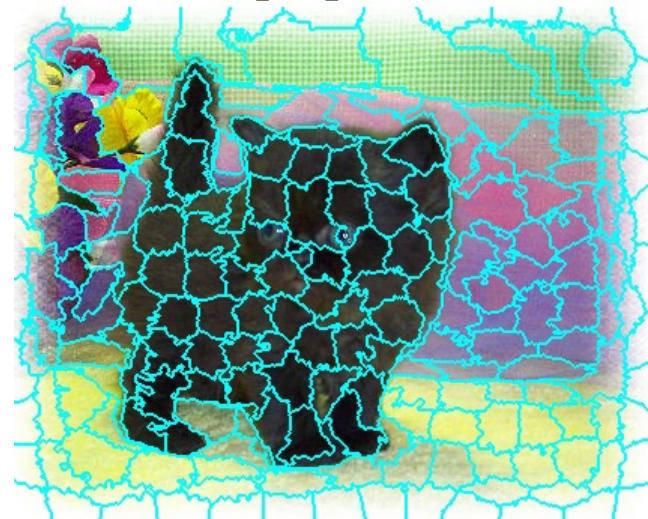


Exercise lab 3 - Grabcat

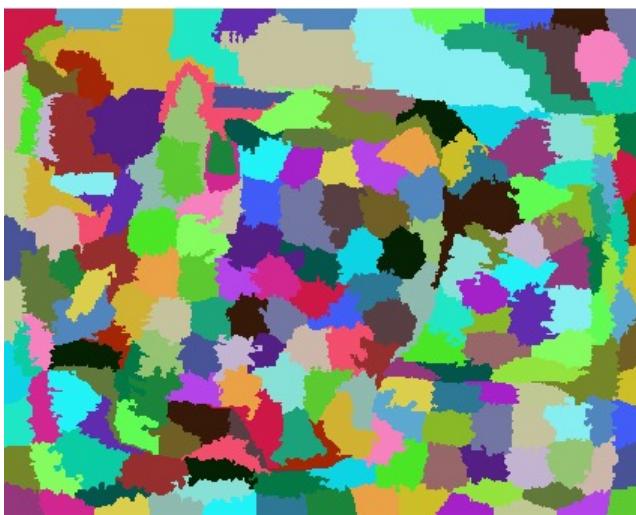
Input image



Superpixels



Superpixels color labeling



Result obtained using superpixels



SLIC Superpixels Compared to State-of-the-art Superpixel Methods

Radhakrishna Achanta, Appu Shaji, Kevin Smith,
Aurelien Lucchi, Pascal Fua, and Sabine Süsstrunk

Abstract—Computer vision applications have come to rely increasingly on superpixels in recent years, but it is not always clear what constitutes a good superpixel algorithm. In an effort to understand the benefits and drawbacks of existing methods, we empirically compare five state-of-the-art superpixel algorithms for their ability to adhere to image boundaries, speed, memory efficiency, and their impact on segmentation performance. We then introduce a new superpixel algorithm, *simple linear iterative clustering* (SLIC), which adapts a k -means clustering approach to efficiently generate superpixels. Despite its simplicity, SLIC adheres to boundaries as well as or better than previous methods. At the same time, it is faster and more memory efficient, improves segmentation performance, and is straightforward to extend to supervoxel generation.

regular lattice, such as [23], is probably a better choice. While it is difficult to define what constitutes an ideal approach for all applications, we believe the following properties are generally desirable:

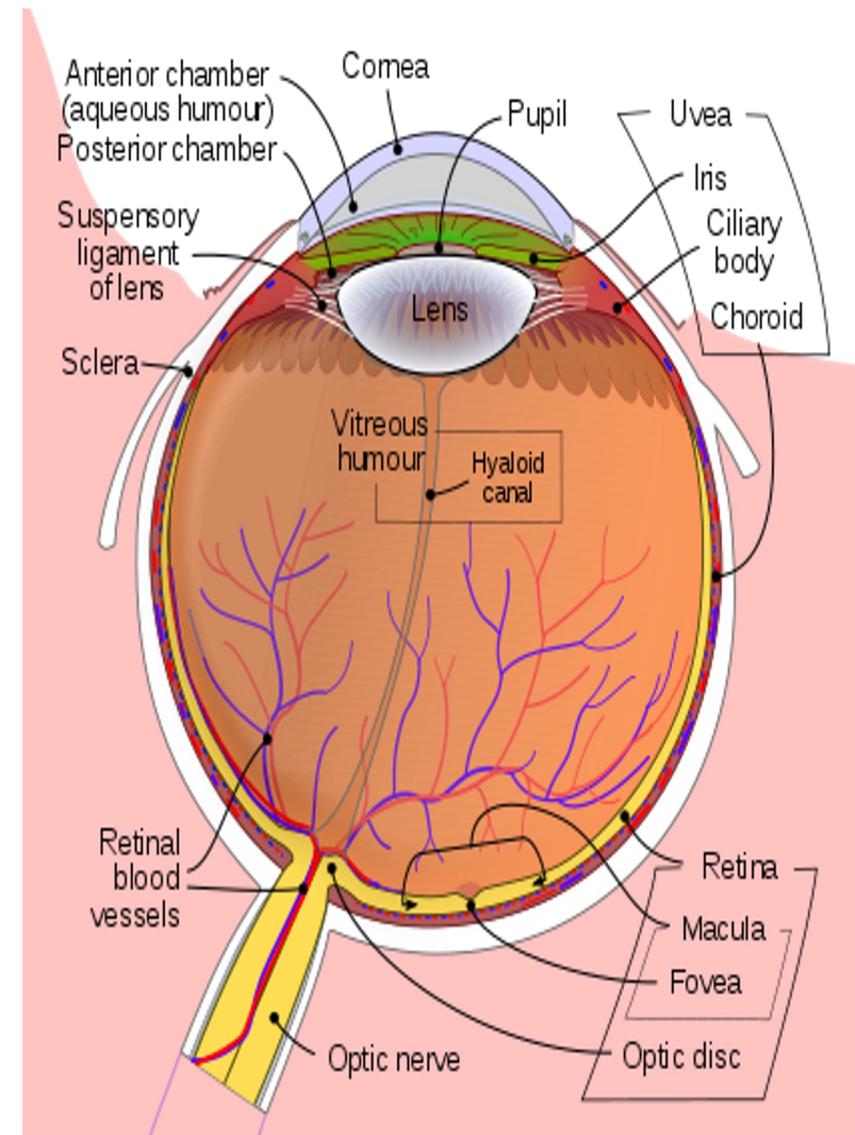
- 1) Superpixels should adhere well to image boundaries.
- 2) When used to reduce computational complexity as a pre-processing step, superpixels should be fast to compute, memory efficient, and simple to use.
- 3) When used for segmentation purposes, superpixels should both increase the speed and improve the quality of the results.



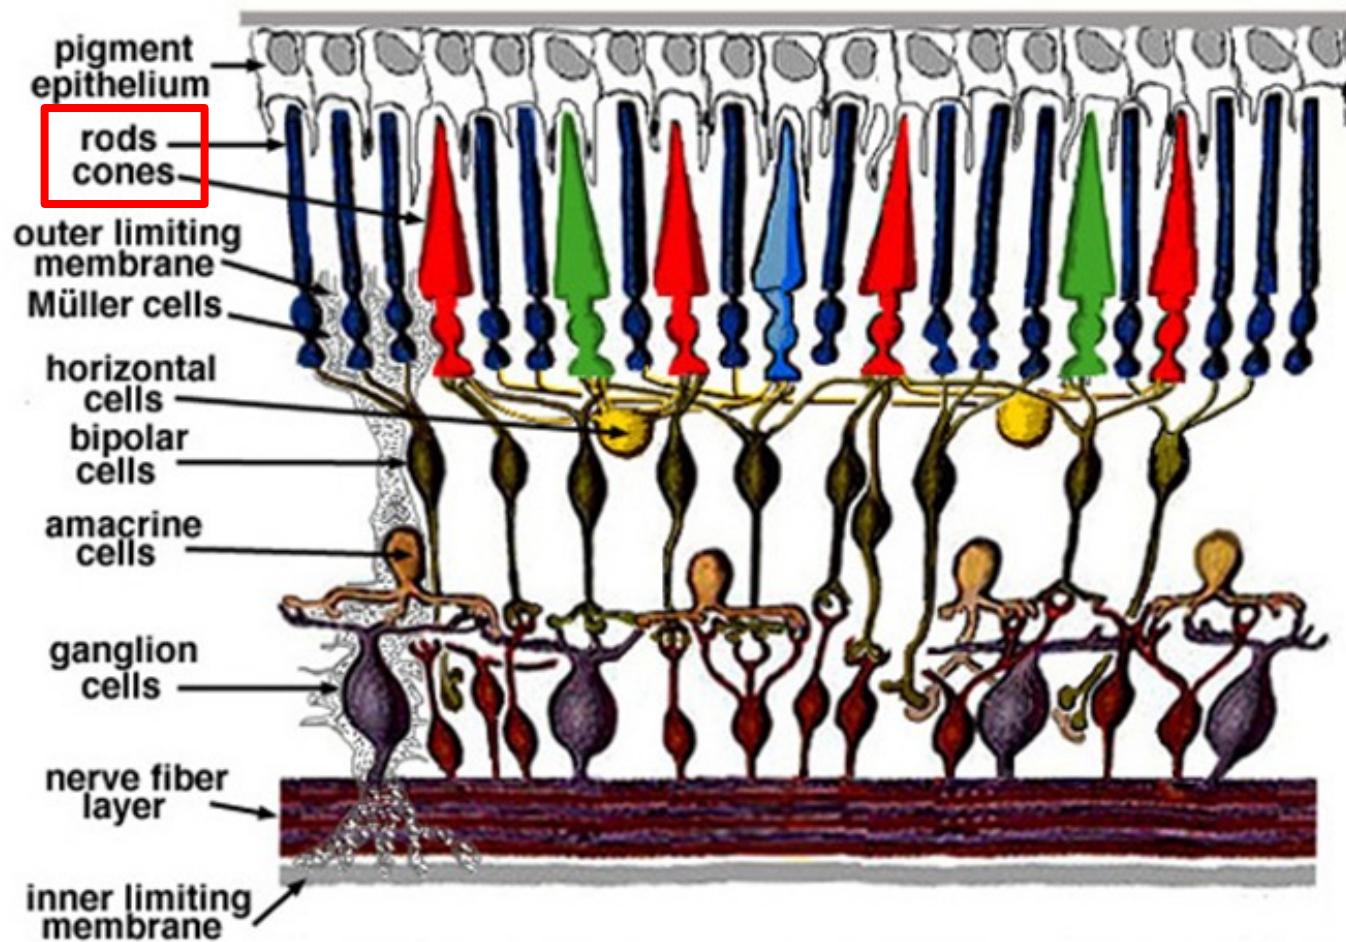
Fig. 1: Images segmented using SLIC into superpixels of size 64, 256, and 1024 pixels (approximately).

So how do human eyes work?

- Complex!
- Light passes through
 - Cornea, humours, lens refract light to focus
- Hit the retina
- Absorbed by photosensitive cells
- Info transmitted through optic nerve, processed by visual cortex

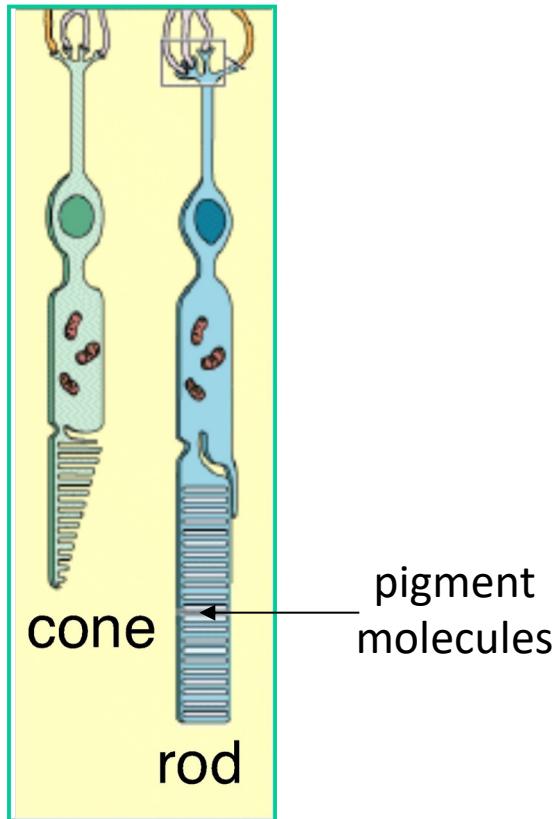


Retina



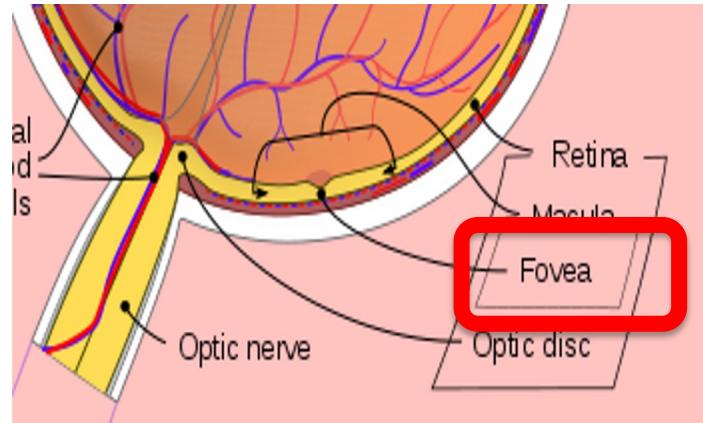
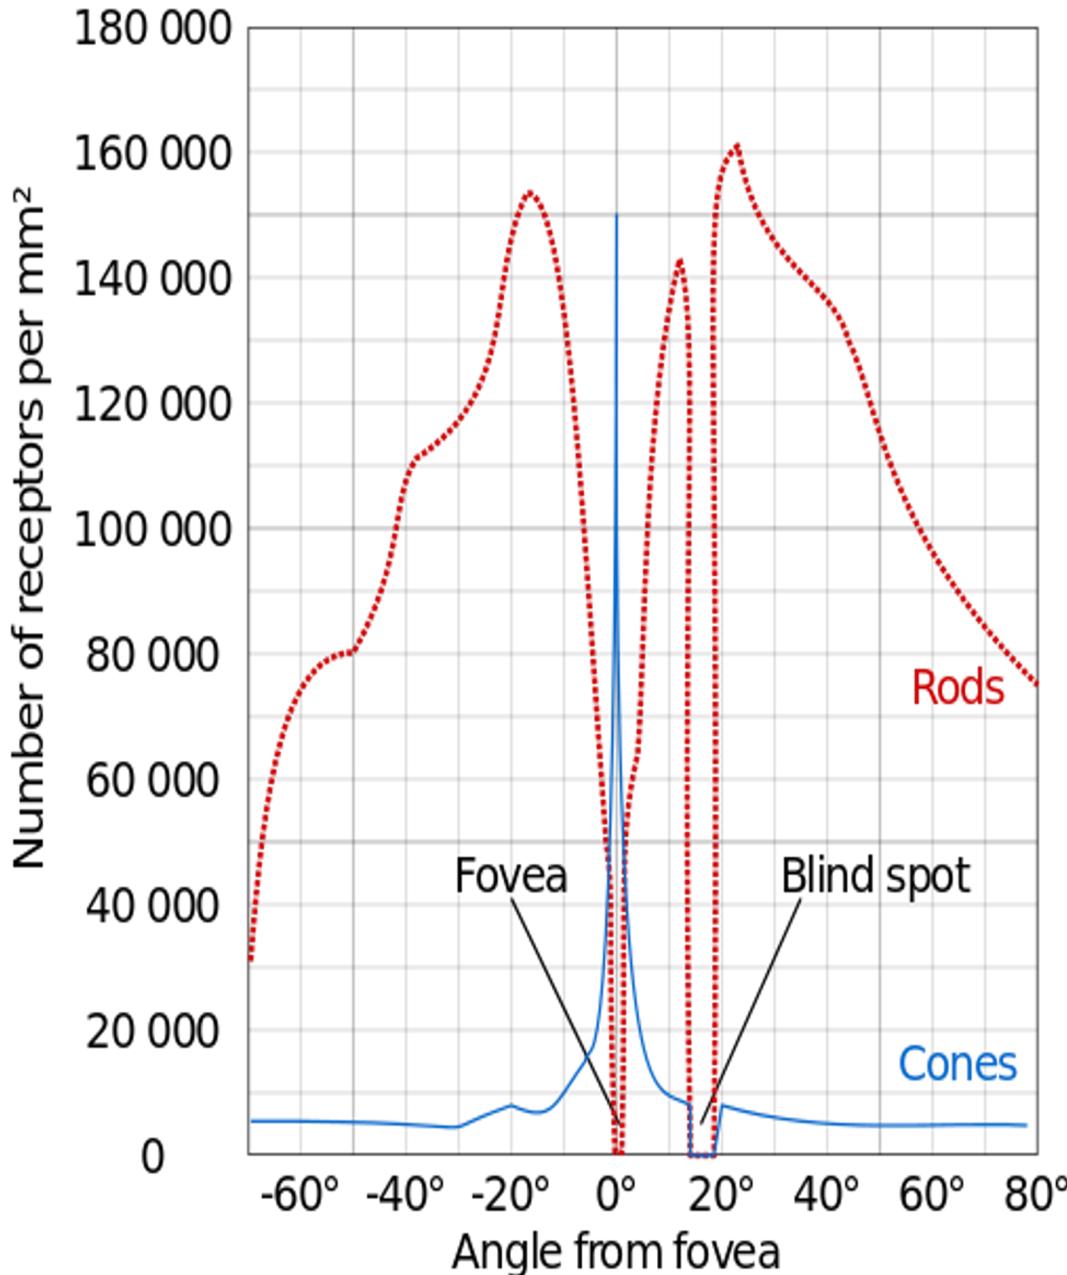
Simple diagram of the organization of the retina.

Photoreceptive cells - rods and cones



- Cones (RO: conuri)
 - around 6 million
 - need many photons to activate, bright light
 - fast response time
 - fine details
 - fast changes over time
 - responsible for most daytime vision (color vision)
 - mostly packed into one region: Fovea
- Rods (RO: bastonaşe)
 - ~120 million
 - sensitive to 1 photon
 - can pool responses
 - slow response time
 - only operate in low-light conditions
 - saturate quickly in lots of light
 - take ~7 minutes to adjust (night vision)

Fovea: where it's all happening



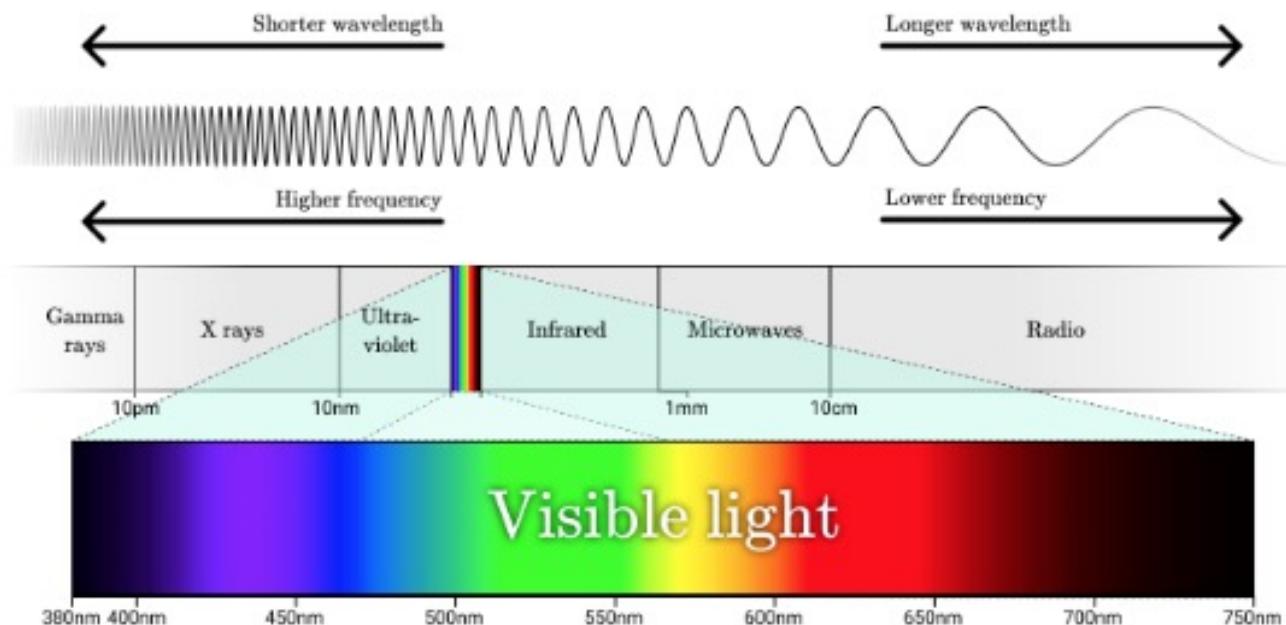
Fovea

- Small circle on the retina, 1.5mm^2
- Densely packed with cones
 - $200,000 \text{ cones/mm}^2$
- Highest visual acuity
- Reading: move your eyes so the text is centered in fovea
- Responsible for sharp central vision;

What is color?

Color = the result of interaction between physical light in the environment and our visual system

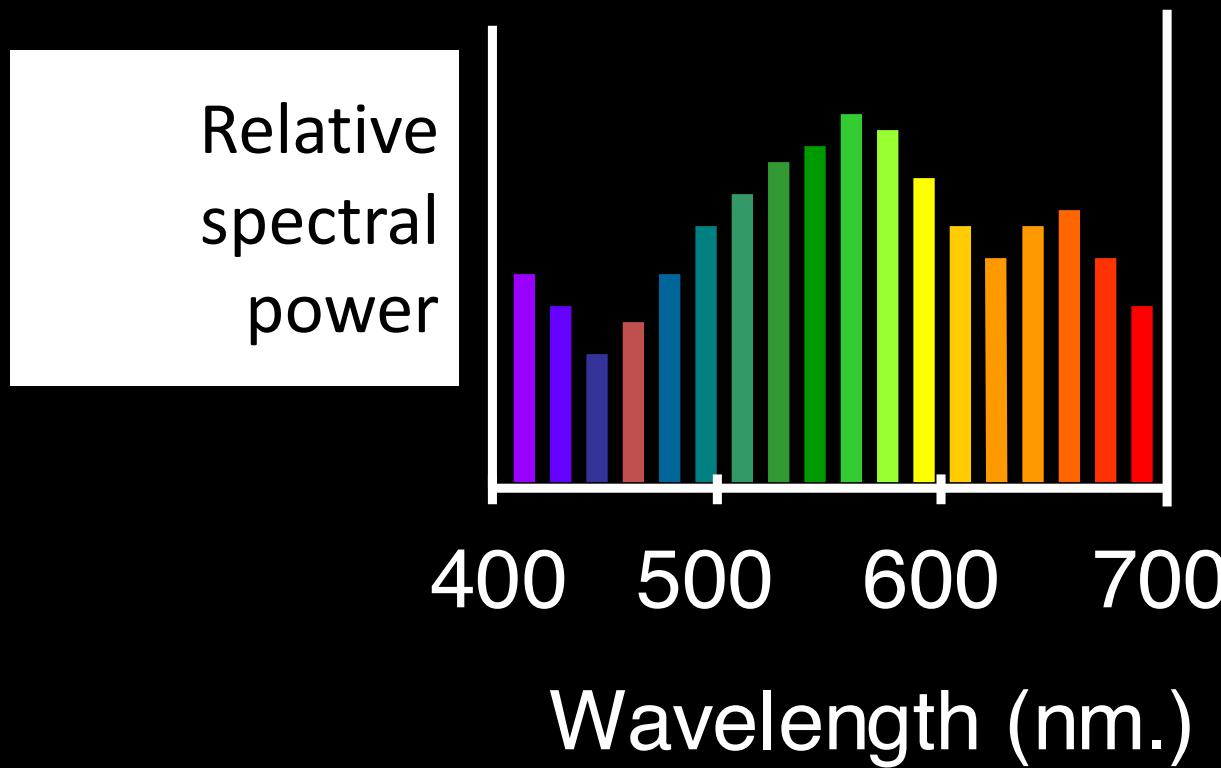
Color = electromagnetic wave detected by the eye and interpreted by the brain. The visible light is composed of a relatively narrow band of the electromagnetic spectrum. The visible spectrum is in the wavelength range: 380nm - 780nm.



The electromagnetic spectrum

The Physics of Light

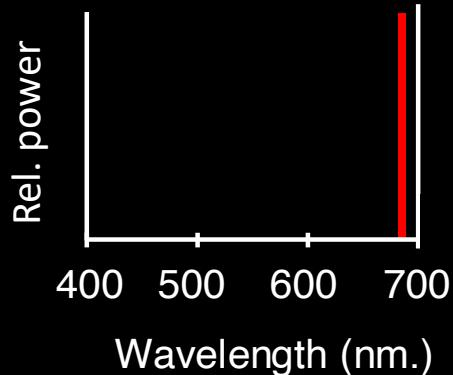
Any source of light can be completely described physically by its spectrum: the amount of energy emitted (per time unit) at each wavelength 400 - 700 nm.



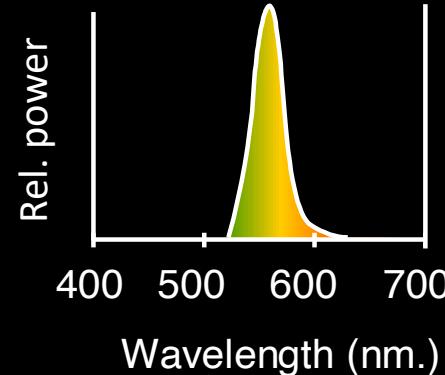
Spectra of Light Sources

Some examples of the spectra of light sources

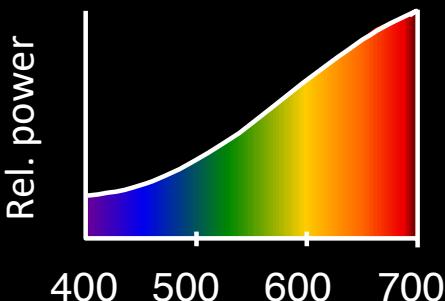
A. Ruby Laser



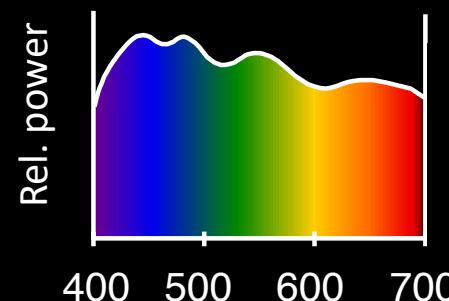
B. Gallium Phosphide Crystal



C. Tungsten Lightbulb

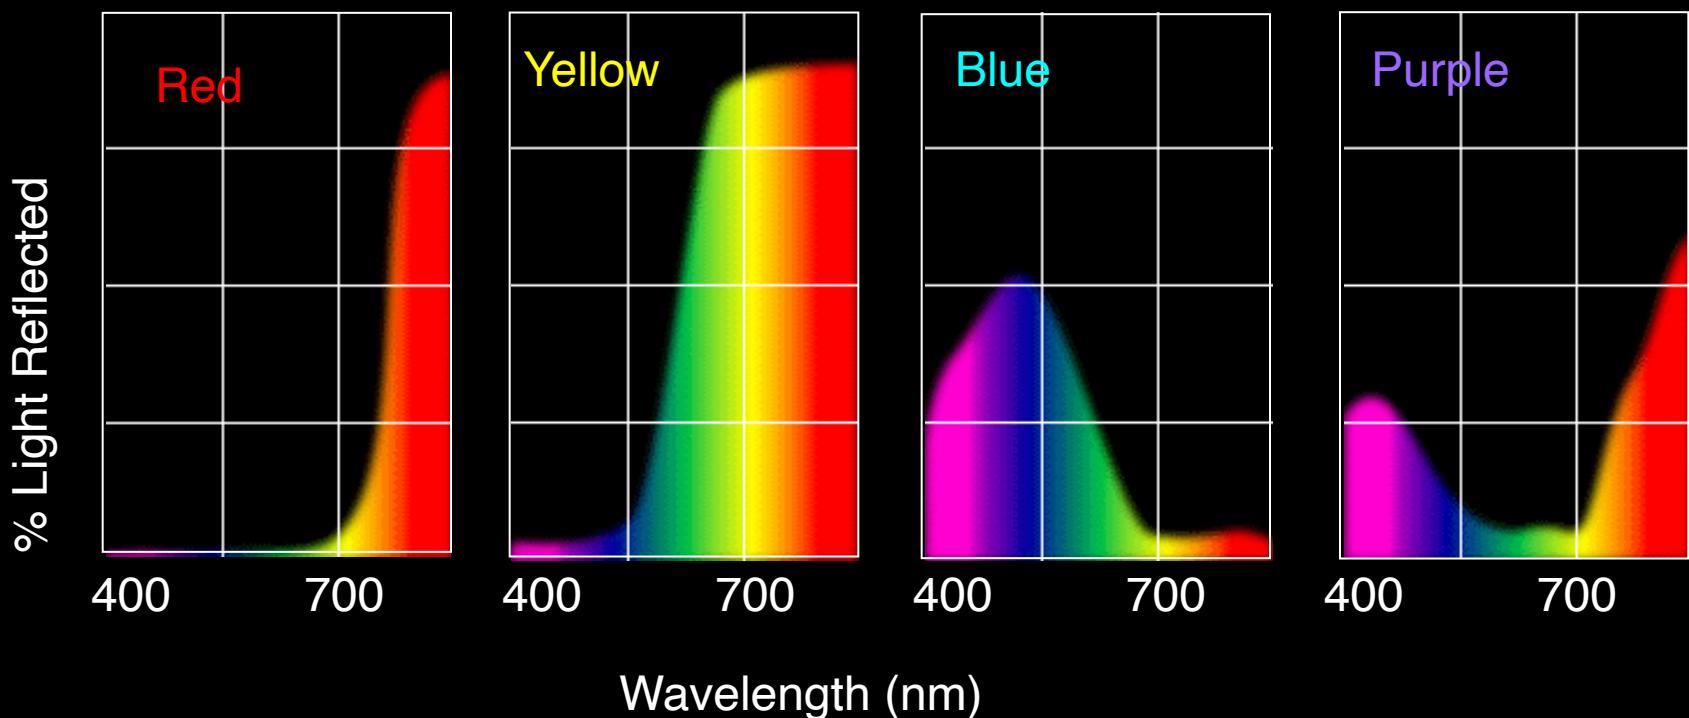


D. Normal Daylight



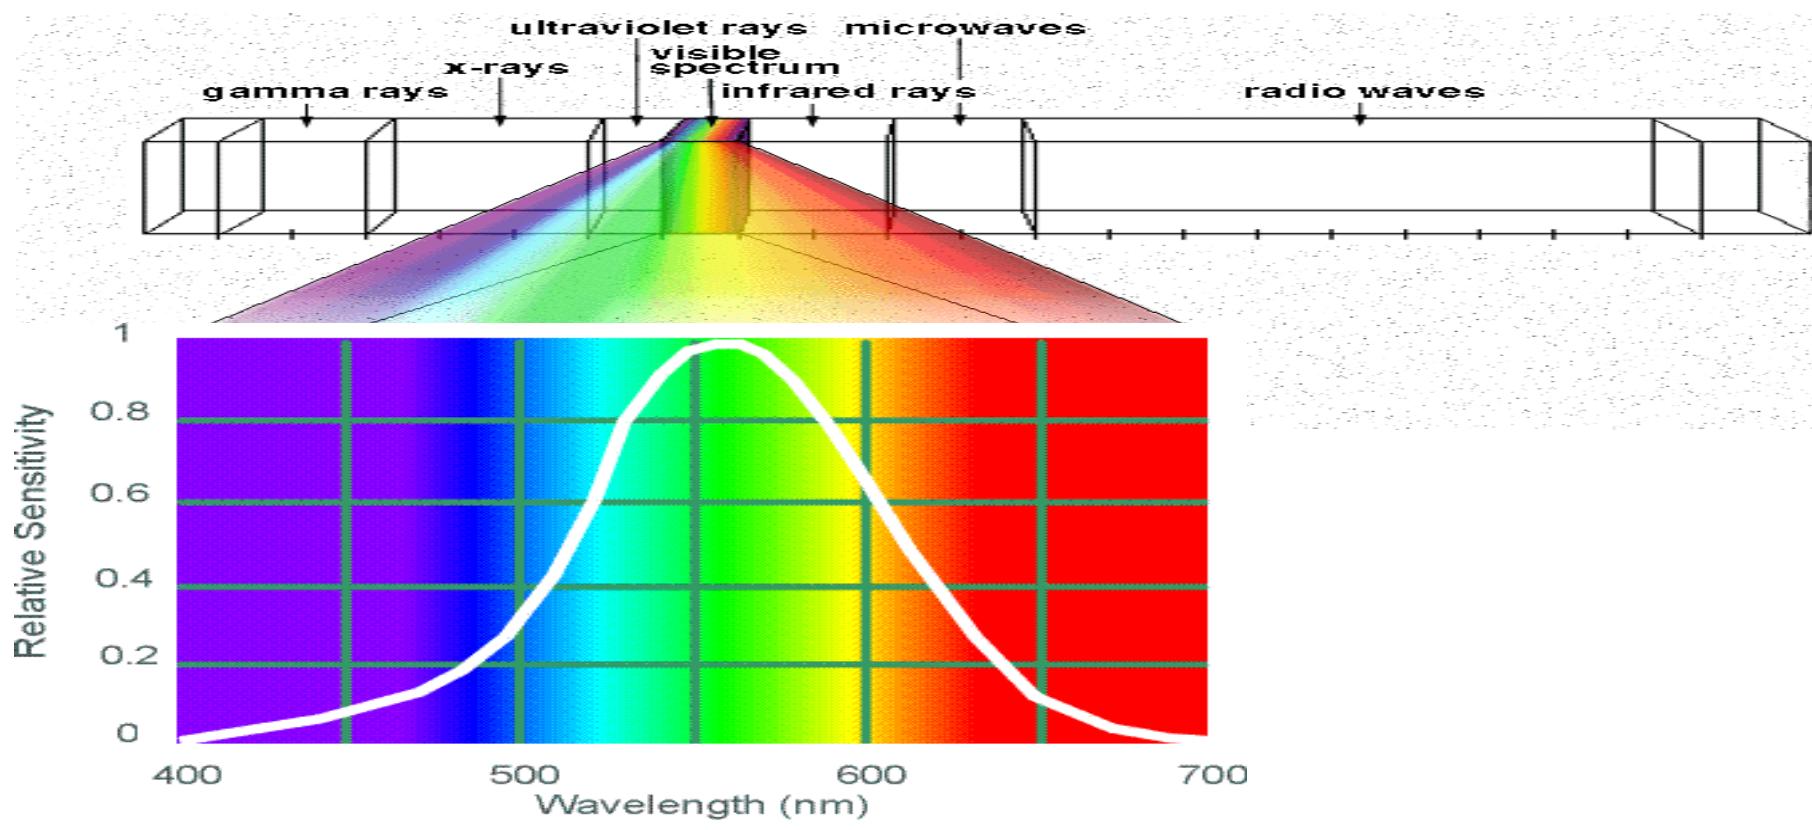
Reflectance Spectra of Surfaces

Some examples of the reflectance spectra of surfaces



Human Luminance Sensitivity Function

Light is composed of a spectrum of wavelengths. If we measure a human's eye sensitivity to every wavelength we get the standard human luminance sensitivity function



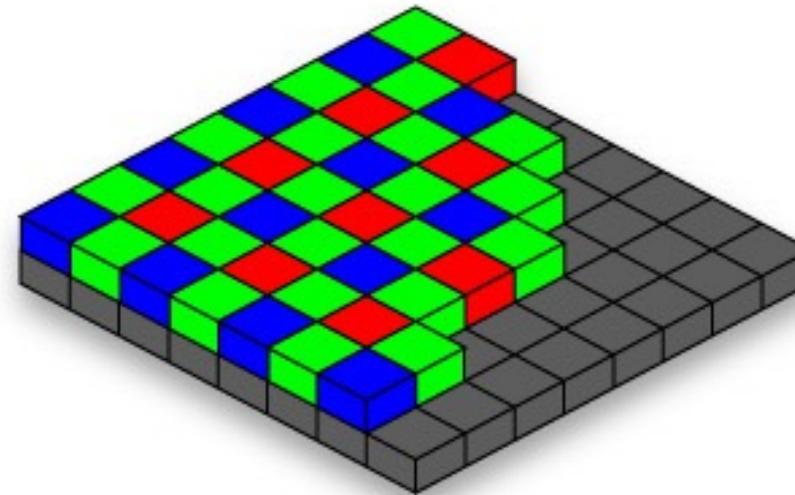
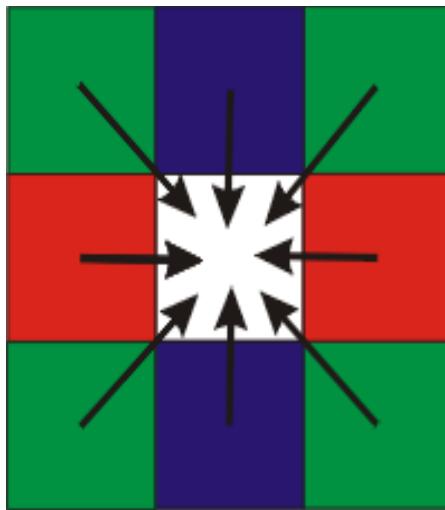
Human Luminance Sensitivity Function

Slide Credit: A. Efros

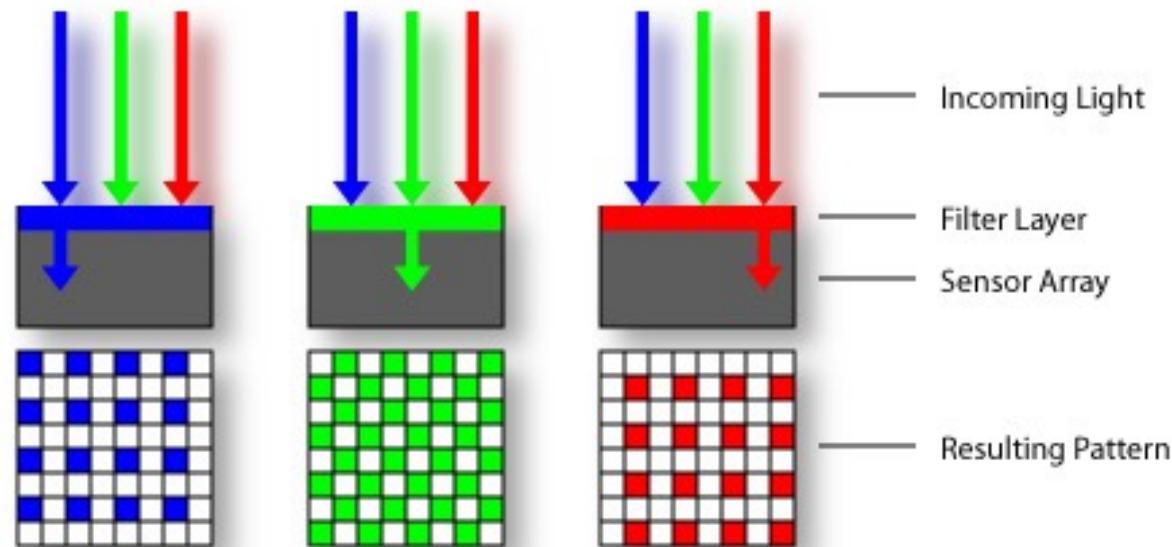
Why RGB?

If light is a spectrum, why are images RGB?

Color Sensing: Bayer Grid

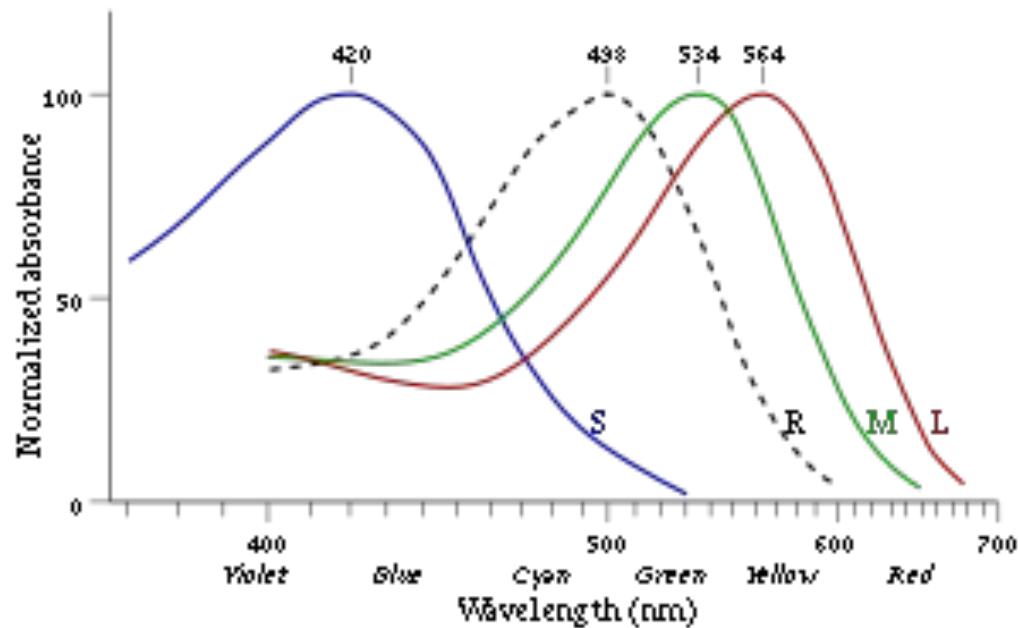


- Estimate RGB at each cell from neighboring values



Photoreceptors and light

- Each receptor has a responsiveness curve
- Receptors more responsive to some wavelengths, less responsive to others
- Rods: peak around 498 nm
- Cones: 3 kinds
 - Short: peak around 420 nm
 - Medium: peak around 530 nm
 - Long: peak around 560 nm



Cones and color

Our perception of color comes from cones

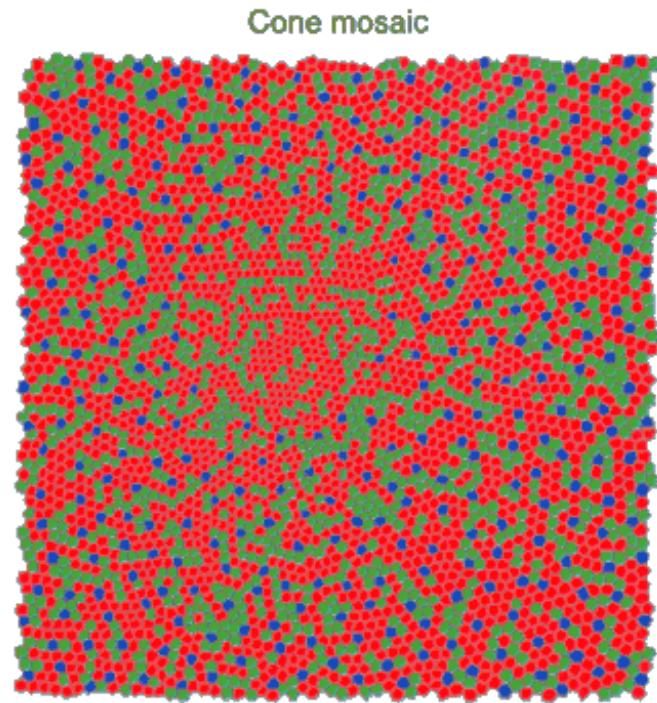
Different waveforms provoke different responses

Each cone has essentially one “output”

To calculate:

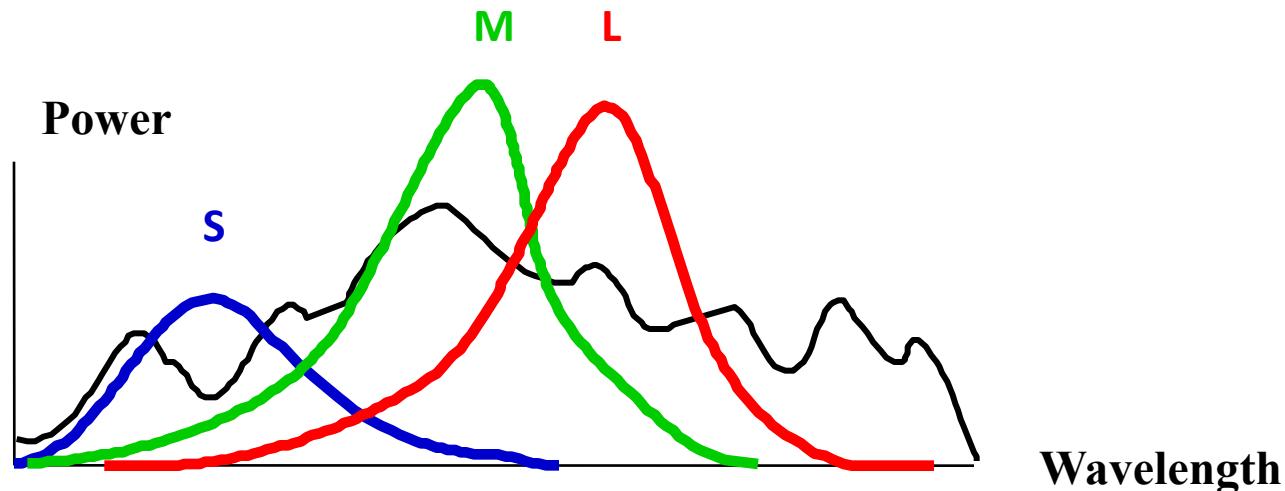
- Multiply input waveform by response curve
- Integrate area under the curve

The “color” we see is the relative activation of the 3 kinds of cones



- Ratio of Long to Medium to Short cones: approx. 10:5:1
- Almost no S cones in the center of the fovea

Color perception

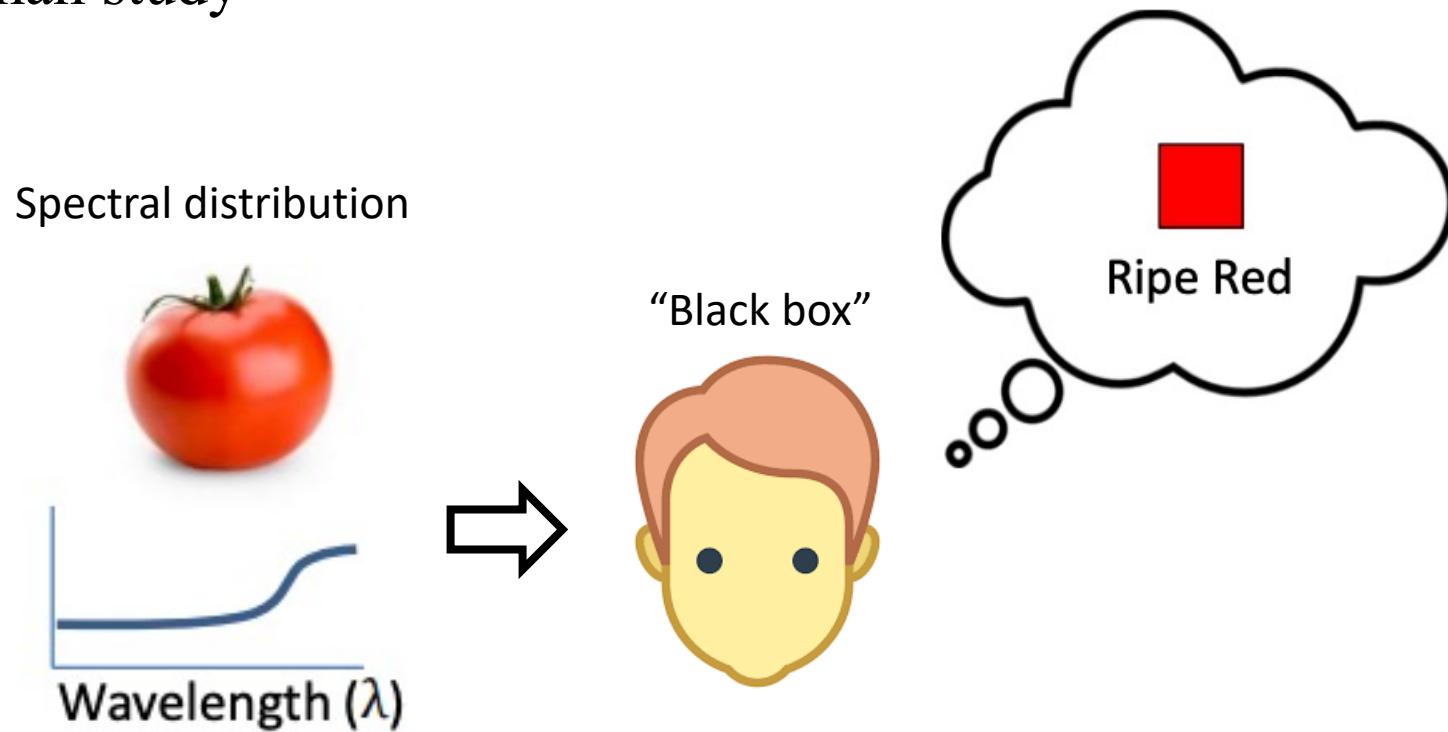


Rods and cones act as *filters* on the spectrum

- To get the output of a filter, multiply its response curve by the spectrum, integrate over all wavelengths
 - Each cone yields one number
 - How can we represent an entire spectrum with three numbers? We can't! Most of the information is lost
 - As a result, two different spectra may appear indistinguishable (such spectra are known as **metamers**)

Quantifying color

- Spectral distributions go through a “black box” (human visual system) and are perceived as color
- The only way to quantify the “black box” is to perform a human study

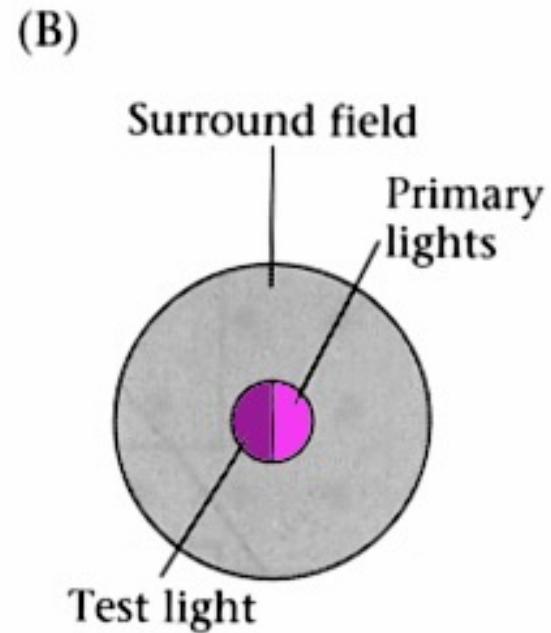
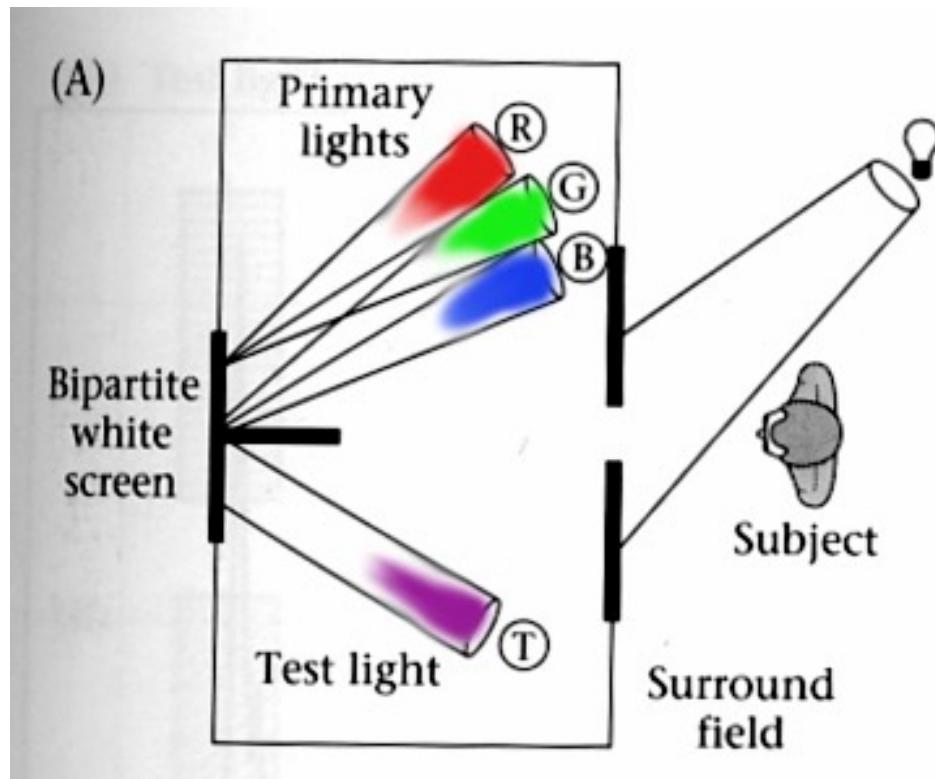


Color matching experiments

- We would like to understand which spectra produce the same color sensation in people under similar viewing conditions.
- Late 1920s experimented with colors! (and people)
 - Subjects get controls to 3 “primary” lights
 - Show them a light
 - Subject adjusts their lights to match the given light

Color matching experiments

- Late 1920s experimented with colors! (and people)
 - Subjects get controls to 3 “primary” lights
 - Show them a light
 - Subject adjusts their lights to match the given light



Trichromatic color theory

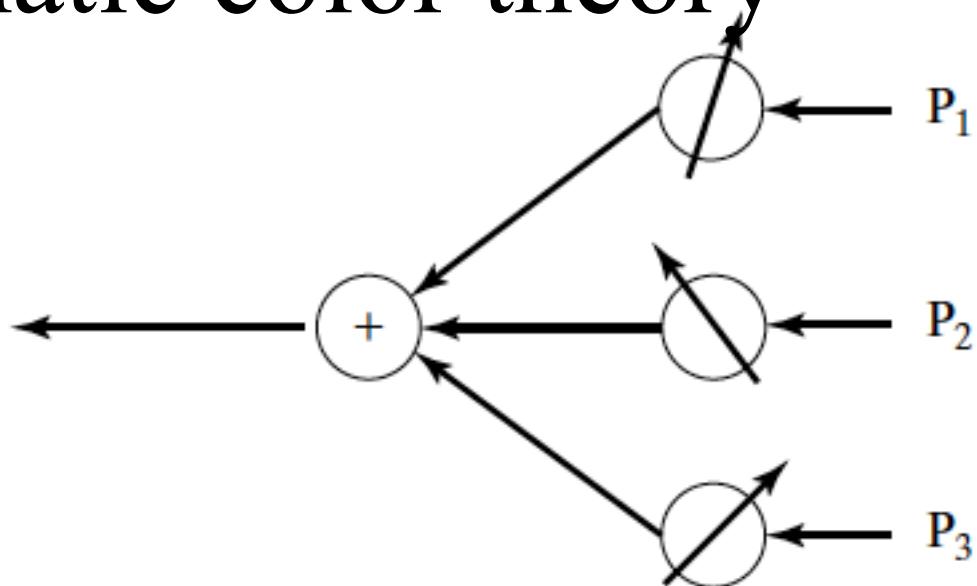
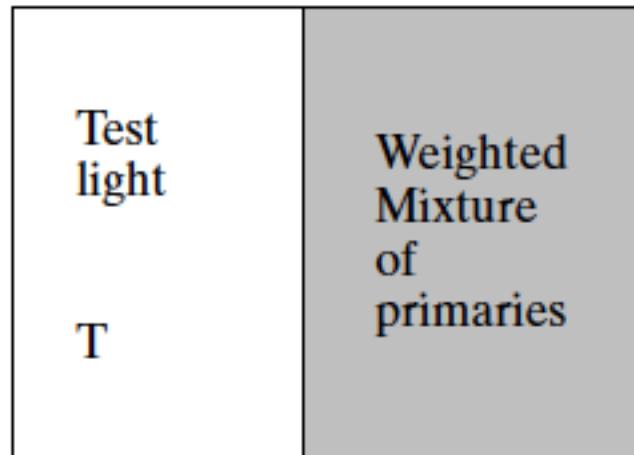
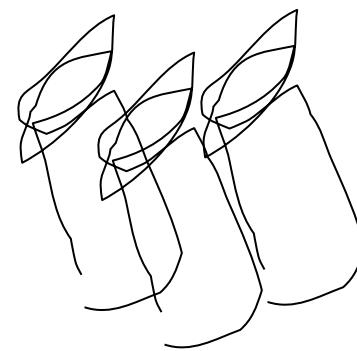
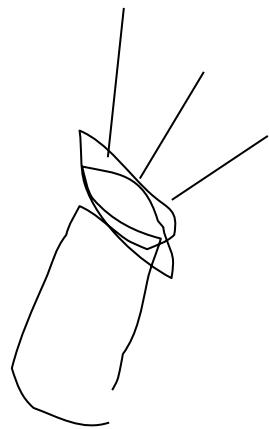
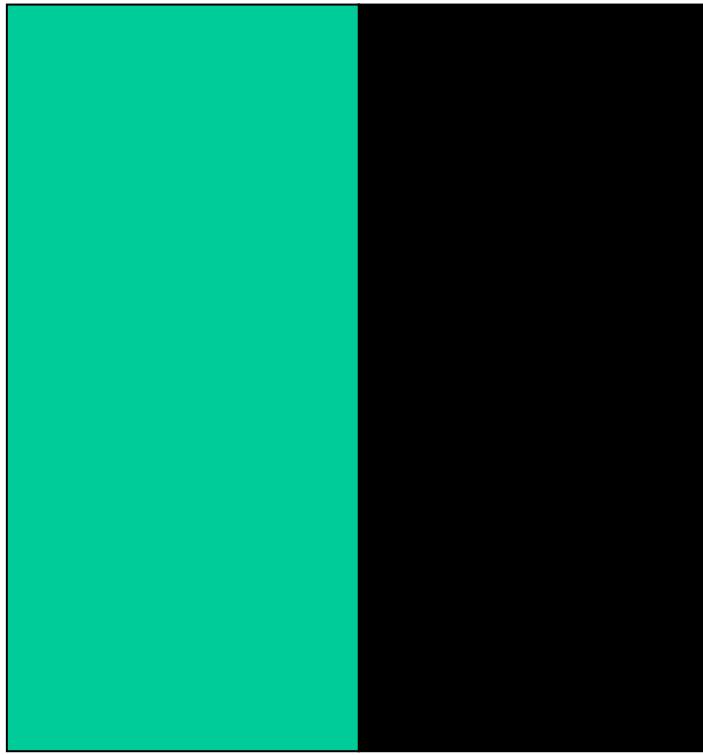
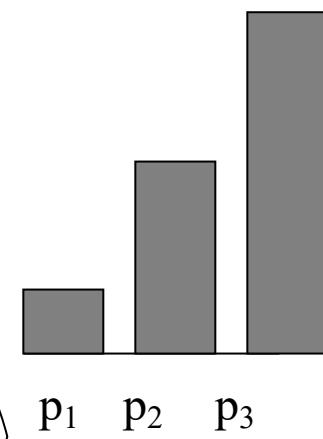
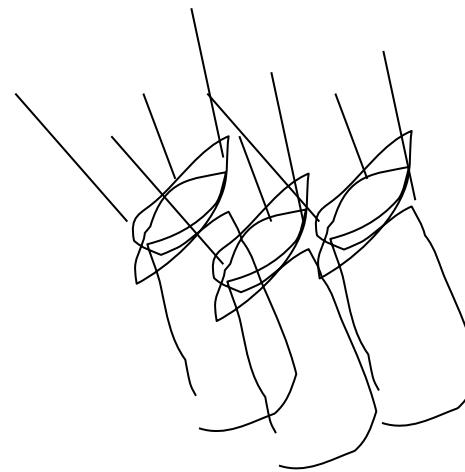
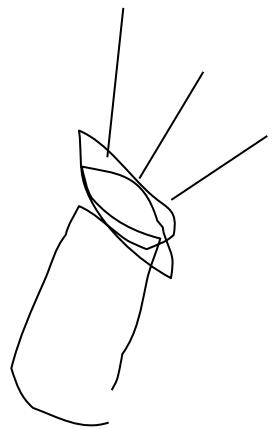
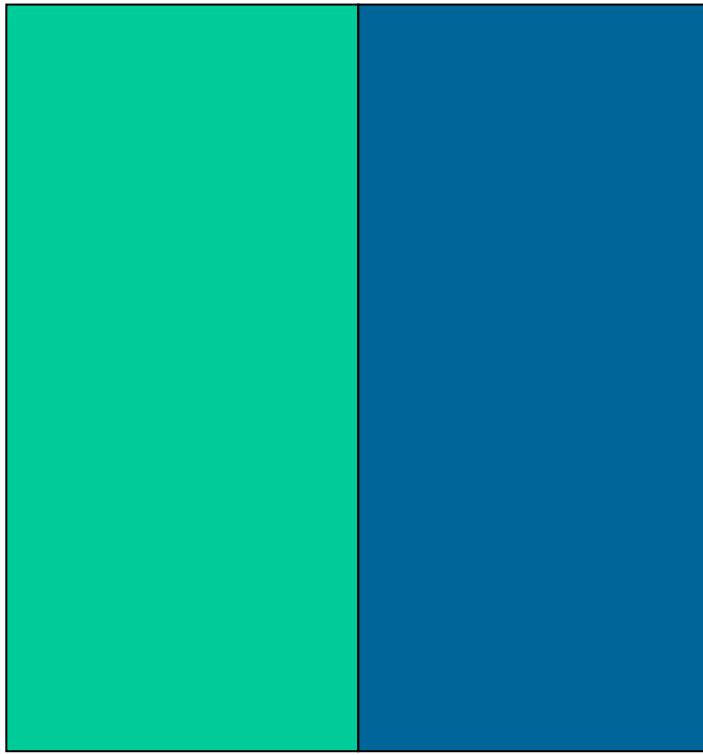


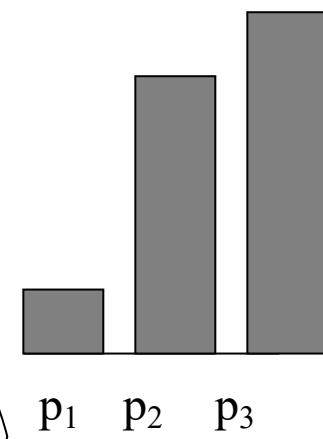
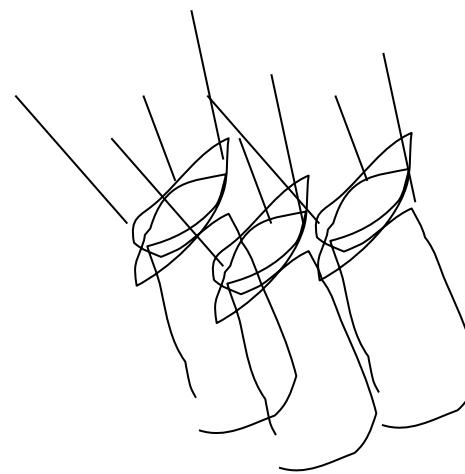
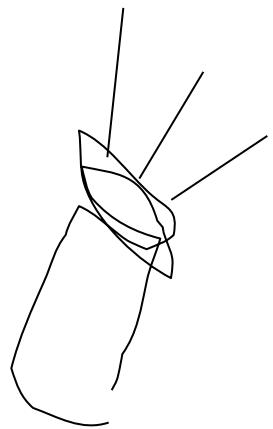
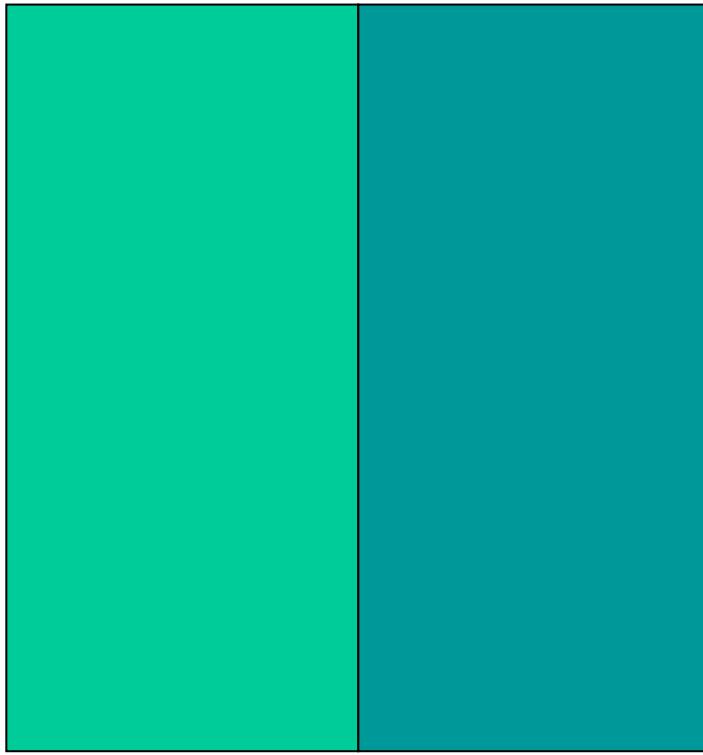
FIGURE 3.2: Human perception of color can be studied by asking observers to mix colored lights to match a test light shown in a split field. The drawing shows the outline of such an experiment. The observer sees a test light T and can adjust the amount of each of three primaries in a mixture displayed next to the test light. The observer is asked to adjust the amounts so that the mixture looks the same as the test light. The mixture of primaries can be written as $w_1 P_1 + w_2 P_2 + w_3 P_3$; if the mixture matches the test light, then we write $T = w_1 P_1 + w_2 P_2 + w_3 P_3$. It is a remarkable fact that for most people three primaries are sufficient to achieve a match for many colors, and three primaries are sufficient for all colors if we allow subtractive matching (i.e., some amount of some of the primaries is mixed with the test light to achieve a match). Some people require fewer primaries. Furthermore, most people choose the same mixture weights to match a given test light.

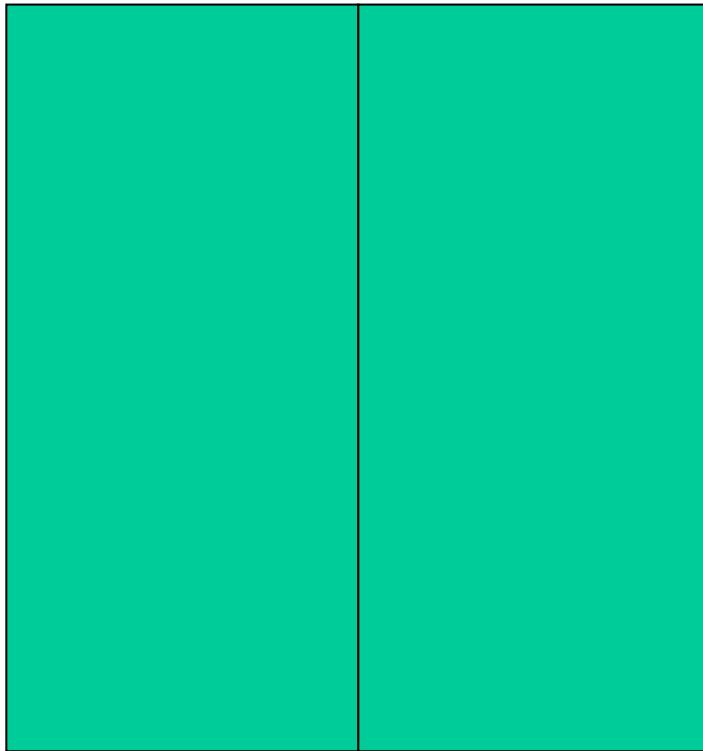
Trichromatic color theory

- In color matching experiments, most people can match any given light with three primaries
 - Primaries must be *independent*
- For the same light and same primaries, most people select the same weights
 - Exception: color blindness
- Trichromatic color theory
 - Three numbers seem to be sufficient for encoding color
 - Dates back to 18th century (Thomas Young)

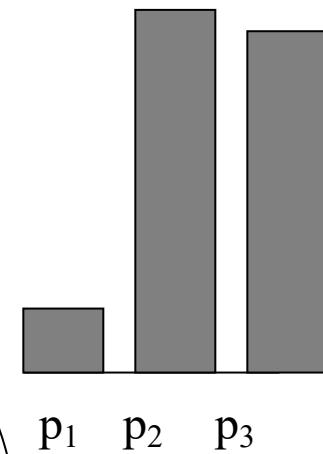
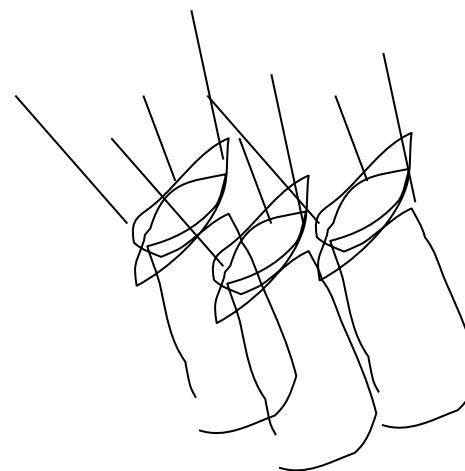
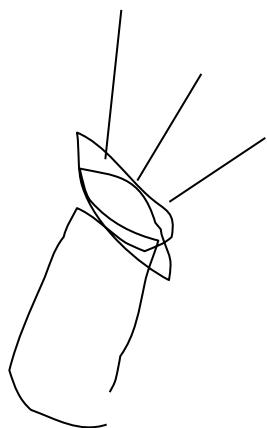


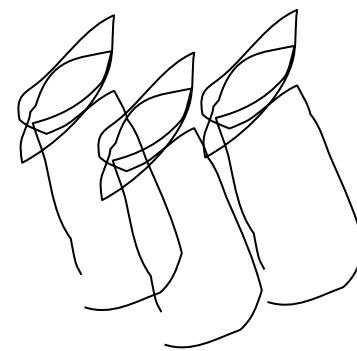
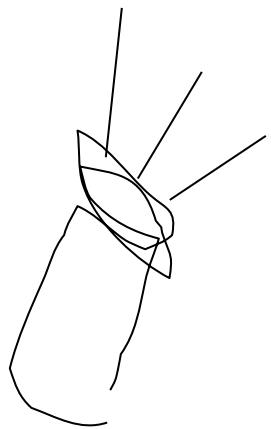
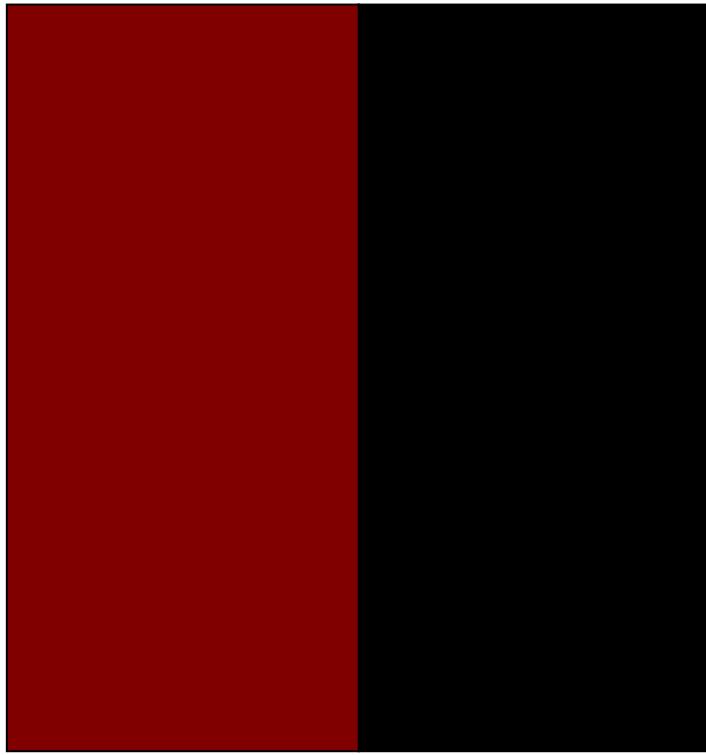


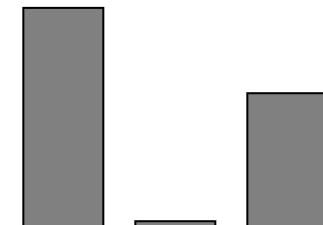
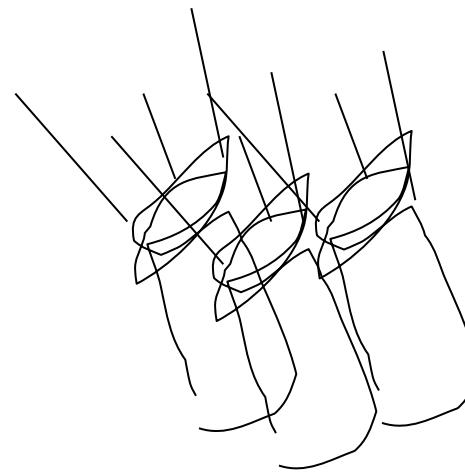
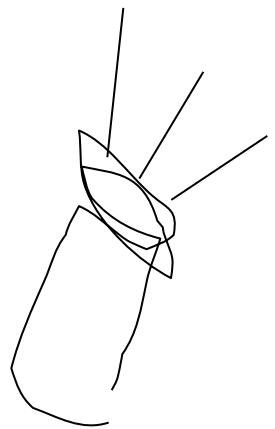
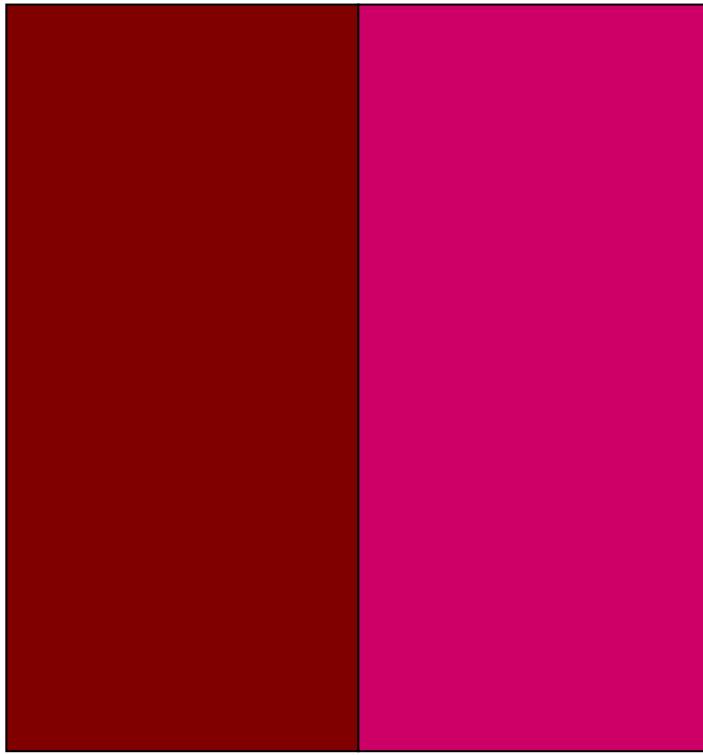




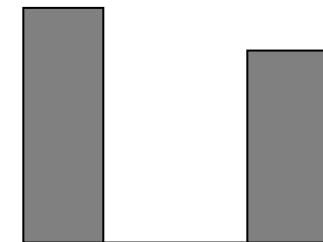
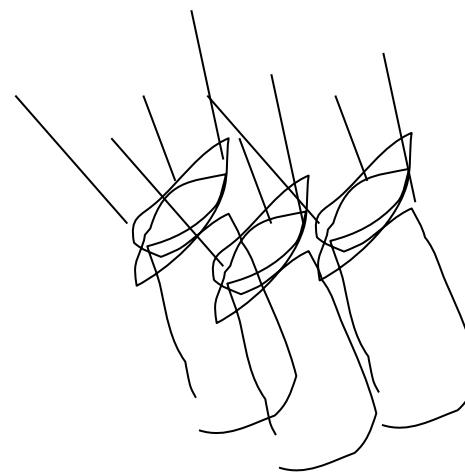
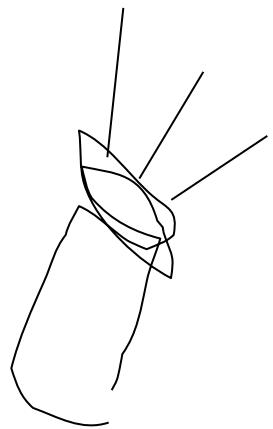
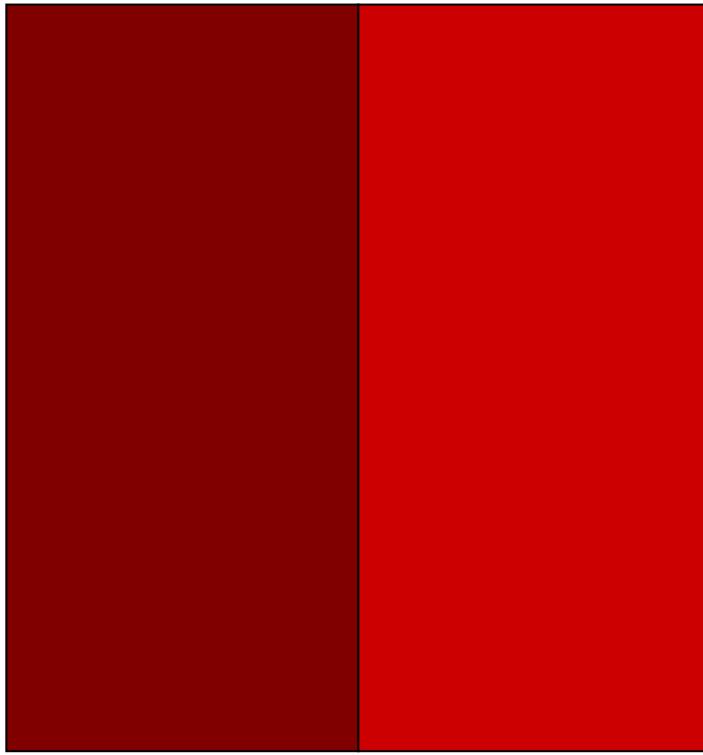
**The primary color
amounts needed
for a match**





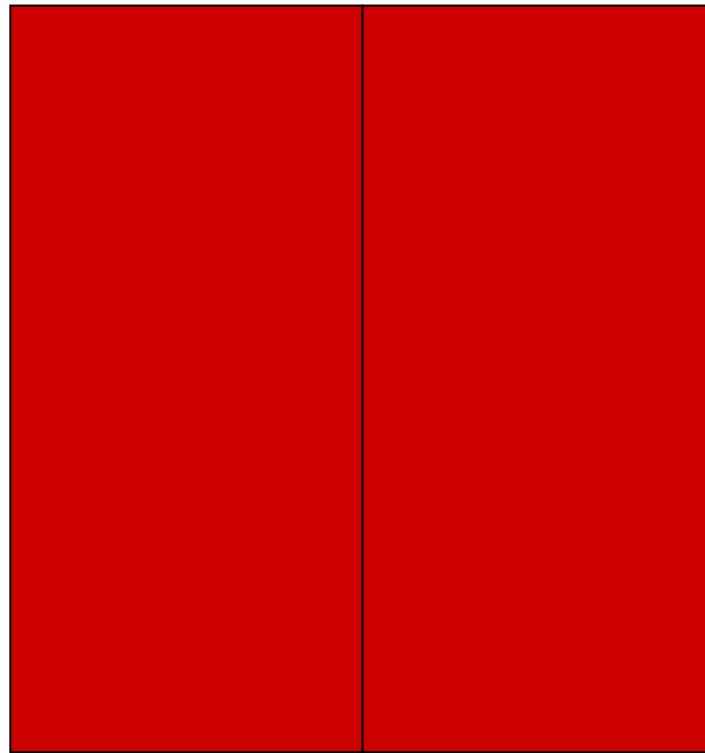


$p_1 \quad p_2 \quad p_3$

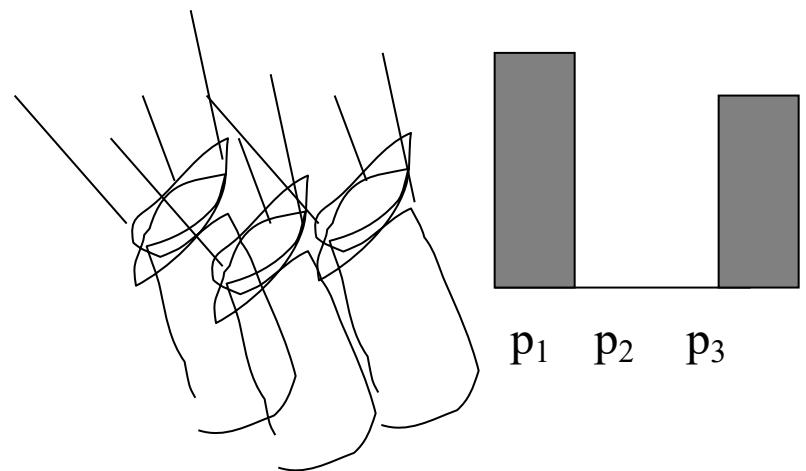
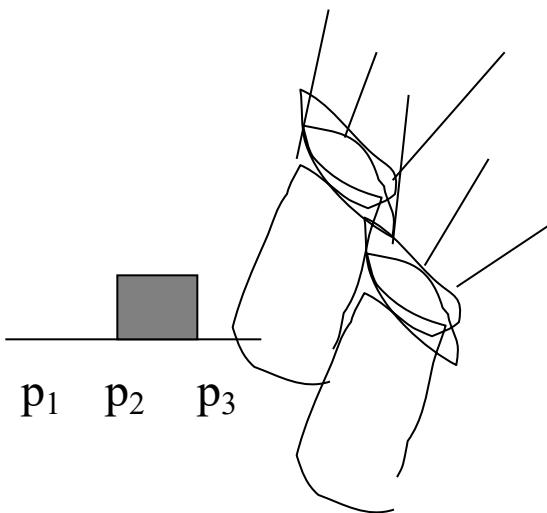
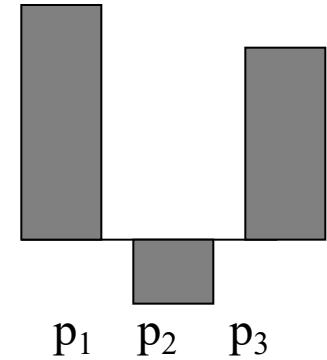


$p_1 \quad p_2 \quad p_3$

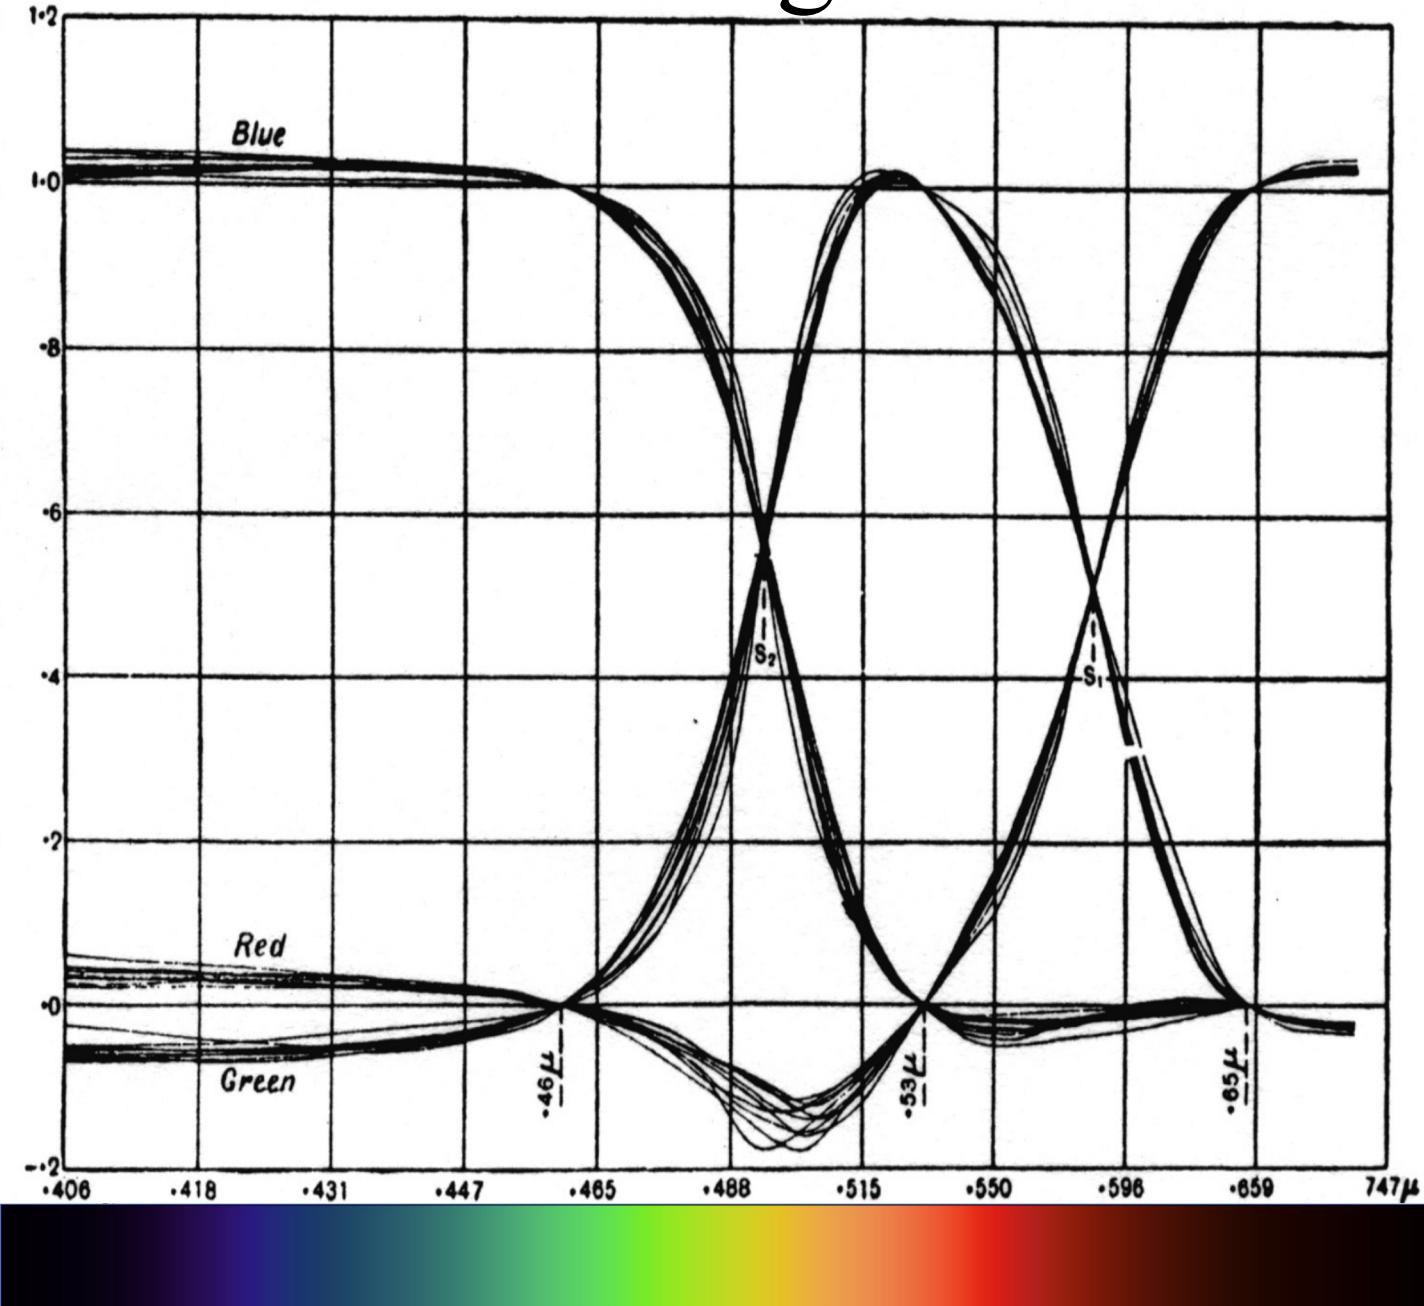
We say a
“negative” amount
of p_2 was needed to
make the match,
because we added
it to the test color’s
side.



The primary color amounts needed for a match:



RGB matching functions



RGB matching functions

