

Computer Vision

Bogdan Alexe

bogdan.alexe@fmi.unibuc.ro

University of Bucharest, 2nd semester, 2023-2024

Course structure

1. Features and filters: low-level vision

Linear filters, color, texture, edge detection, template matching

2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

4. Object Recognition: high – level vision

Object classification, object detection, part based models, bovw models

5. Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

Background subtraction

- given an image (video frame) we want to identify the *foreground objects* in that image.
- in most cases, objects are of interest and not the scene



Background subtraction

- simple techniques can do ok with static camera
- ...but hard to do perfectly
- widely used:
 - traffic monitoring (counting vehicles, detecting & tracking vehicles, pedestrians),
 - human action recognition (run, walk, jump, squat),
 - human-computer interaction
 - object tracking

Thresholding

- Given a grayscale image or an intermediate matrix → threshold to create a binary output.

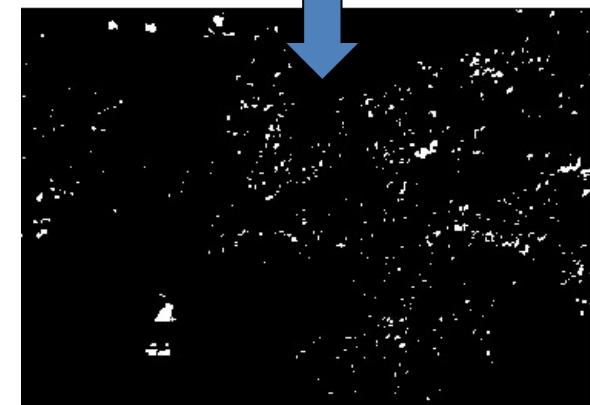
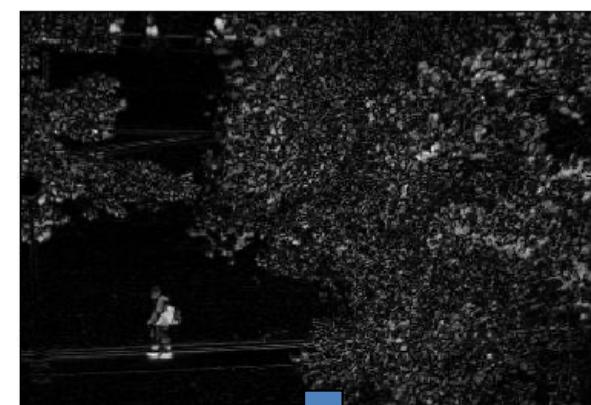
Example: background subtraction



-



=



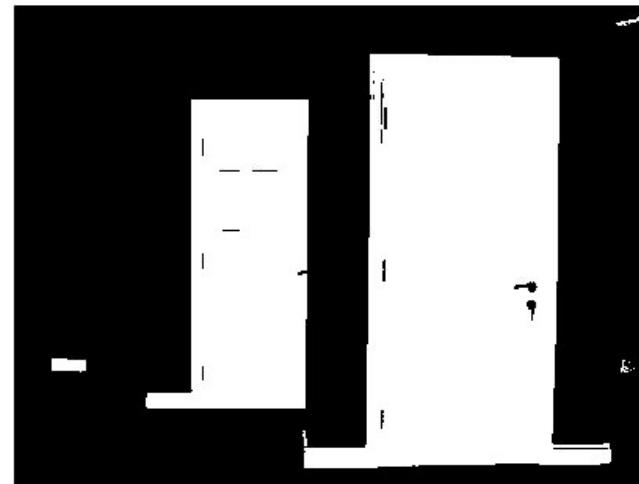
Looking for pixels that differ significantly from the “empty” background.

```
fg_pix = find(diff > t);
```

Thresholding

- Given a grayscale image or an intermediate matrix → threshold to create a binary output.

Example: intensity-based detection

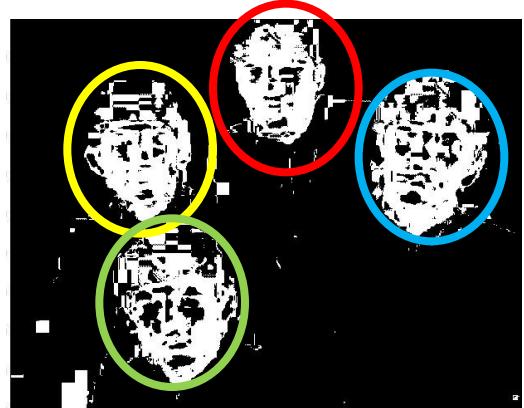


```
fg_pix = find(im < 65);
```

Looking for dark pixels

Issues

- What to do with “noisy” binary outputs?
 - Holes
 - Extra small fragments
- How to delimitate multiple regions of interest?
 - Count objects
 - Compute further features per object

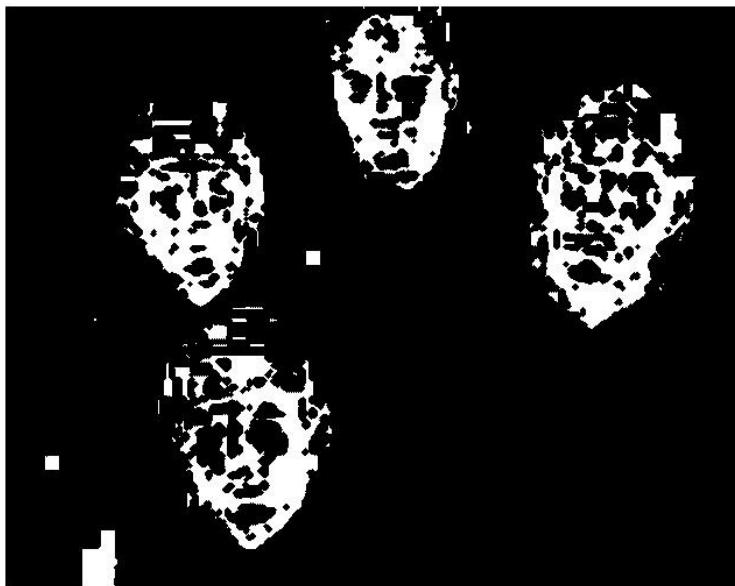


Morphological operators

- Change the shape of the foreground regions via intersection/union operations between a scanning structuring element and binary image.
- Useful to clean up result from thresholding
- Basic operators are:
 - Dilation
 - Erosion

Dilation

- Expands connected components
- Grow features
- Fill holes



Before dilation



After dilation

Erosion

- Erode connected components
- Shrink features
- Remove bridges, branches, noise



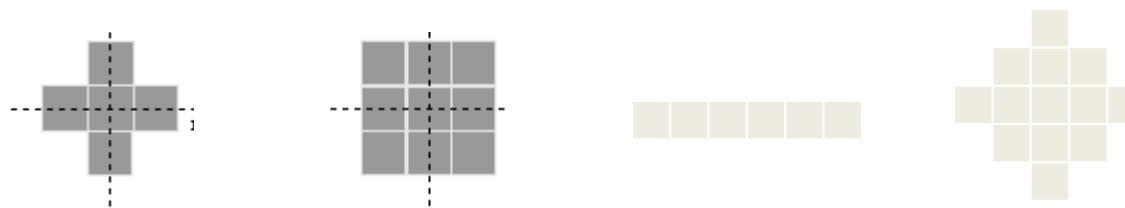
Before erosion



After erosion

Structuring elements

- **Masks** of varying shapes and sizes used to perform morphology, for example:



- Scan mask across foreground pixels to transform the binary image

Dilation vs. Erosion

At each position:

- **Dilation:** if current pixel is foreground OR the center of structuring element covers a foreground pixel in the input image.

Example for Dilation (1D)

Input

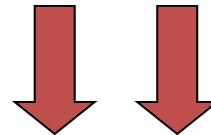
| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|

$$O(x) = I(x) \oplus ES$$



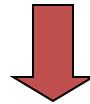
Output

| | | | | | | | | | |
|---|---|--|--|--|--|--|--|--|--|
| 1 | 1 | | | | | | | | |
|---|---|--|--|--|--|--|--|--|--|

Example for Dilation (1D)

Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



**Structuring
Element**

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|

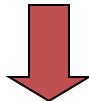
Output

| | | | | | | | | | |
|---|---|--|--|--|--|--|--|--|--|
| 1 | 1 | | | | | | | | |
|---|---|--|--|--|--|--|--|--|--|

Example for Dilation (1D)

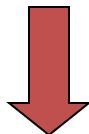
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



**Structuring
Element**

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|--|--|--|--|--|--|--|
| 1 | 1 | 0 | | | | | | | |
|---|---|---|--|--|--|--|--|--|--|

Example for Dilation (1D)

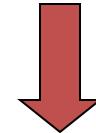
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



**Structuring
Element**

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|--|--|--|--|--|--|
| 1 | 1 | 0 | 0 | | | | | | |
|---|---|---|---|--|--|--|--|--|--|

Example for Dilation (1D)

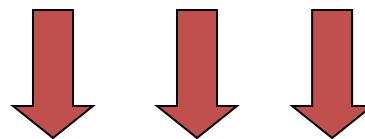
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|---|---|--|--|--|--|
| 1 | 1 | 0 | 1 | 1 | 1 | | | | |
|---|---|---|---|---|---|--|--|--|--|

Example for Dilation (1D)

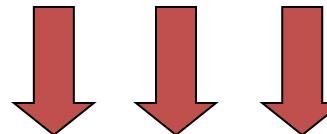
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|--|--|--|
| 1 | 1 | 0 | 1 | 1 | 1 | 1 | | | |
|---|---|---|---|---|---|---|--|--|--|

Example for Dilation (1D)

Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Structuring
Element

Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|--|--|
| 1 | 1 | 0 | 1 | 1 | 1 | 1 | 1 | | |
|---|---|---|---|---|---|---|---|--|--|

Example for Dilation (1D)

Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|

Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|--|--|
| 1 | 1 | 0 | 1 | 1 | 1 | 1 | 1 | | |
|---|---|---|---|---|---|---|---|--|--|

Example for Dilation (1D)

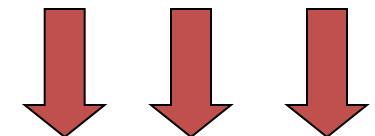
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 1 | 0 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|

Note that the object gets bigger and holes are filled

2D example for dilation

| | | | | | | | |
|---|---|---|---|---|---|---|--|
| | | | | | | | |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | |
| | | | 1 | 1 | 1 | 1 | |
| | | | 1 | 1 | 1 | 1 | |
| | | 1 | 1 | 1 | 1 | 1 | |
| | | 1 | 1 | 1 | 1 | 1 | |
| | 1 | 1 | | | | | |
| | | | | | | | |

(a) Binary image \mathbf{B}

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

(b) Structuring element \mathbf{S}

Dilation vs. Erosion

At each position:

- **Dilation:** if **current pixel** is foreground OR the center of structuring element covers foreground pixel in the input image.
- **Erosion:** if **current pixel** is foreground OR if **every pixel** under the structuring element's nonzero entries is foreground

Example for Erosion (1D)

Input

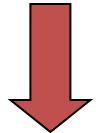
| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|

$$O(x) = I(x) \ominus ES$$



Output

| | | | | | | | | | |
|---|--|--|--|--|--|--|--|--|--|
| 0 | | | | | | | | | |
|---|--|--|--|--|--|--|--|--|--|

Example for Erosion (1D)

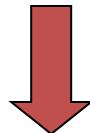
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|--|--|--|--|--|--|--|--|
| 0 | 0 | | | | | | | | |
|---|---|--|--|--|--|--|--|--|--|

Example for Erosion (1D)

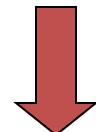
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



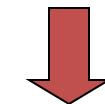
Output

| | | | | | | | | | |
|---|---|---|--|--|--|--|--|--|--|
| 0 | 0 | 0 | | | | | | | |
|---|---|---|--|--|--|--|--|--|--|

Example for Erosion (1D)

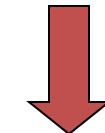
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|

Structuring
Element



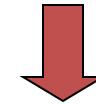
Output

| | | | | | | | | | |
|---|---|---|---|--|--|--|--|--|--|
| 0 | 0 | 0 | 0 | | | | | | |
|---|---|---|---|--|--|--|--|--|--|

Example for Erosion (1D)

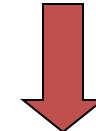
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|---|--|--|--|--|--|
| 0 | 0 | 0 | 0 | 0 | | | | | |
|---|---|---|---|---|--|--|--|--|--|

Example for Erosion (1D)

Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



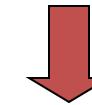
Output

| | | | | | | | | | |
|---|---|---|---|---|---|--|--|--|--|
| 0 | 0 | 0 | 0 | 0 | 1 | | | | |
|---|---|---|---|---|---|--|--|--|--|

Example for Erosion (1D)

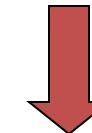
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|--|--|--|
| 0 | 0 | 0 | 0 | 0 | 1 | 0 | | | |
|---|---|---|---|---|---|---|--|--|--|

Example for Erosion (1D)

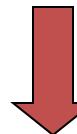
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|--|--|
| 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | | |
|---|---|---|---|---|---|---|---|--|--|

Example for Erosion (1D)

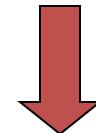
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



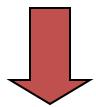
Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|--|
| 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | |
|---|---|---|---|---|---|---|---|---|--|

Example for Erosion (1D)

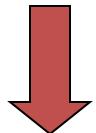
Input

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 1 | 0 | 0 | 0 | 1 | 1 | 1 | 0 | 1 | 1 |
|---|---|---|---|---|---|---|---|---|---|



Structuring
Element

| | | |
|---|---|---|
| 1 | 1 | 1 |
|---|---|---|



Output

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 1 |
|---|---|---|---|---|---|---|---|---|---|

Note that the object gets smaller

2D example for erosion

| | | | | | | | |
|---|---|---|---|---|---|---|--|
| | | | | | | | |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | |
| | | 1 | 1 | 1 | 1 | 1 | |
| | 1 | 1 | 1 | 1 | 1 | 1 | |
| | 1 | 1 | 1 | 1 | 1 | 1 | |
| | 1 | 1 | 1 | 1 | 1 | 1 | |
| | 1 | 1 | | | | | |
| | | | | | | | |

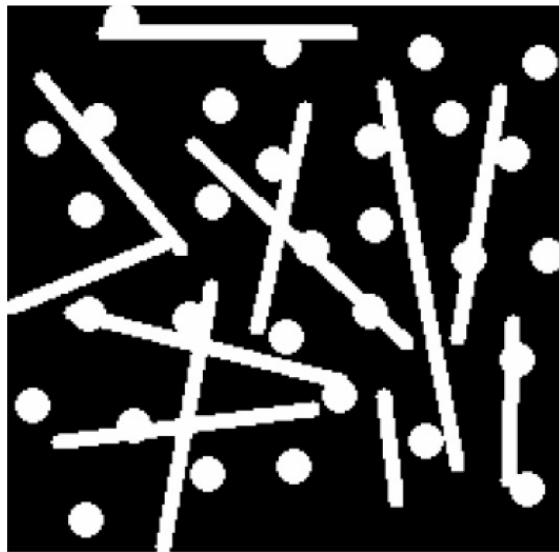
(a) Binary image \mathbf{B}

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

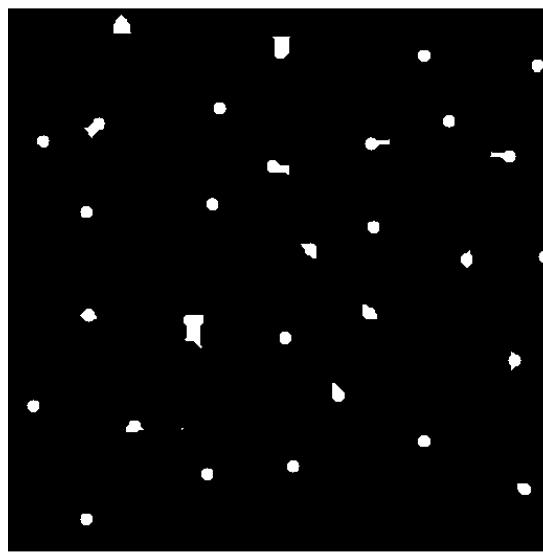
(b) Structuring element \mathbf{S}

Opening

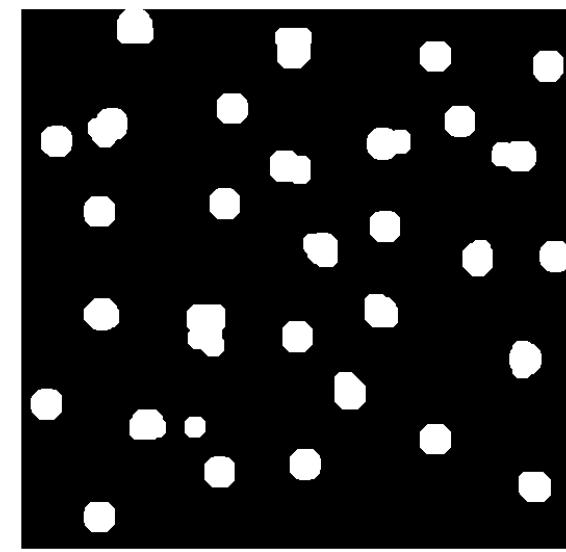
- Erode, then dilate
- Remove small objects, keep original shape



Initial image



Erosion



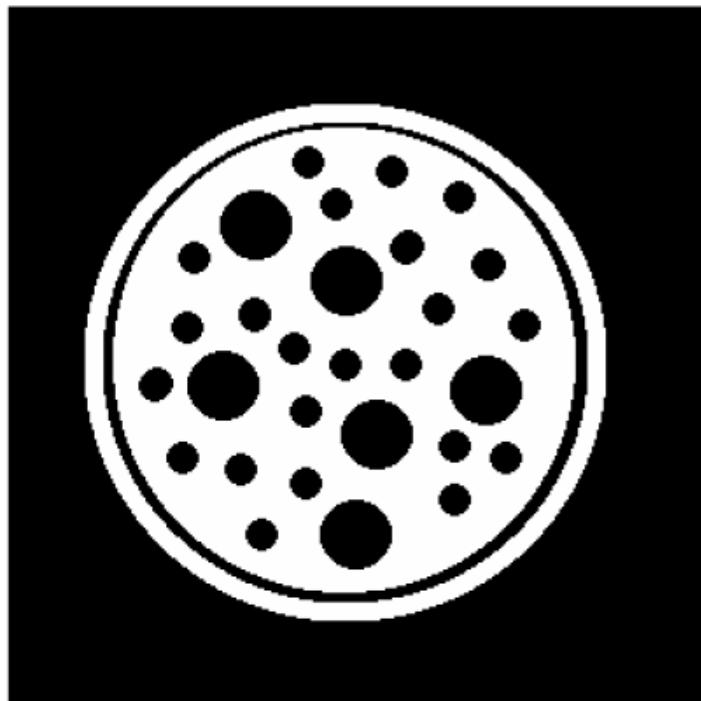
Erosion + Dilation

Structural element

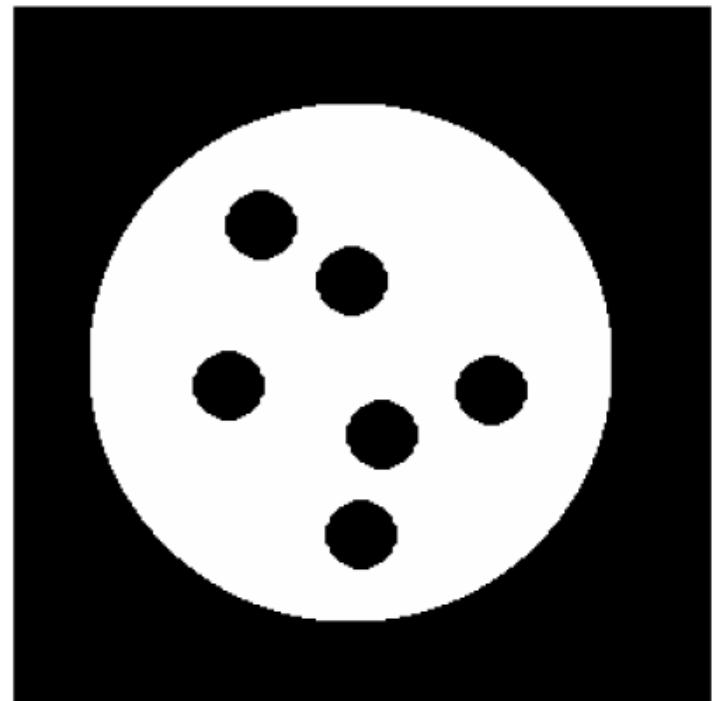


Closing

- Dilate, then erode
- Fill holes, but keep original shape

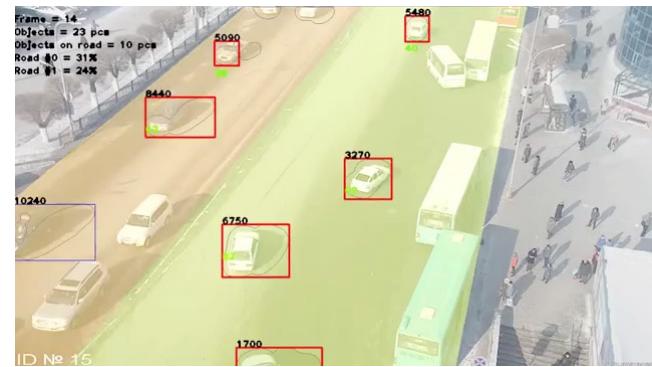


Before closing



After closing

Object tracking



traffic



body



eyes



soccer



snooker



face

Further applications

- Motion capture
- Augmented Reality
- Action Recognition
- Security, traffic monitoring
- Video Compression
- Video Summarization
- Medical Screening
- Sport Analysis

Object tracking

Optical flow is used to determine how each point is moving from frame to frame

Object tracking

Optical flow is used to determine how each point is moving from frame to frame

Different problem: track the location of target objects in each frame of a video sequence.

Object tracking

Optical flow is used to determine how each point is moving from frame to frame

Different problem: track the location of target objects in each frame of a video sequence.

We want the tracking algorithm to be insensitive to changes in:

- illumination
- scale
- rotations
- possible occlusions

Object tracking

Different problem: track the location of target objects in each frame of a video sequence.

We can build tracks for each of the object of interest and understand how objects are moving with respect to each other wrt 3D scene.

Object tracking

Different problem: track the location of target objects in each frame of a video sequence.

We can build tracks for each of the object of interest and understand how objects are moving with respect to each other wrt 3D scene.

Topics:

- (1) Change Detection (detect meaningful changes in the scene)
- (2) Gaussian Mixture Model
- (3) Tracking by Feature Detection
- (4) Object Tracking using Particle Filter

Change detection

Given: static camera observing scene (room, street, station, etc.)

Find: meaningful changes (moving objects, people, etc)



Want to build algorithms which are insensitive to uninteresting changes.
Robust and real-time classification of each pixel as foreground
(motion/change) or background (static)

Change detection: challenges

Ignore uninteresting changes:

- Background fluctuation
- Image noise



Change detection: challenges

Ignore uninteresting changes:

- Background fluctuation
- Image noise
- Rain, snow, turbulence



Change detection: challenges

Ignore uninteresting changes:

- Background fluctuation
- Image noise
- Rain, snow, turbulence
- Illumination changes & shadows
- Camera shake



PHO East (30 fps 8.35 Mb/s) 2007-08-08 16:53:28.11

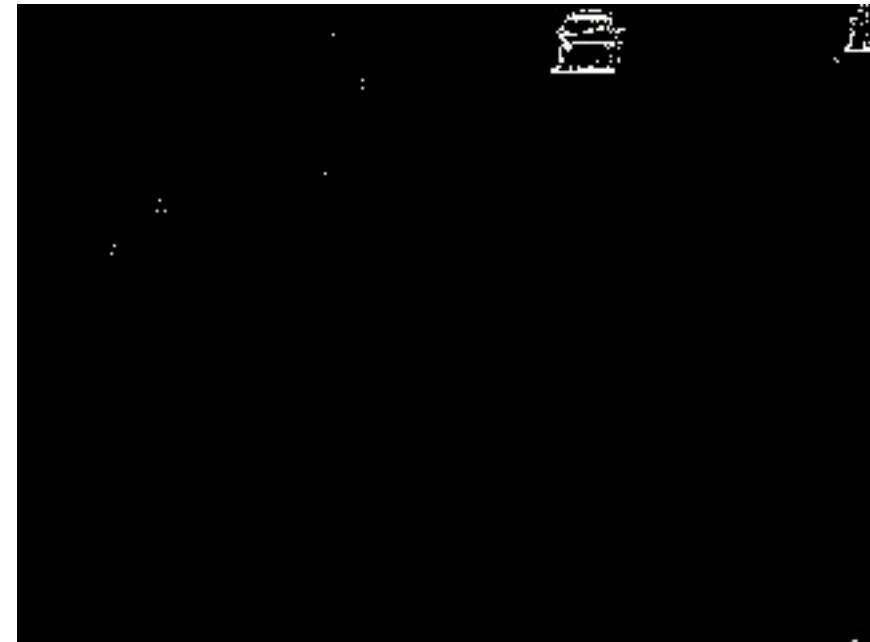
Simple frame difference

Label significant difference between current and previous frame as foreground.

$$F_t = |I_t - I_{t-1}| > T, \quad T: \text{threshold}$$



Input video sequence



Frame difference

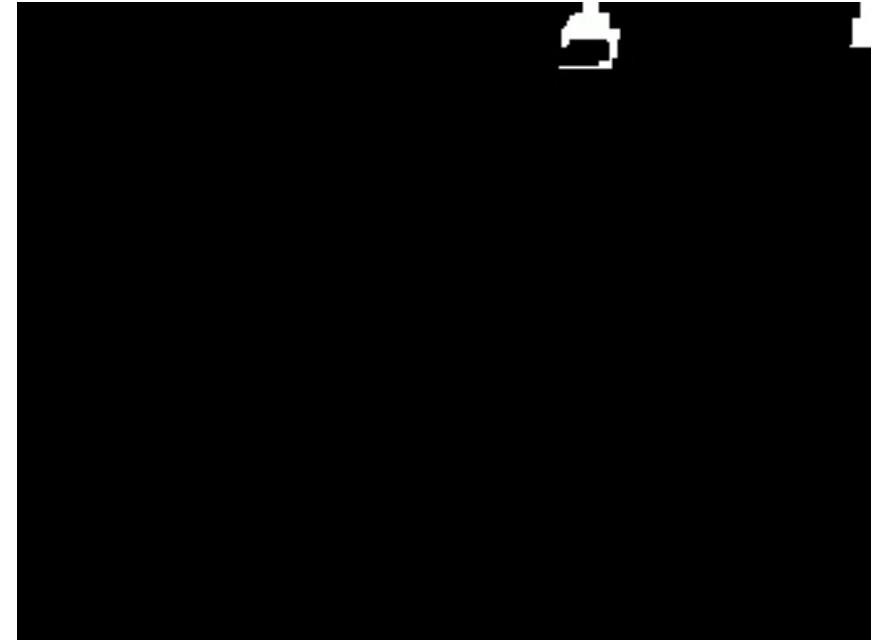
Simple frame difference

Label significant difference between current and previous frame as foreground.

$$F_t = |I_t - I_{t-1}| > T, \quad T: \text{threshold}$$



Input video sequence



Frame difference + dilation + erosion

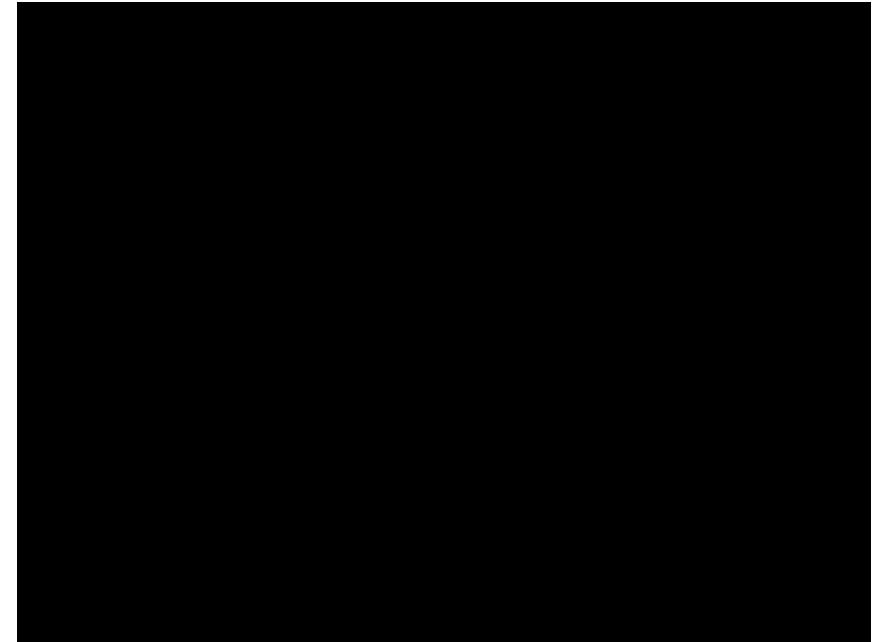
Simple frame difference

Label significant difference between current and previous frame as foreground.

$$F_t = |I_t - I_{t-1}| > T, \quad T: \text{threshold}$$



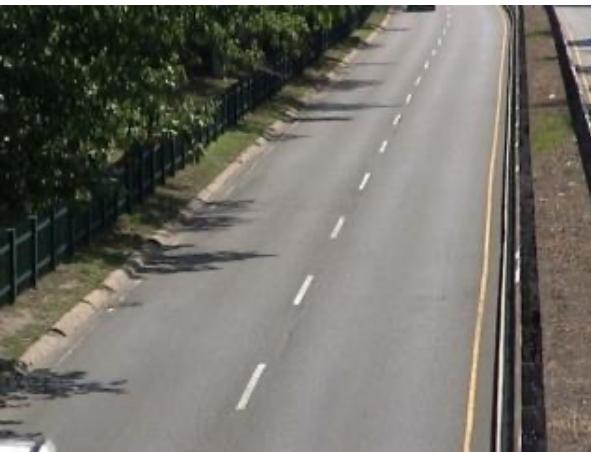
Input video sequence



Frame difference + erosion + dilation

Background modelling: Average

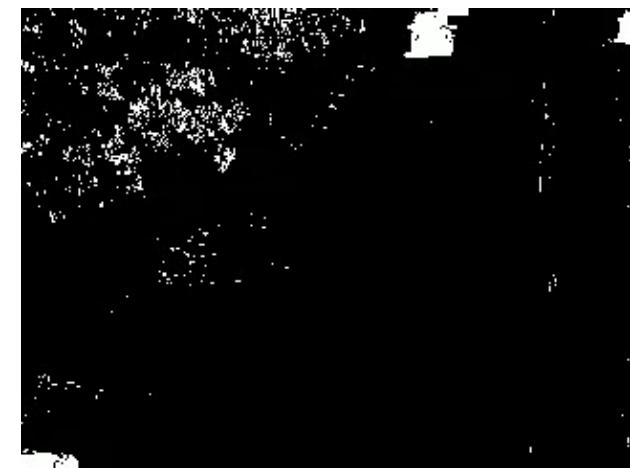
Build simple model of background before classification



Background B
average{I₁, I₂, ..., I_k}
(first K frames)



Input frame I_t

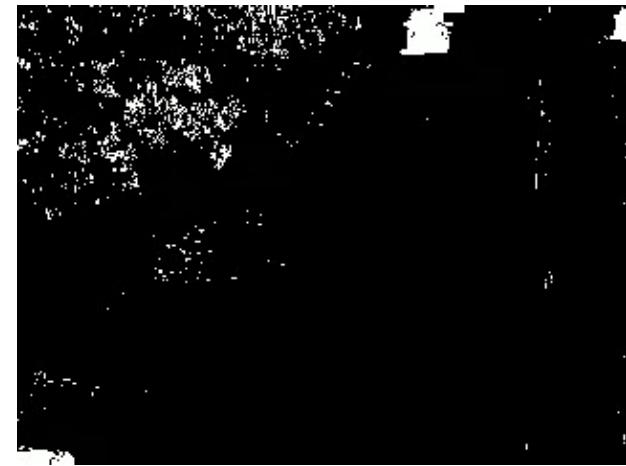


Foreground F_t
 $F_t = |I_t - B| > T$

Cannot handle change in background, lighting, etc.

Background modelling: Median

Build simple model of background before classification



Background B

median{ I_1, I_2, \dots, I_k }
(first K frames)

Input frame I_t

Foreground F_t

$$F_t = |I_t - B| > T$$

Cannot handle change in background, lighting, etc.

Background modelling: Moving Median

Build simple **adaptive** model of background over time.



Background B

median{ $I_{t-1}, I_{t-2}, \dots, I_{t-k}$ }
(last K frames)

Input frame I_t

Foreground F_t

$$F_t = |I_t - B| > T$$

Cannot handle significant pixel fluctuations

Mixture Model

Intensity distribution at each pixel over time:



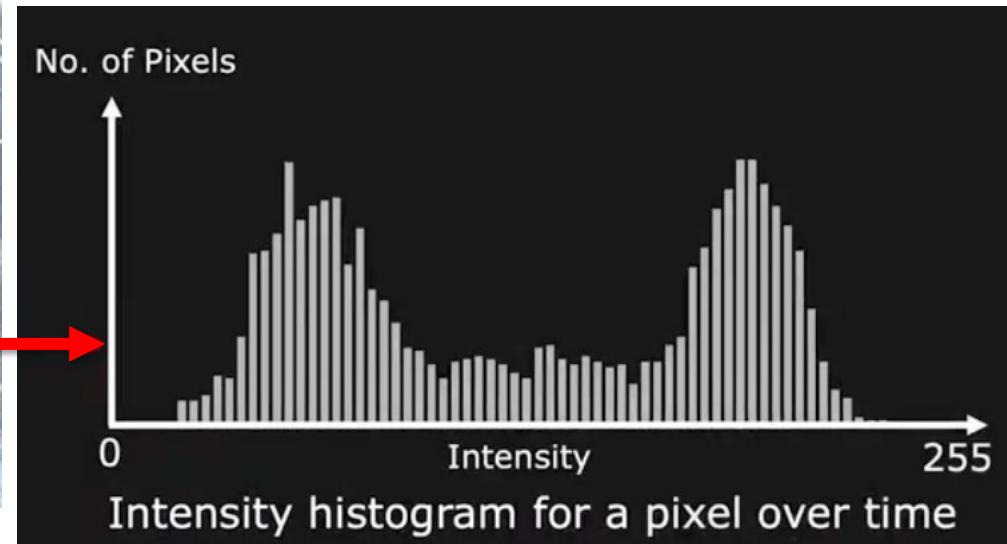
Input video sequence

Mixture Model

Intensity distribution at each pixel over time:



Input video sequence

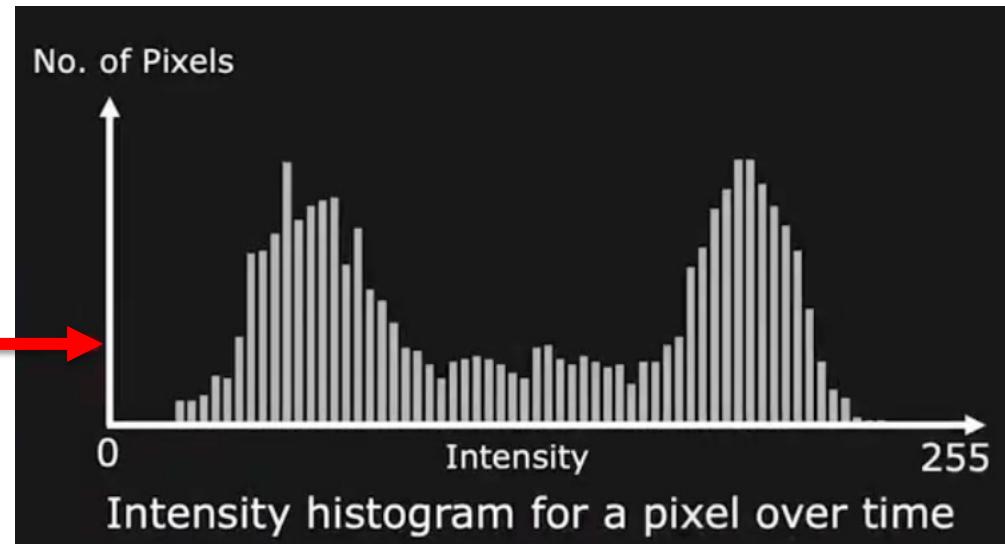


Mixture Model

Intensity distribution at each pixel over time:



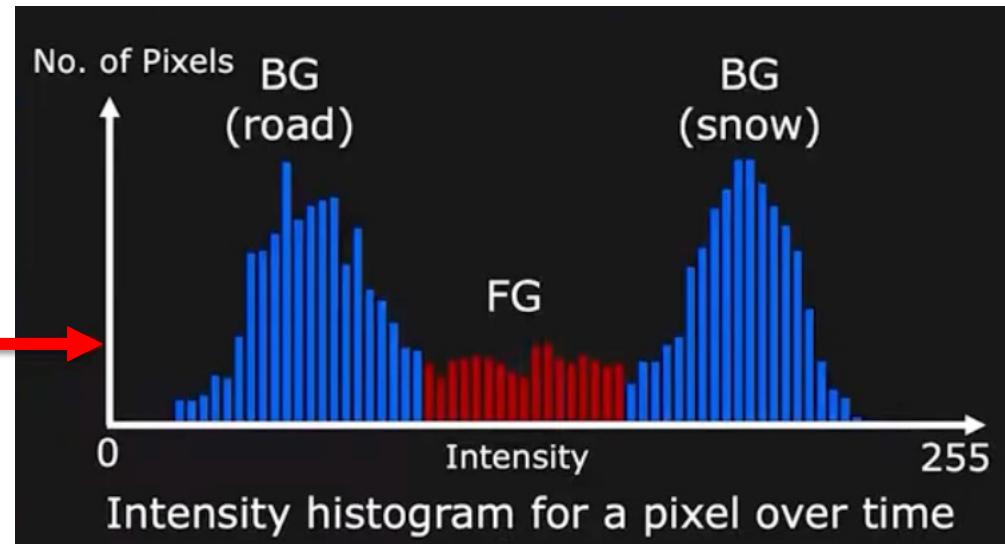
Input video sequence



Intensity variations due to static scene (road), noise (snow) and occasional moving objects (vehicles).

Mixture Model

Intensity distribution at each pixel over time:



Input video sequence

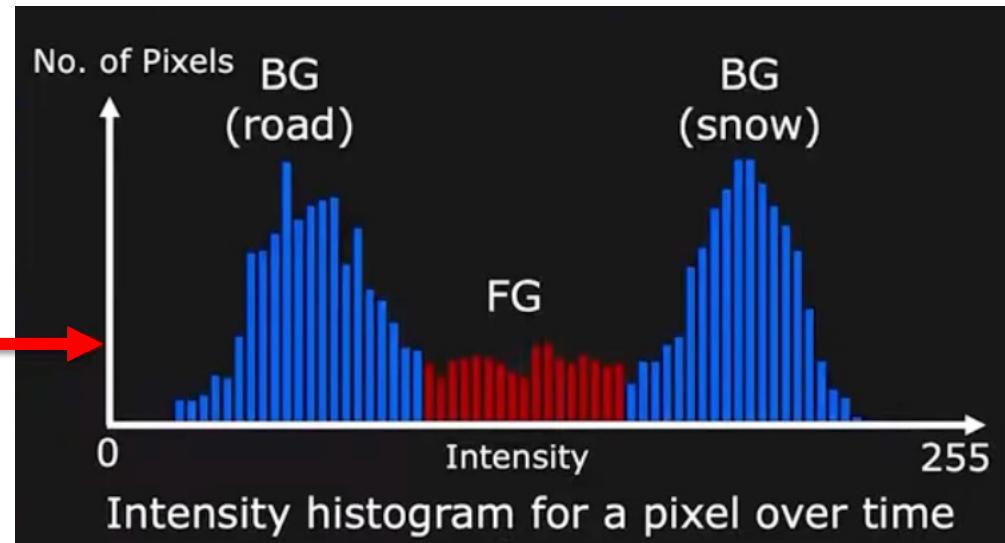
Intensity variations due to static scene (**road**), noise (**snow**) and occasional moving objects (**vehicles**).

Mixture Model

Intensity distribution at each pixel over time:



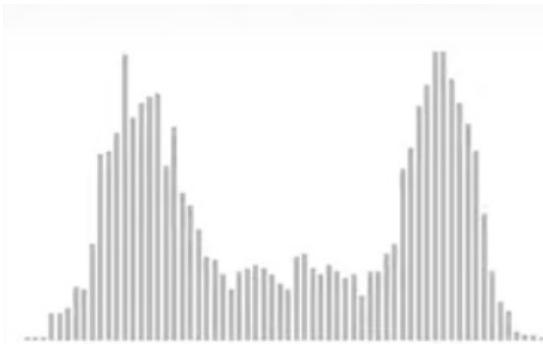
Input video sequence



Intensity variations due to static scene (**road**), noise (**snow**) and occasional moving objects (**vehicles**).

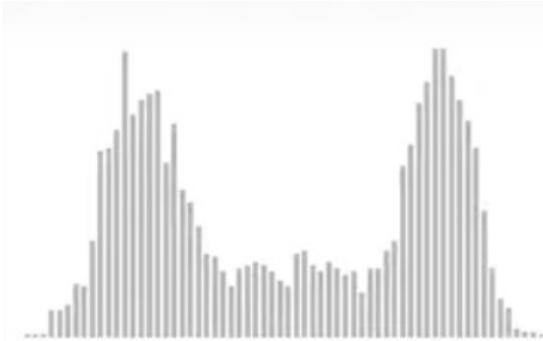
Intuition: Pixels are background most of the time

Gaussian Model

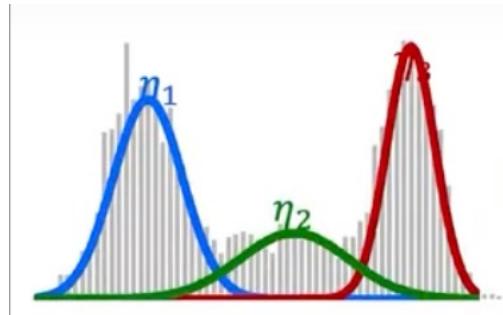


Probability Distribution
 $P(x)$ (x : pixel intensity)

Mixture of Gaussians



Probability Distribution
 $P(x)$ (x : pixel intensity)



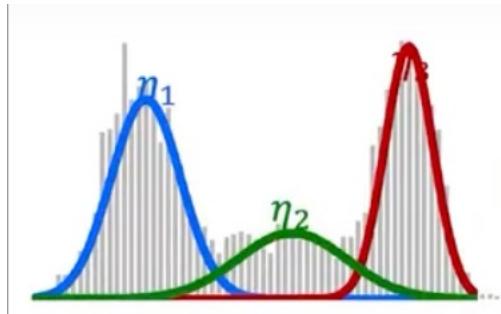
Mixture of Gaussians
 $\omega_k, \eta_k(x, \mu_k, \sigma_k)$

Assume $P(x)$ is made of K different Gaussians.

Gaussian Mixture Model (GMM)

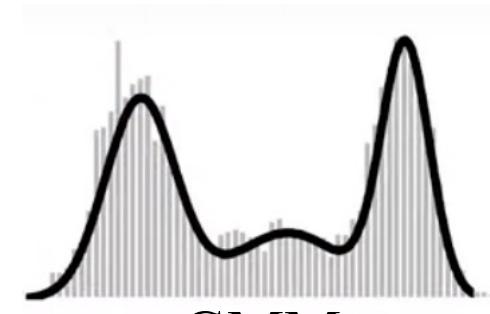


Probability Distribution
 $P(x)$ (x : pixel intensity)



Mixture of Gaussians

$$\omega_k, \eta_k(x, \mu_k, \sigma_k) \quad P(x) = \sum_{k=1}^K \omega_k \eta_k(x, \mu_k, \sigma_k)$$



GMM

Gaussian Mixture Model Distribution : weighted sum of K

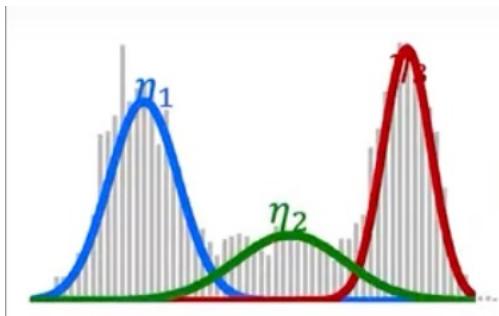
$$P(x) \cong \sum_{k=1}^K \omega_k \eta_k(x, \mu_k, \sigma_k) \quad \text{Gaussians.}$$

$$\text{such that } \sum_{k=1}^K \omega_k = 1$$

Gaussian Mixture Model (GMM)

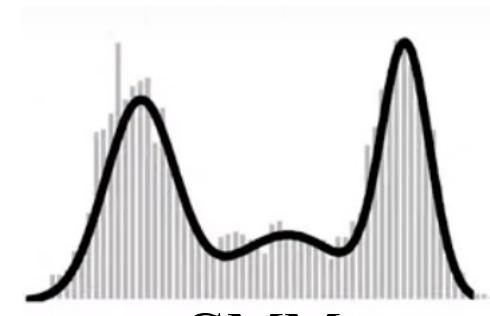


Probability Distribution
 $P(x)$ (x : pixel intensity)



Mixture of Gaussians

$$\omega_k, \eta_k(x, \mu_k, \sigma_k) \quad P(x) = \sum_{k=1}^K \omega_k \eta_k(x, \mu_k, \sigma_k)$$



GMM

Gaussian Mixture Model Distribution : weighted sum of K

$$P(x) \cong \sum_{k=1}^K \omega_k \eta_k(x, \mu_k, \sigma_k) \quad \text{such that} \quad \sum_{k=1}^K \omega_k = 1$$

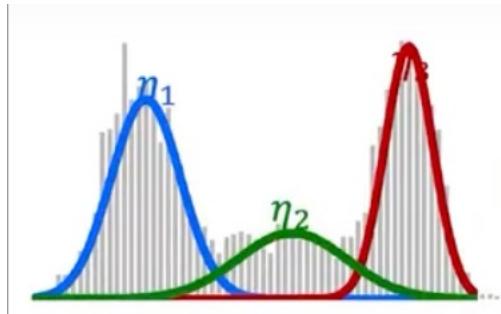
Gaussians.

We can extend this model to high dimensional GMM, to model pixel colors (r,g,b), go from 1D to 3D

Gaussian Mixture Model (GMM)

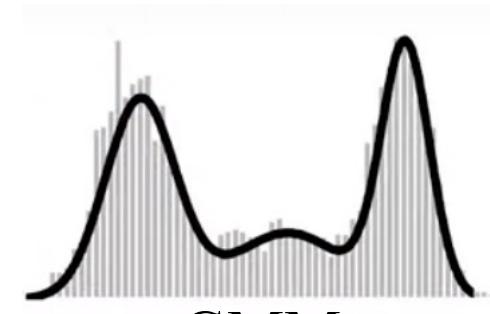


Probability Distribution
 $P(x)$ (x : pixel intensity)



Mixture of Gaussians

$$\omega_k, \eta_k(x, \mu_k, \sigma_k) \quad P(x) = \sum_{k=1}^K \omega_k \eta_k(x, \mu_k, \sigma_k)$$



GMM

Gaussian Mixture Model Distribution : weighted sum of K

$$P(x) \cong \sum_{k=1}^K \omega_k \eta_k(x, \mu_k, \sigma_k) \quad \text{Gaussians.}$$

$$\text{such that } \sum_{k=1}^K \omega_k = 1$$

GMM can be estimated from $P(x)$

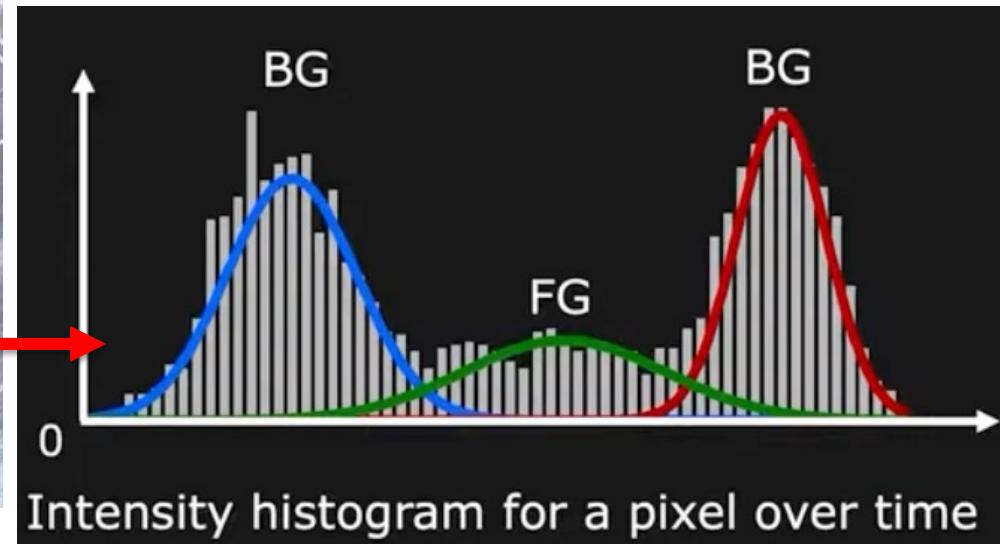
Background Modeling with GMM

Given: A GMM for intensity/color variation at a pixel over time

Classify: Individual Gaussians as foreground/background



Input video sequence

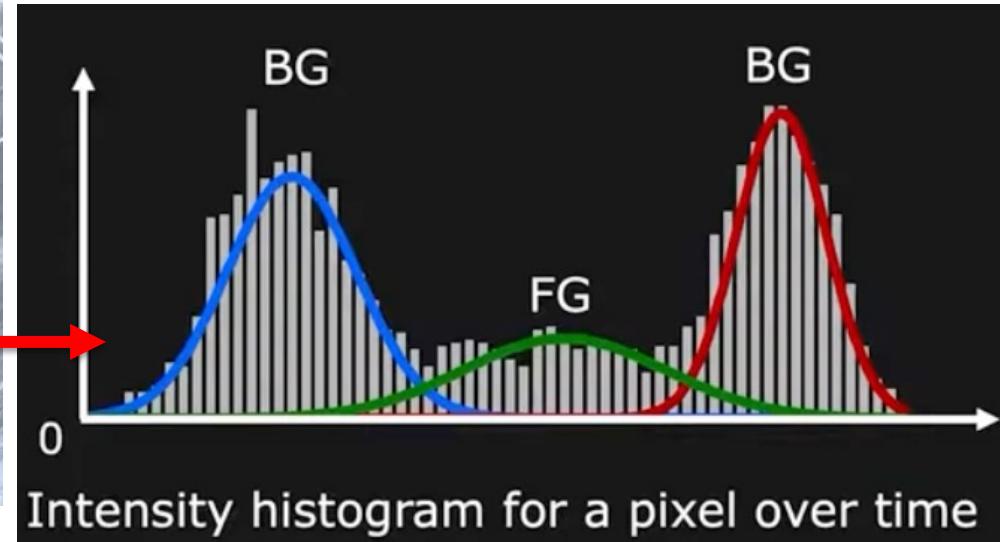


Intuition: Pixels are background most of the time. That is, Gaussians with large supporting evidence ω and small σ .

Background Modeling with GMM

Given: A GMM for intensity/color variation at a pixel over time

Classify: Individual Gaussians as foreground/background



Input video sequence

Intuition: Pixels are background most of the time. That is, Gaussians with large supporting evidence ω and small σ .

Large $\frac{\omega}{\sigma}$: Background

Small $\frac{\omega}{\sigma}$: Foreground

Change Detection using GMM

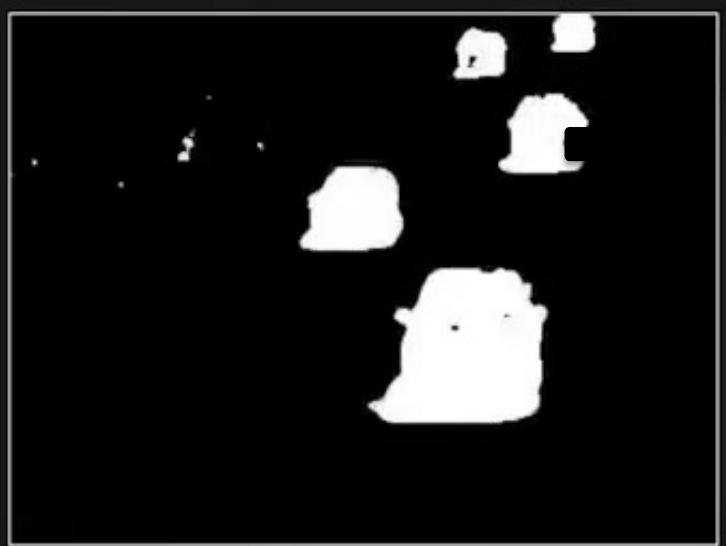
For each pixel:

1. Compute pixel color histogram H using first N frames.
2. Normalize histogram H : $\bar{H} = \frac{H}{\|H\|}$
3. Model \bar{H} as a mixture of K (3 to 5) Gaussians.
4. For each subsequent frame:
 1. The pixel value X belongs to Gaussian k in GMM for which $\|X - \mu_k\|$ is minimum but also $\|X - \mu_k\| < 2.5\sigma_k$
 2. If $\frac{\omega_k}{\sigma_k}$ is large classify pixel as background. Else classify as foreground.
 3. Update histogram H using new pixel intensity
 4. If \bar{H} and $\frac{H}{\|H\|}$ differ a lot refit GMM

Adaptive GMM based Change Detection using GMM



Input video



Foreground

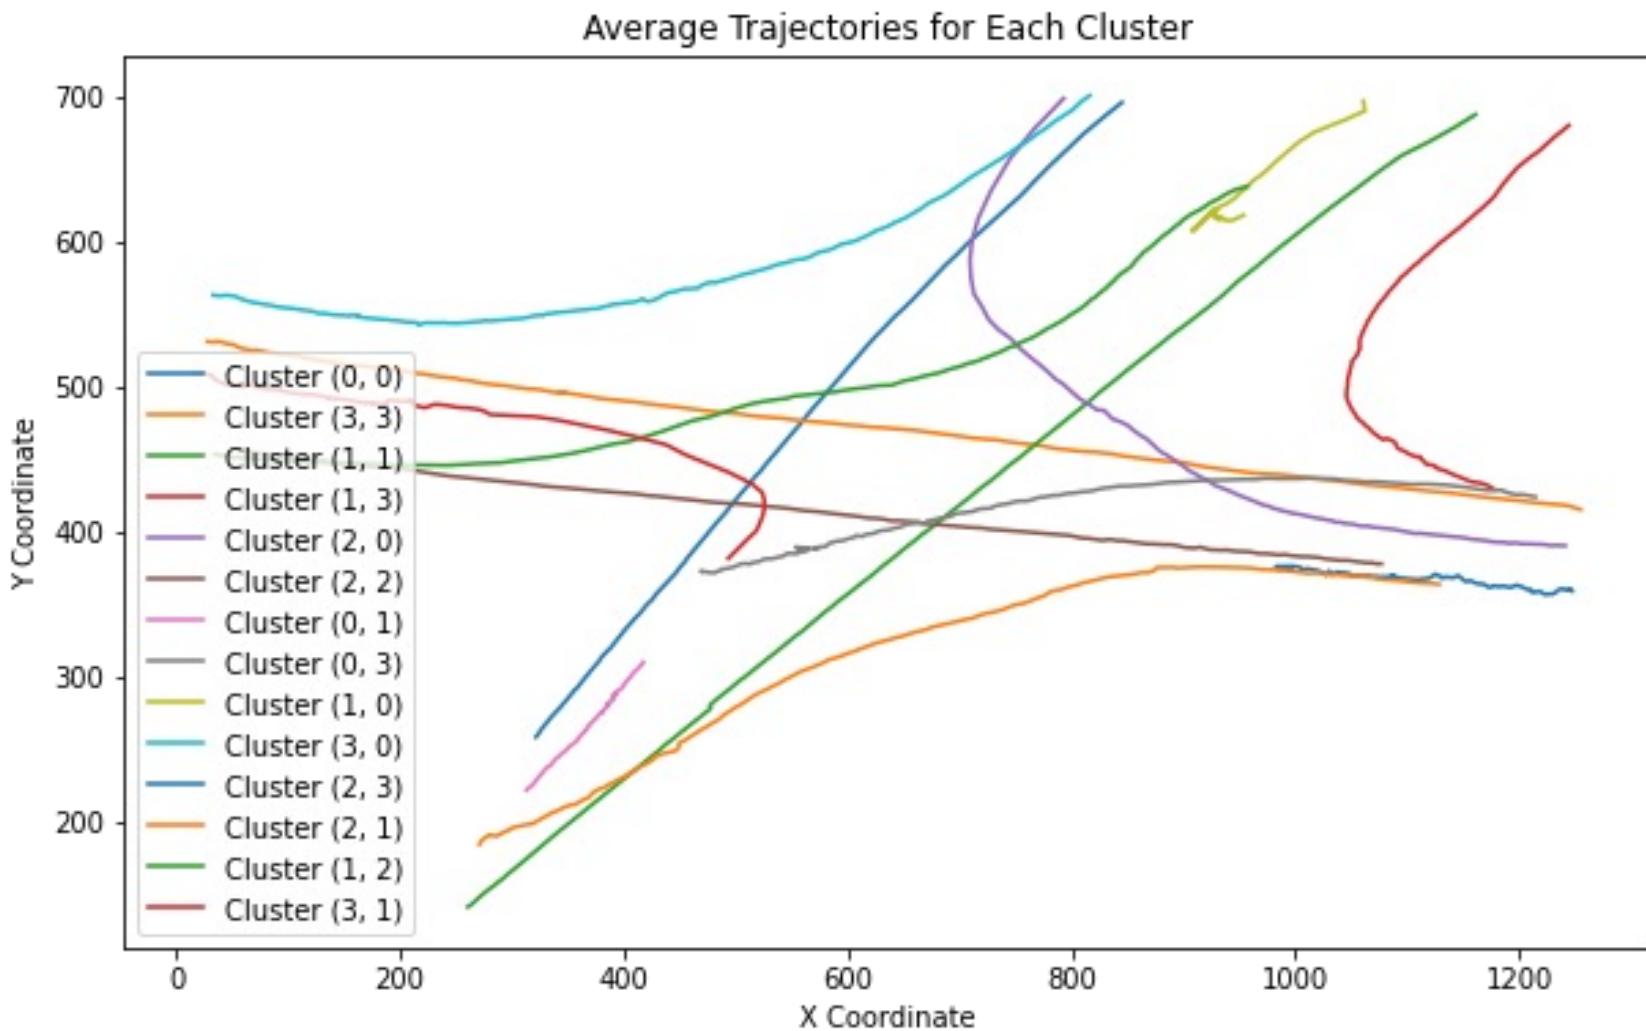
Strategies for tracking

- **Feature-based tracking**
 - extract visual features (corners, textured areas) and track them over multiple frames
 - sparse motion fields, but more robust tracking
 - figure out which features can be tracked, background vs foreground model
 - some points may change appearance over time (e.g., due to rotation, moving into shadows, etc.)
- **Tracking by repeated detection**
 - works well if object is easily detectable (e.g., face, ball or colored glove) and there is only one
 - need some way to link up detections (obtain tracks)
 - best you can do, if you can't predict motion
- **Tracking with dynamics**
 - based on a model of expected motion, predict where objects will occur in next frame, before even seeing the image
 - restrict search for the object
 - measurement noise is reduced by trajectory smoothness

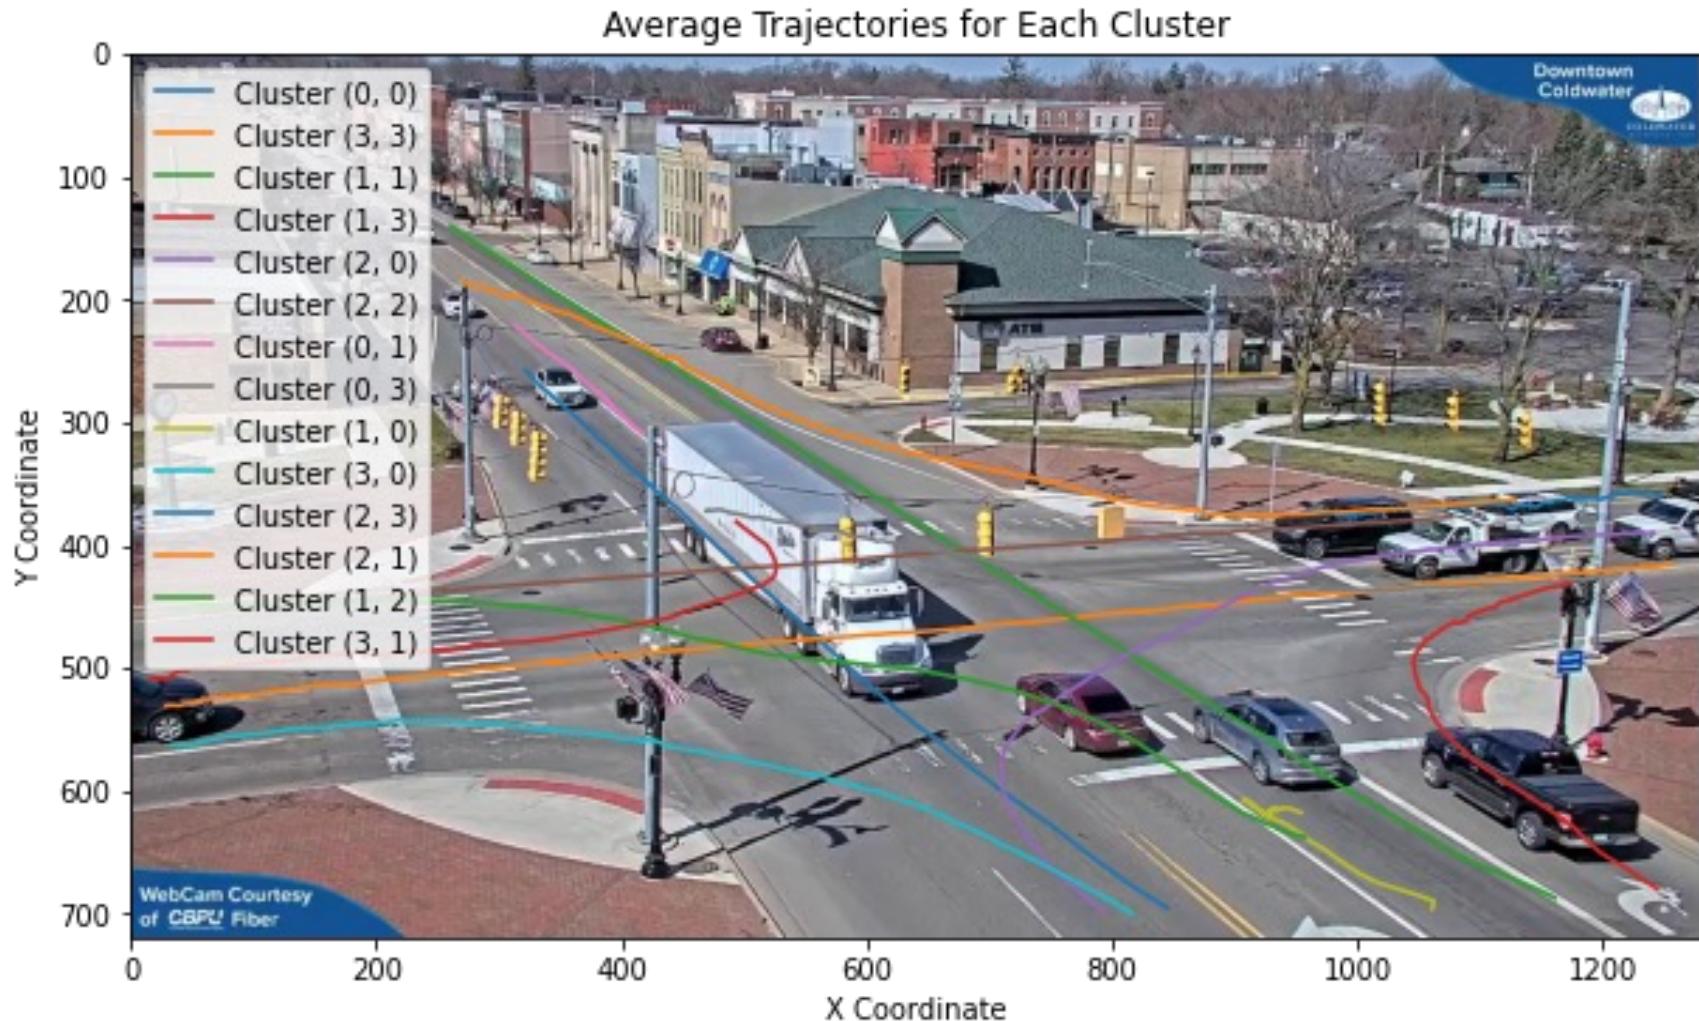
Lab 7 – Extracting trajectorie in traffic surveillance



Lab 7 – Extracting trajectorie in traffic surveillance

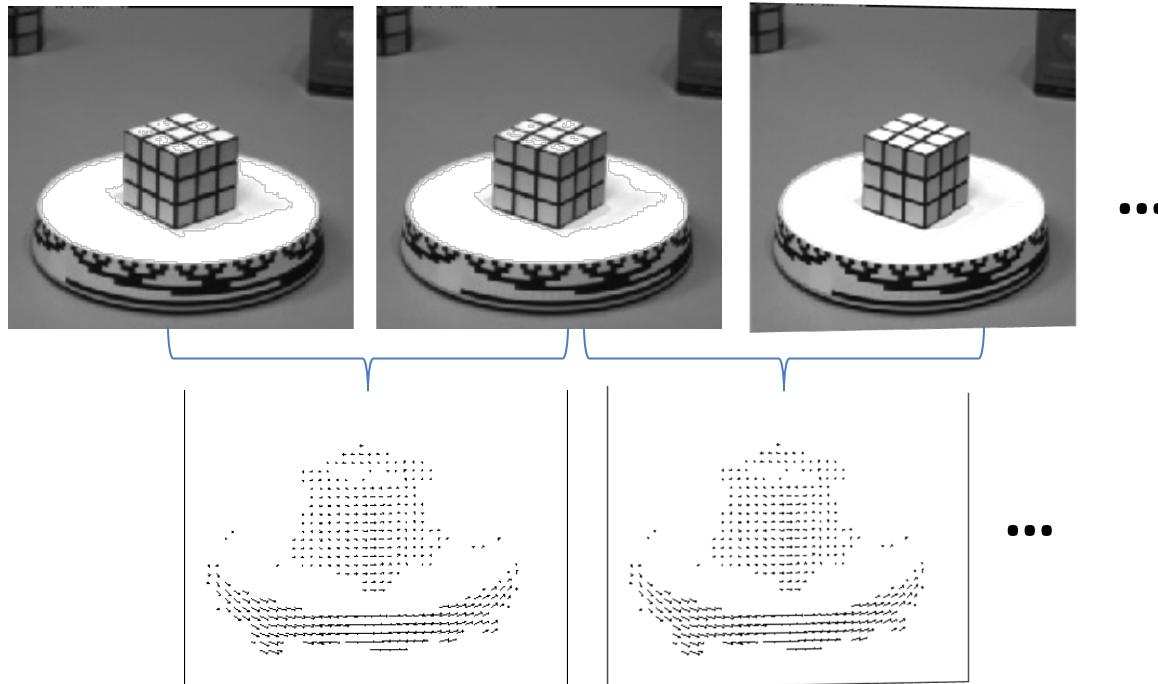


Lab 7 – Extracting trajectorie in traffic surveillance



Optical flow for tracking?

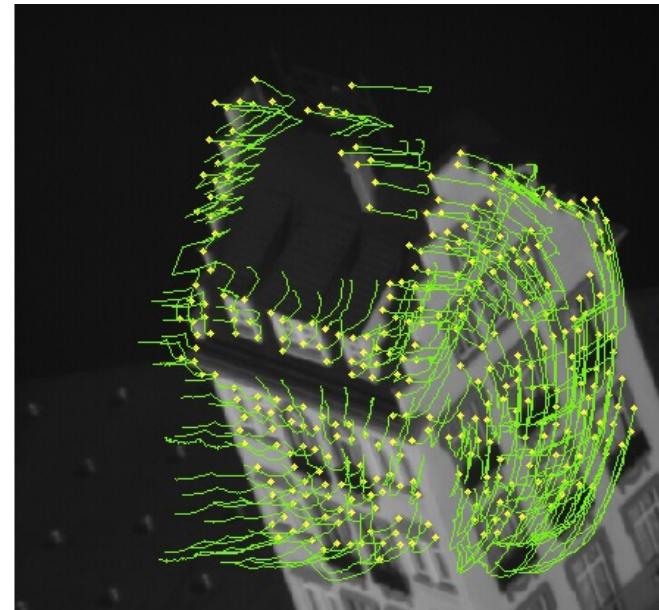
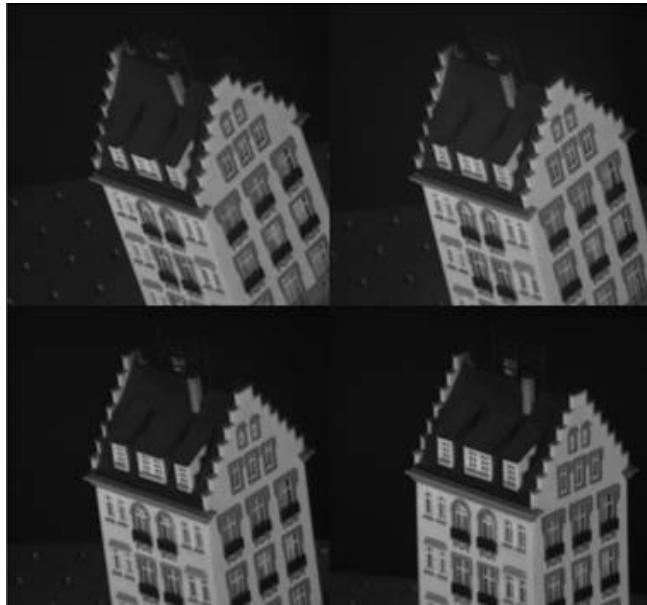
If we have more than just a pair of frames, we could compute flow from one to the next:



But flow only reliable for small motions, and we may have occlusions, textureless regions that yield bad estimates anyway...

Feature tracking

- If we have more than two images, we can track a feature from one frame to the next by following the optical flow
- Challenges
 - Finding good features to track
 - Adding and deleting tracks (trajectories of objects in video)

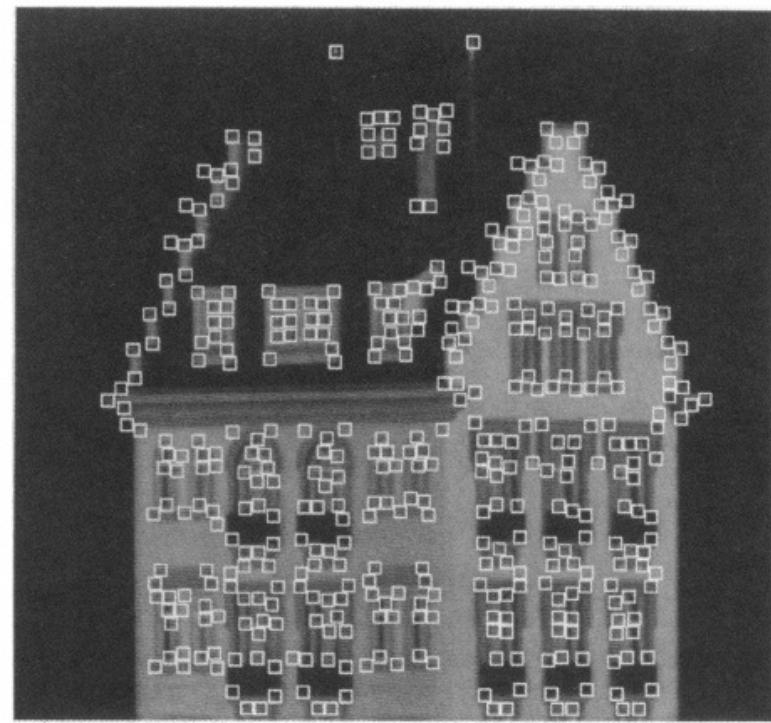
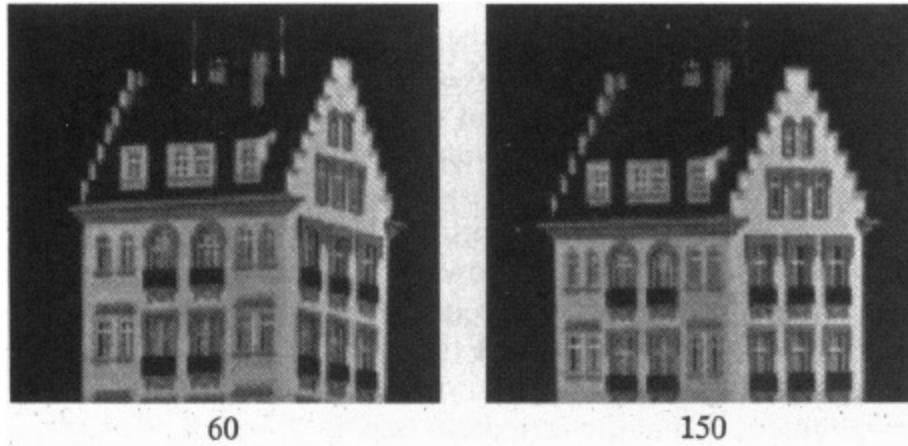


Feature tracking

- Challenges
 - Figure out which features can be tracked
 - Efficiently track across frames
 - Some points may change appearance over time (e.g., due to rotation, moving into shadows, etc.)
 - Drift: small errors can accumulate as appearance model is updated
 - Points may appear or disappear: need to be able to add/delete tracked points

Lucas-Kanade feature tracker

- Same as Lucas-Kanade optical flow, but instead for all pixels just for selected features (e.g. Harris corners)



Shi-Tomasi feature tracker

- Find good features using eigenvalues of second-moment matrix
 - Key idea: “good” features to track are the ones whose motion can be estimated reliably (Harris corner)
- From frame to frame, track with Lucas-Kanade
 - This amounts to assuming a translation model for frame-to-frame feature movement
- Check consistency of tracks by *affine* registration to the first observed instance of the feature
 - Affine model is more accurate for larger displacements
 - Comparing to the first frame helps to minimize drift

Tracking example



Figure 1: Three frame details from Woody Allen's *Manhattan*. The details are from the 1st, 11th, and 21st frames of a subsequence from the movie.



Figure 2: The traffic sign windows from frames 1,6,11,16,21 as tracked (top), and warped by the computed deformation matrices (bottom).

Detection vs. tracking



t=1



t=2

...

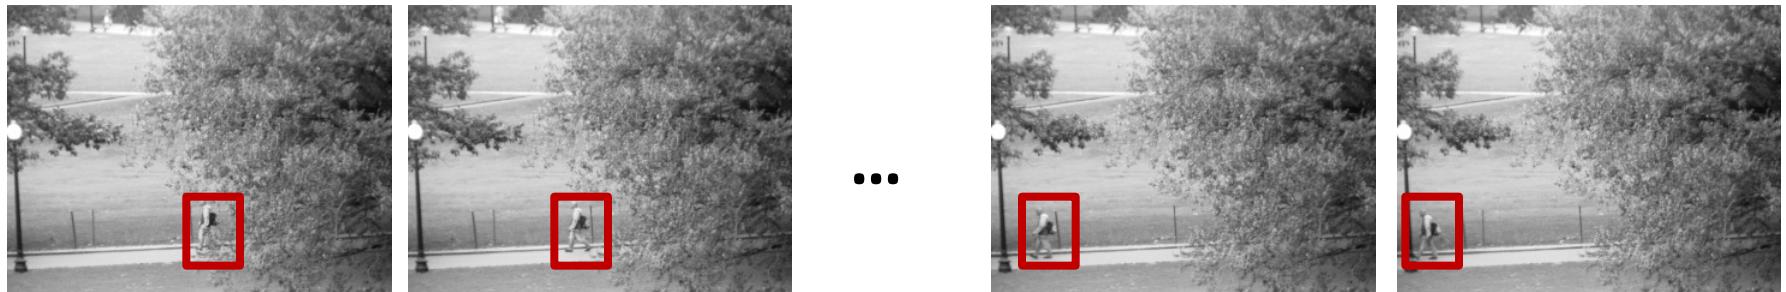


t=20



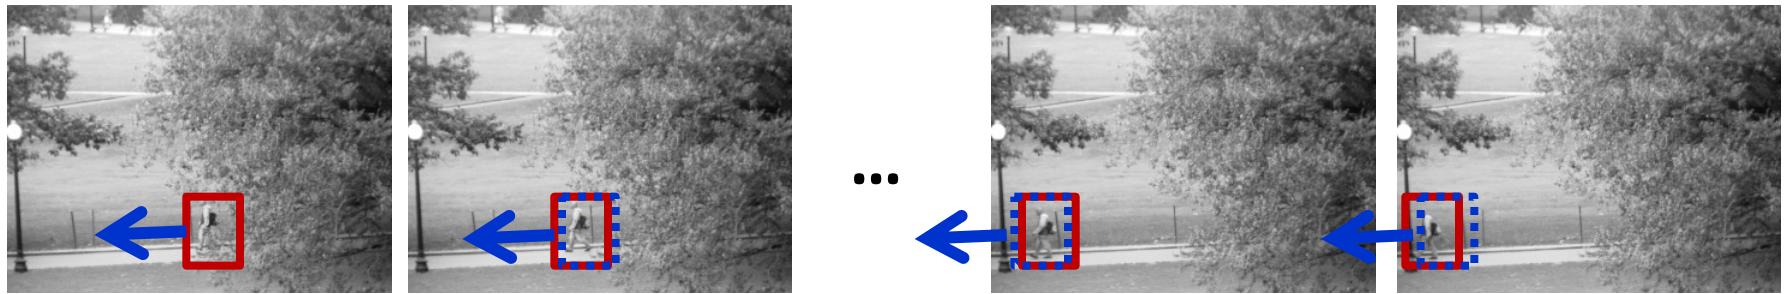
t=21

Detection vs. tracking



Detection: We detect the object independently in each frame and can record its position over time, e.g., based on blob's centroid or detection window coordinates

Detection vs. tracking



Tracking with *dynamics*: We use image measurements to estimate position of object, but also incorporate position predicted by dynamics, i.e., our expectation of object's motion pattern.

Tracking with dynamics

- Use model of expected motion to *predict* where objects will occur in next frame, even before seeing the image.
- **Intent:**
 - Do less work looking for the object, restrict the search.
 - Get improved estimates since measurement noise is tempered by smoothness, dynamics priors.
- **Assumption:** continuous motion patterns:
 - Camera is not moving instantly to new viewpoint
 - Objects do not disappear and reappear in different places in the scene
 - Gradual change in pose between camera and scene

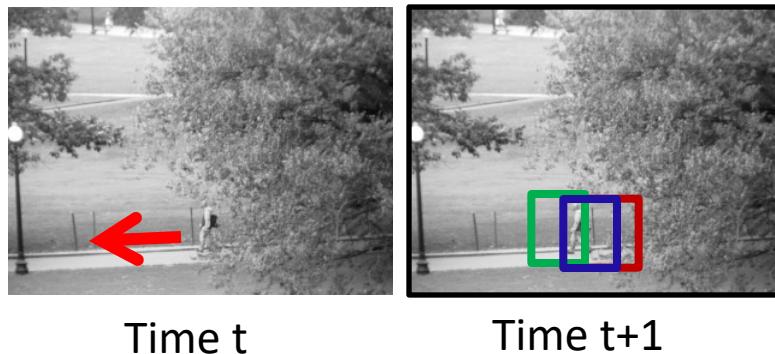
Tracking as inference

- The *hidden state* consists of the true parameters we care about, denoted X (e.g. a vector containing the position, velocity, acceleration).
- The *measurement* is our noisy observation that results from the underlying state, denoted Y . Could be the position of some image points, the position of some image regions or pretty much anything else.
- At each time step, state changes (from X_{t-1} to X_t) and we get a new observation Y_t .

Tracking as inference

- The *hidden state* consists of the true parameters we care about, denoted X .
- The *measurement* is our noisy observation that results from the underlying state, denoted Y .
- At each time step, state changes (from X_{t-1} to X_t) and we get a new observation Y_t .
- Our goal: recover most likely state X_t given
 - All observations seen so far.
 - Knowledge about dynamics of state transitions.

Tracking as inference: intuition

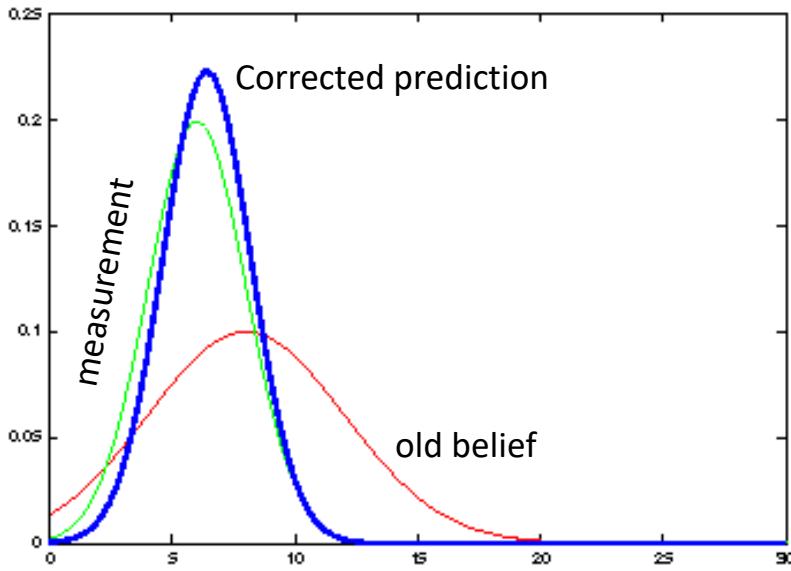
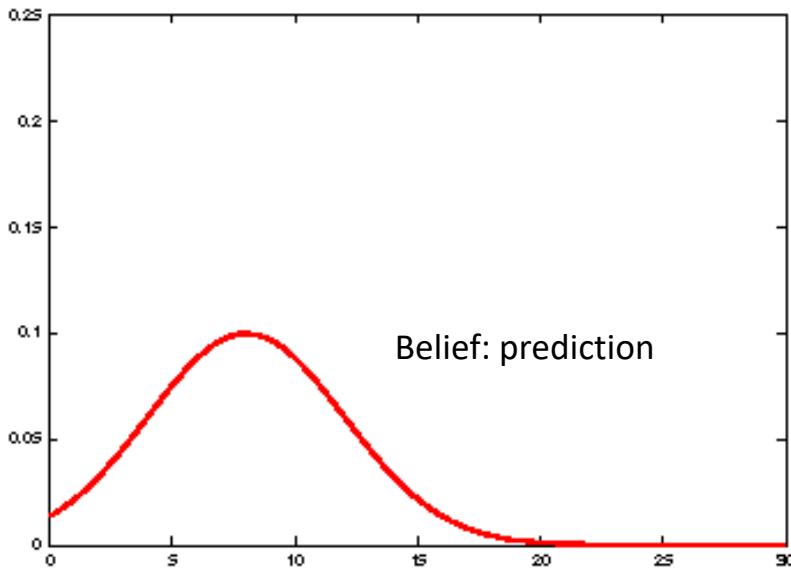


Belief

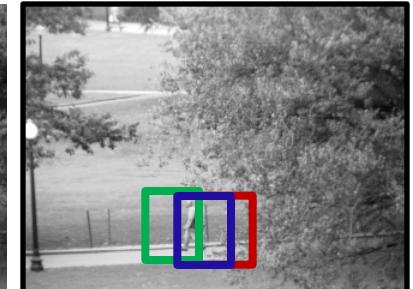
Measurement

Corrected prediction

Tracking as inference: intuition



Time t



Time $t+1$

Independence assumptions

- Only immediate past state influences current state

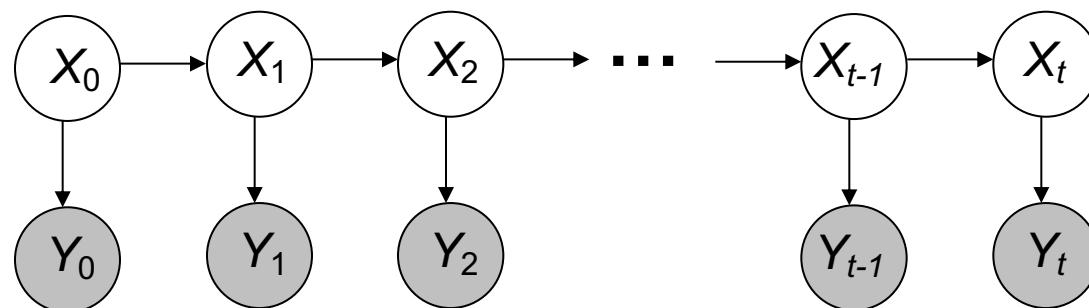
$$P(X_t | X_0, \dots, X_{t-1}) = P(X_t | X_{t-1})$$

dynamics model

- Measurement at time t depends on current state

$$P(Y_t | X_0, Y_0, \dots, X_{t-1}, Y_{t-1}, X_t) = P(Y_t | X_t)$$

observation model



Hidden
Markov
model

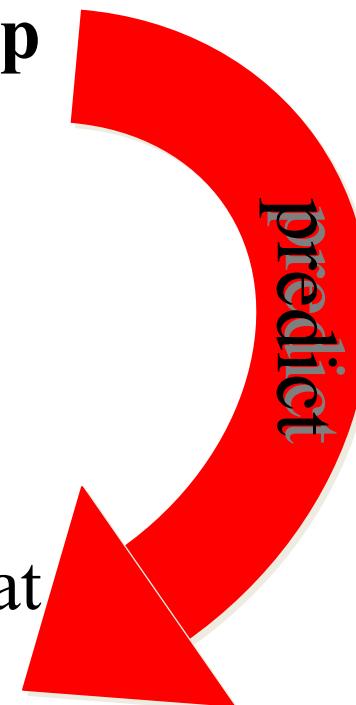
Tracking as inference

- Prediction:
 - Given the measurements we have seen up to this point, what state should we predict?

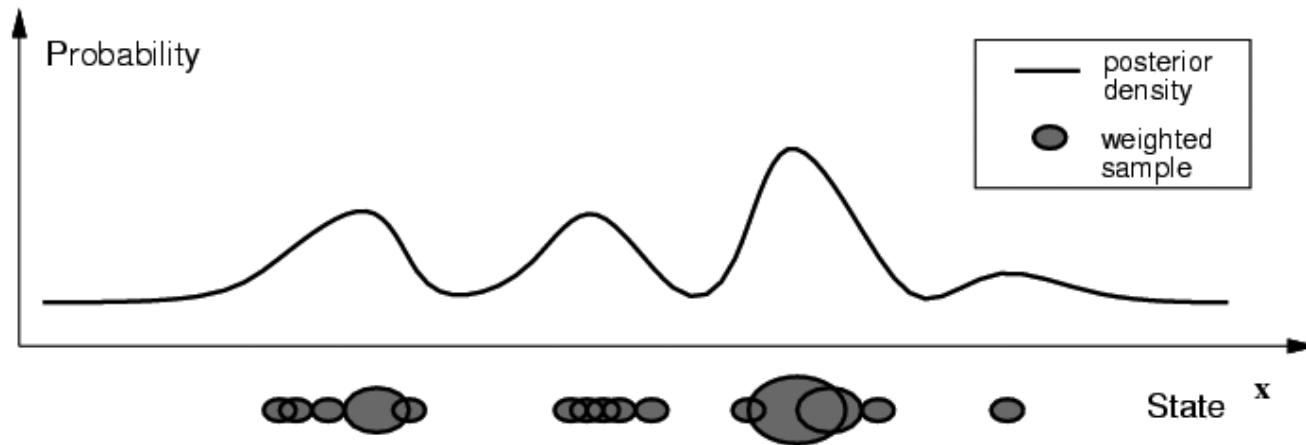
$$P(X_t | y_0, \dots, y_{t-1})$$

- Correction:
 - Now given the **current** measurement, what state should we predict?

$$P(X_t | y_0, \dots, y_t)$$



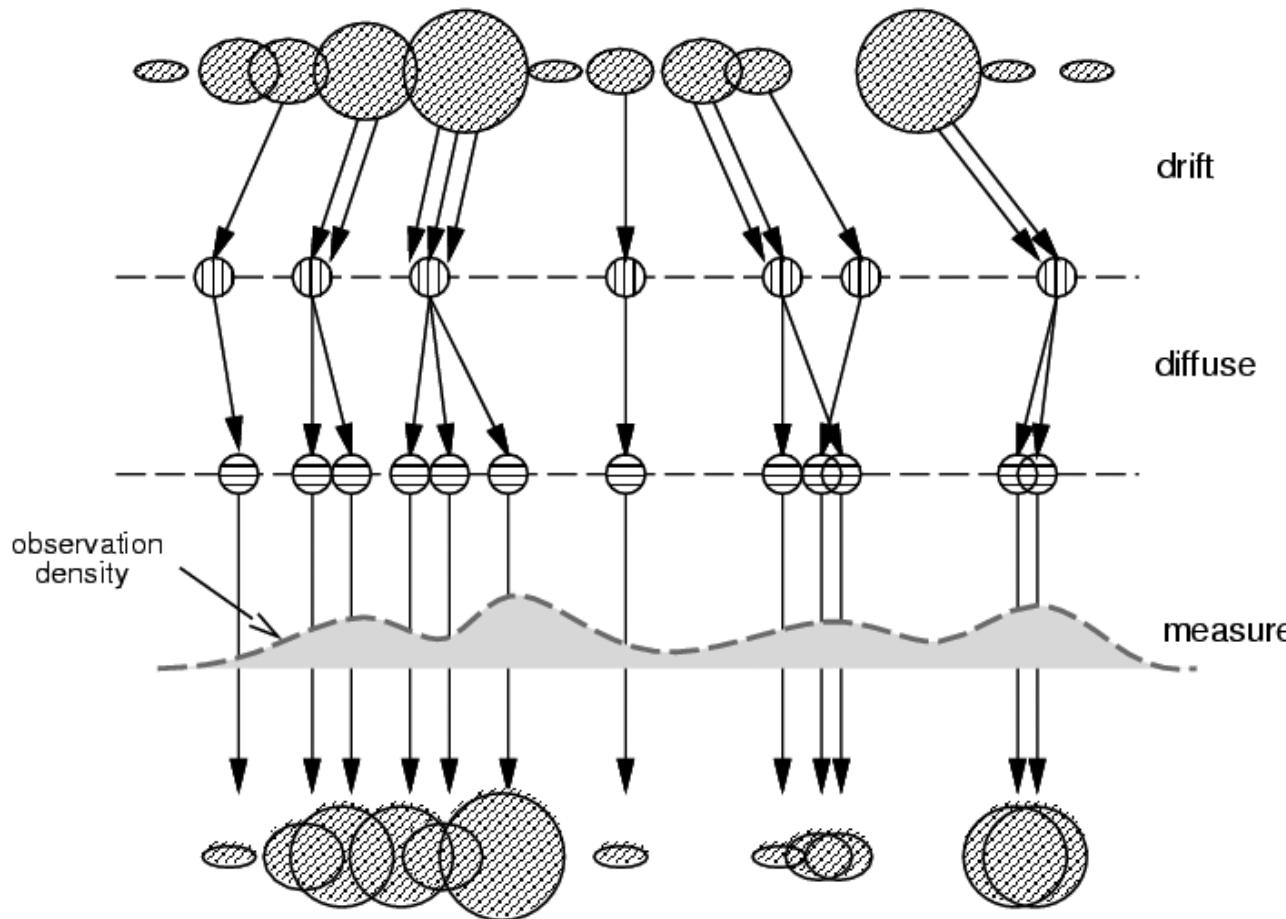
Particle filtering



Represent the state distribution non-parametrically

- Prediction: Sample possible values X_{t-1} for the previous state and shift according to the dynamics model
- Correction: Compute likelihood of X_t based on weighted samples and $P(y_t|X_t)$

Particle filtering



Start with weighted samples from previous time step

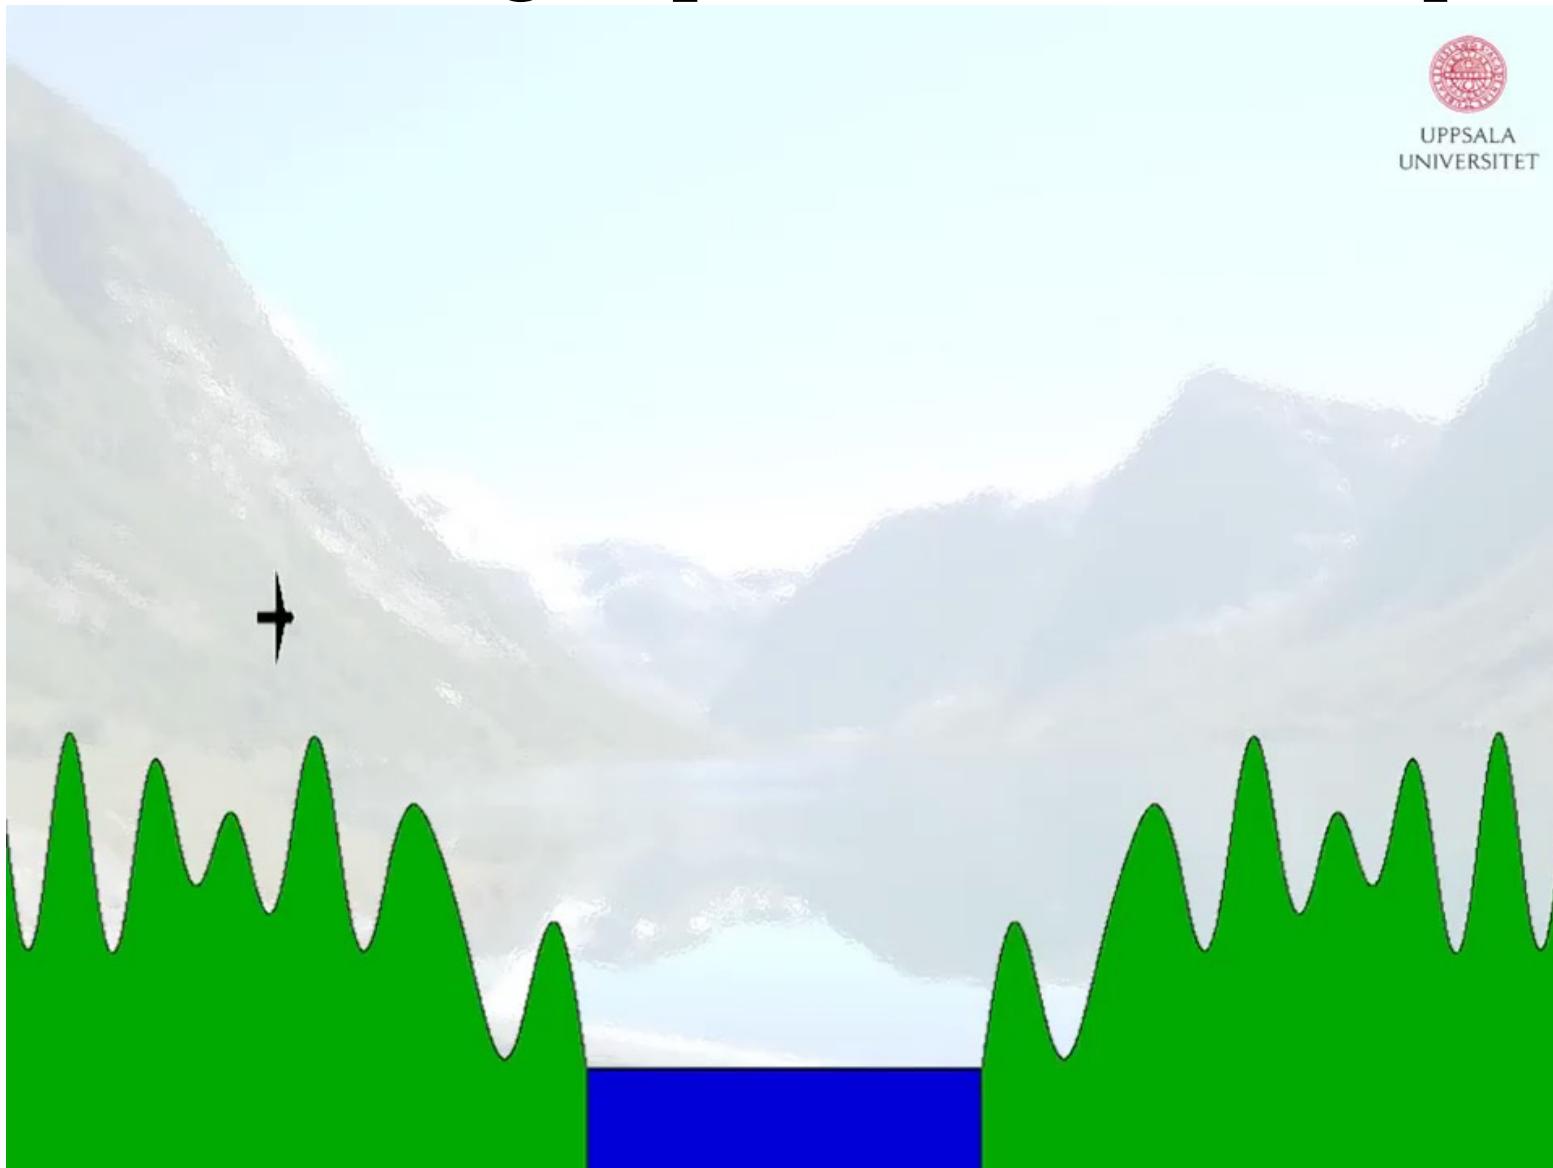
Sample and shift according to dynamics model

Spread due to randomness; this is predicted density $P(X_t | Y_{t-1})$

Weight the samples according to observation density

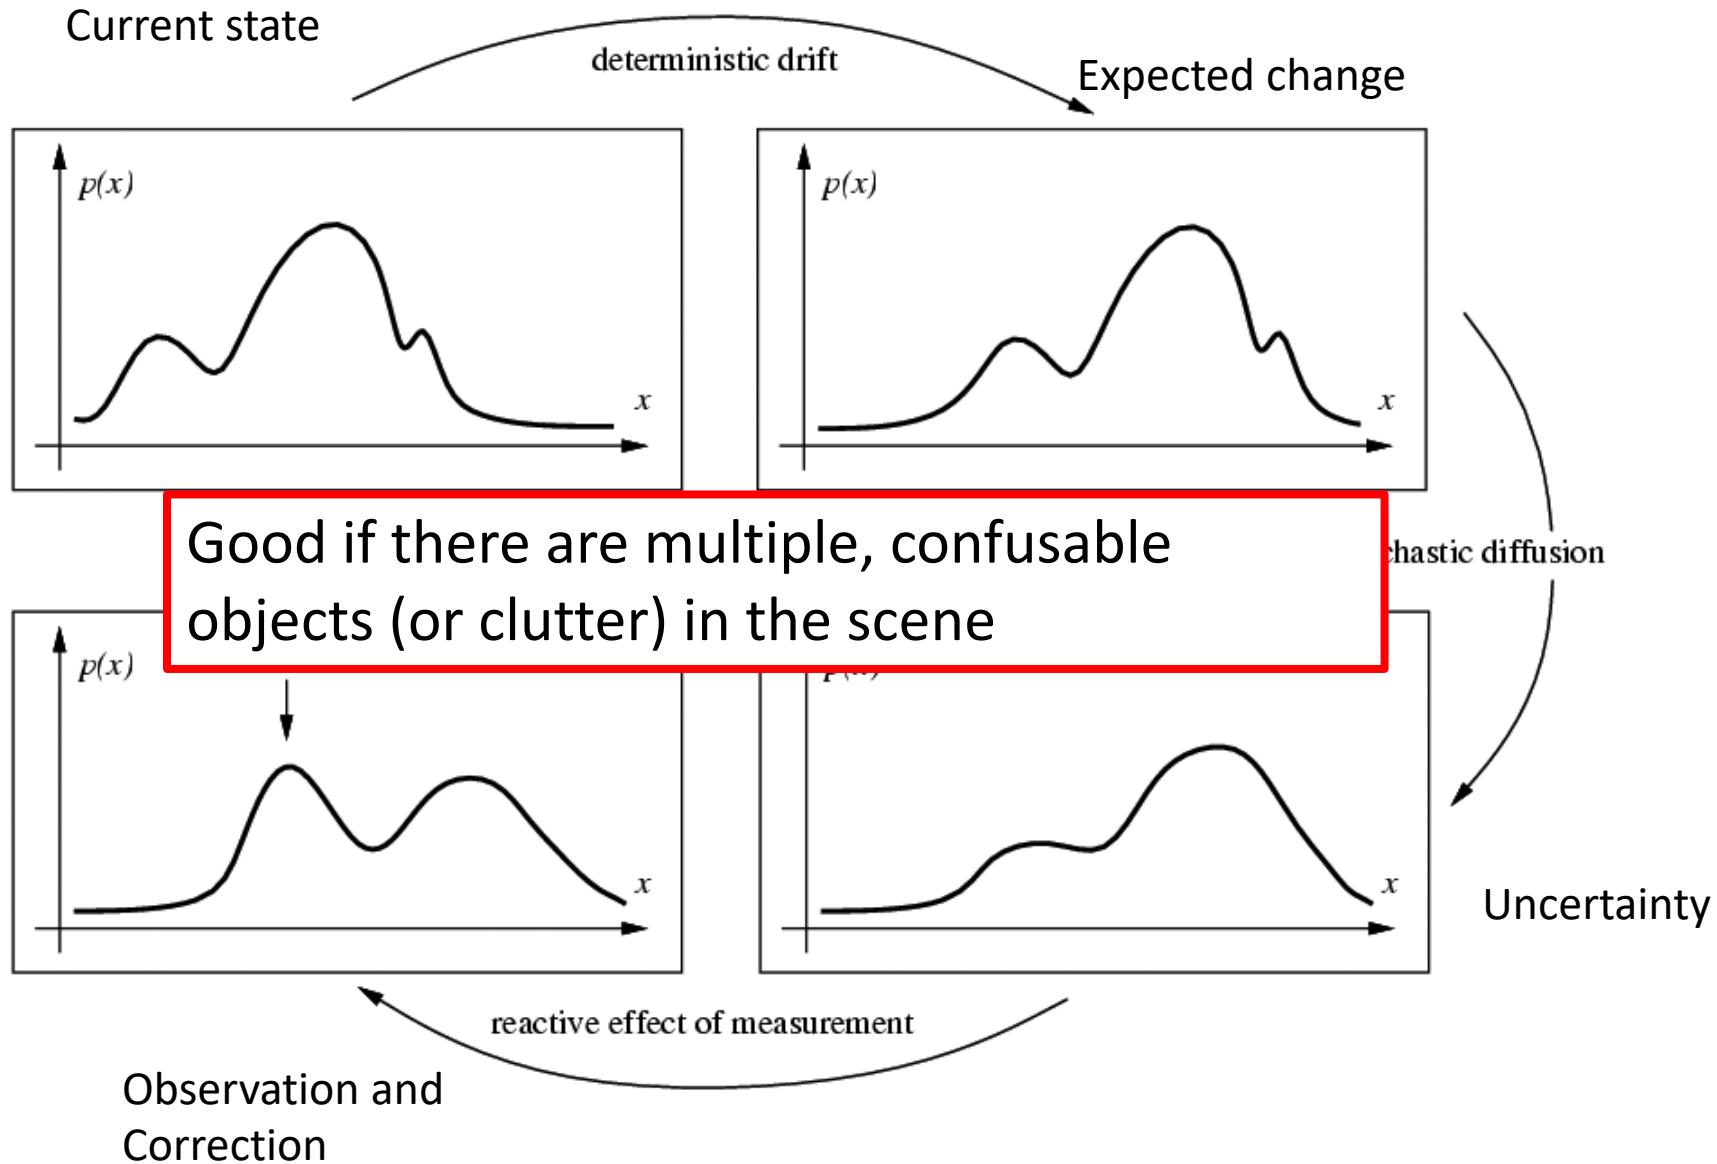
Arrive at corrected density estimate $P(X_t | Y_t)$

Particle filtering explained without equations



<https://www.youtube.com/watch?v=aUkBa1zMkv4>

Propagation of non-parametric densities



Particle filtering results

**Particle Filter combined with
Gentle Adaboost Online
Classifier**

**Yellow box = PF Estimate
Green box = Ground Truth
Blue = the particles**

Particle filtering results



<https://www.youtube.com/watch?v=tljuflnUqZM>

Tracking issues

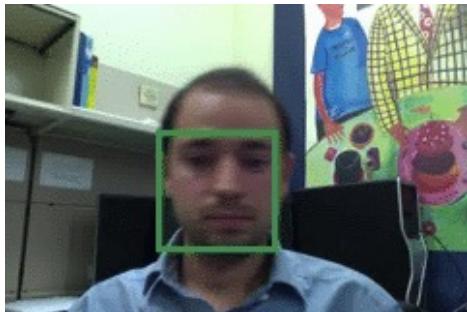
- Initialization
 - Manual
 - Background subtraction
 - Detection

Tracking issues

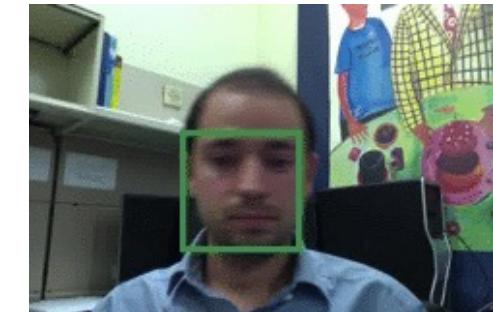
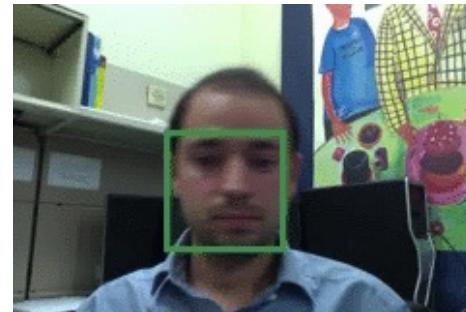
- Initialization
- Getting observation and dynamics models
 - Observation model: match a template or use a trained detector
 - Dynamics model: usually specify using domain knowledge

Tracking issues

- Initialization
- Obtaining observation and dynamics model
- Uncertainty of prediction vs. correction
 - If the dynamics model is too strong, will end up ignoring the data
 - If the observation model is too strong, tracking is reduced to repeated detection



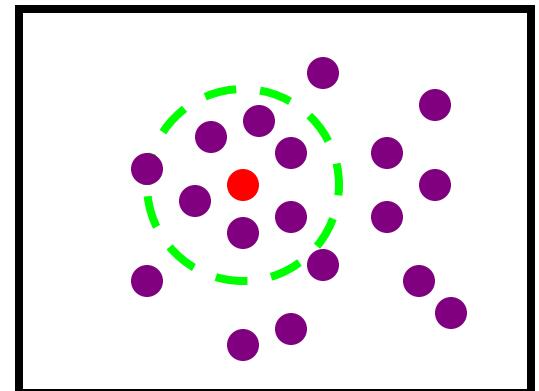
Too strong dynamics model



Too strong observation model

Tracking issues

- Initialization
- Getting observation and dynamics models
- Prediction vs. correction
- Data association
 - When tracking multiple objects, need to assign right objects to right tracks (particle filters good for this)



Tracking issues

- Initialization
- Getting observation and dynamics models
- Prediction vs. correction
- Data association
- Drift
 - Errors can accumulate over time

Drift



D. Ramanan, D. Forsyth, and A. Zisserman. [Tracking People by Learning their Appearance](#). PAMI 2007.